

# Introduction to Computer Vision

---

**Kaveh Fathian**

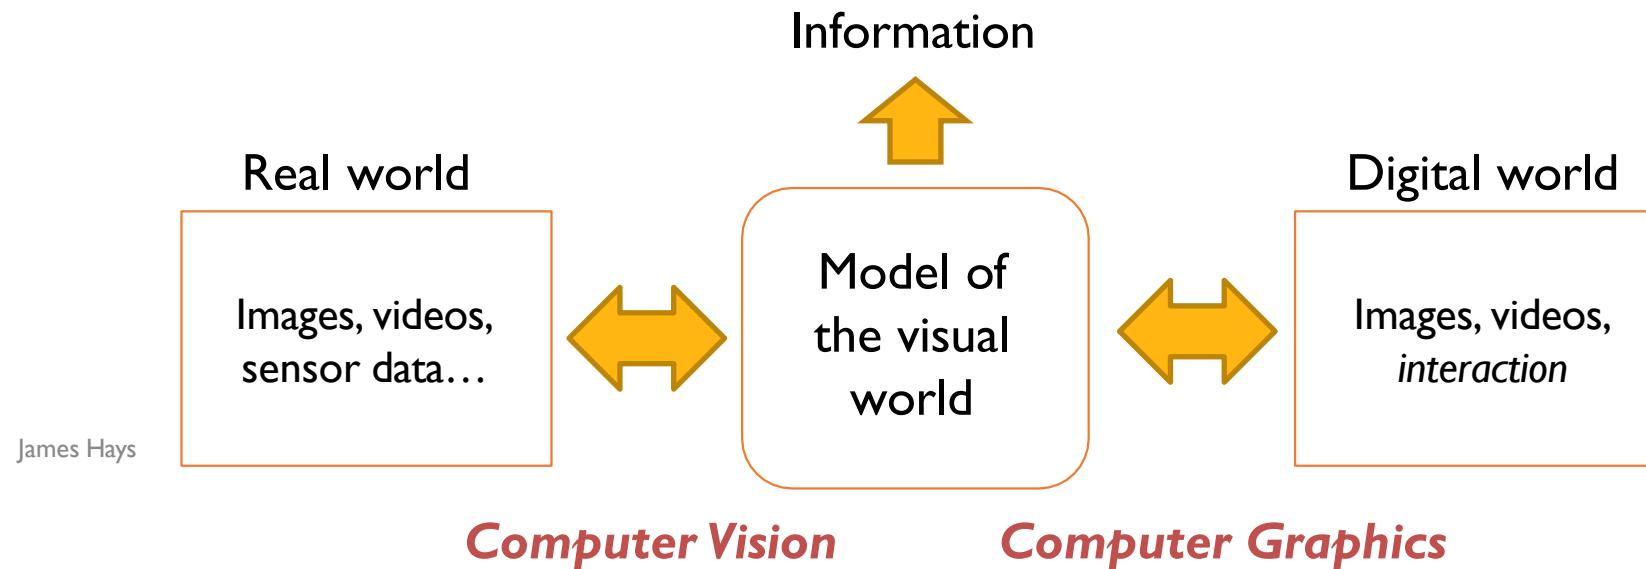
Assistant Professor

Computer Science Department

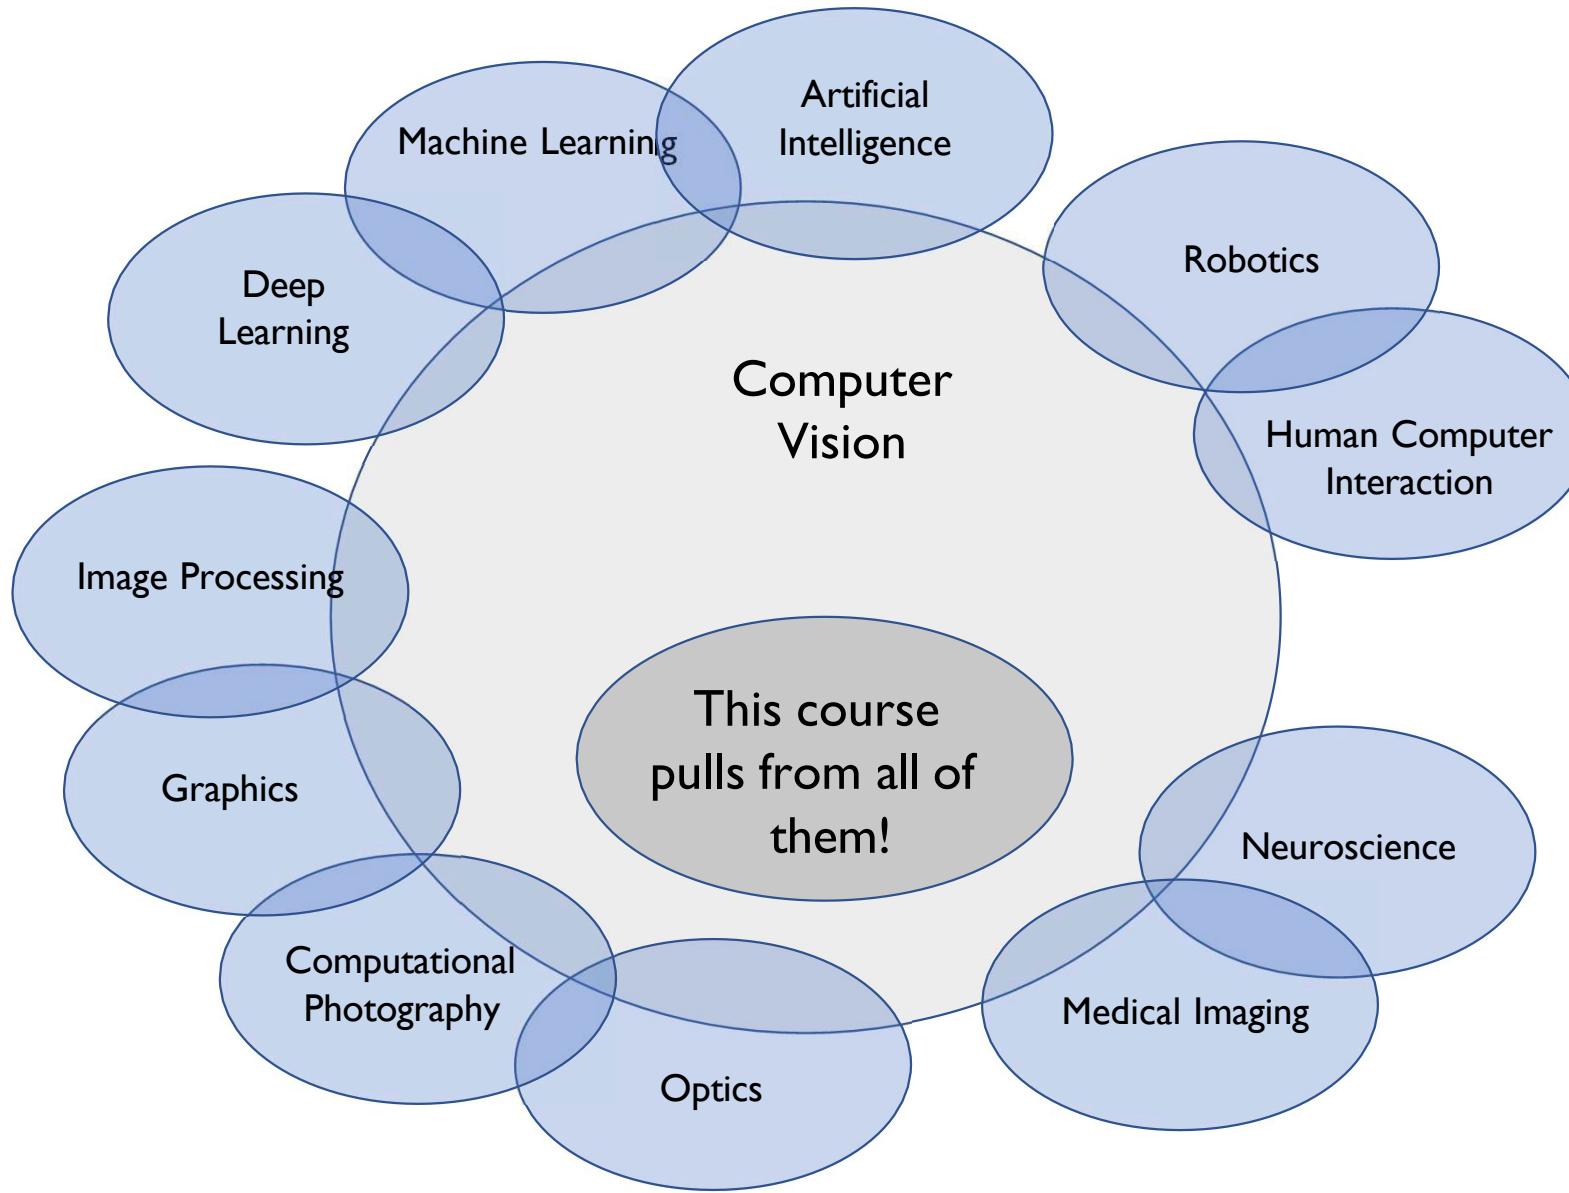
Colorado School of Mines

**Lecture 2**

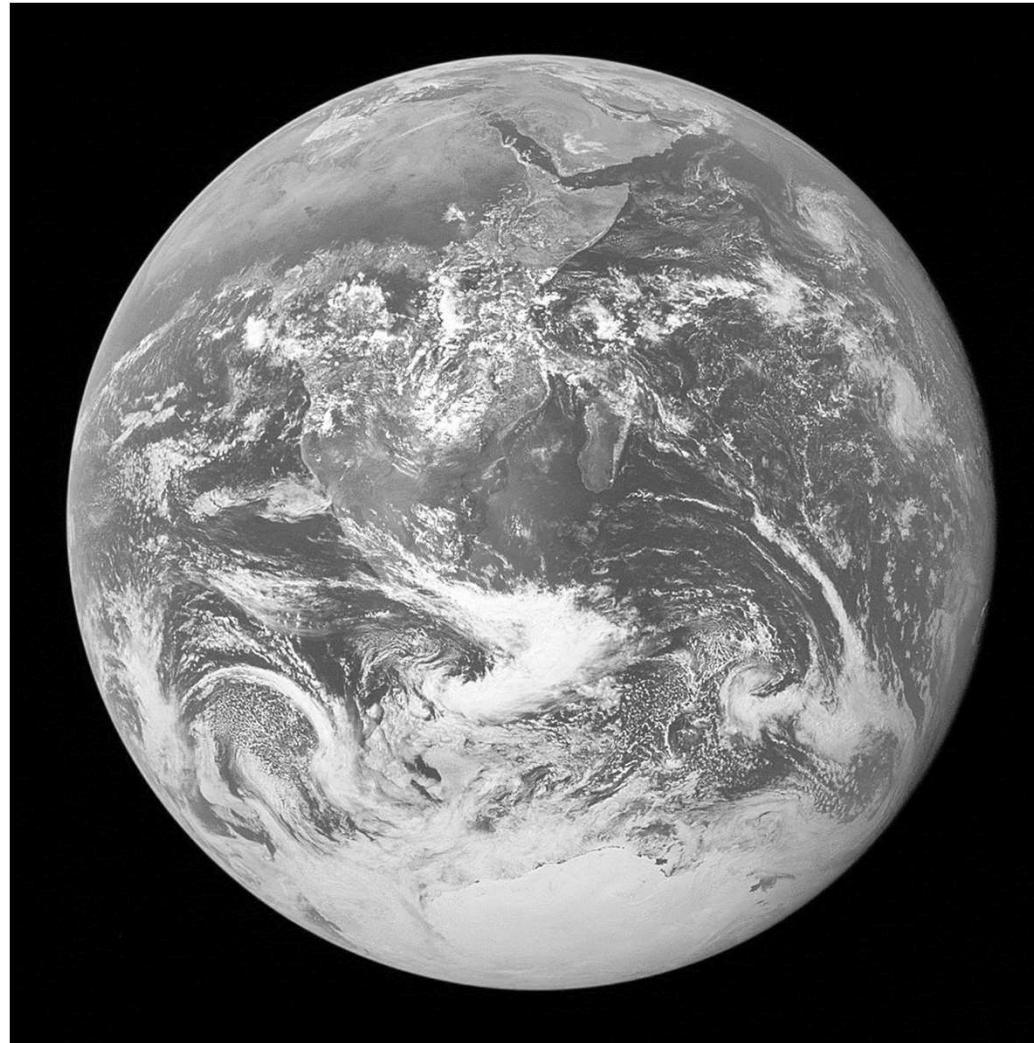
# Computer Vision and Nearby Fields



# Scope of Computer Vision

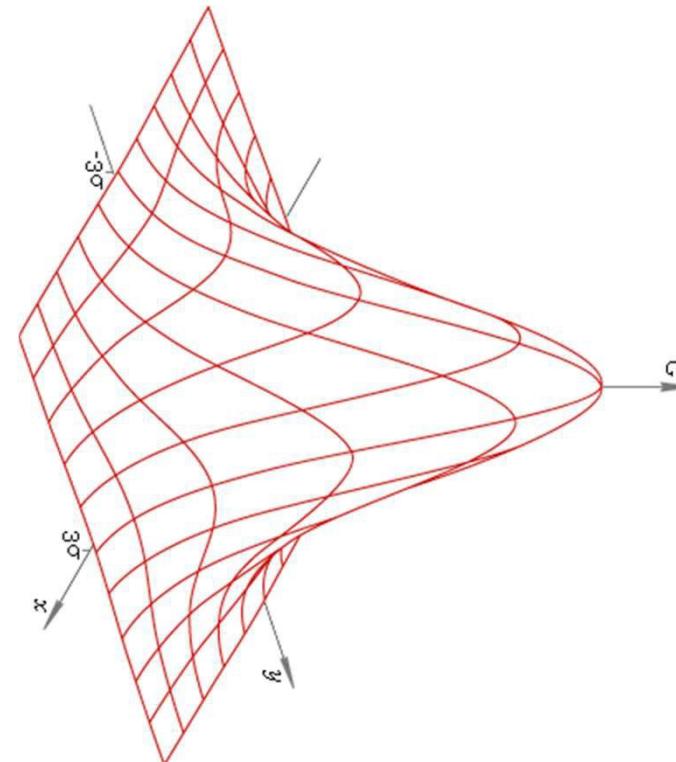
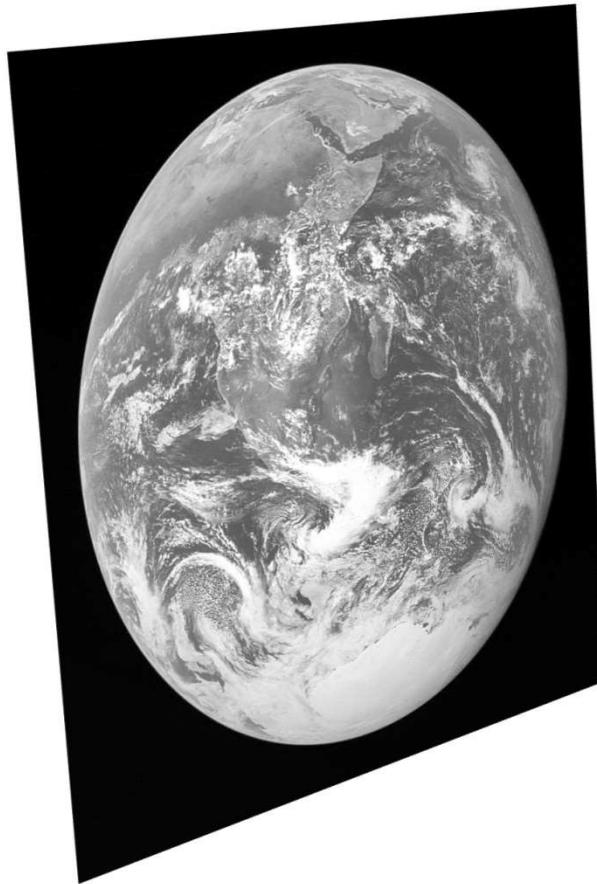


# Image



Blue Marble. NASA | Apollo 17

# Image

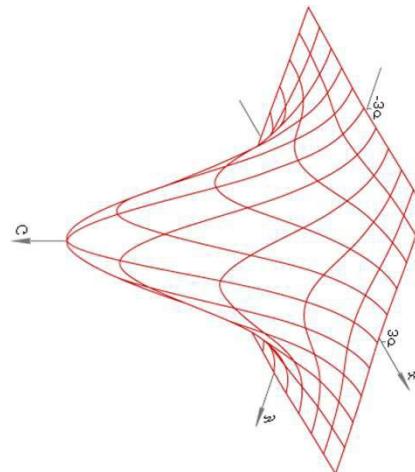
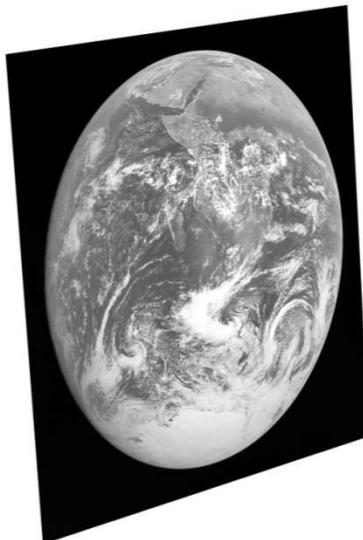


Blue Marble. NASA | Apollo 17

# Signal

## Definition

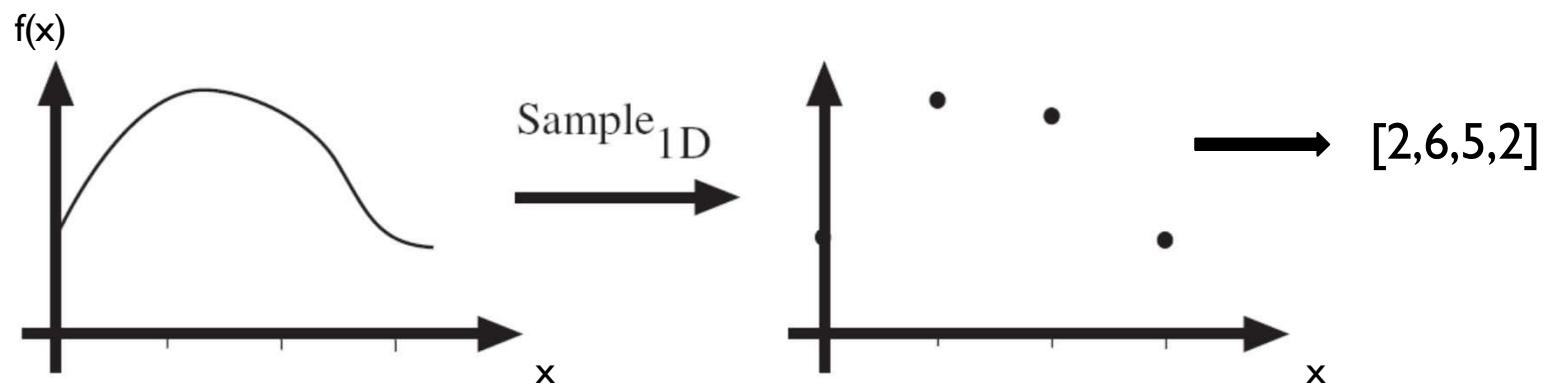
**Signal:** A (multi-dimensional) function that contains information about a phenomenon.



- Signals can be
  - Continuous: light
  - Discrete: measurement of a light
  - Sampling: reduction of continuous signal to a discrete signal
- Any phenomenon
  - Light
  - Heat
  - Gravity

# Sampling

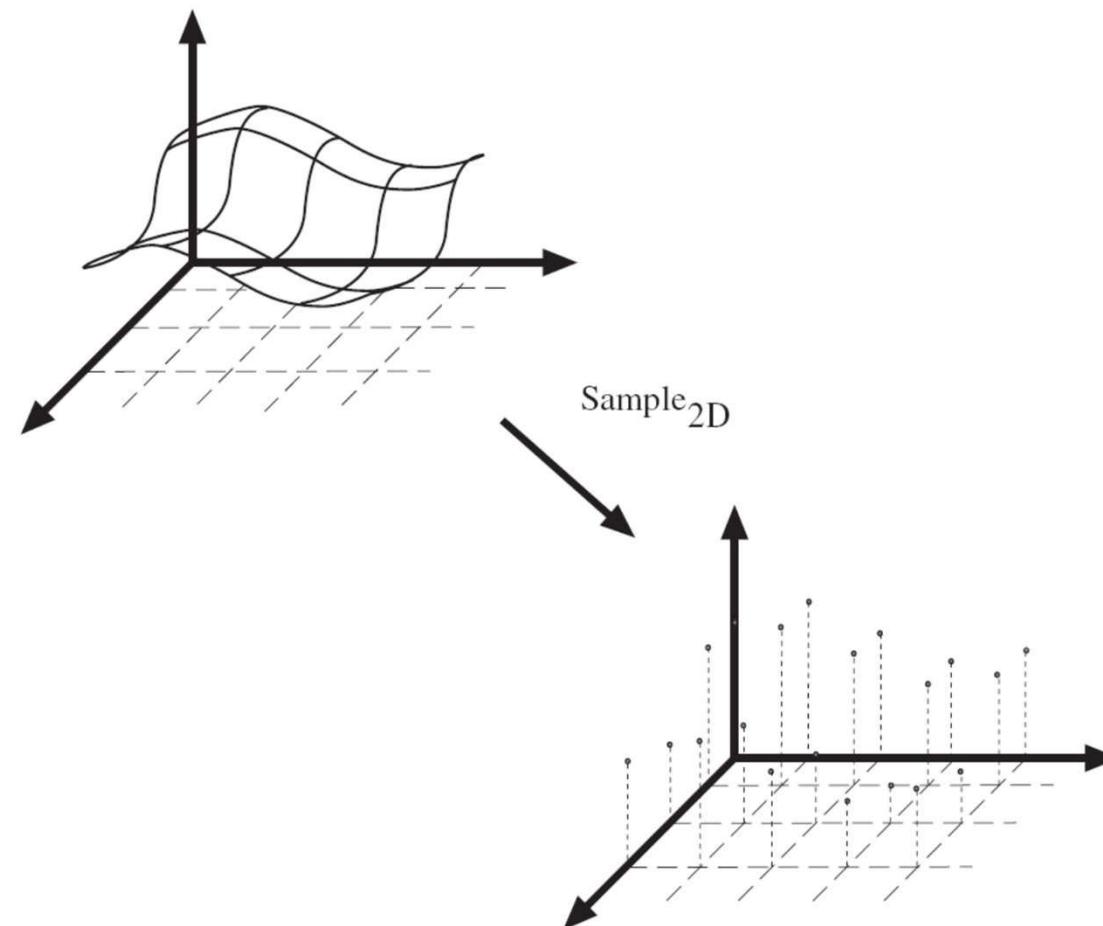
- Sampling in 1D takes a function and returns a vector whose elements are values of that function at the sample points.



Danny Alexander

# Sampling

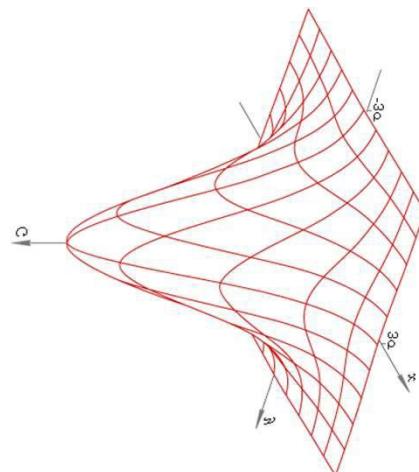
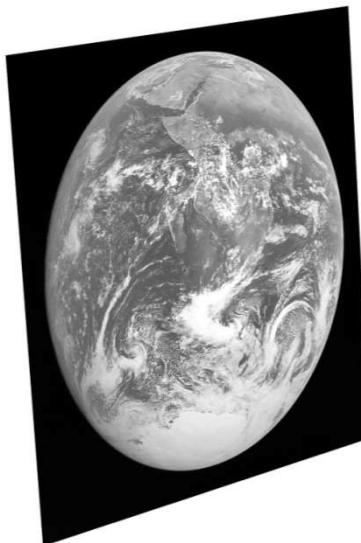
- Sampling in 2D takes a function and returns a matrix



# 2D Image

## Definition

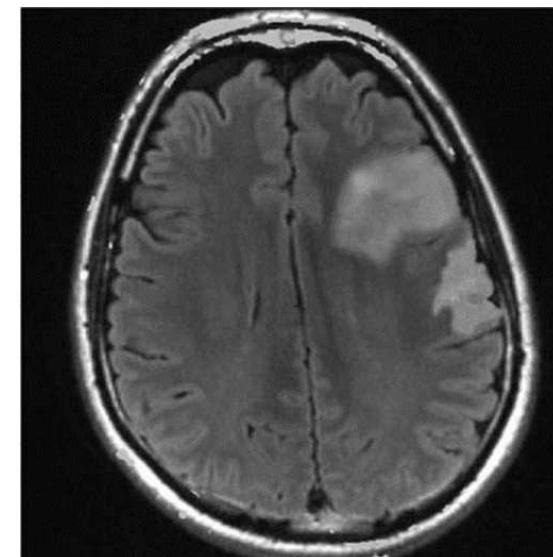
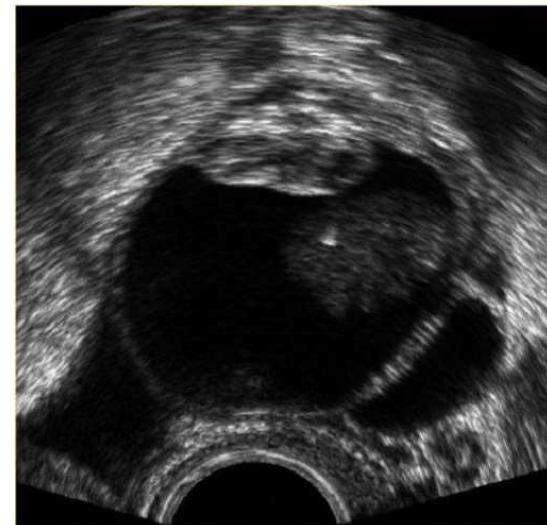
**Image:** A sampling of a function that contains information about a 2D\* signal.



\* or a 2D projection of a multi-dimensional signal

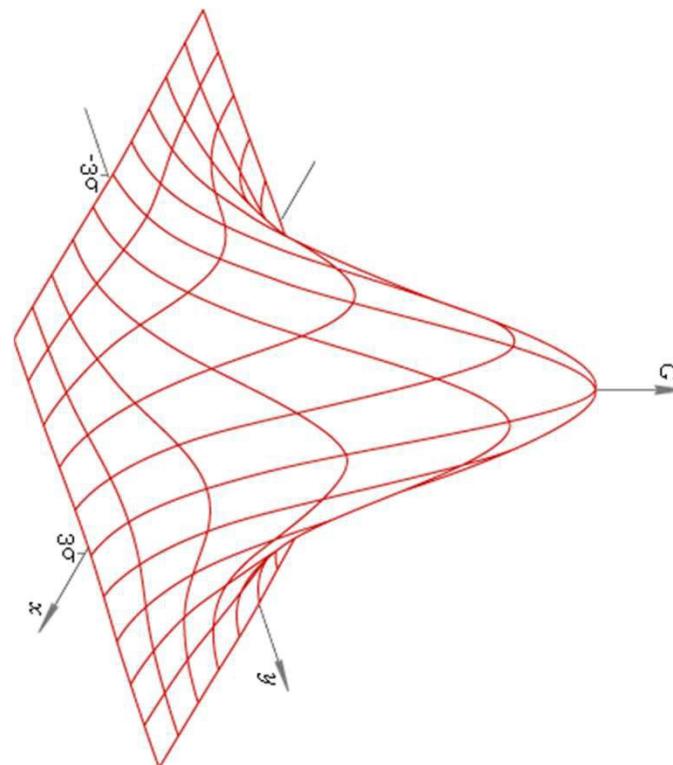
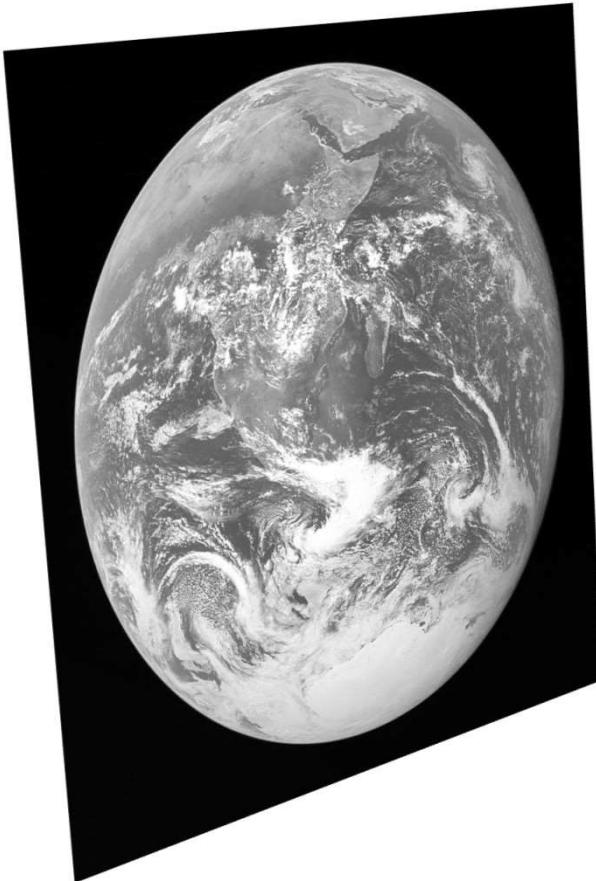
- Function stores ‘brightness’
- 2D signals are special for us
  - Brightness along x and y dimensions
- Video: xy-coordinates + time
  - Time-varying 2D signal

# Example of 2D Images

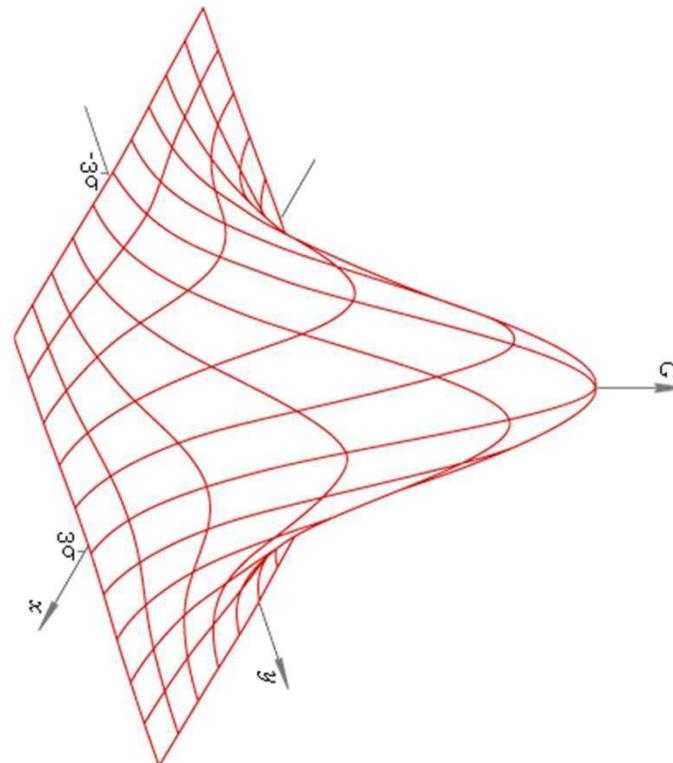
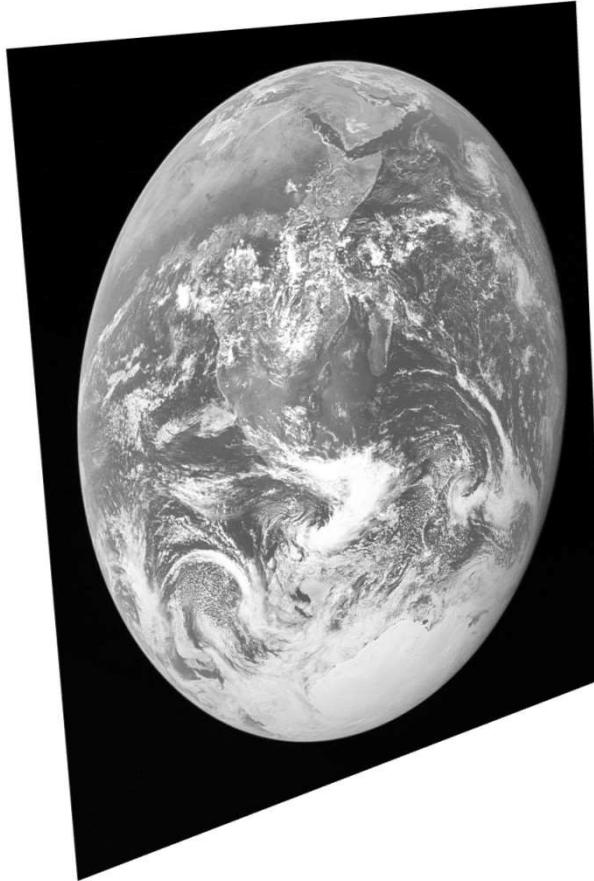


[http://radiologykey.com/wp-content/uploads/2016/01/9781604063325\\_c014\\_f004.jpg](http://radiologykey.com/wp-content/uploads/2016/01/9781604063325_c014_f004.jpg)

# Sampling in Practice

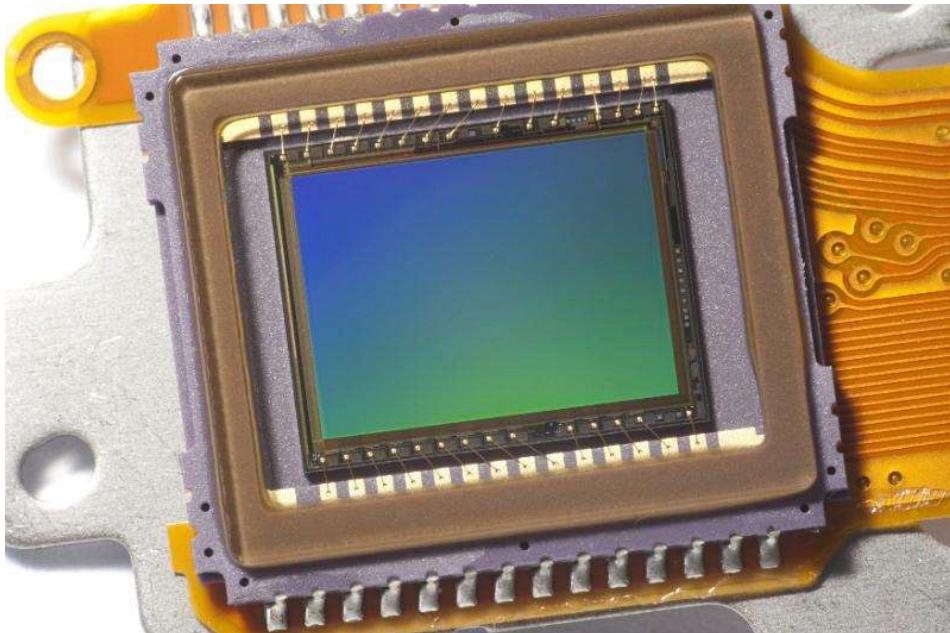


# Sampling in Practice: Digital Image

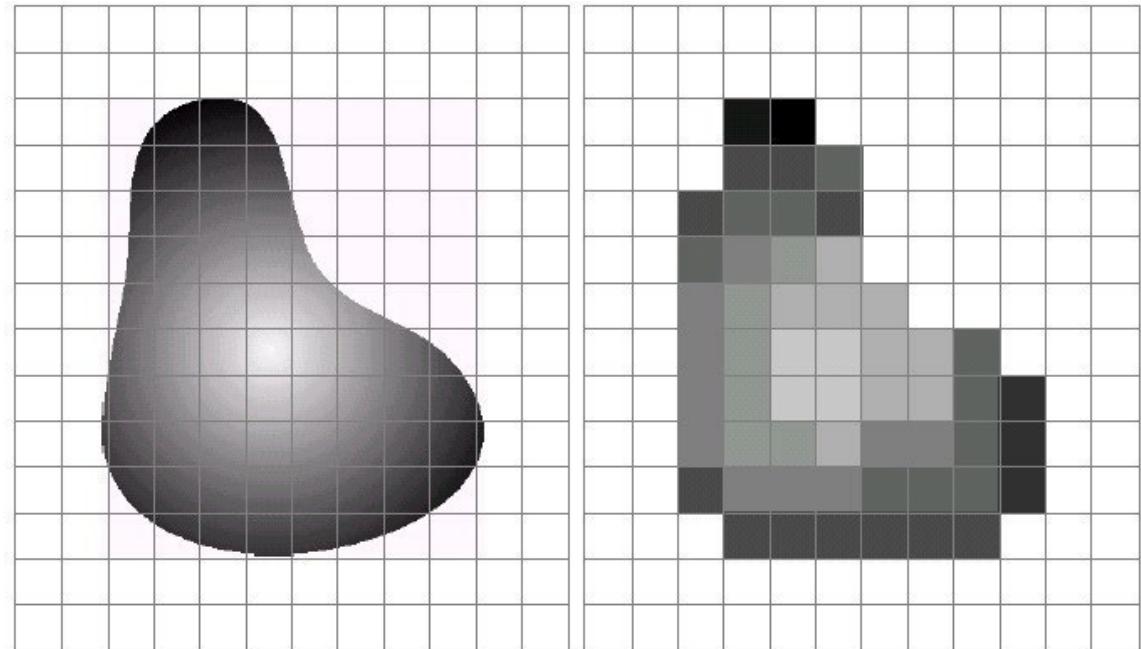


CMOS / CCD

# Sensor Array



CMOS sensor



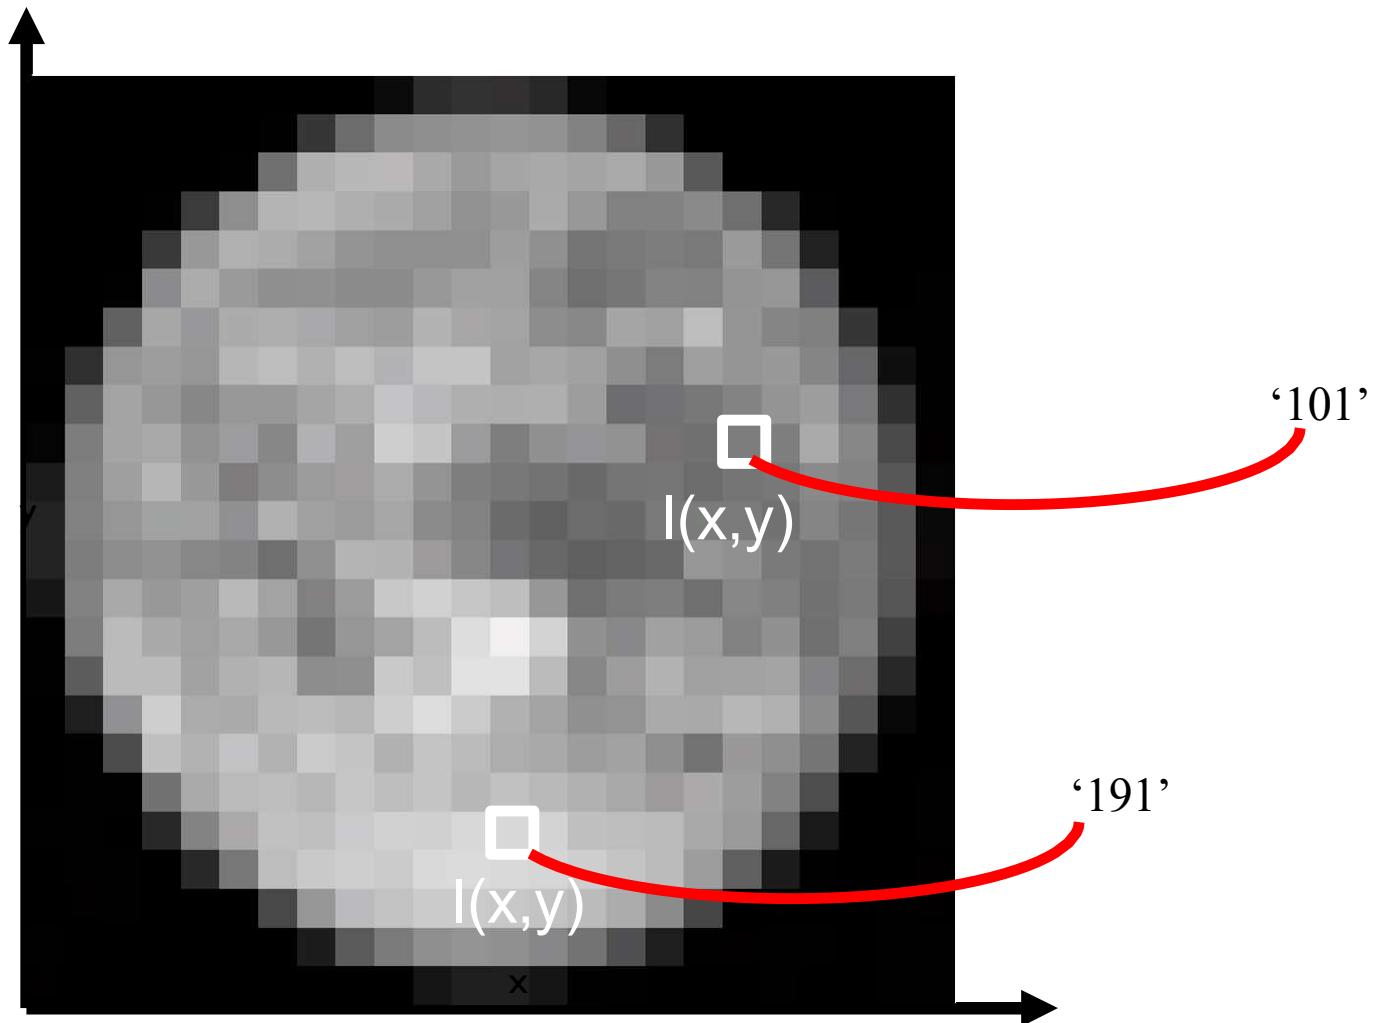
a b

**FIGURE 2.17** (a) Continuous image projected onto a sensor array. (b) Result of image sampling and quantization.

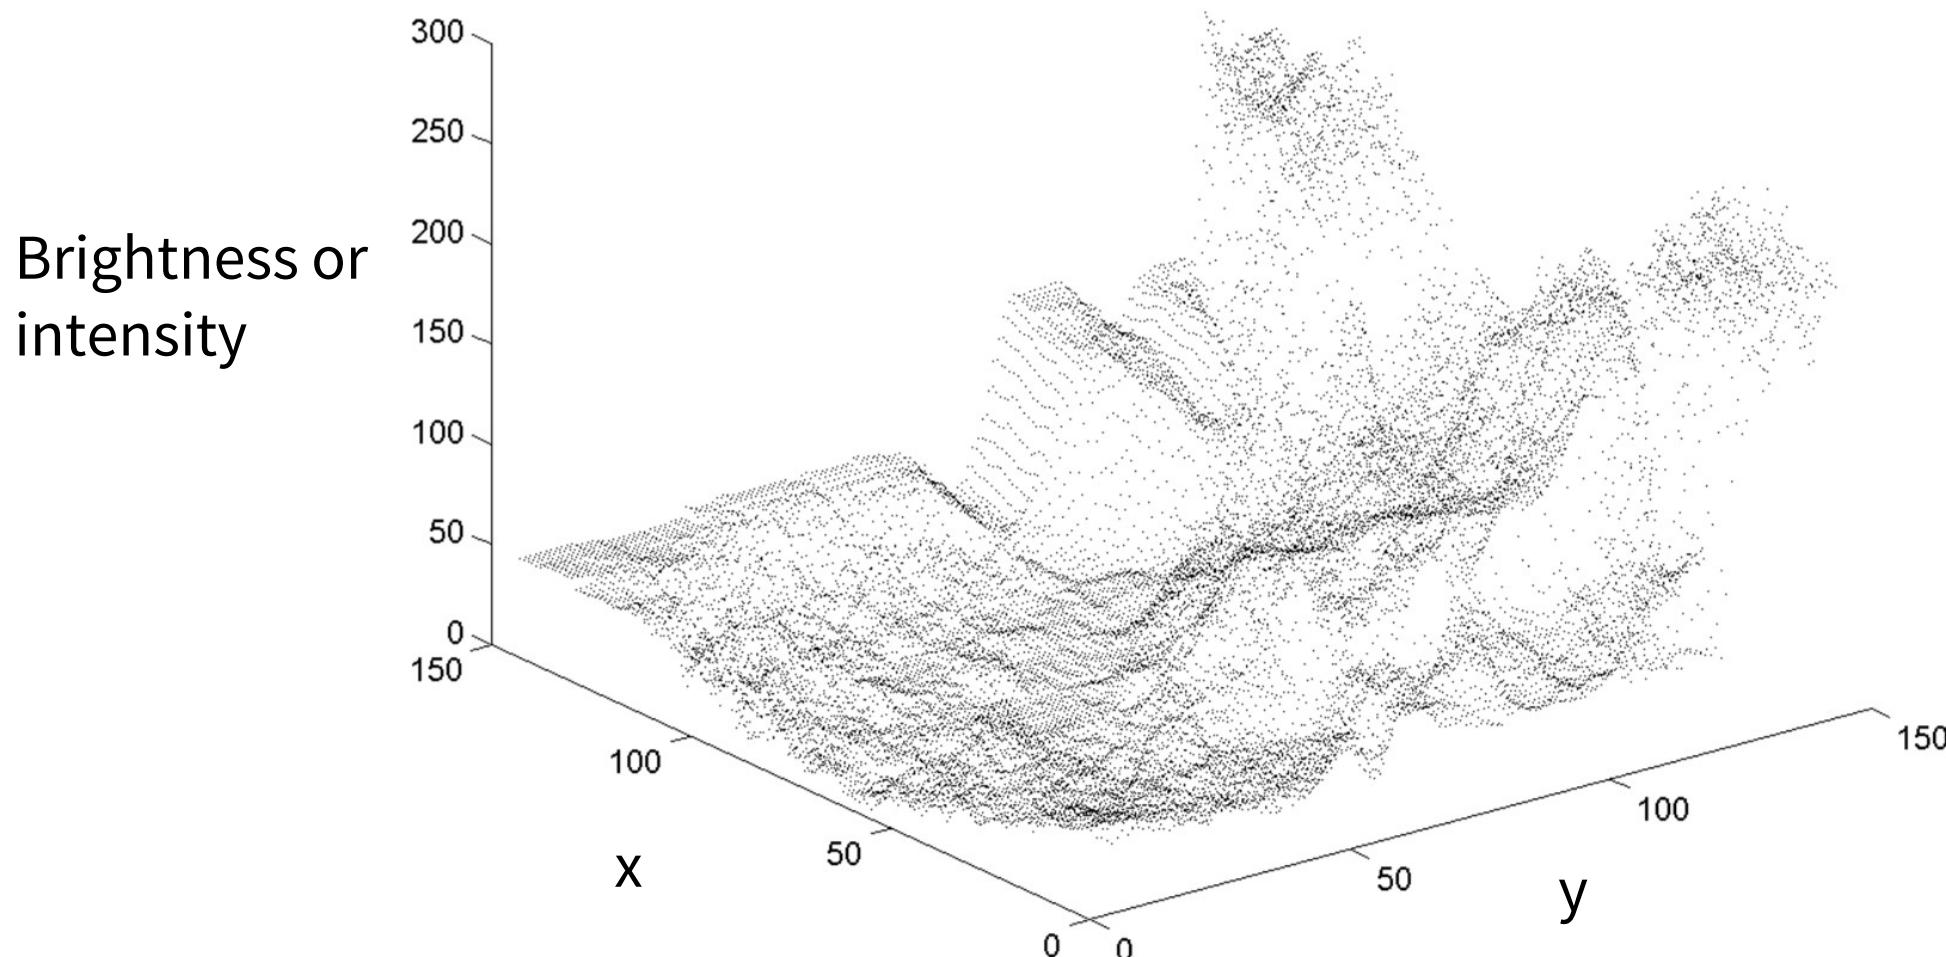
James Hays

# Elements of a Digital Image

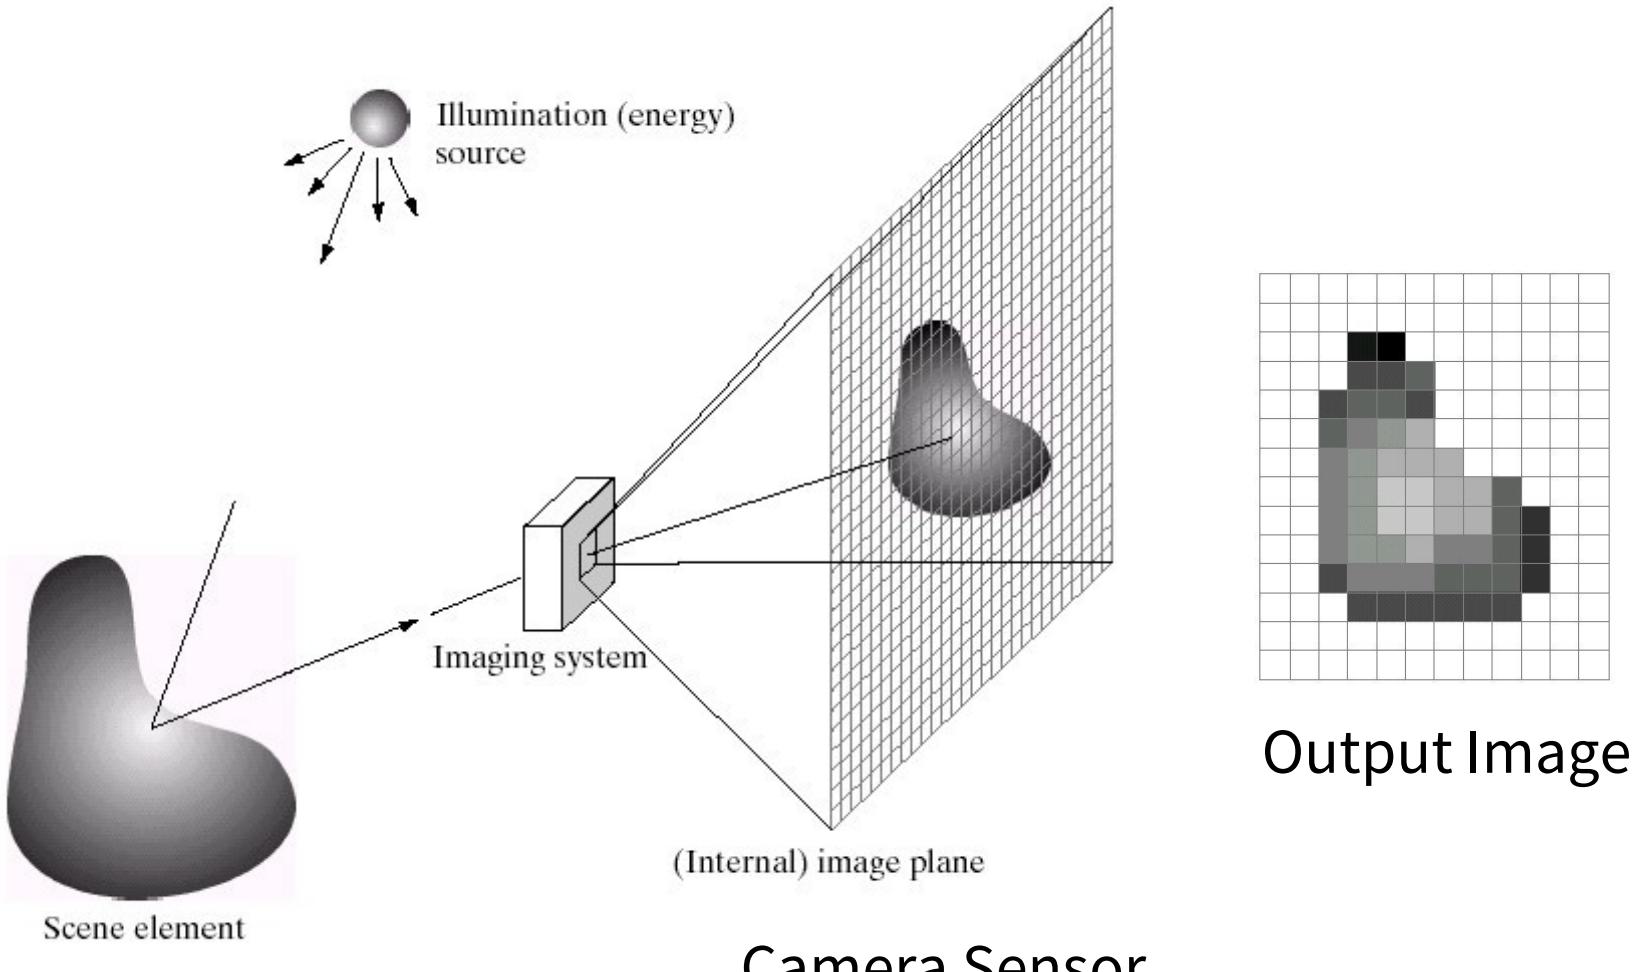
**Pixel:** picture element



# Digital Image is a 2D Signal

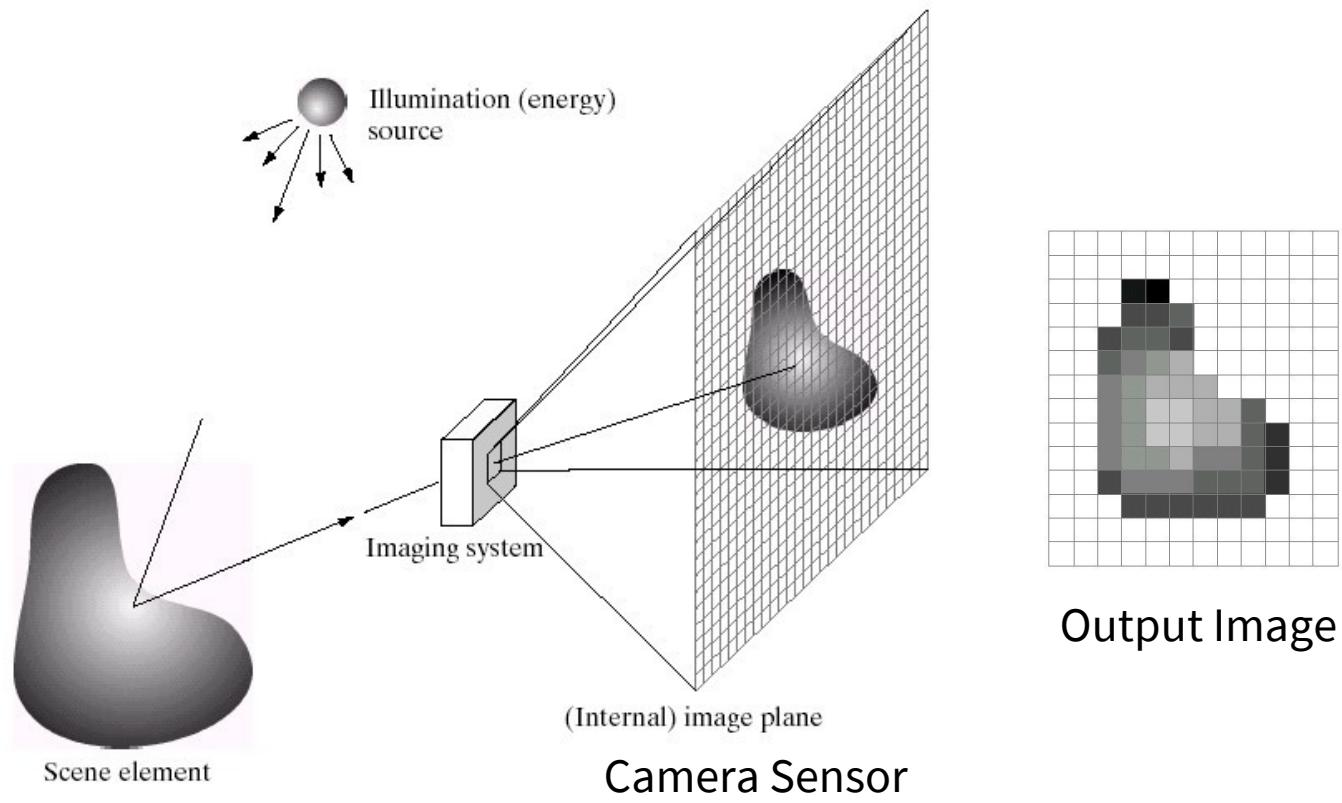


# Light Integration Over the “Frustum”

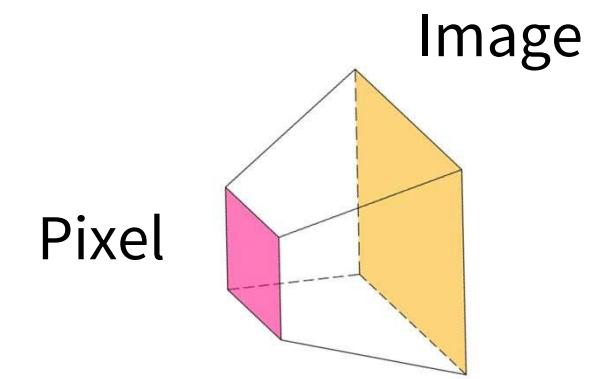


James Hays

# Light Integration Over the “Frustum”



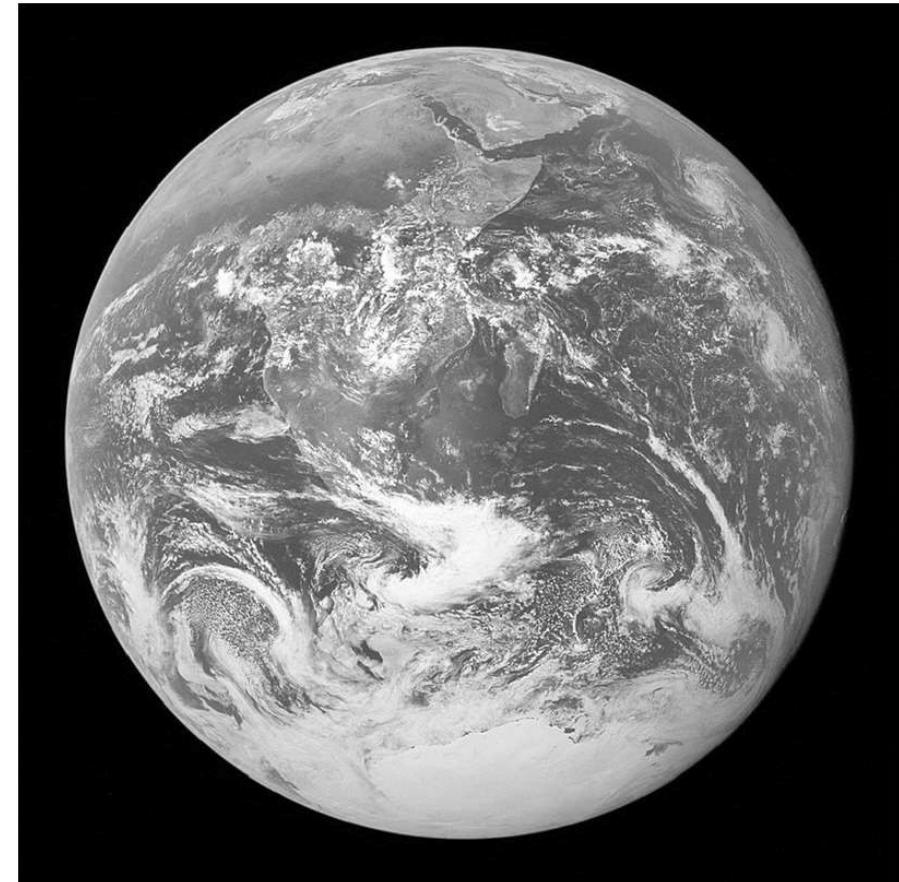
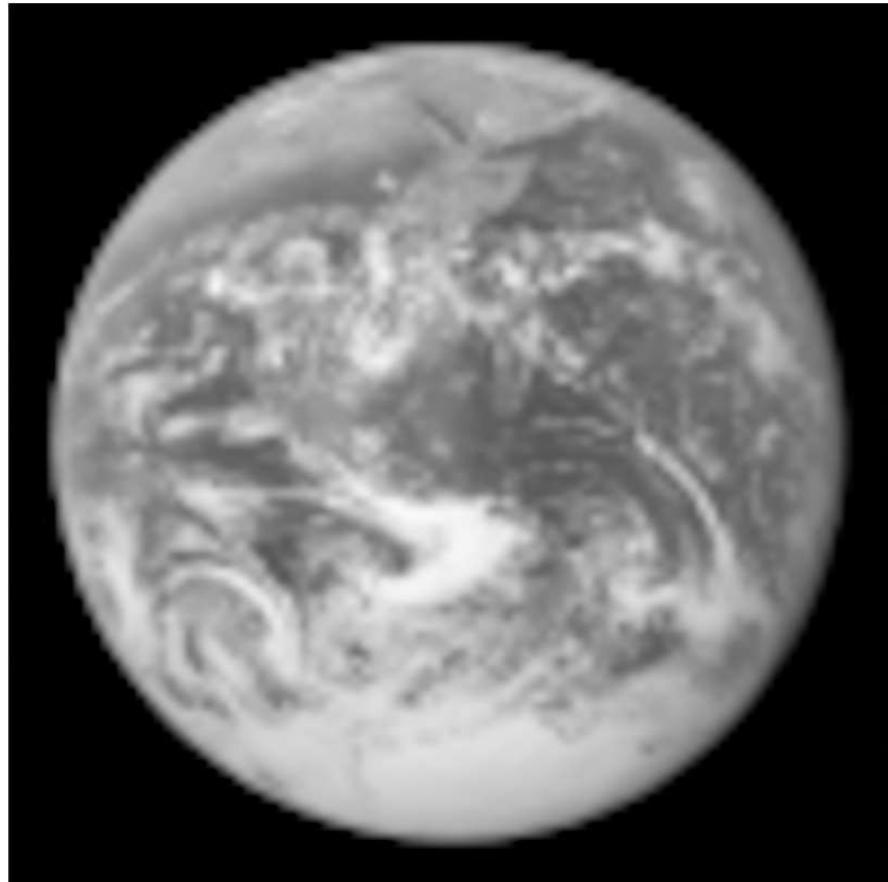
James Hays



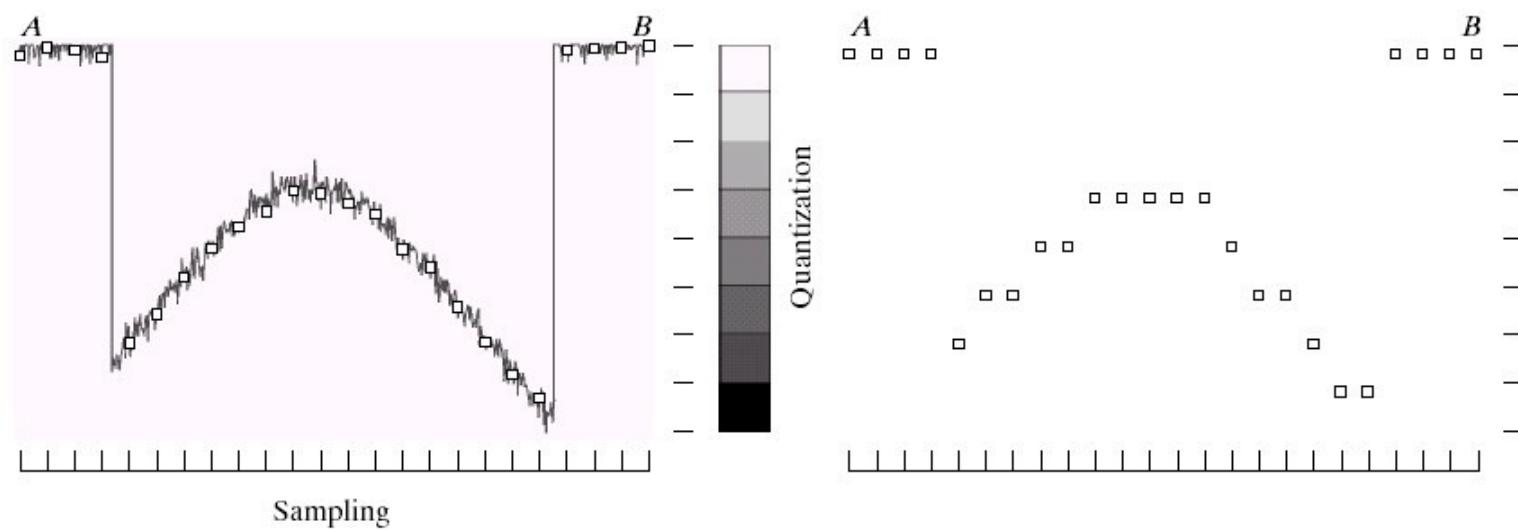
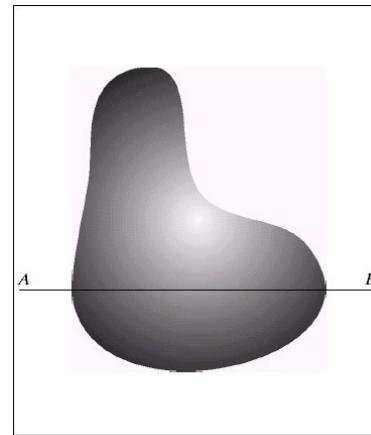
A frustum approximates the ray space that a pixel samples

# Resolution: geometric vs. spatial

Both images are 1000x1000 pixels:



# Quantization



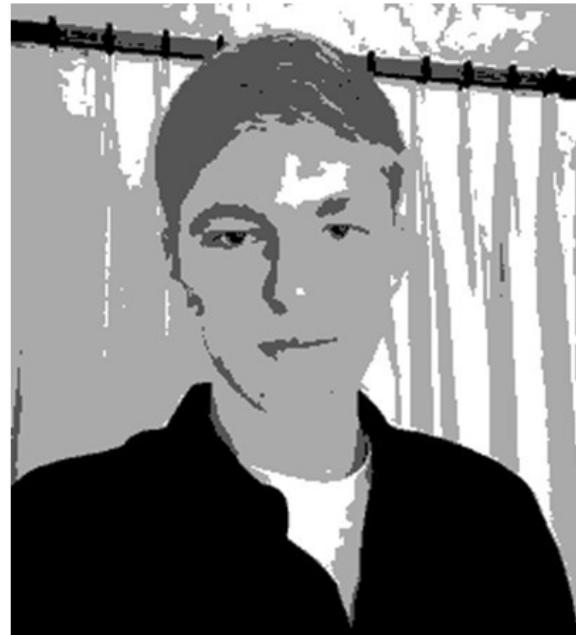
# Quantization Effects – Radiometric Resolution



8 bit – 256 levels



4 bit – 16 levels



2 bit – 4 levels



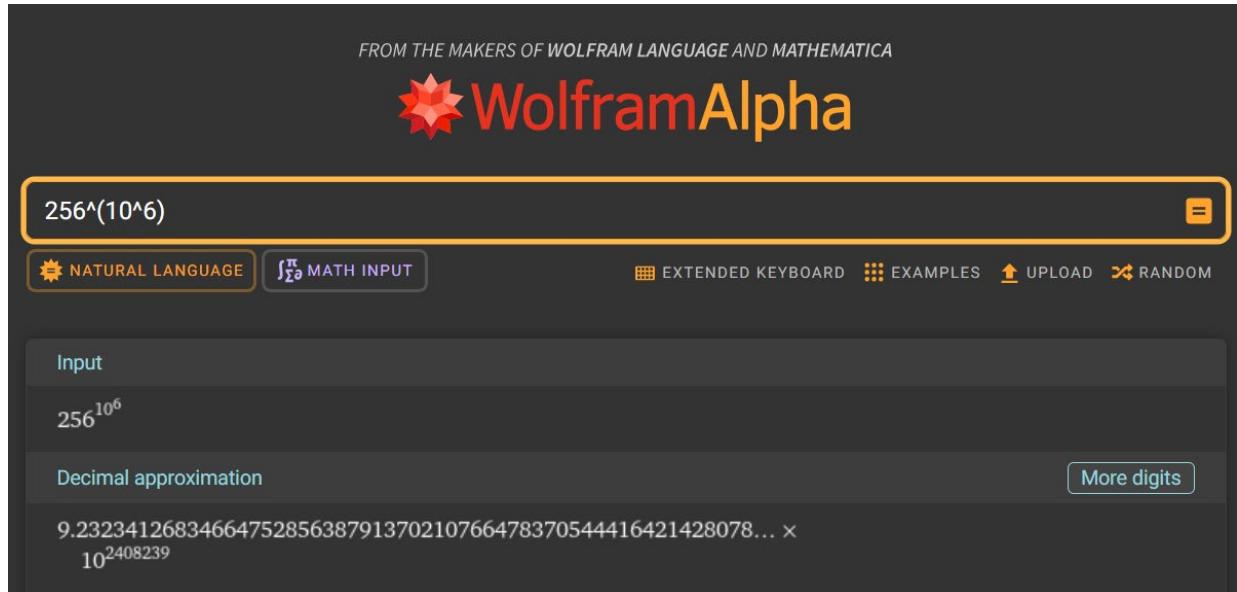
1 bit – 2 levels

We often call this ***bit depth***.

For photography, this is also related to ***dynamic range***.

# Dimensionality of an Image

- An image of size 1000x1000 with 8-bit quantization per pixel
  - $= 256 \text{ values}^{\wedge} (1000 \times 1000)$
- All images ever created
  - of size 1000x1000



\* it is estimated that there may be  $10^{78}$  to  $10^{82}$  atoms in the known universe!

- Computer vision as making sense of an extremely high-dimensional data
  - 'natural' images are a subspace

# Images in Python (import numpy)

$N \times M$  grayscale image “im”

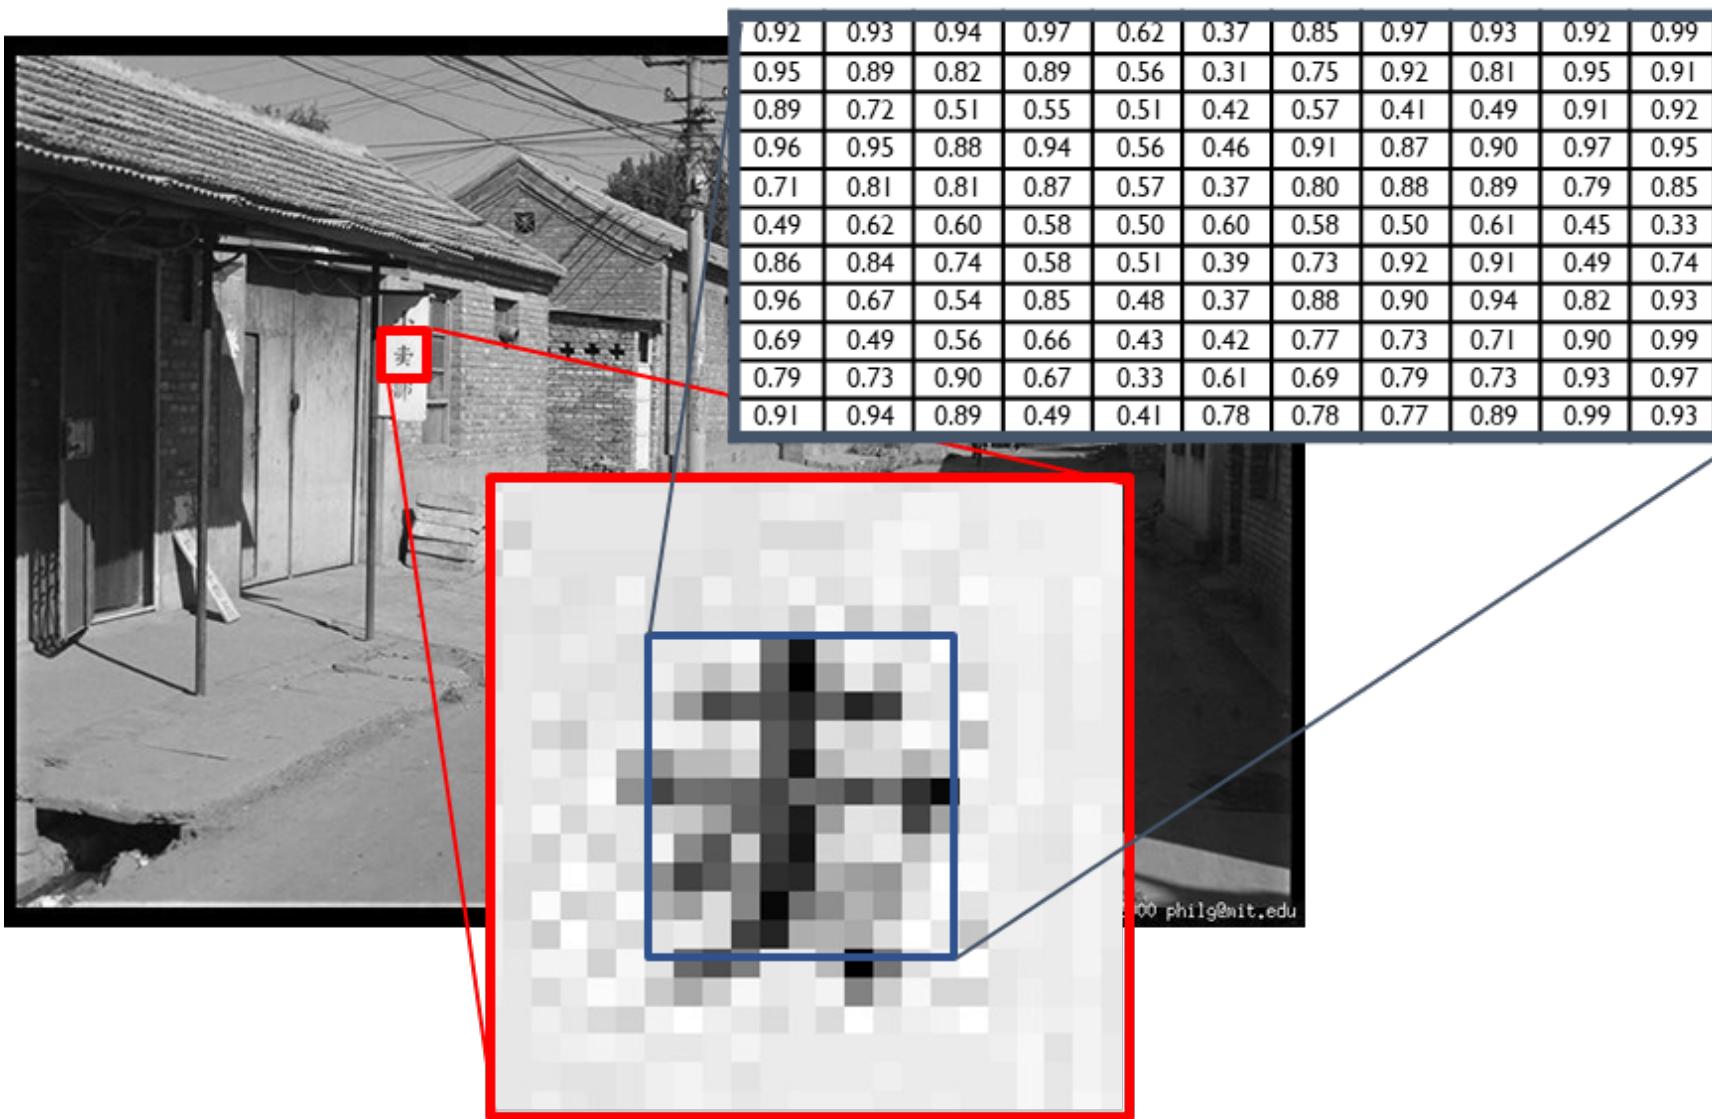
- $im[0, 0]$  = top-left pixel value
- $im[y, x]$  =  $y$  pixels down,  $x$  pixels to right
- $im[N-1, M-1]$  = bottom-right pixel

Row  Column 

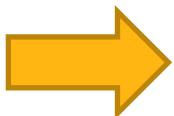
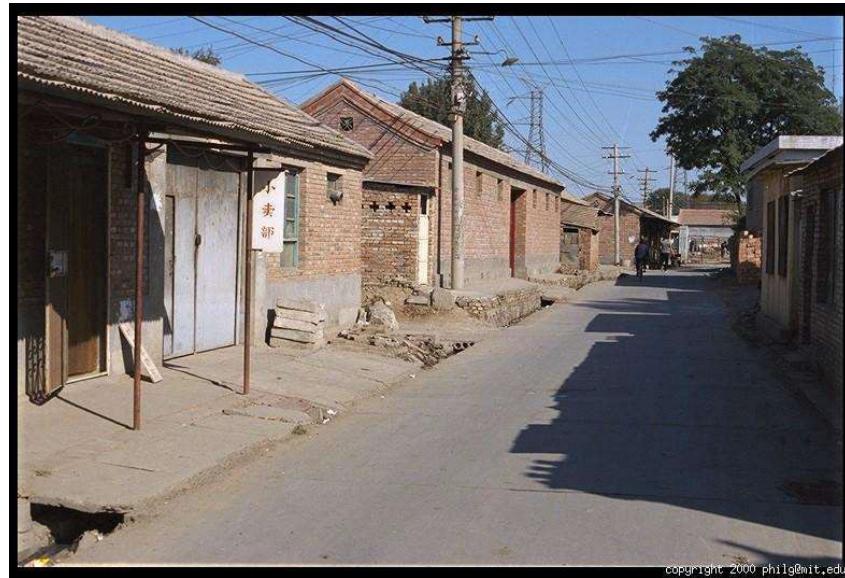
0.92	0.93	0.94	0.97	0.62	0.37	0.85	0.97	0.93	0.92	0.99
0.95	0.89	0.82	0.89	0.56	0.31	0.75	0.92	0.81	0.95	0.91
0.89	0.72	0.51	0.55	0.51	0.42	0.57	0.41	0.49	0.91	0.92
0.96	0.95	0.88	0.94	0.56	0.46	0.91	0.87	0.90	0.97	0.95
0.71	0.81	0.81	0.87	0.57	0.37	0.80	0.88	0.89	0.79	0.85
0.49	0.62	0.60	0.58	0.50	0.60	0.58	0.50	0.61	0.45	0.33
0.86	0.84	0.74	0.58	0.51	0.39	0.73	0.92	0.91	0.49	0.74
0.96	0.67	0.54	0.85	0.48	0.37	0.88	0.90	0.94	0.82	0.93
0.69	0.49	0.56	0.66	0.43	0.42	0.77	0.73	0.71	0.90	0.99
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93

James Hays

# Grayscale Intensity



# Color



Red intensity



Green

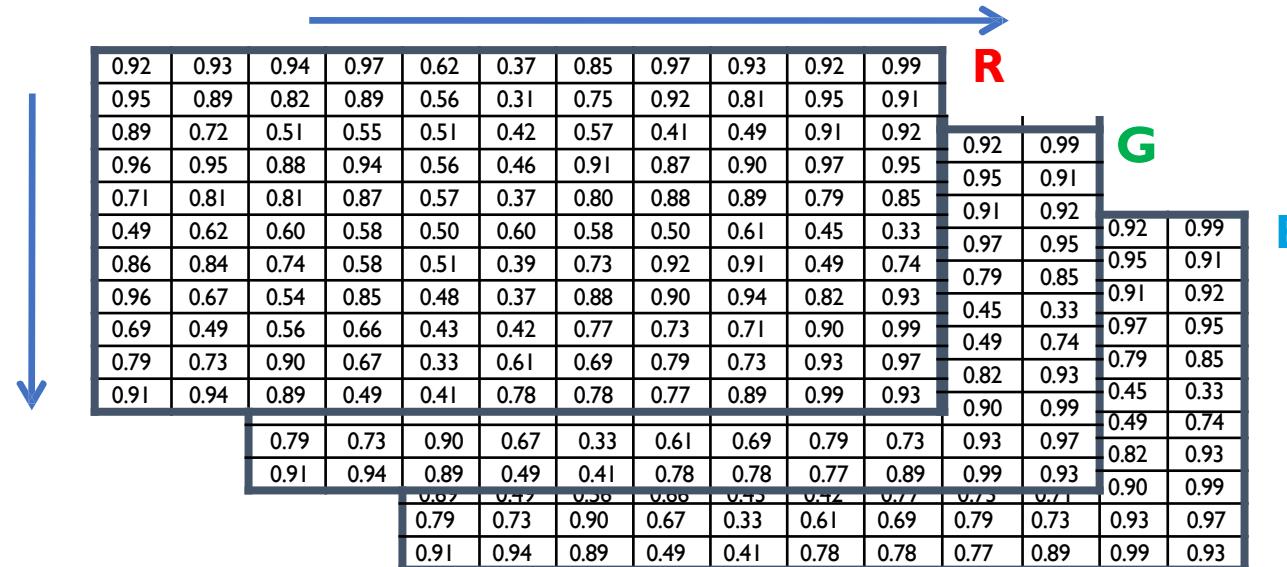


Blue

# Images in Python (import numpy)

$N \times M$  grayscale image “im”

- $im[0, 0, 0]$  = top-left pixel value, **red channel**
- $im[y, x, 1]$  =  $y$  pixels down,  $x$  pixels to right, **green channel**
- $im[N-1, M-1, 2]$  = bottom-right pixel, **blue channel**

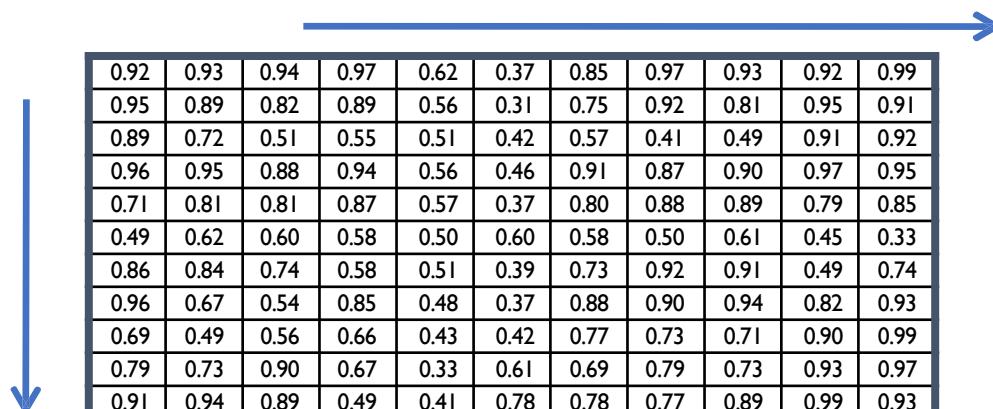


James Hays

# Images in Python (import numpy, scikit)

Take care of types!

- uint8 (values 0 to 255) - `io.imread("file.jpg")`
- float32 (values 0 to 255) - `io.imread("file.jpg").astype(np.float32)`
- float32 (values 0 to 1) - `img_as_float32(io.imread("file.jpg"))`



0.92	0.93	0.94	0.97	0.62	0.37	0.85	0.97	0.93	0.92	0.99
0.95	0.89	0.82	0.89	0.56	0.31	0.75	0.92	0.81	0.95	0.91
0.89	0.72	0.51	0.55	0.51	0.42	0.57	0.41	0.49	0.91	0.92
0.96	0.95	0.88	0.94	0.56	0.46	0.91	0.87	0.90	0.97	0.95
0.71	0.81	0.81	0.87	0.57	0.37	0.80	0.88	0.89	0.79	0.85
0.49	0.62	0.60	0.58	0.50	0.60	0.58	0.50	0.61	0.45	0.33
0.86	0.84	0.74	0.58	0.51	0.39	0.73	0.92	0.91	0.49	0.74
0.96	0.67	0.54	0.85	0.48	0.37	0.88	0.90	0.94	0.82	0.93
0.69	0.49	0.56	0.66	0.43	0.42	0.77	0.73	0.71	0.90	0.99
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93

James Hays

# Images in Python

```
>>> from numpy import random as r  
>>> I = r.rand(256,256)
```

- What is this? What does it look like?
- Which values does it have?
- How many are there?

# Images in Python

```
>>> from matplotlib import pyplot as p  
>>> I = r.rand(256,256)  
>>> p.imshow(I, cmap='gray')  
>>> p.show()
```

Is it an image?

