

# Introduction to Computer Vision

---

**Kaveh Fathian**

Assistant Professor

Computer Science Department

Colorado School of Mines

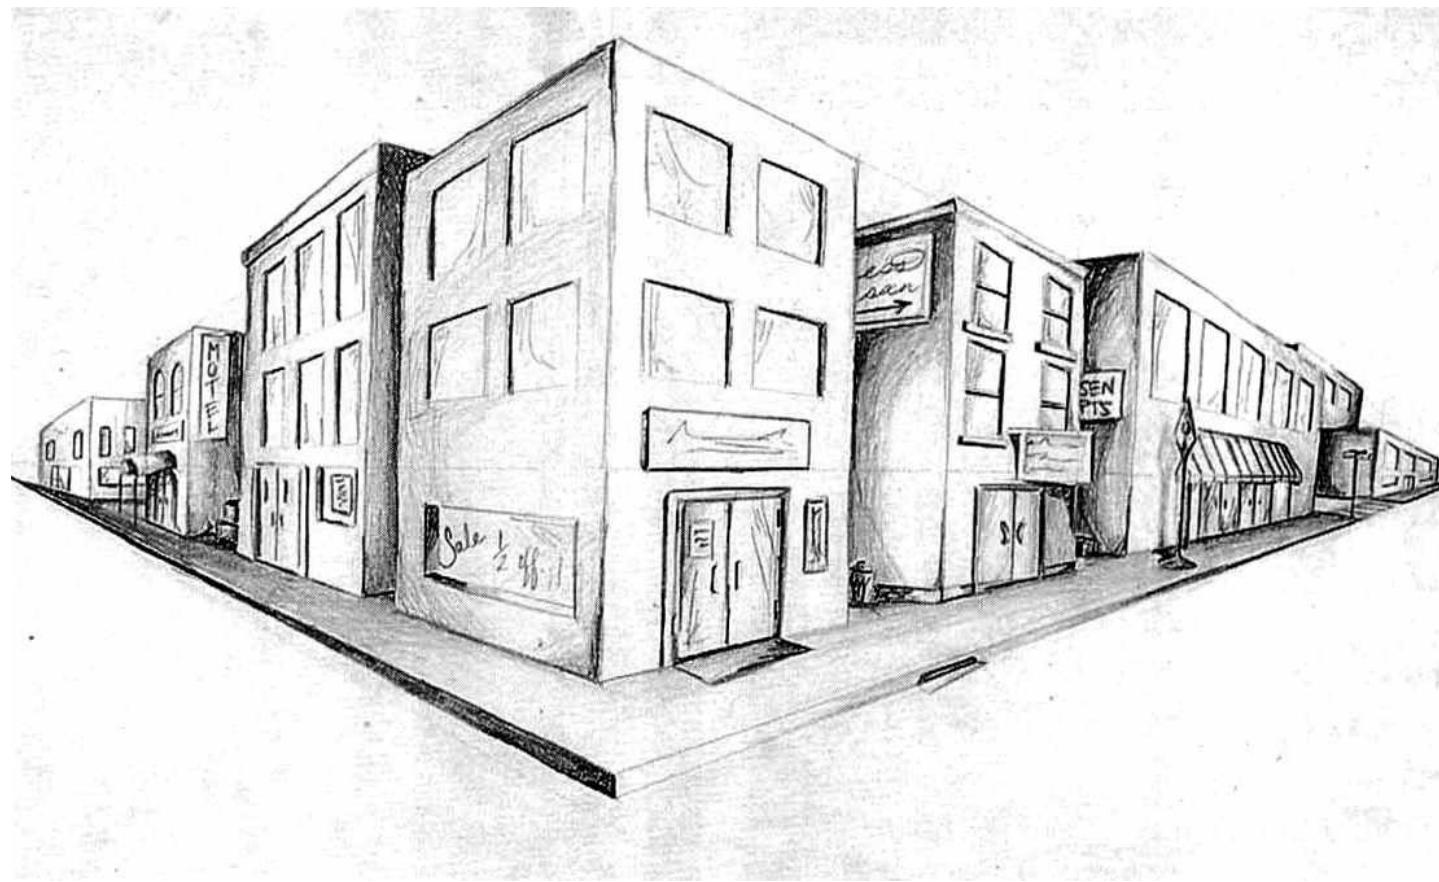
**Lecture 6**

# Feature Points

## Learning outcomes:

- Features
- Detection, Description, & Matching
- Distinctiveness, Repeatability, & Compactness/Efficiency
- Geometric Transformations
- Photometric Transformations
- Corners
- Corner Detection by Auto-Correlation
- Harris Corner Detection
- Taylor Series Expansion & Second Moment Matrix
- Second Moment Matrix
- Invariance and Covariance
- Templates and Histograms
- Scale-Invariant Feature Transform (SIFT)
- SIFT Descriptors

# Detecting Corners



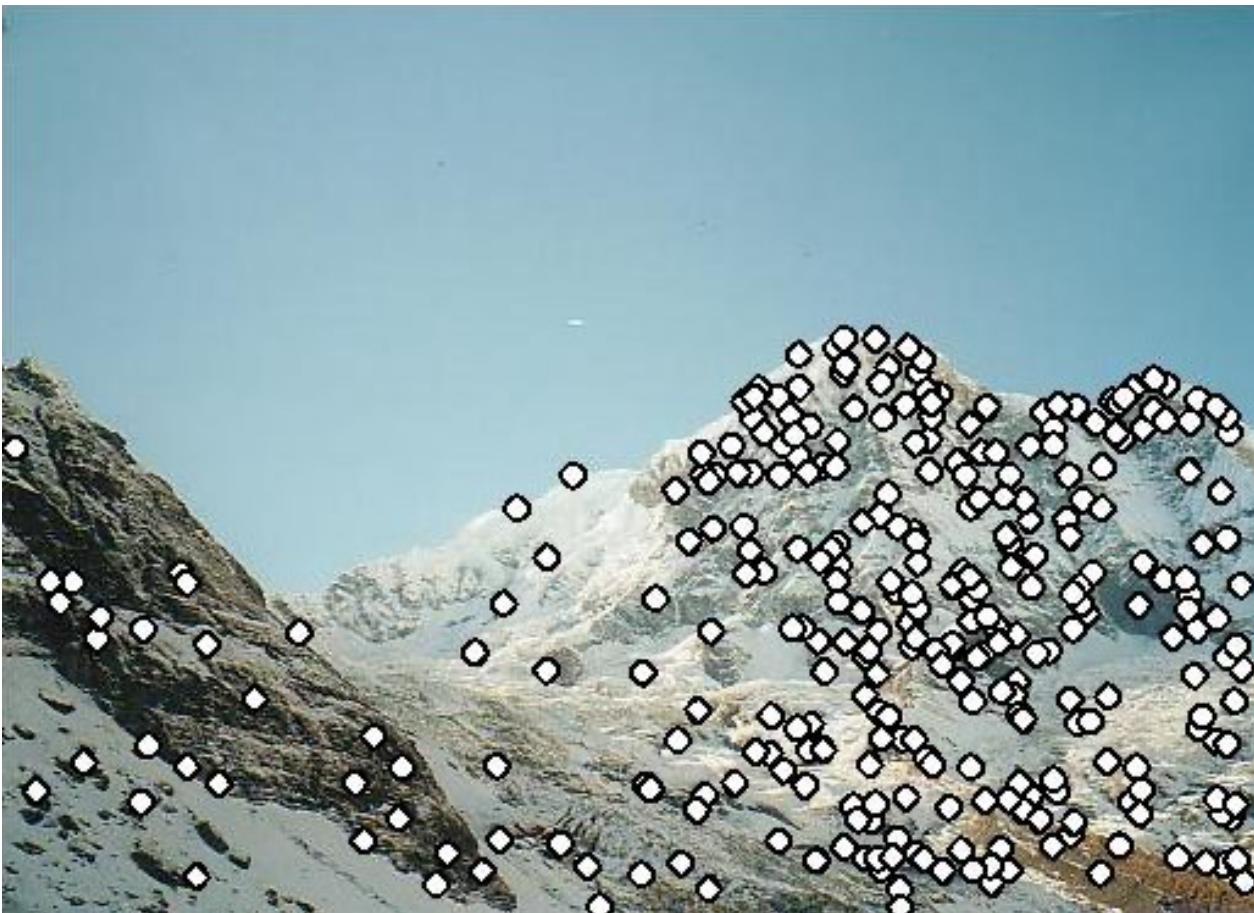
# Application: Image Alignment

- We have two images – how do we overlay them?
- How would you do it by eye?



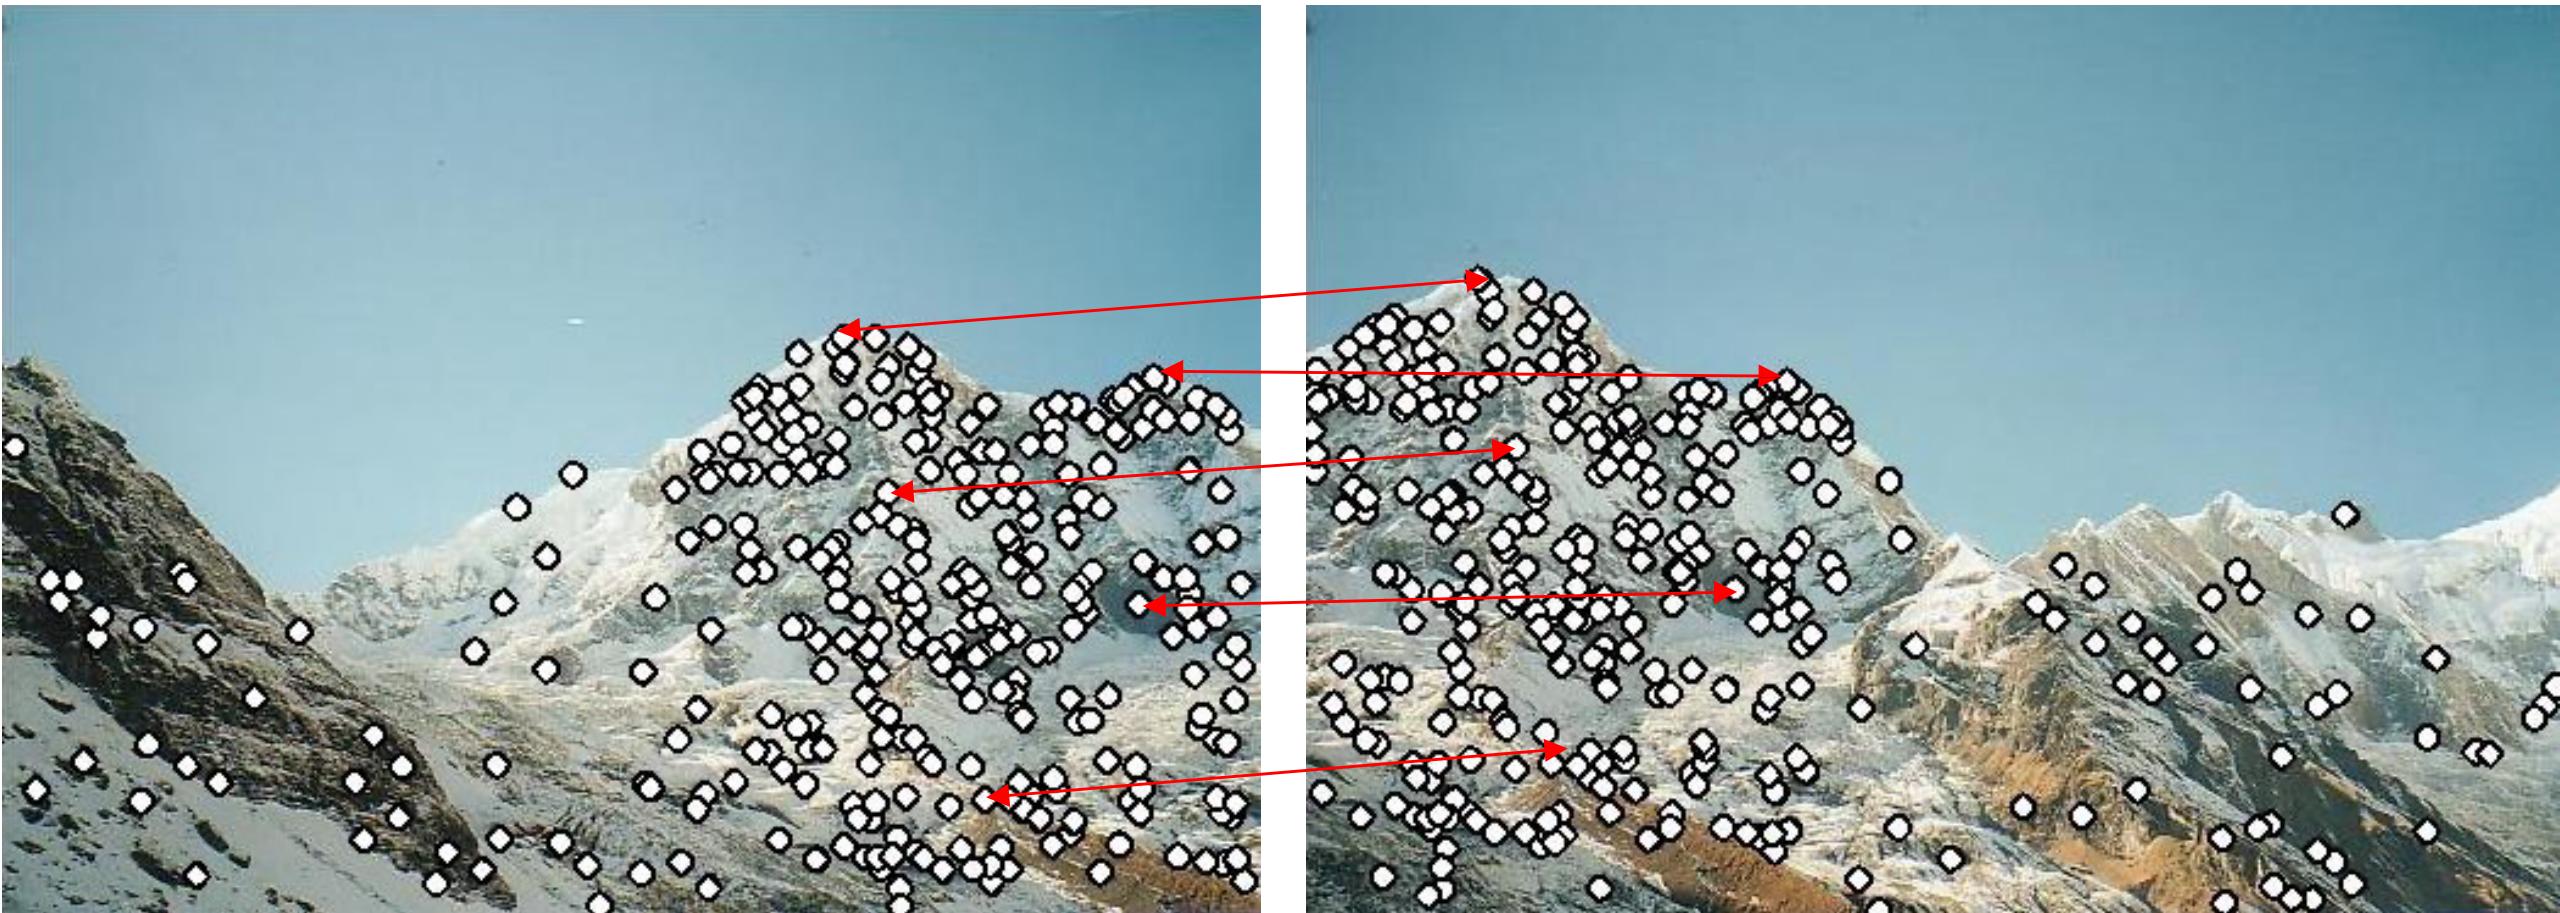
# Application: Image Alignment

- Detect feature points in both images



# Application: Image Alignment

- Detect feature points in both images
- Find corresponding pairs



# Application: Image Alignment

- Detect feature points in both images
- Find corresponding pairs
- Use these pairs to align images



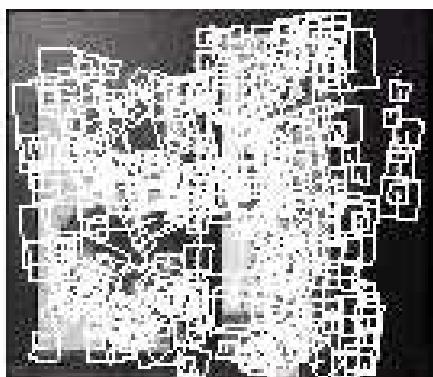
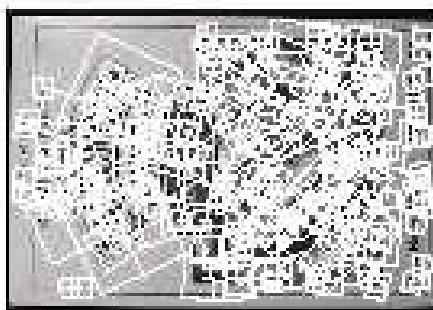
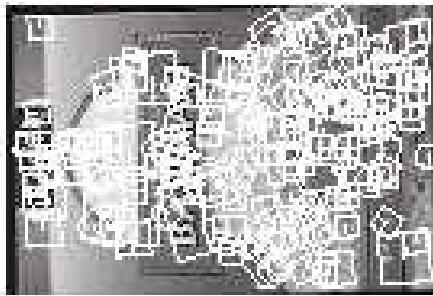
# Why Detect Corners?

Corners are fundamental to CV applications:

- Image alignment
- 3D reconstruction
- Motion tracking
- Indexing and database retrieval
- Object recognition



# Object Recognition

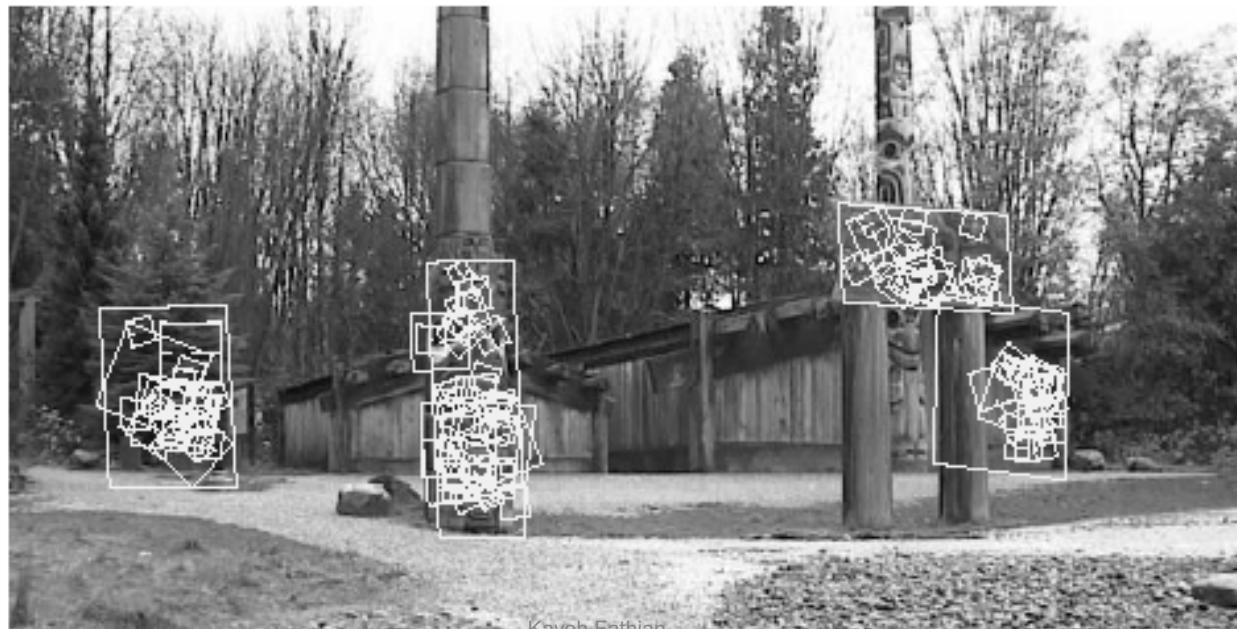
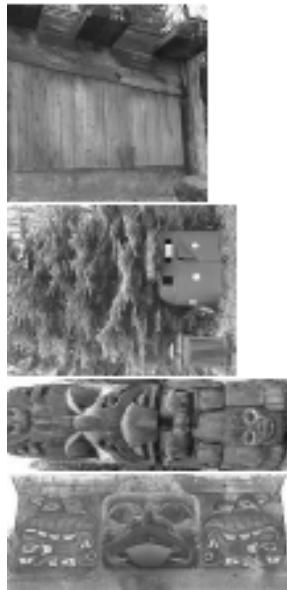


Database of objects



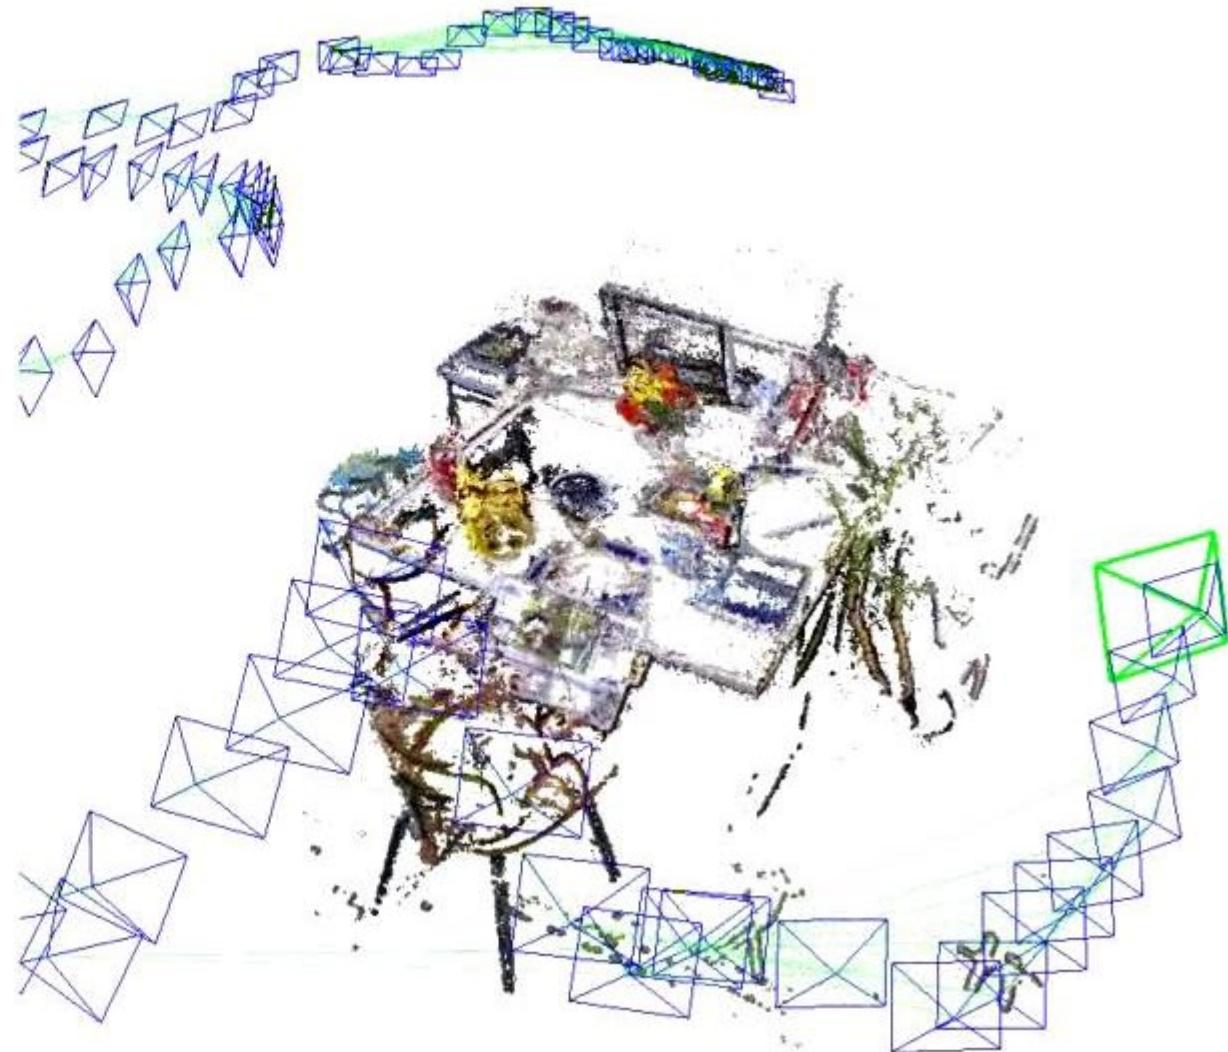
Instance recognition

# Location Recognition



# Localization & Mapping

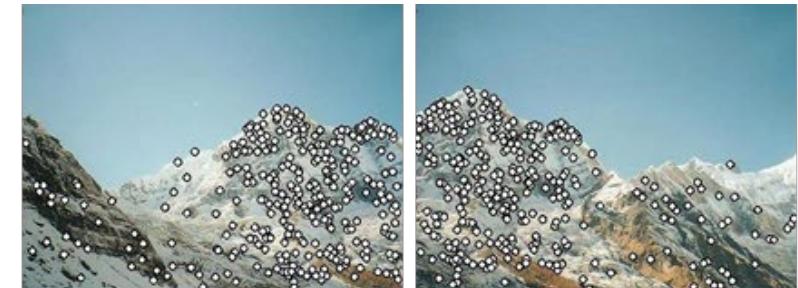
<https://youtu.be/ufvPS5wJAx0?t=77>



# Local Feature Matching: Main Components

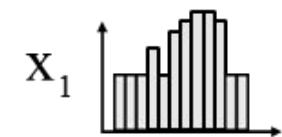
## 1) Detection:

Find a set of distinctive features (e.g., key points)

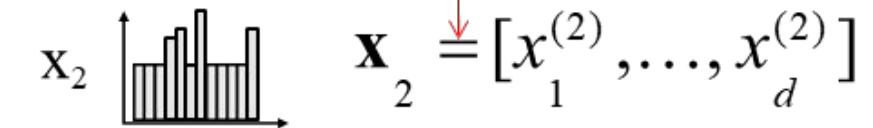
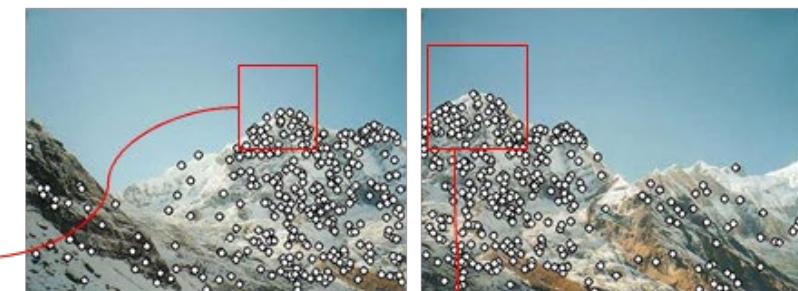


## 2) Description:

Extract feature descriptor around each interest point as vector.



$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$



$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$

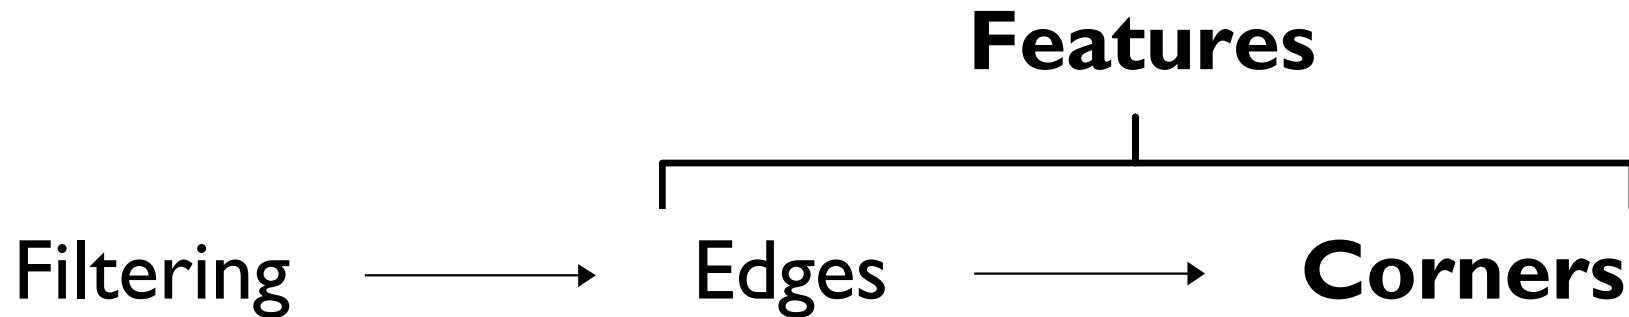
## 3) Matching:

Compute distance between feature vectors to find correspondence.

$$d(x_1, x_2) < T$$



# Features

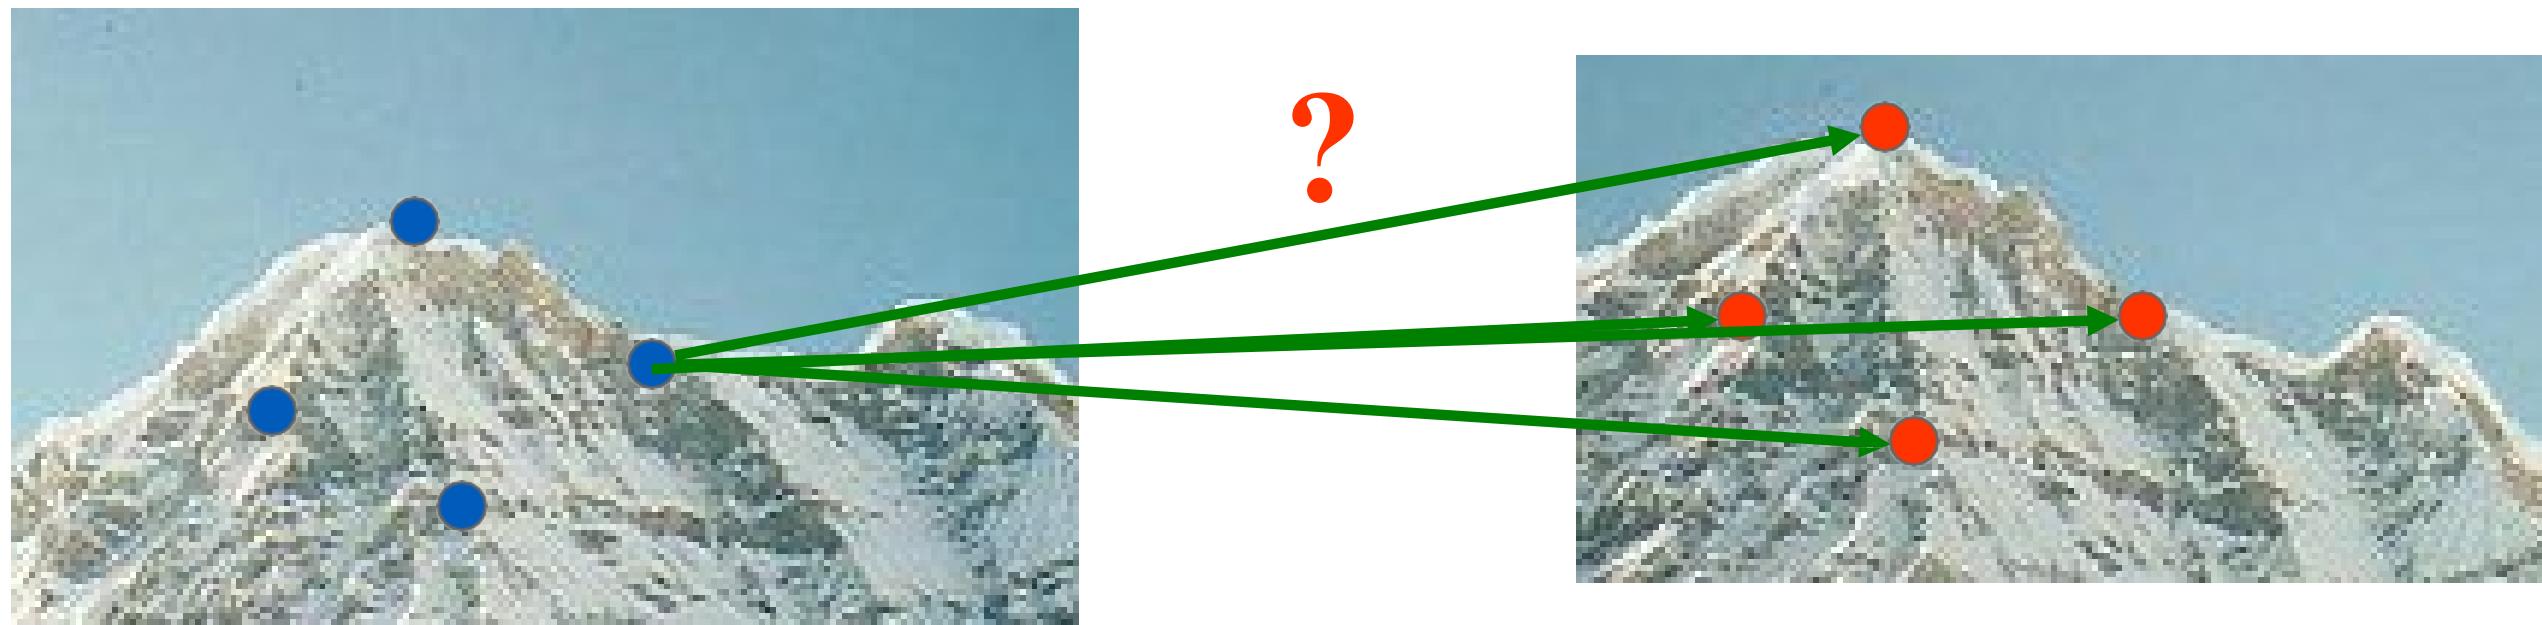


- Corners are also called **interest points** or **key points**
- “Local features”

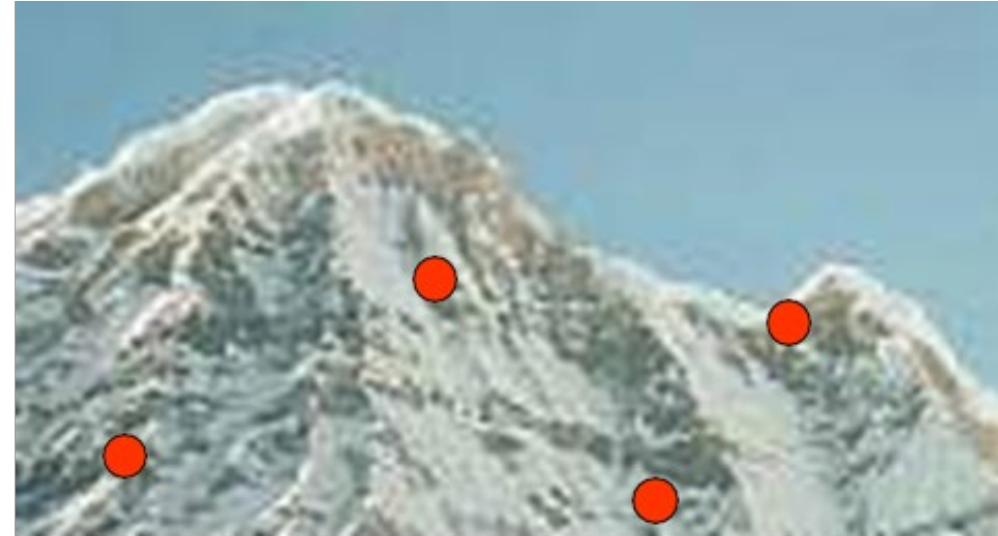
# Good Features

## Goal: Distinctiveness

- We want to be able to reliably determine which point goes with which
- May be difficult in structured environments with repeated elements



# Good Features



With these points, there's no chance to find true matches!

## Goal: Repeatability

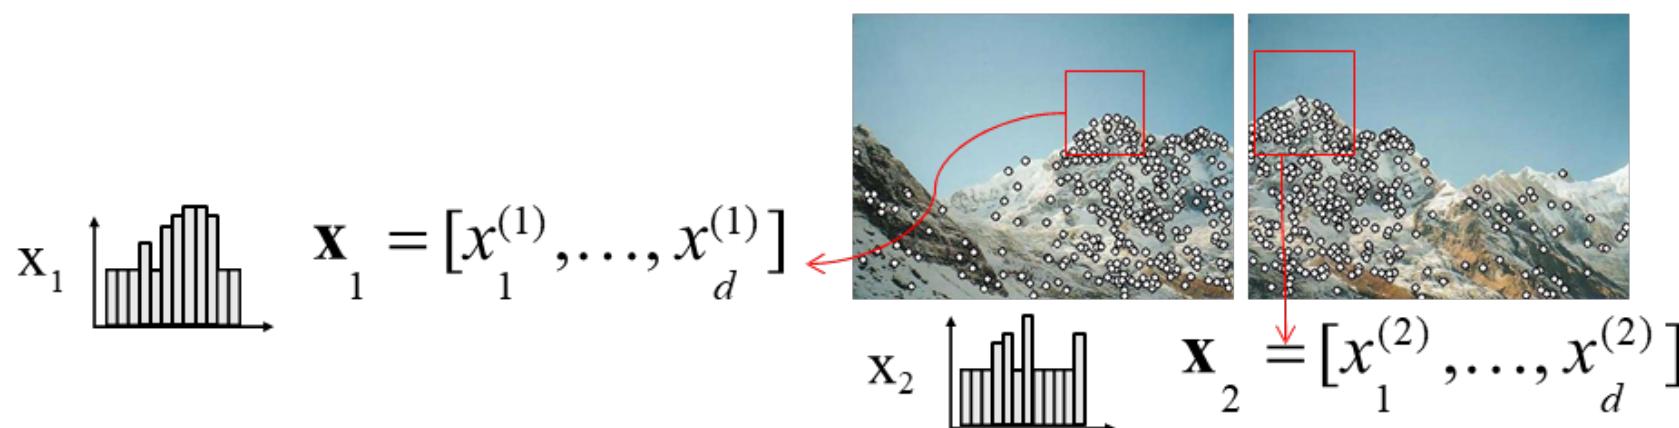
- We want to detect (at least some of) the same points in both images
- Under geometric and photometric variations



# Good Features

## Goal: Compactness and Efficiency

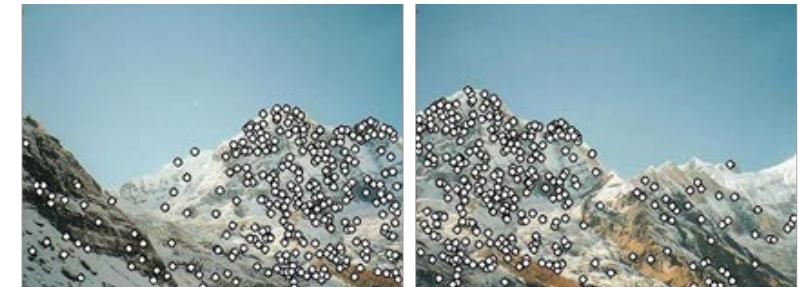
- We want the representation to be as small and as fast as possible
  - Much smaller than a whole image
- Sometimes, we'd like to run the detection procedure independently per image
  - Match just the compact descriptors for speed
  - Difficult! We don't get to see 'the other image' at match time, e.g., object detection



# Local Feature Matching: Main Components

## 1) Detection:

Find a set of distinctive features (e.g., key points)

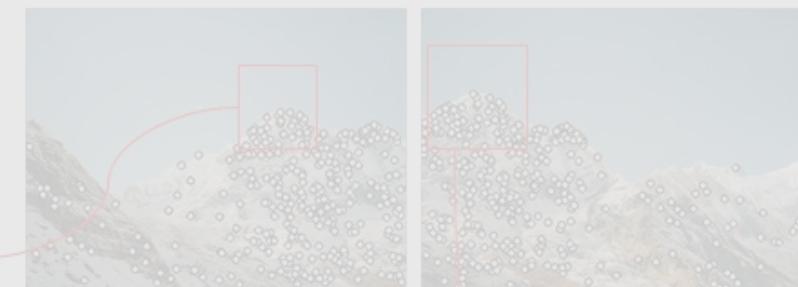


## 2) Description:

Extract feature descriptor around each interest point as vector.



$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$



$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$

## 3) Matching:

Compute distance between feature vectors to find correspondence.

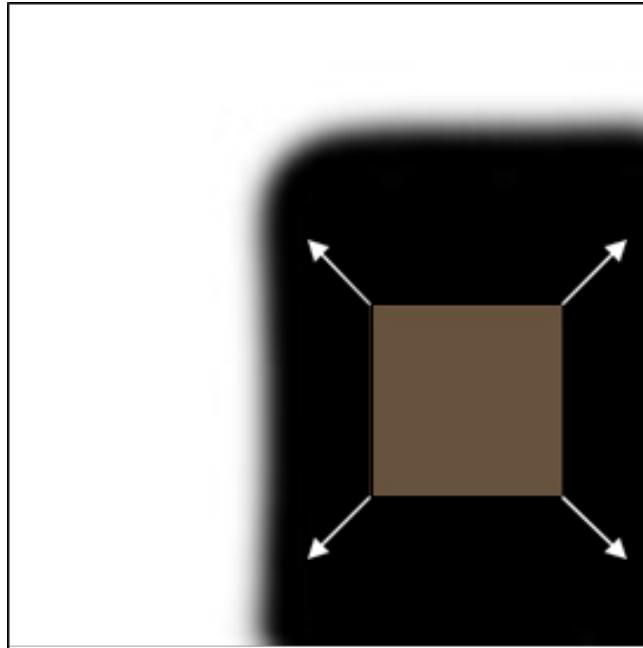
$$d(x_1, x_2) < T$$



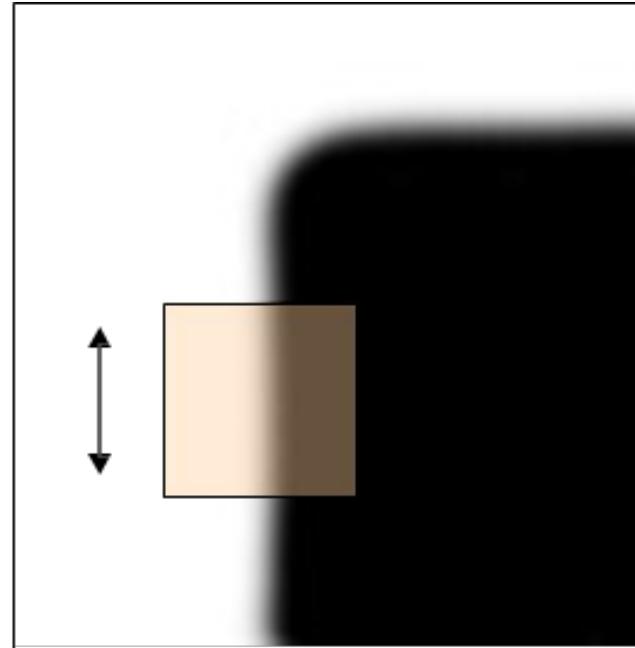
K. Grauman, B. Leibe

# Corner Detection: Basic Idea

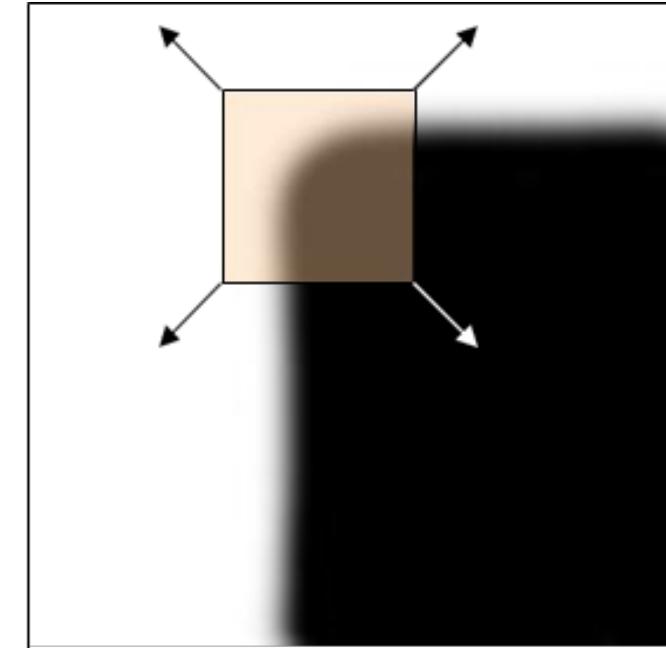
- Recognize corners by looking at small window
- We want a window shift in any direction to give a large change in intensity



**“Flat” region:**  
no change in all  
directions



**“Edge”:**  
no change along  
the edge direction



**“Corner”:**  
significant change in  
all directions

# Corner Detection by Auto-correlation

Change in appearance of window  $w(x, y)$  for shift  $[u, v]$ :

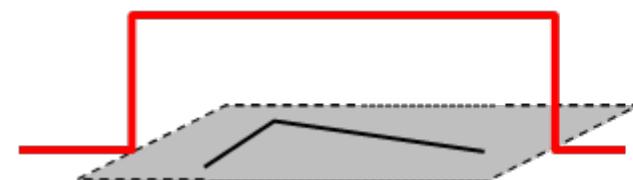
$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Diagram illustrating the components of the auto-correlation function  $E(u, v)$ :

- Window function:  $w(x, y)$
- Shifted intensity:  $I(x + u, y + v)$
- Intensity:  $I(x, y)$

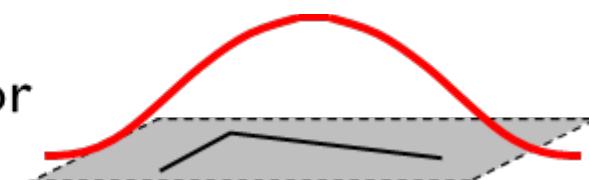
Also called 'sum of squared differences' or SSD

Window function  $w(x, y) =$



1 in window, 0 outside

or

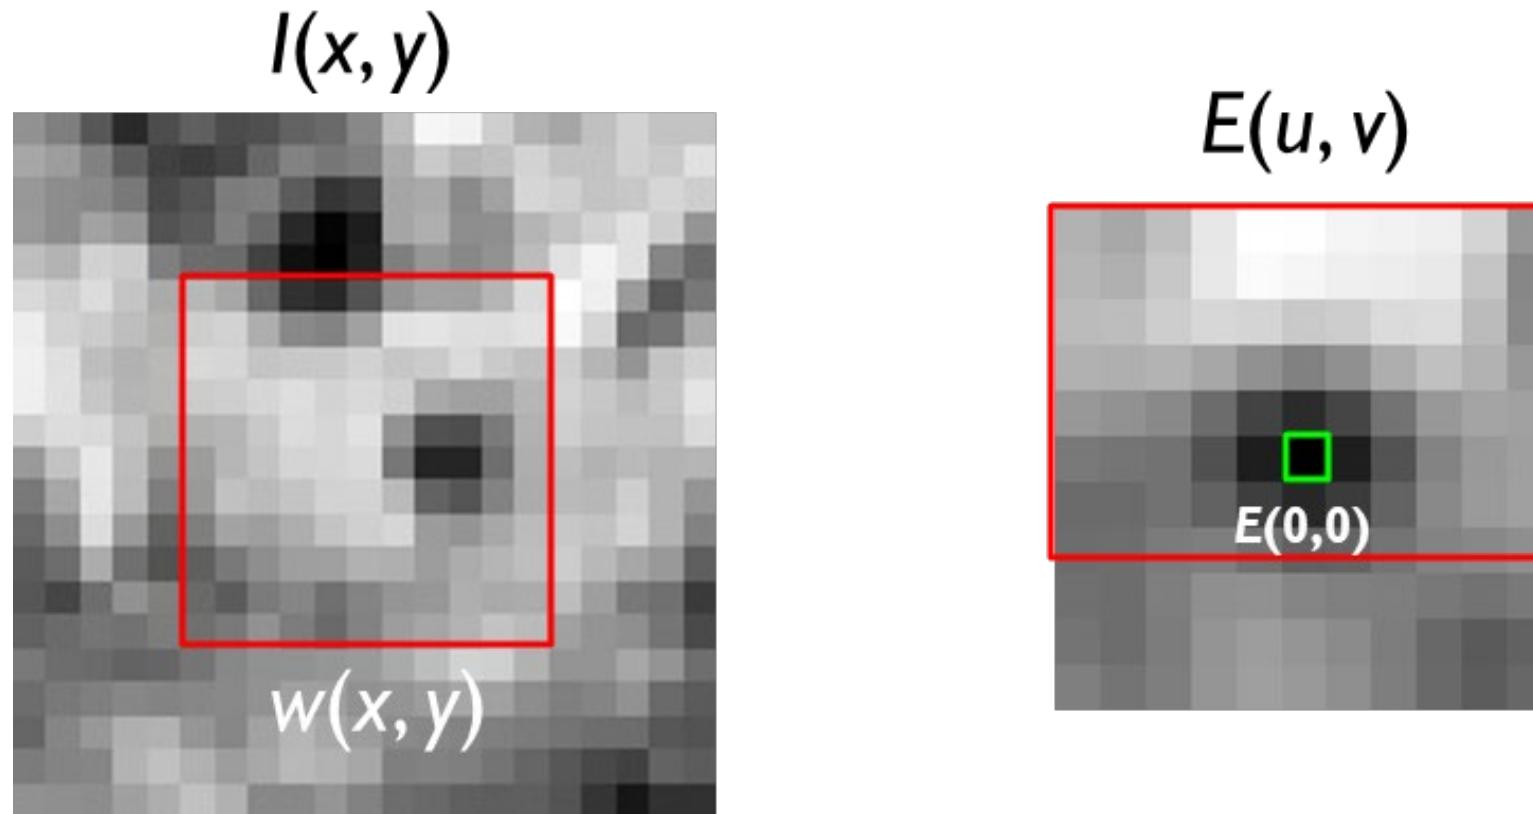


Gaussian

# Corner Detection by Auto-correlation

Change in appearance of window  $w(x, y)$  for shift  $[u, v]$ :

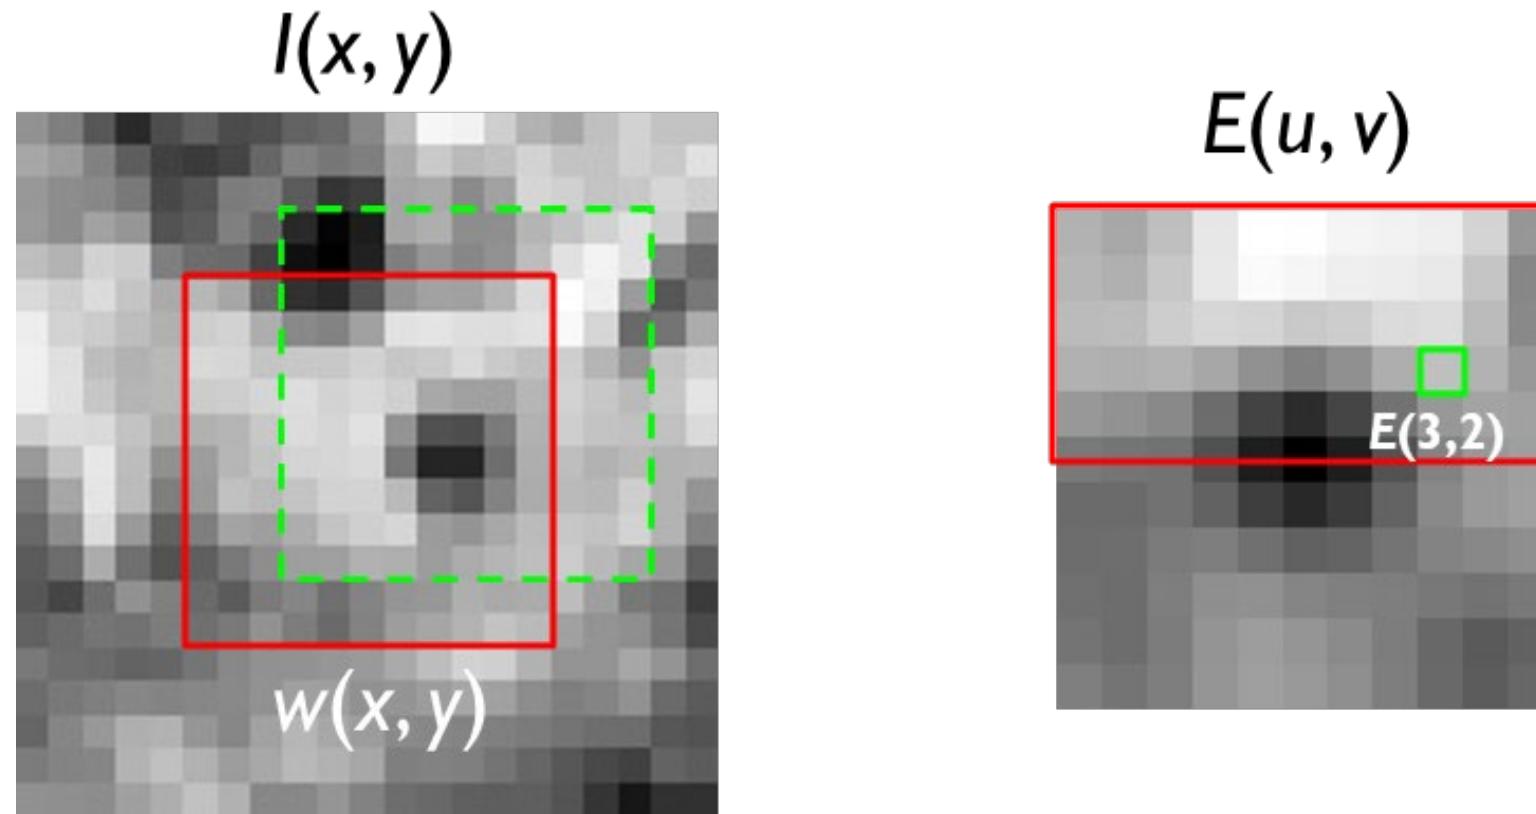
$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$



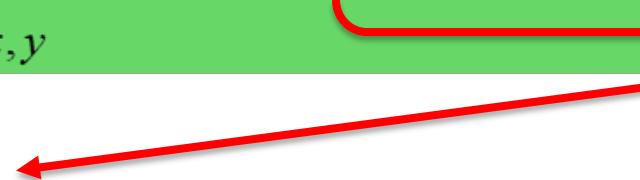
# Corner Detection by Auto-correlation

Change in appearance of window  $w(x, y)$  for shift  $[u, v]$ :

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$



# Auto-correlation

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$


**Autocorrelation** is the correlation of the image with itself

- Windows centered on an edge point will have autocorrelation that falls off slowly in the direction along the edge and rapidly in the direction across (perpendicular to) the edge.
- Windows centered on a corner point will have autocorrelation that falls off rapidly in all directions.

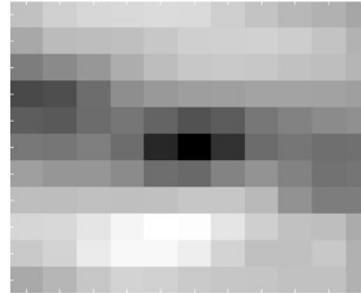
# Exercise

Correspond the three red crosses to (a,b,c)

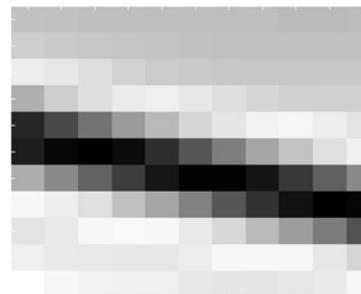
$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$



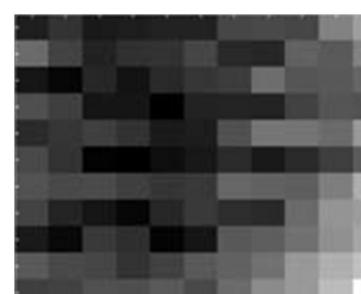
$E(u, v)$



(a)

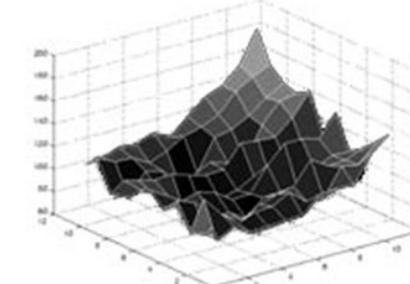
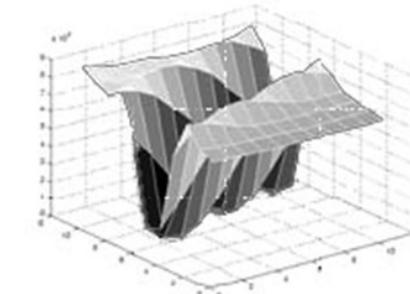
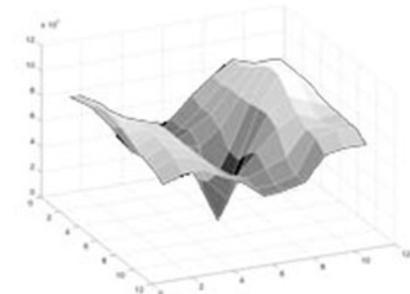


(b)



(c)

$E(u, v)$  as a surface

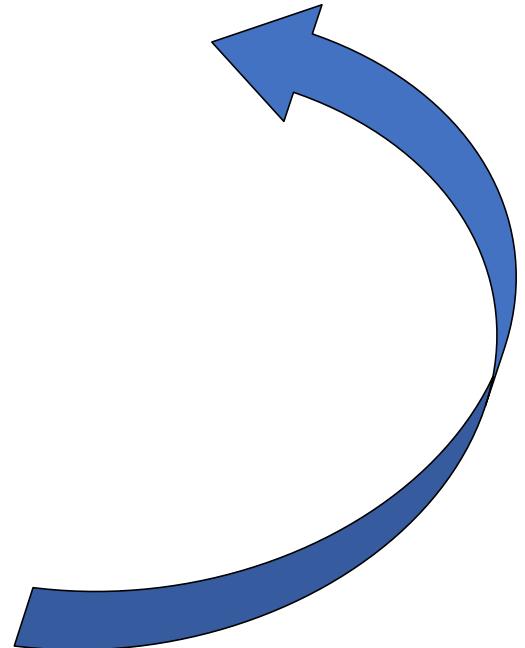


# Corner Detection by Auto-correlation

- Change in appearance of window  $w(x, y)$  for shift  $[u, v]$ :

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

- We want to discover how  $E$  behaves for small shifts
- But this is very slow to compute naively:  
 $O(\text{window\_width}^2 * \text{shift\_range}^2 * \text{image\_width}^2)$
- $O(11^2 * 11^2 * 600^2) = 5.2$  billion of these  
14.6k ops per image pixel

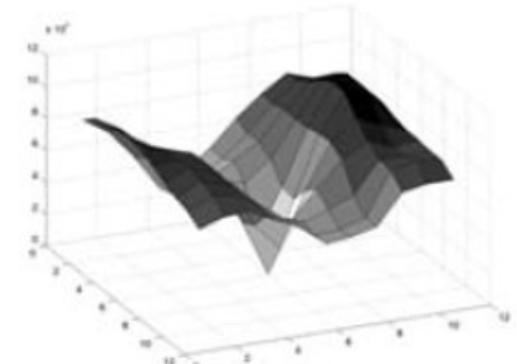
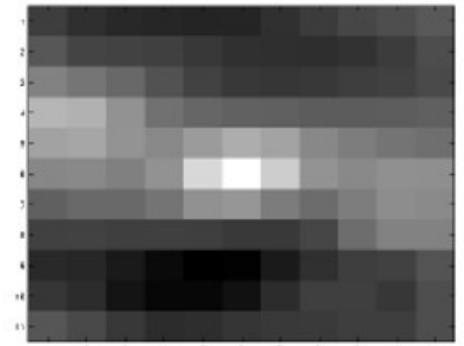


# Corner Detection by Auto-correlation

- Change in appearance of window  $w(x, y)$  for shift  $[u, v]$ :

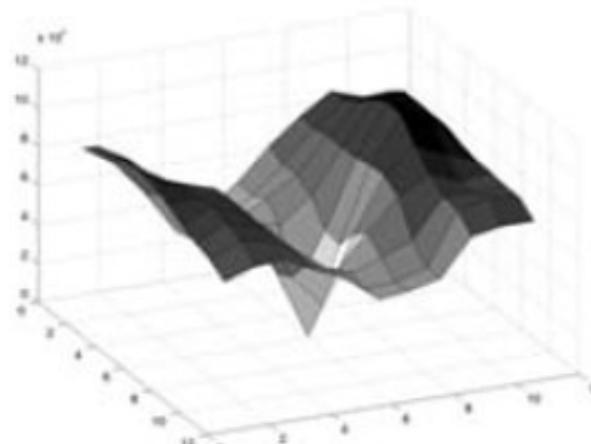
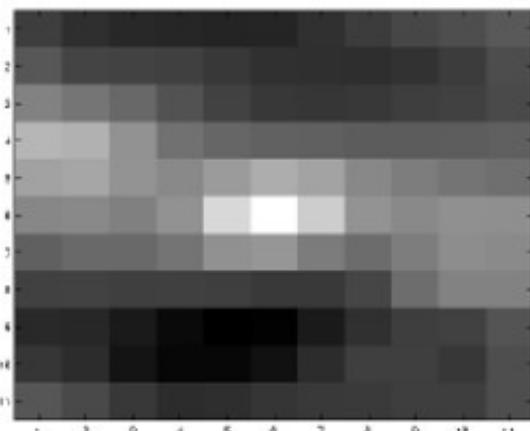
$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

- We want to discover how  $E$  behaves for small shifts
- But we know the response in  $E$  that we are looking for
  - strong peak

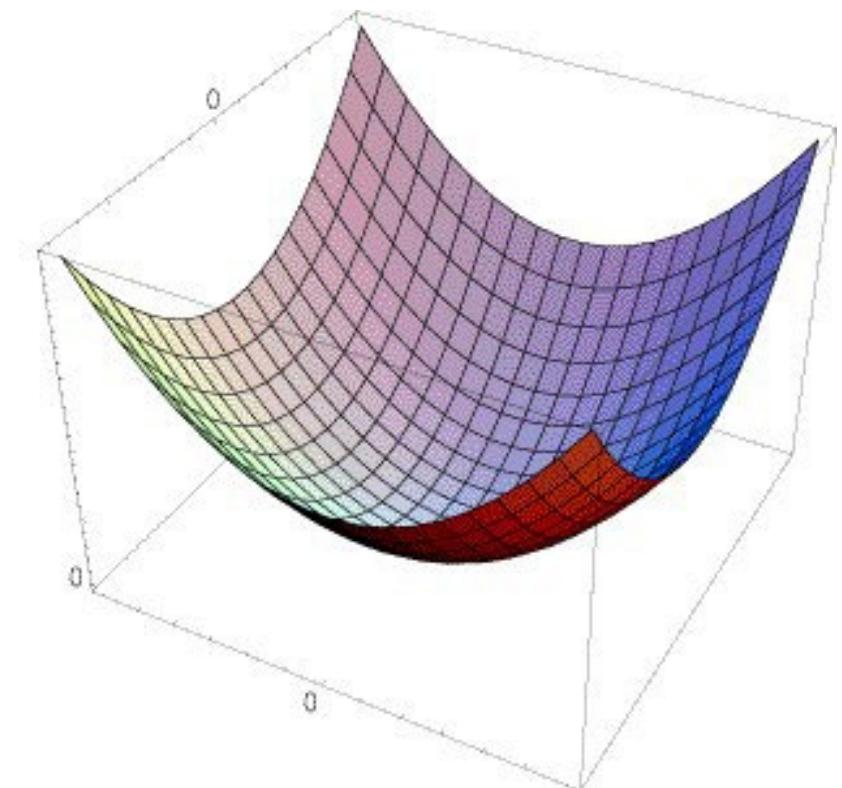


# Strategy

Approximate  $E(u, v)$  locally by a quadratic surface, and look for that instead



≈



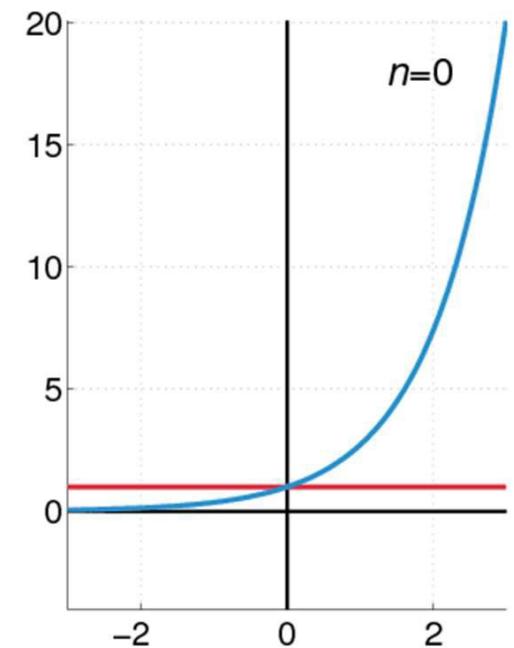
# Recall: Taylor Series Expansion

A function  $f$  can be represented by an infinite series of its derivatives at a single point  $a$ :

$$f(a) + \frac{f'(a)}{1!}(x - a) + \frac{f''(a)}{2!}(x - a)^2 + \frac{f'''(a)}{3!}(x - a)^3 + \dots$$

As we care about window centered, we set  $a = 0$   
(MacLaurin series)

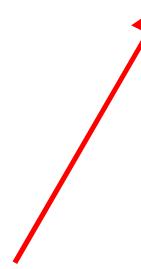
Approximation of  
 $f(x) = e^x$   
centered at  $f(0)$



# Quadratic Approximation

Local quadratic approximation of  $E(u, v)$  in the neighborhood of  $(0,0)$  is given by the second-order Taylor expansion:

$$E(u, v) \approx E(0,0) + [u \ v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$



Notation: partial derivative

# Quadratic Approximation

Local quadratic approximation of  $E(u, v)$  in the neighborhood of  $(0,0)$  is given by the second-order Taylor expansion:

$$E(u, v) \approx E(0,0) + [u \ v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$



Ignore function  
value; set to 0



Ignore first  
derivative, set  
to 0



Just look at shape of  
second derivative  
(2D quadratic surface)

# Additional Slide

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Second-order Taylor expansion of  $E(u, v)$  about (0,0):

$$E(u, v) \approx E(0,0) + [u \ v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E_u(u, v) = \sum_{x, y} 2w(x, y) [I(x + u, y + v) - I(x, y)] I_x(x + u, y + v)$$

$$E_{uu}(u, v) = \sum_{x, y} 2w(x, y) I_x(x + u, y + v) I_x(x + u, y + v) \\ + \sum_{x, y} 2w(x, y) [I(x + u, y + v) - I(x, y)] I_{xx}(x + u, y + v)$$

$$E_{uv}(u, v) = \sum_{x, y} 2w(x, y) I_y(x + u, y + v) I_x(x + u, y + v) \\ + \sum_{x, y} 2w(x, y) [I(x + u, y + v) - I(x, y)] I_{xy}(x + u, y + v)$$

# Additional Slide

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Second-order Taylor expansion of  $E(u, v)$  about (0,0):

$$E(u, v) \approx E(0,0) + [u \ v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E(0,0) = 0$$

$$E_u(0,0) = 0$$

$$E_v(0,0) = 0$$

$$E_{uu}(0,0) = \sum_{x, y} 2w(x, y) I_x(x, y) I_x(x, y)$$

$$E_{vv}(0,0) = \sum_{x, y} 2w(x, y) I_y(x, y) I_y(x, y)$$

$$E_{uv}(0,0) = \sum_{x, y} 2w(x, y) I_x(x, y) I_y(x, y)$$

# Additional Slide

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Second-order Taylor expansion of  $E(u, v)$  about (0,0):

$$E(u, v) \approx [u \ v] \begin{bmatrix} \sum_{x, y} w(x, y) I_x^2(x, y) & \sum_{x, y} w(x, y) I_x(x, y) I_y(x, y) \\ \sum_{x, y} w(x, y) I_x(x, y) I_y(x, y) & \sum_{x, y} w(x, y) I_y^2(x, y) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E(0,0) = 0$$

$$E_u(0,0) = 0$$

$$E_v(0,0) = 0$$

$$E_{uu}(0,0) = \sum_{x, y} 2w(x, y) I_x(x, y) I_x(x, y)$$

$$E_{vv}(0,0) = \sum_{x, y} 2w(x, y) I_y(x, y) I_y(x, y)$$

$$E_{uv}(0,0) = \sum_{x, y} 2w(x, y) I_x(x, y) I_y(x, y)$$

# Corner Detection: Mathematics

The quadratic approximation simplifies to

$$E(u, v) \approx [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where  $M$  is a second moment matrix/tensor computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

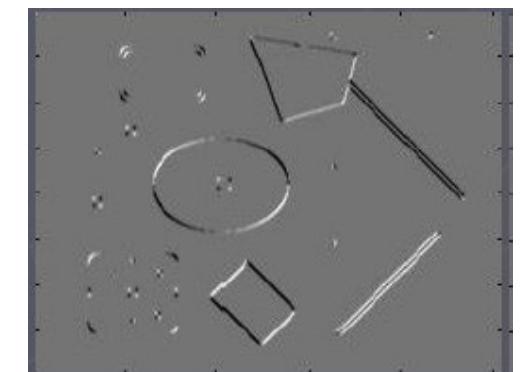
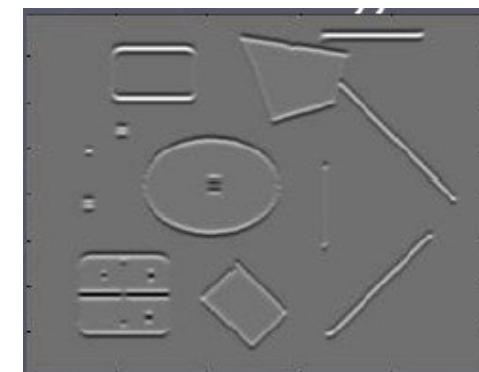
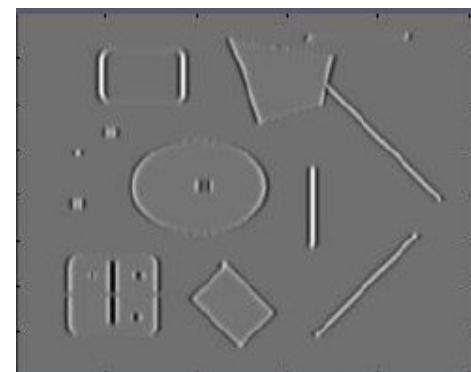
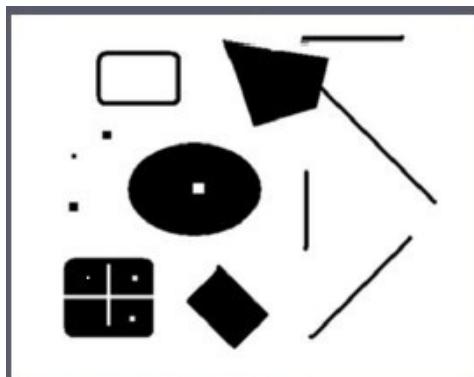
$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y] = \sum \nabla I (\nabla I)^T$$

# Corner Detection

Corners as distinctive interest points

$$M = \sum w(x, y) \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix}$$

2 x 2 matrix of image derivatives (averaged in neighborhood of a point):



Notation:

$$I_x \Leftrightarrow \frac{\partial I}{\partial x}$$

$$I_y \Leftrightarrow \frac{\partial I}{\partial y}$$

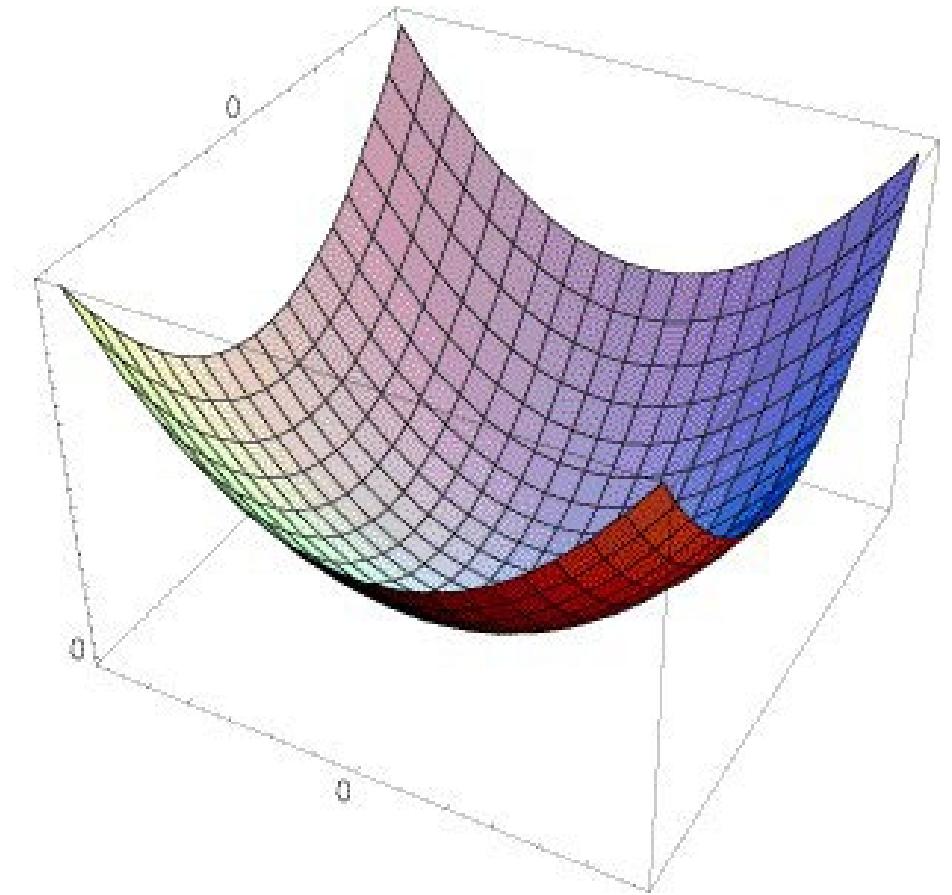
$$I_x I_y \Leftrightarrow \frac{\partial I}{\partial x} \frac{\partial I}{\partial y}$$

# Interpreting the Second Moment Matrix

$E(u, v)$  is locally approximated by a quadratic surface. Let's try to understand how its shape relates to  $M$

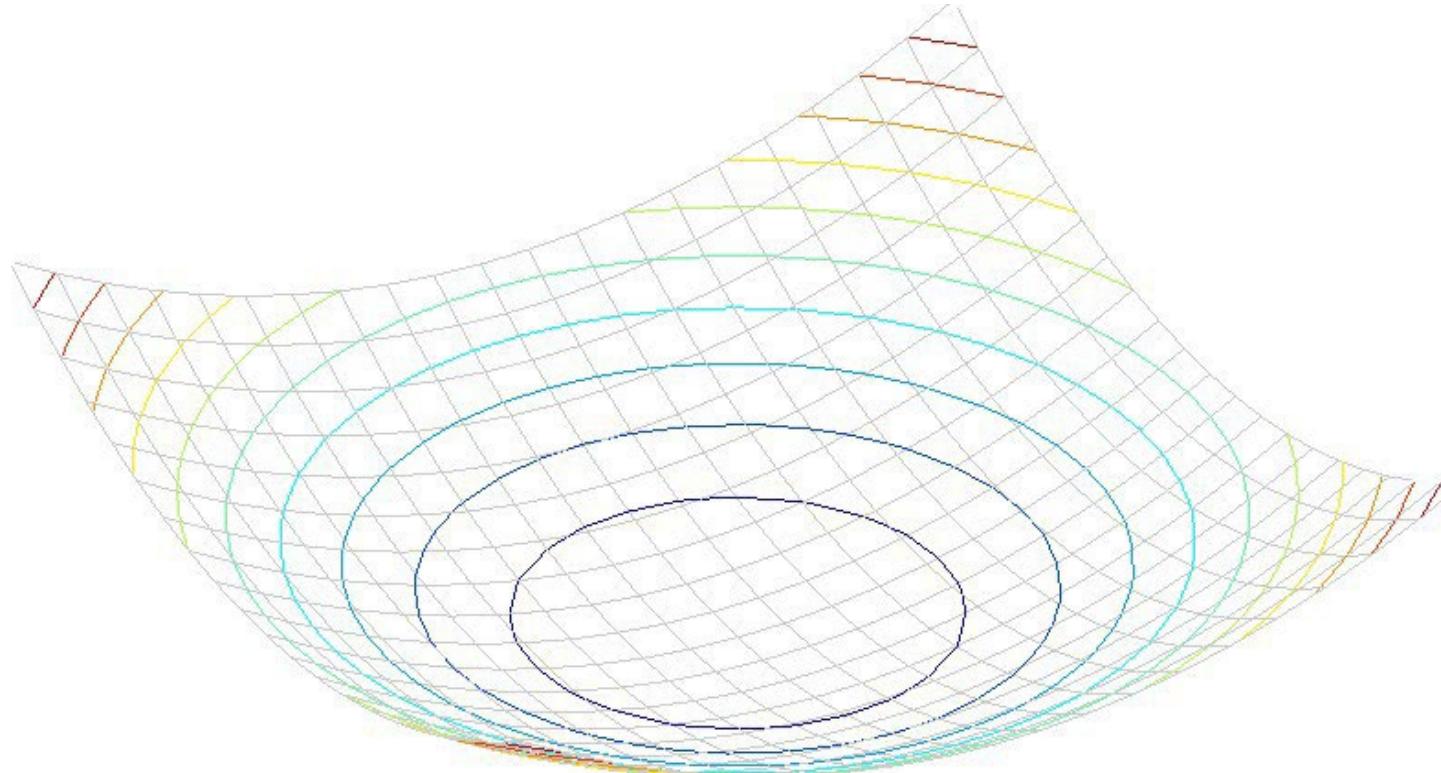
$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$



# Interpreting the Second Moment Matrix

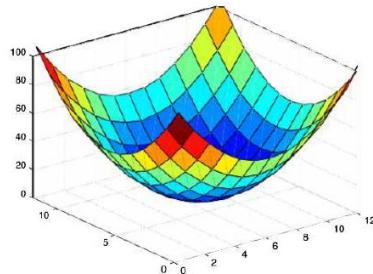
- Let's take horizontal "slices" of our approximation of  $E(u, v)$
- Each colored line in the diagram below is where  $\begin{bmatrix} u & v \end{bmatrix}^T M \begin{bmatrix} u \\ v \end{bmatrix} = \text{constant}$



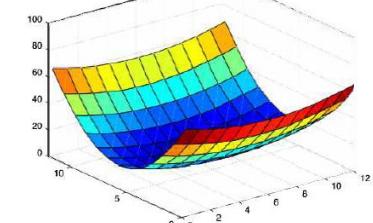
# Visualization of the Second Moment Matrix

Simple checkerboard

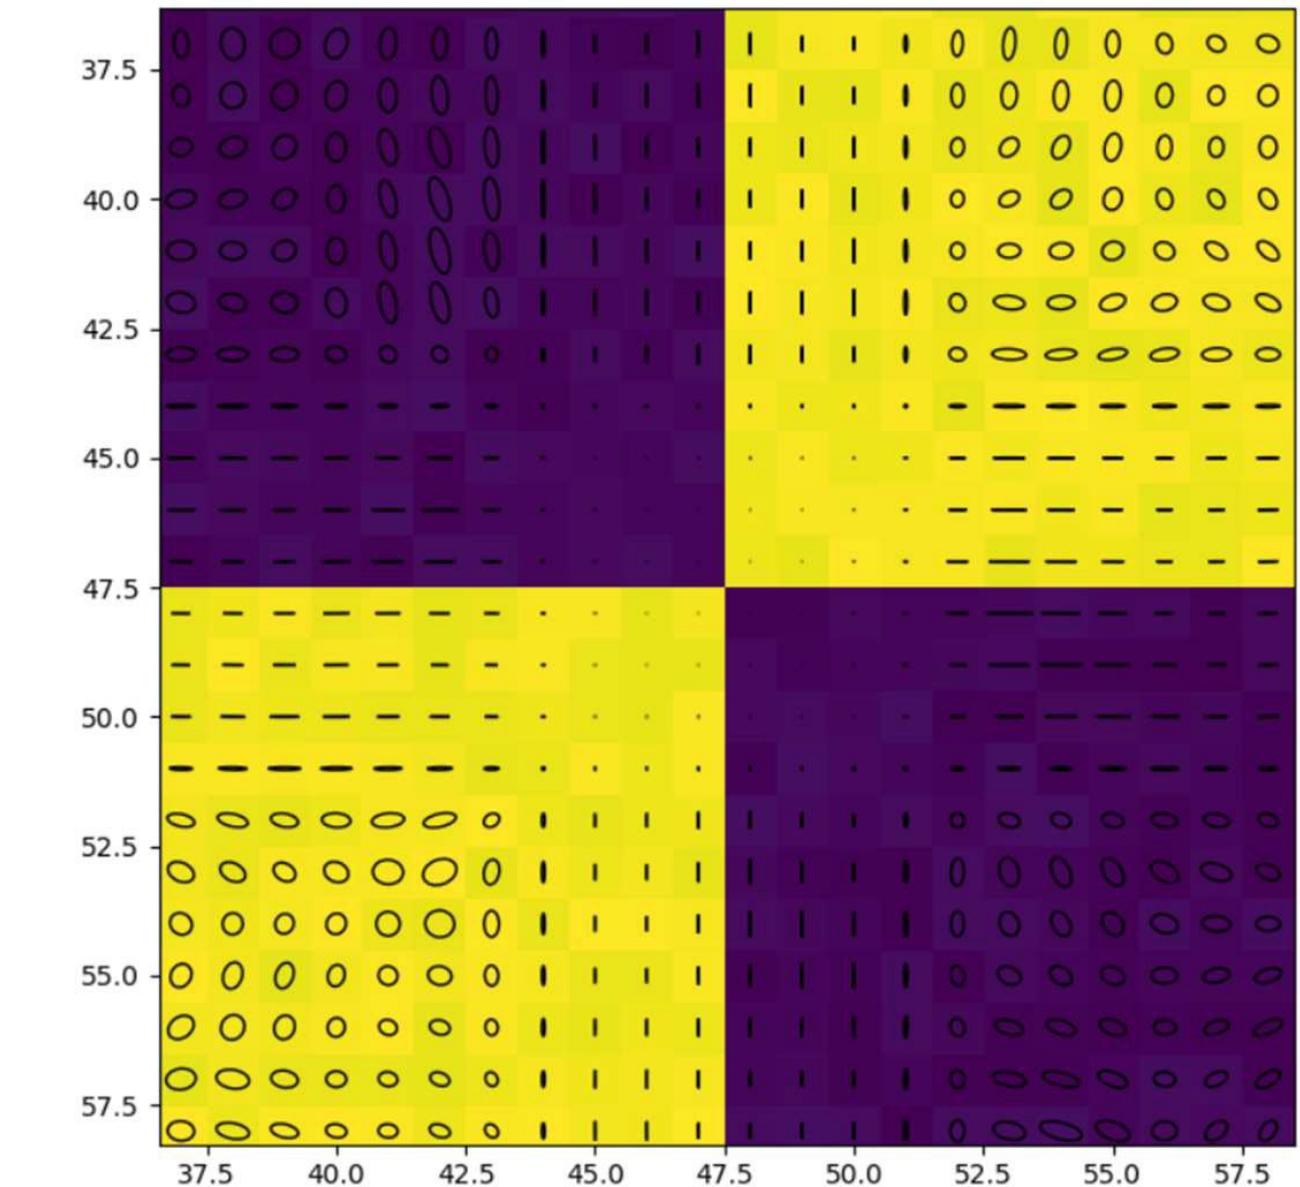
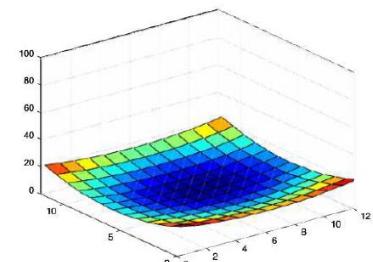
- 1 corner



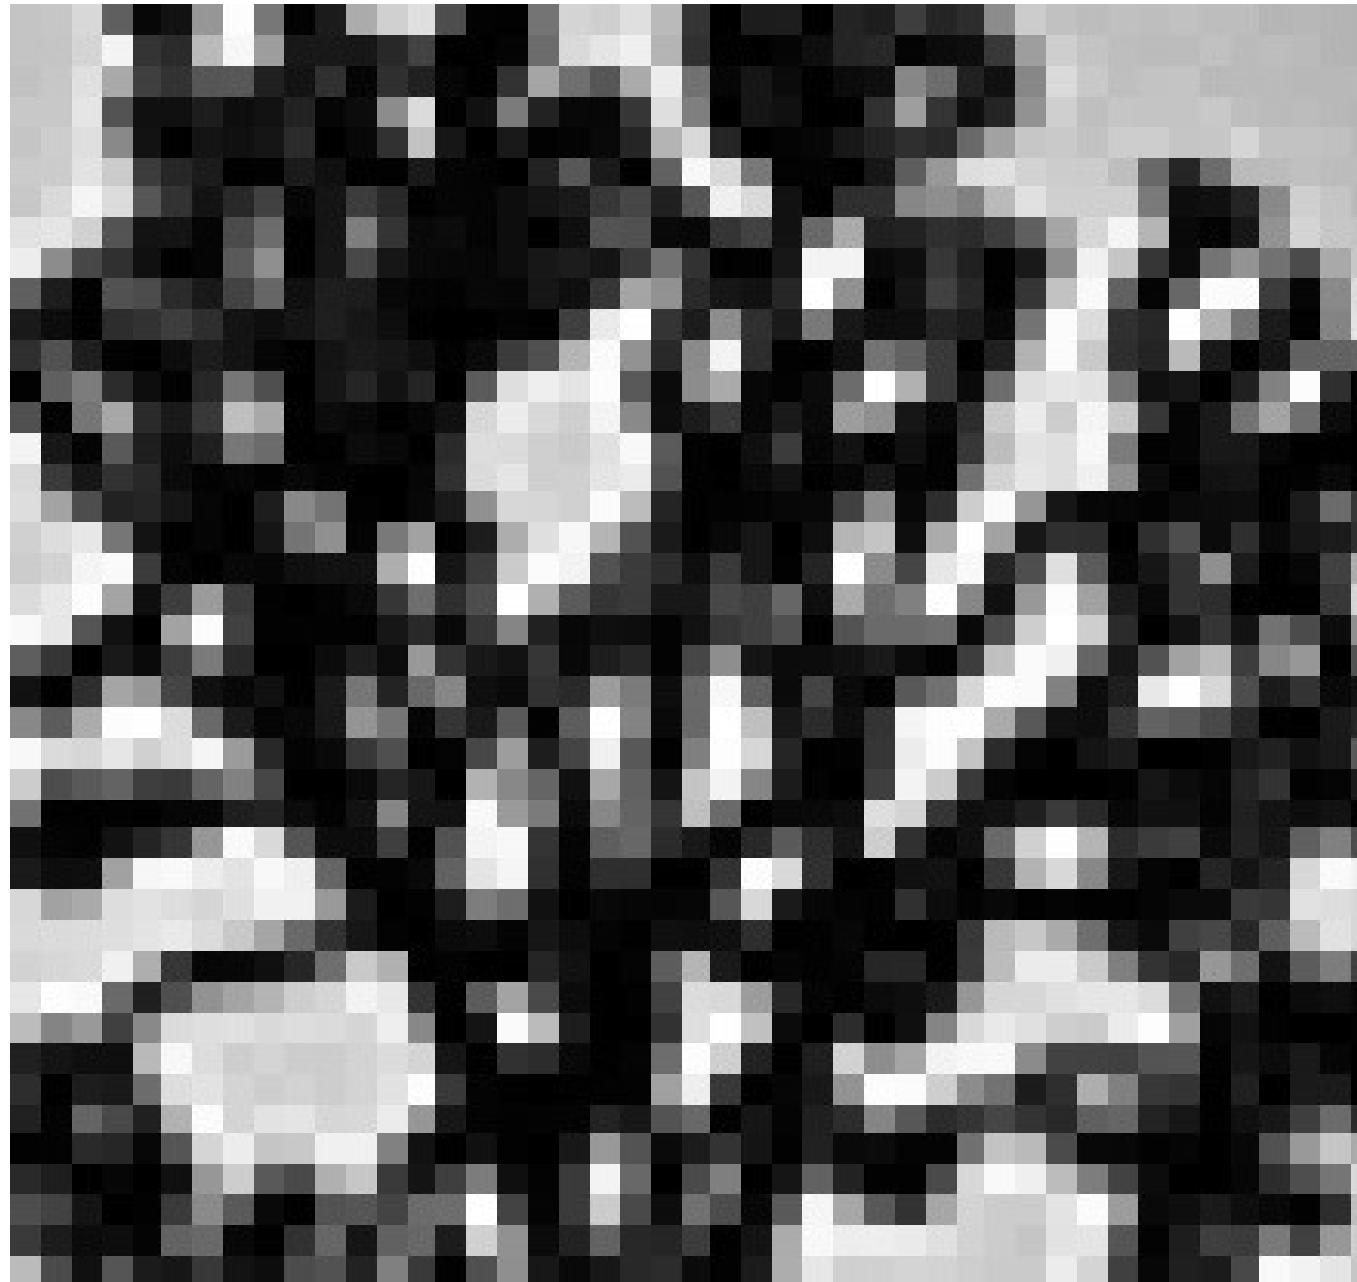
- 4 edges



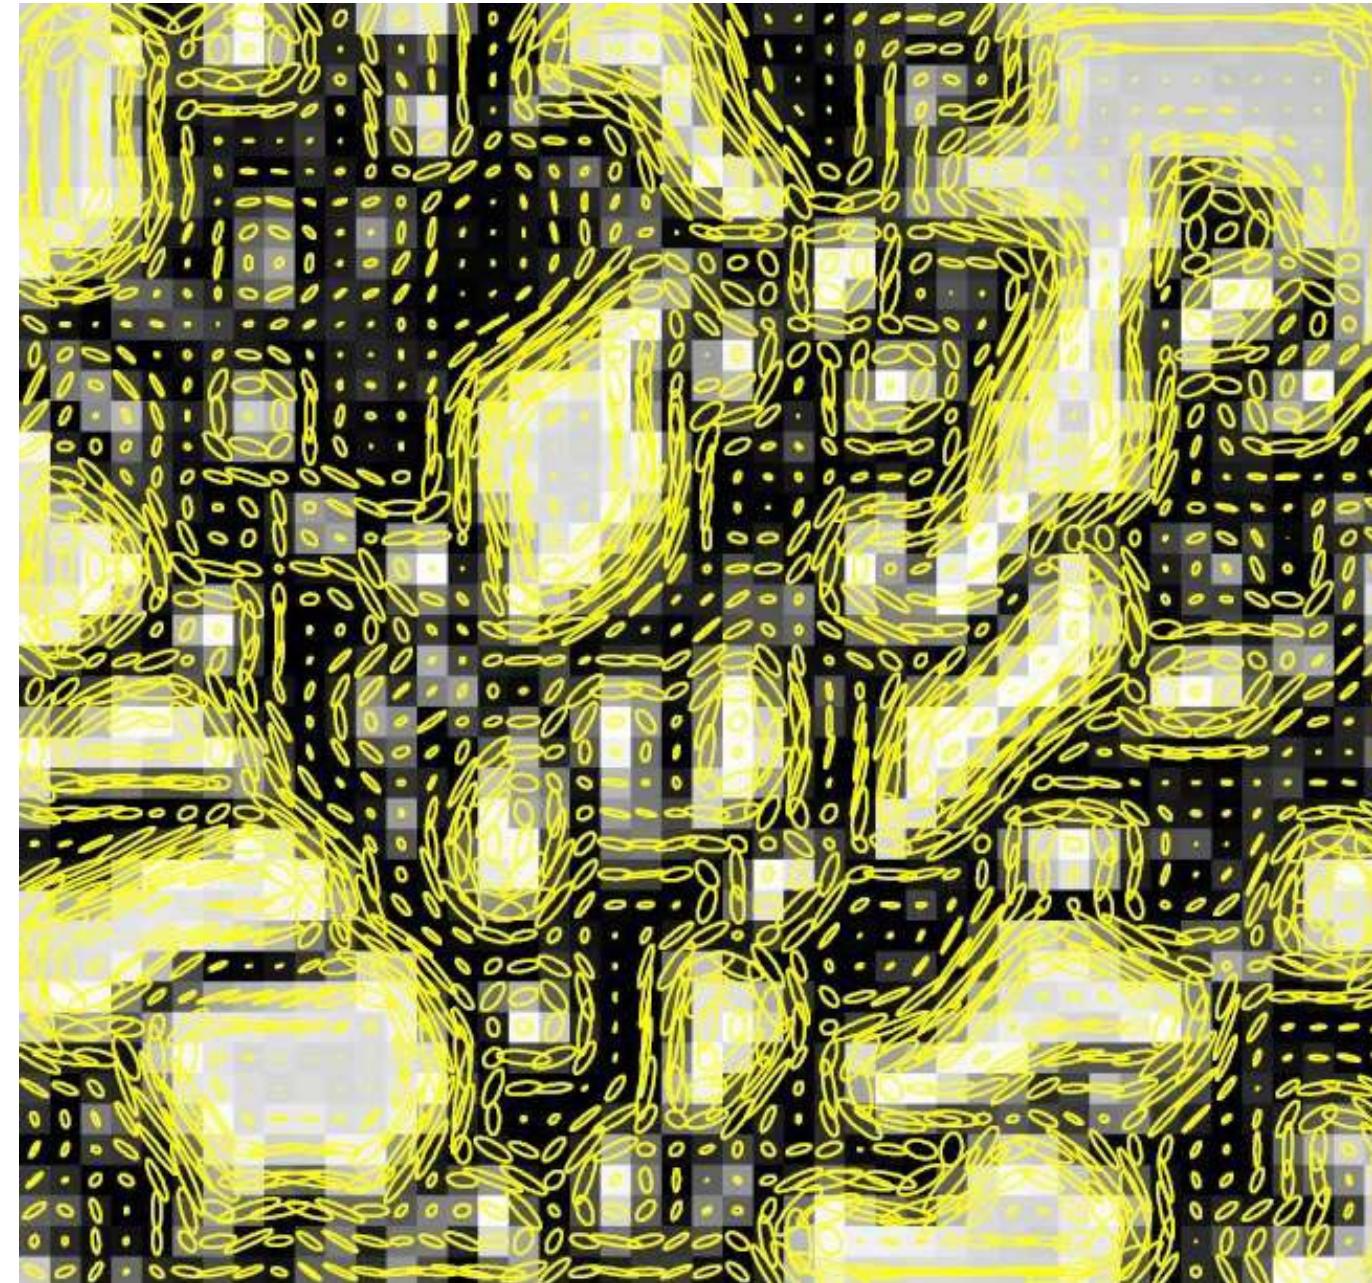
- 4 flat regions



# Visualization of the Second Moment Matrix



# Visualization of the Second Moment Matrix



For “cornerness”, we only care about the ‘steepness’, not the rotation. Can we ignore this somehow?

# Linear Algebra: Review

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- $M$  is **symmetric**. Symmetric matrices have orthogonal eigenvectors (i.e., a basis).
- $M$  is **square**. Square matrices are diagonalizable if some matrix  $P$  exists s.t.  $M = P^{-1}AP$ , where  $A$  has only diagonal entries and  $P$  represents a change of basis (in 2D, a rotation).

# Interpreting the Second Moment Matrix

Diagonalization of  $M$ :

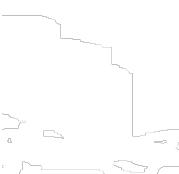
$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$

Consider a horizontal “slice” of  $E(u, v)$ :  
This defines an ellipse.

$$[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$

Form of standard ellipse:

Centered at origin and oriented along the axes


$$\frac{x^2}{a^2} + \frac{y^2}{b^2} = 0$$

$$[x \ y] \begin{bmatrix} \frac{1}{a^2} & 0 \\ 0 & \frac{1}{b^2} \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = 0$$

$$\lambda_{max} = \frac{1}{a^2} \implies a = (\lambda_{max})^{-1/2}$$

$$\lambda_{min} = \frac{1}{b^2} \implies b = (\lambda_{min})^{-1/2}$$

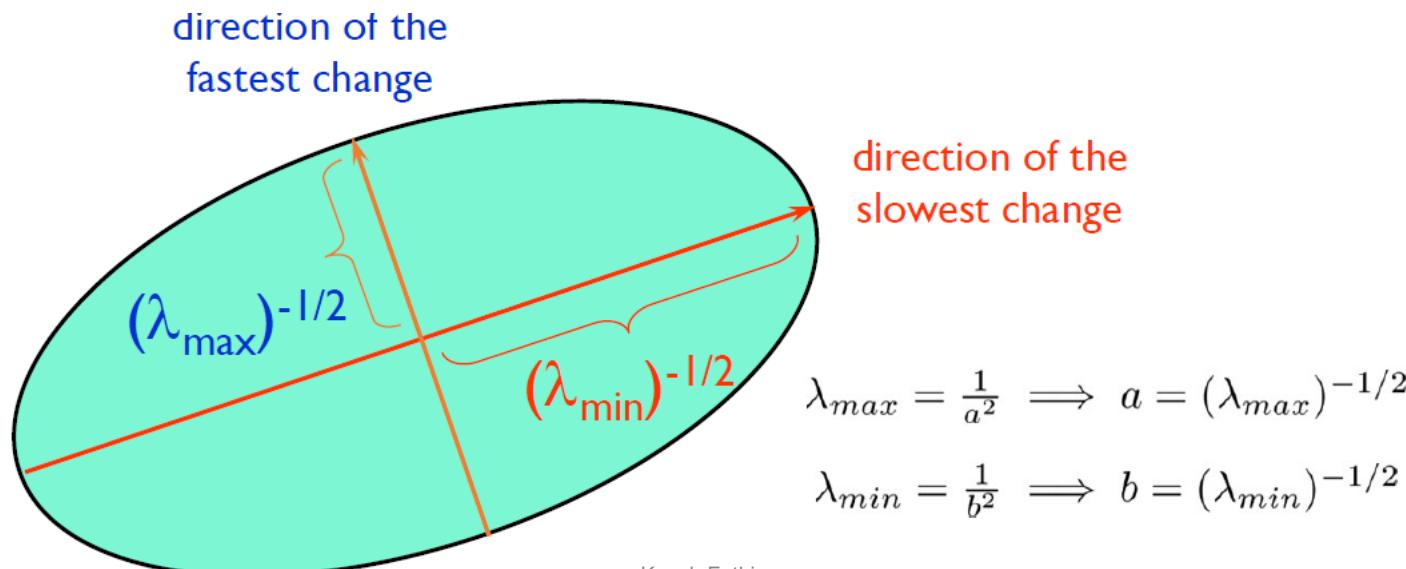
# Interpreting the Second Moment Matrix

Diagonalization of  $M$ :

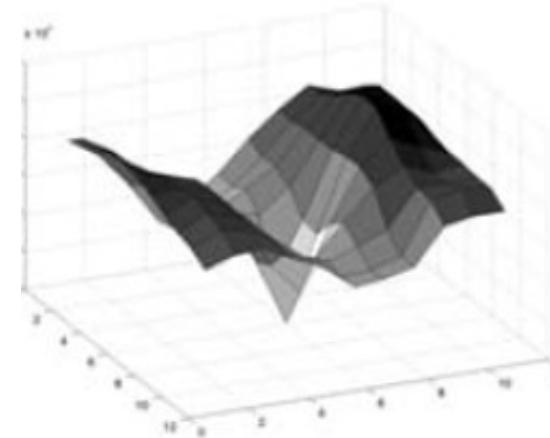
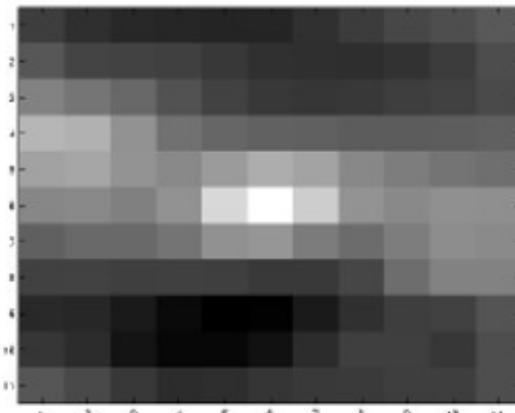
Consider a horizontal “slice” of  $E(u, v)$ :  
This defines an ellipse.

$$[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$

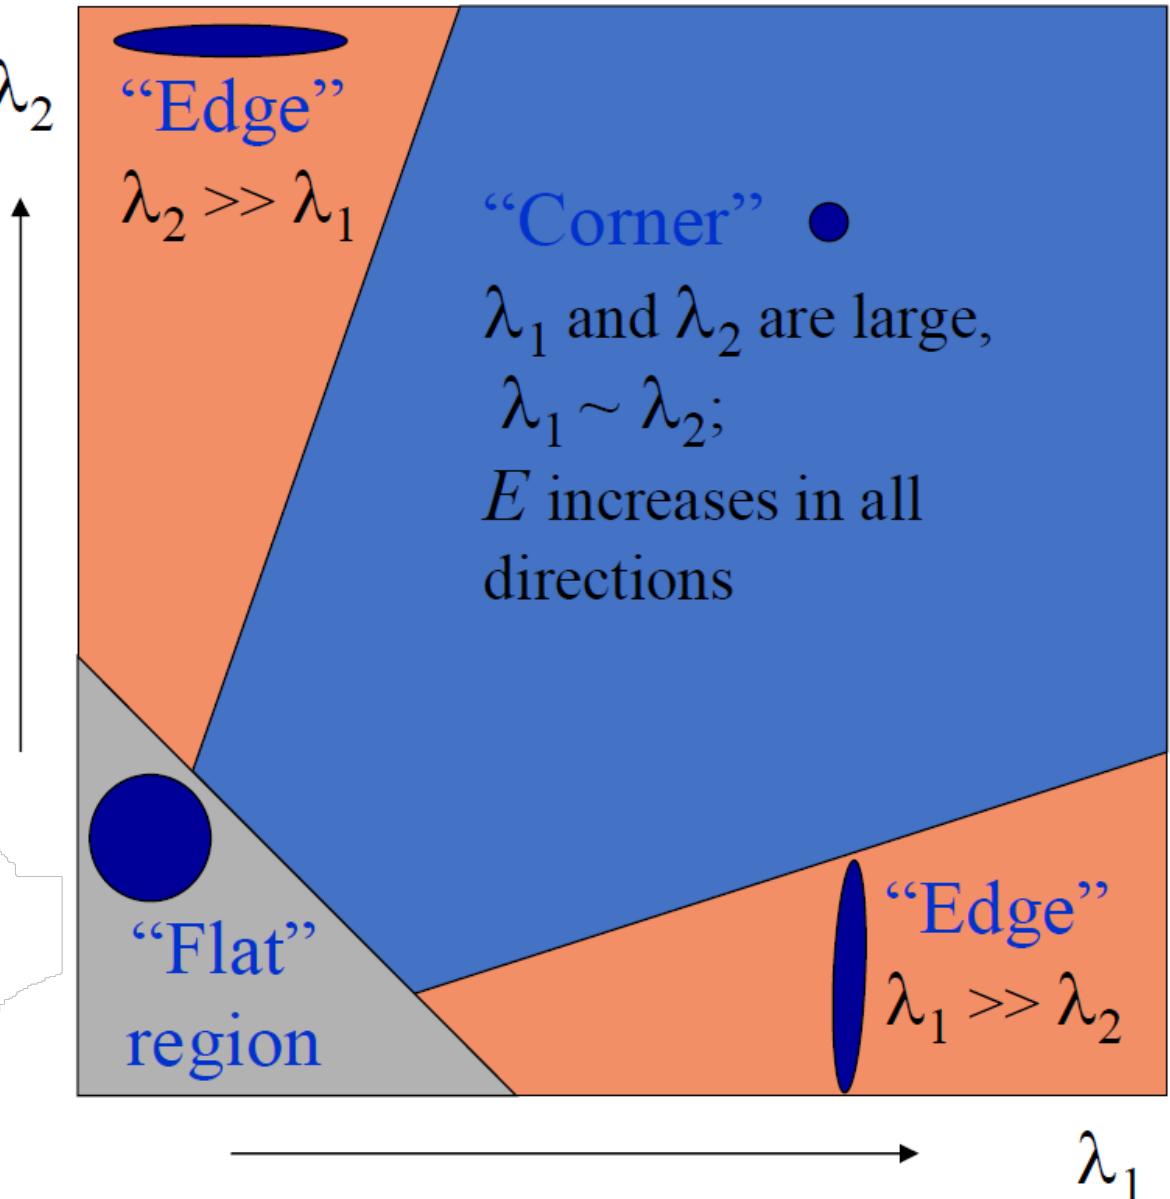
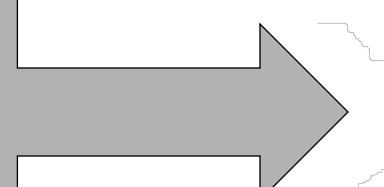
The axis lengths of the ellipse are determined by the eigenvalues, and the orientation is determined by a rotation matrix  $R$



# Classification of image points using eigenvalues of $M$



$\lambda_1$  and  $\lambda_2$  are small;  
 $E$  is almost constant  
in all directions

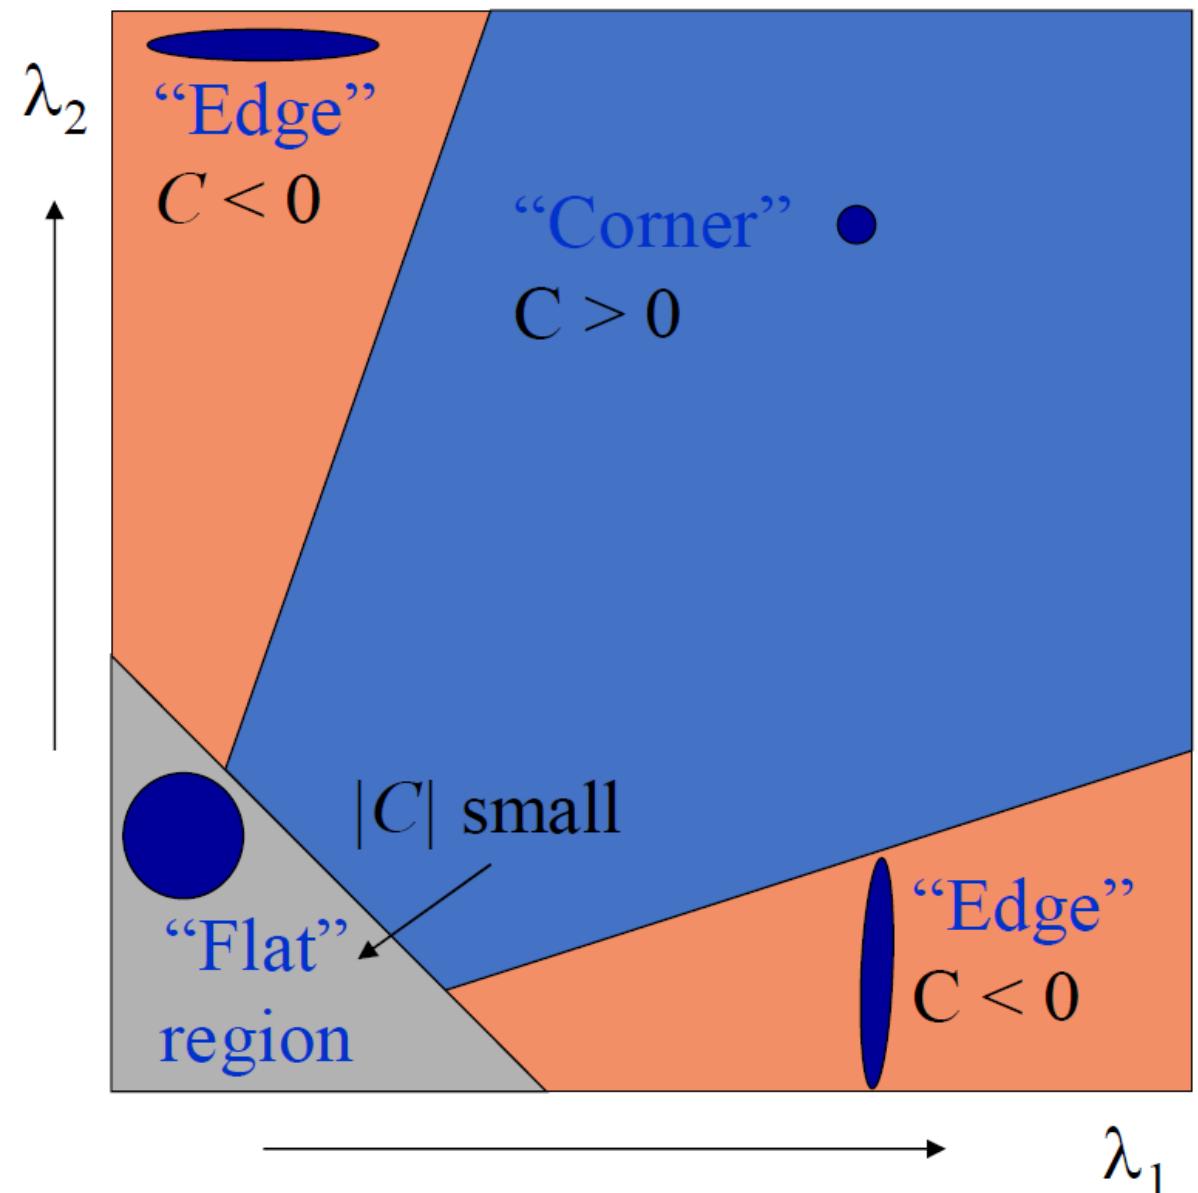


# Classification of image points using eigenvalues of $M$

**Cornerness** score:

$$C = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

$\alpha$ : some constant ( $\sim 0.04$  to  $0.06$ )



# Linear Algebra: Review

$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$

- **Determinant** of a diagonal matrix is the **product** of all **eigenvalues** ( $A$  is diagonal).
- **Trace** of a square matrix is the **sum** of its **diagonal** entries; and is the sum of its **eigenvalues**.

# Classification of image points using eigenvalues of $M$

**Cornerness** score:

$$C = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

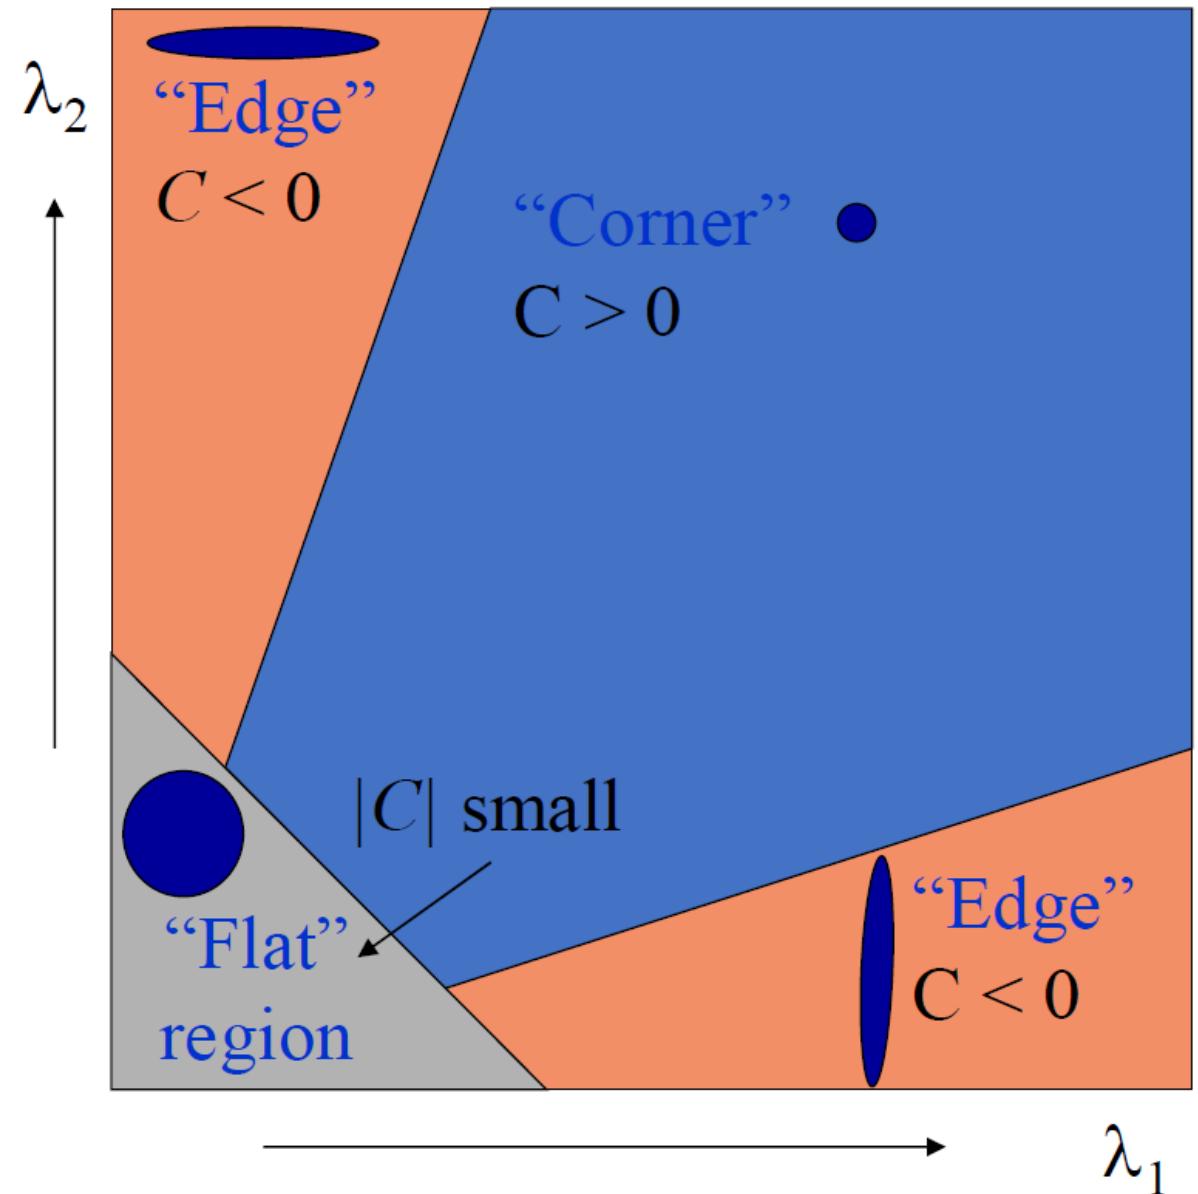
$\alpha$ : some constant ( $\sim 0.04$  to  $0.06$ )

Remember your linear algebra:

- Determinant:  $\det(A) = \prod_{i=1}^n \lambda_i = \lambda_1 \lambda_2 \cdots \lambda_n$
- Trace:  $\text{tr}(A) = \sum_i \lambda_i$

$$C = \det(M) - \alpha \text{trace}(M)^2$$

Why should we compute  $C$  this way?



# Harris Corner Detector

## Algorithm

- 1) Compute  $M$  matrix for each window to recover a cornerness score  $C$ .  
Note: We can find  $M$  purely from the per-pixel image derivatives!
- 2) Threshold to find pixels which give large corner response ( $C >$  threshold).
- 3) Find the local maxima pixels, i.e., non-maxima suppression.

C.Harris and M.Stephens. ["A Combined Corner and Edge Detector."](#) Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

# Harris Corner Detector

## Algorithm

1. Input image: compute  $M$  at each pixel.
2. Compute image derivatives (optionally, blur first).
3. Compute  $M$  components as squares of derivatives.
4. Gaussian filter  $g()$  with width  $s$

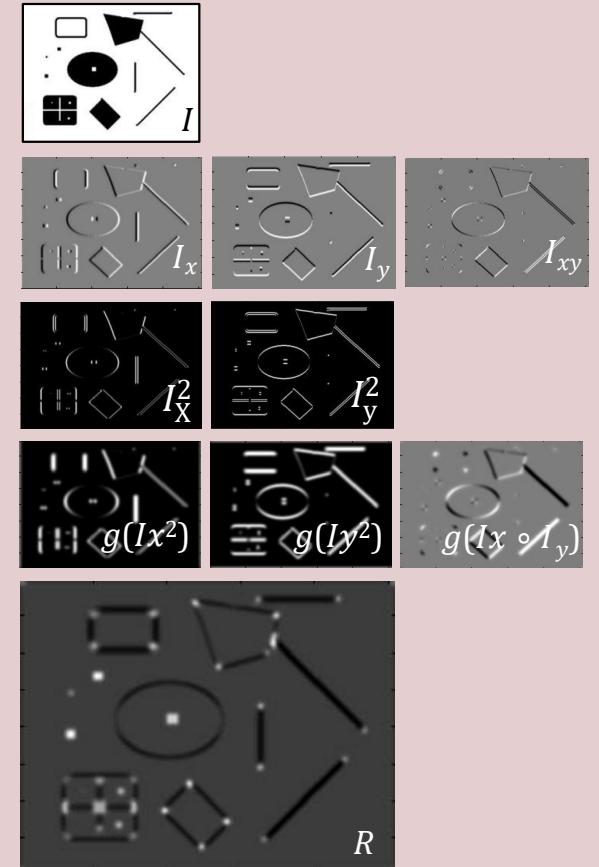
$$g(I_x^2), g(I_y^2), g(I_x \circ I_y)$$

5. Compute cornerness

$$C = \det(M) - \alpha \operatorname{trace}(M)^2 = g(I_x^2) \circ g(I_y^2) - g(I_x \circ I_y)^2 - \alpha [g(I_x^2) + g(I_y^2)]^2$$

Reminder:  $a \circ b$  is  
Hadamard product  
(element-wise  
multiplication)

5. Threshold on  $C$  to pick high cornerness
6. Non-maximal suppression to pick peaks.

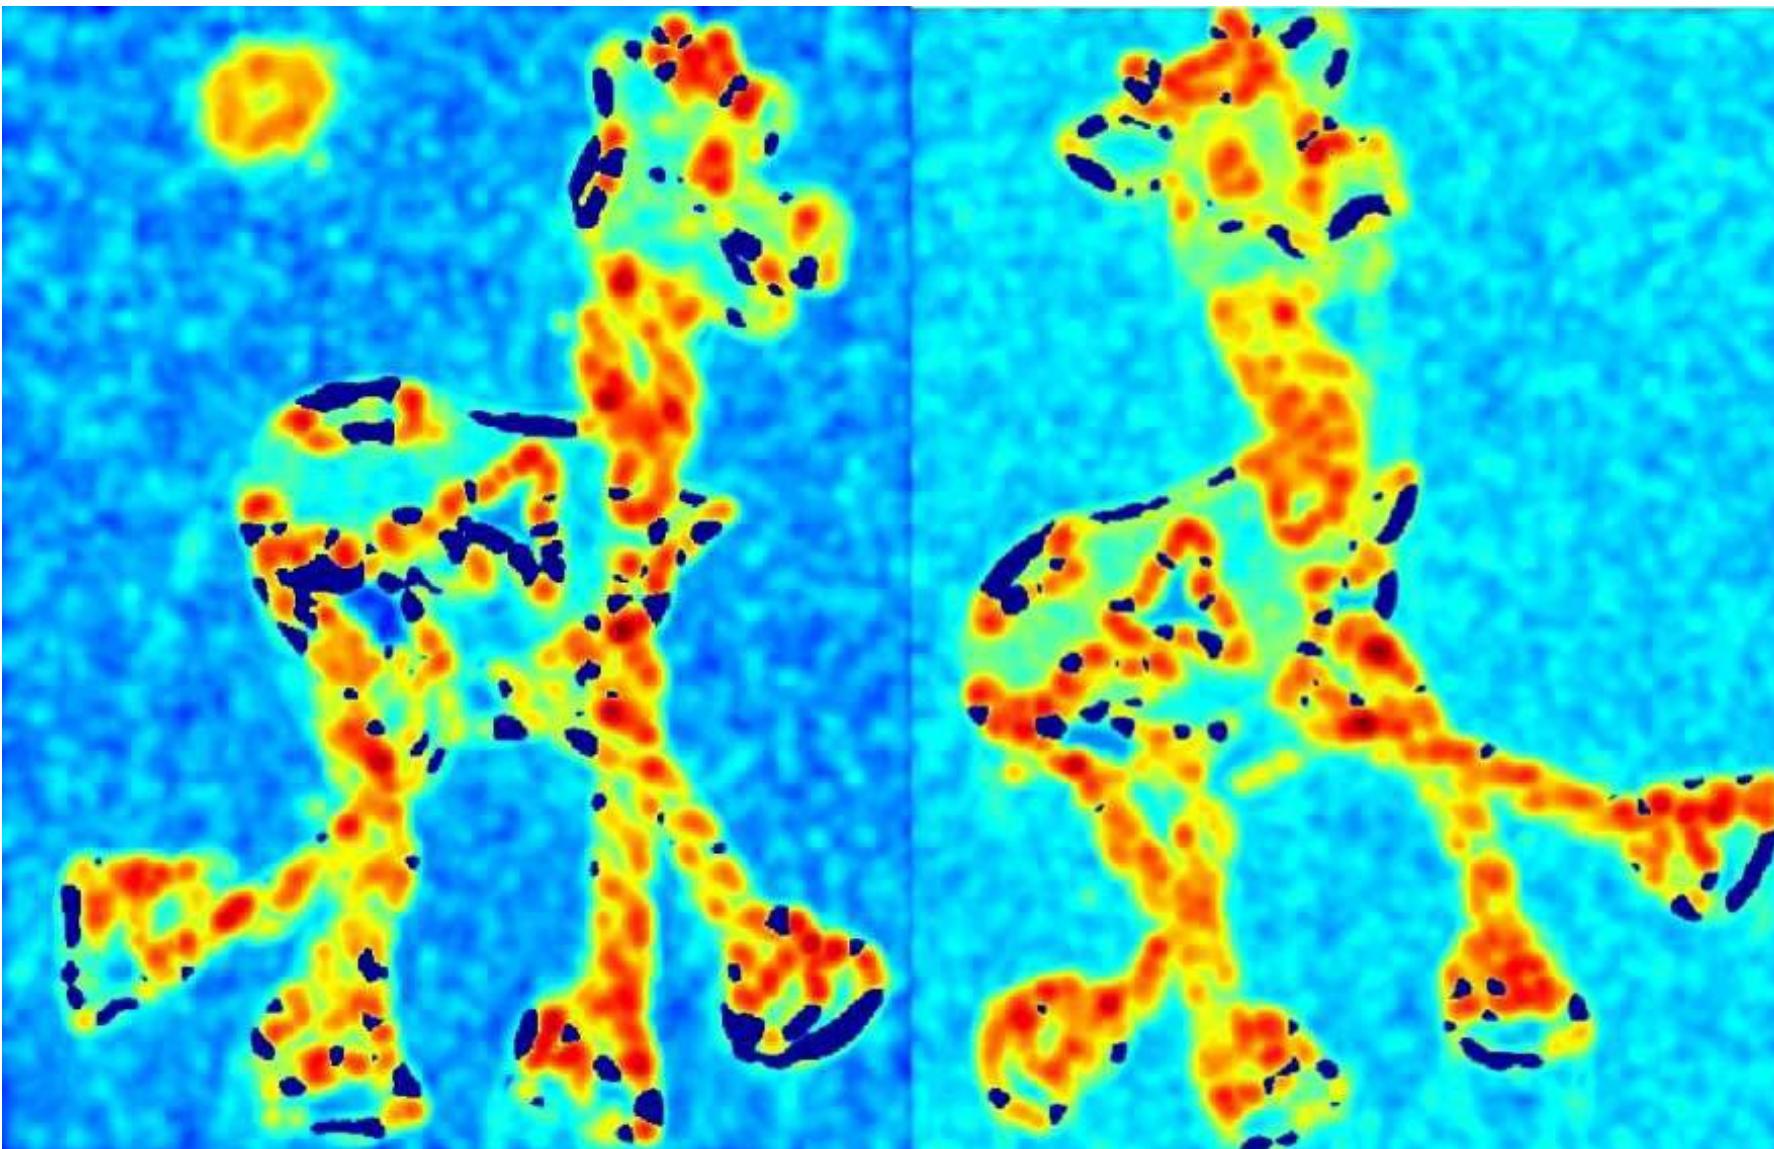


# Harris Detector: Steps



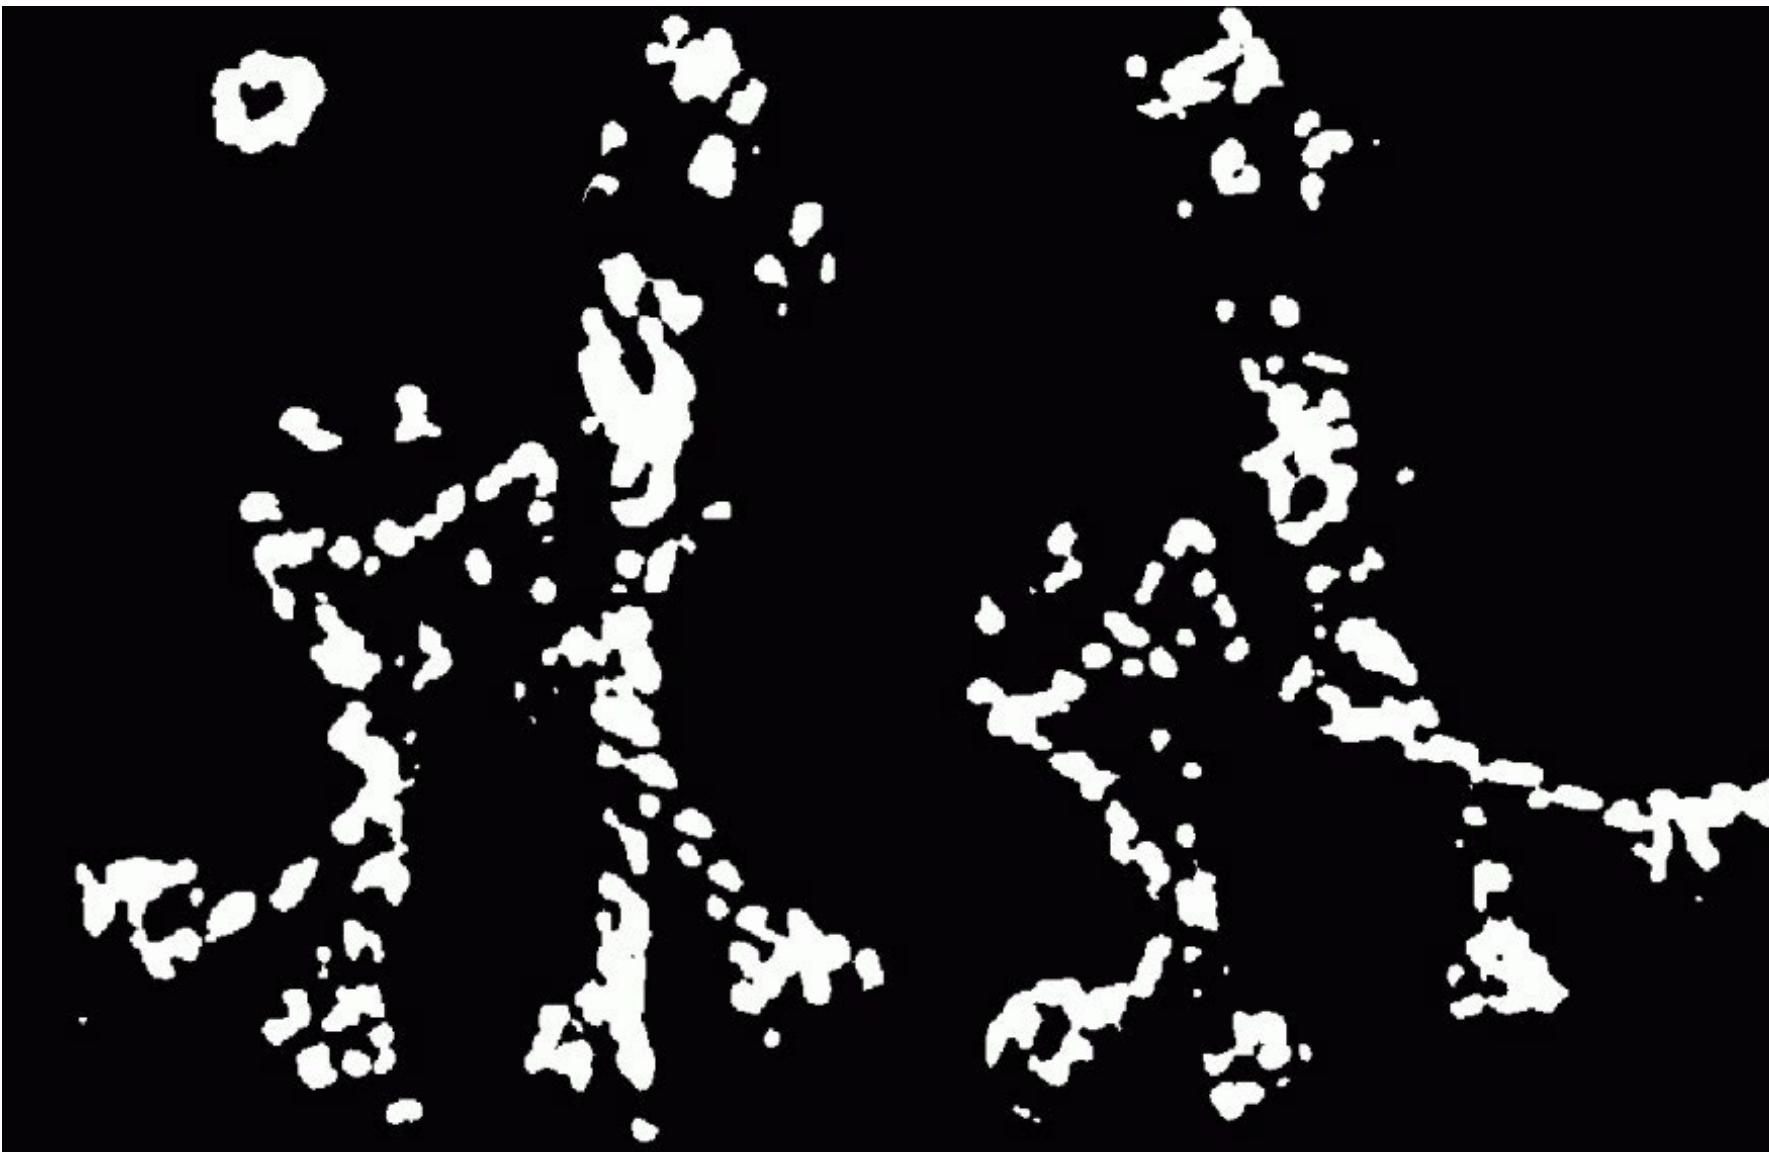
# Harris Detector: Steps

Compute corner response  $C$



# Harris Detector: Steps

Find points with large corner response:  $C > \text{threshold}$



# Harris Detector: Steps

Take only the points of local maxima of  $C$  (non-maximum suppression)



# Harris Detector: Steps



# Invariance and Covariance

# Invariance and Covariance

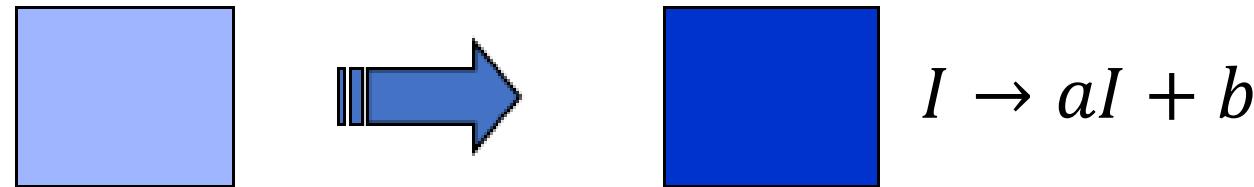
Are locations **invariant** to photometric transformations and **covariant** to geometric transformations?

- **Invariance**: image is transformed and corner locations do not change
- **Covariance / Equivariance**: if we have two transformed versions of the same image, features should be detected in corresponding locations

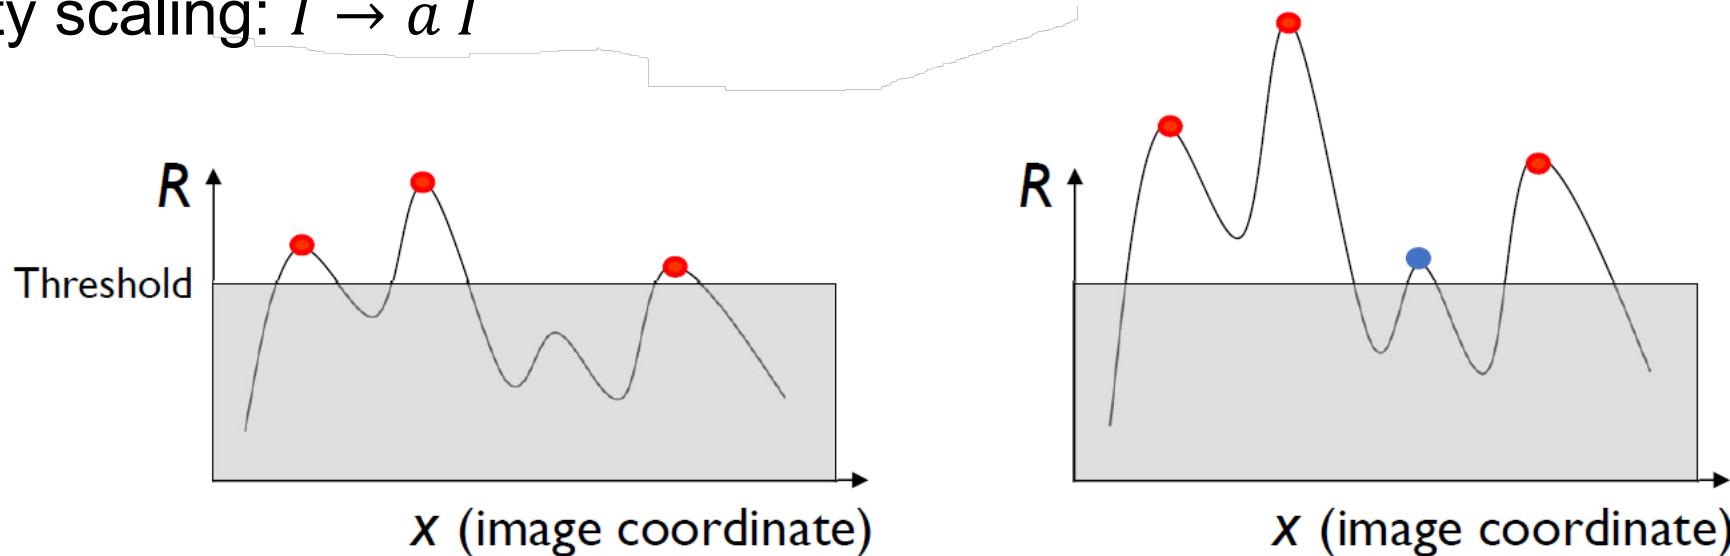


Kaveh Fathian

# Affine Intensity Change

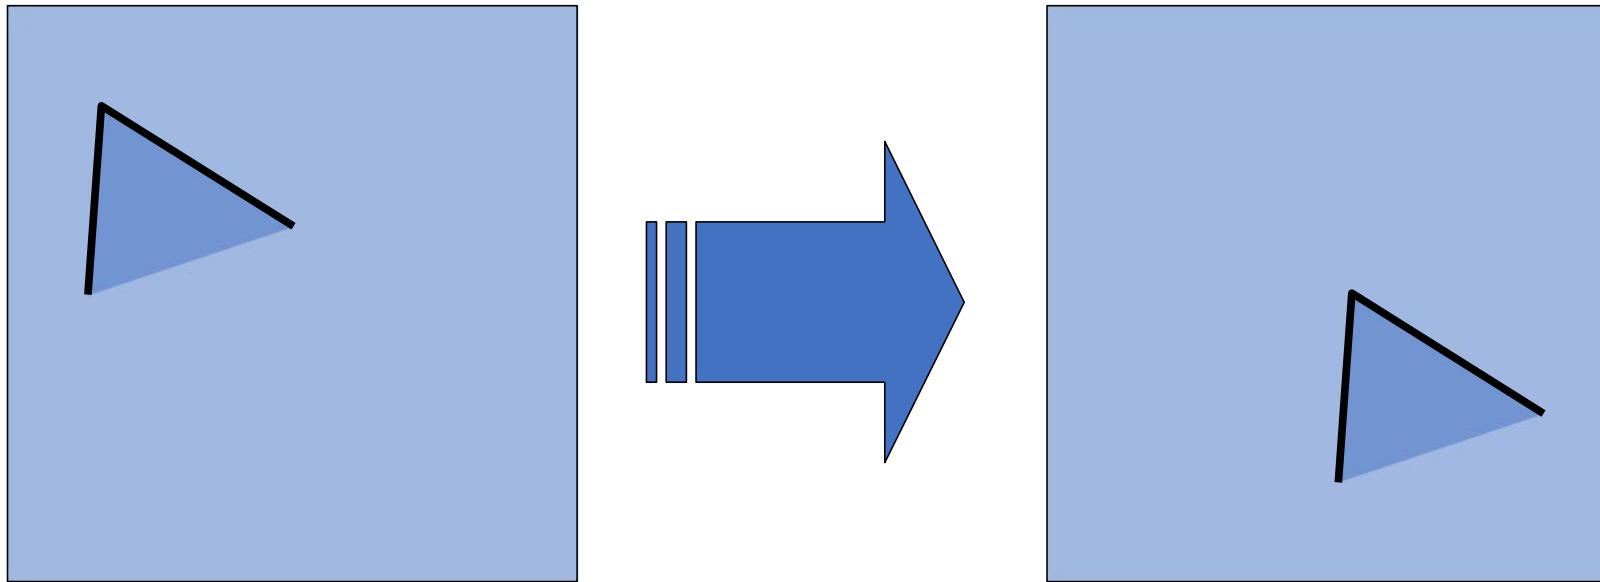


- Only derivatives are used  $\Rightarrow$  invariance to intensity shift  $I \rightarrow I + b$
- Intensity scaling:  $I \rightarrow aI$



**Partially invariant to affine intensity change**

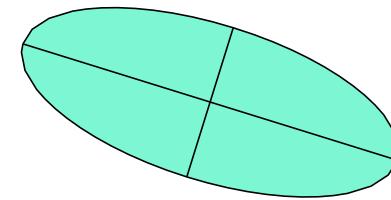
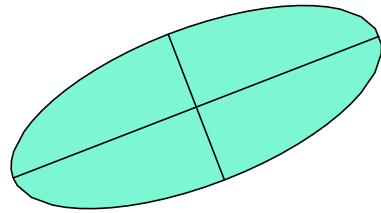
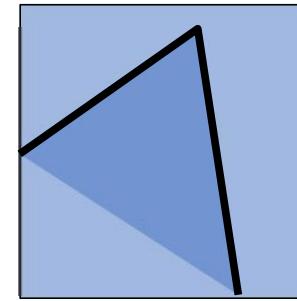
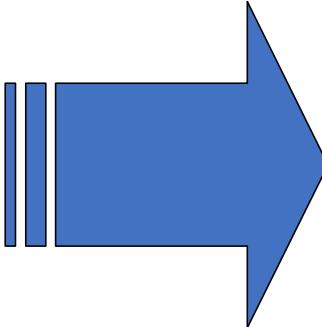
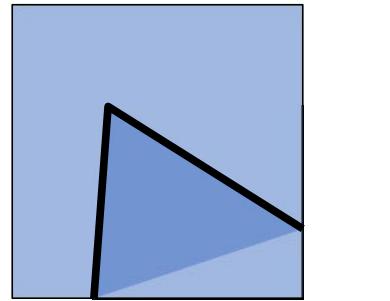
# Image Translation



- Derivatives and window function are shift-invariant

**Corner location is covariant w.r.t. translation**

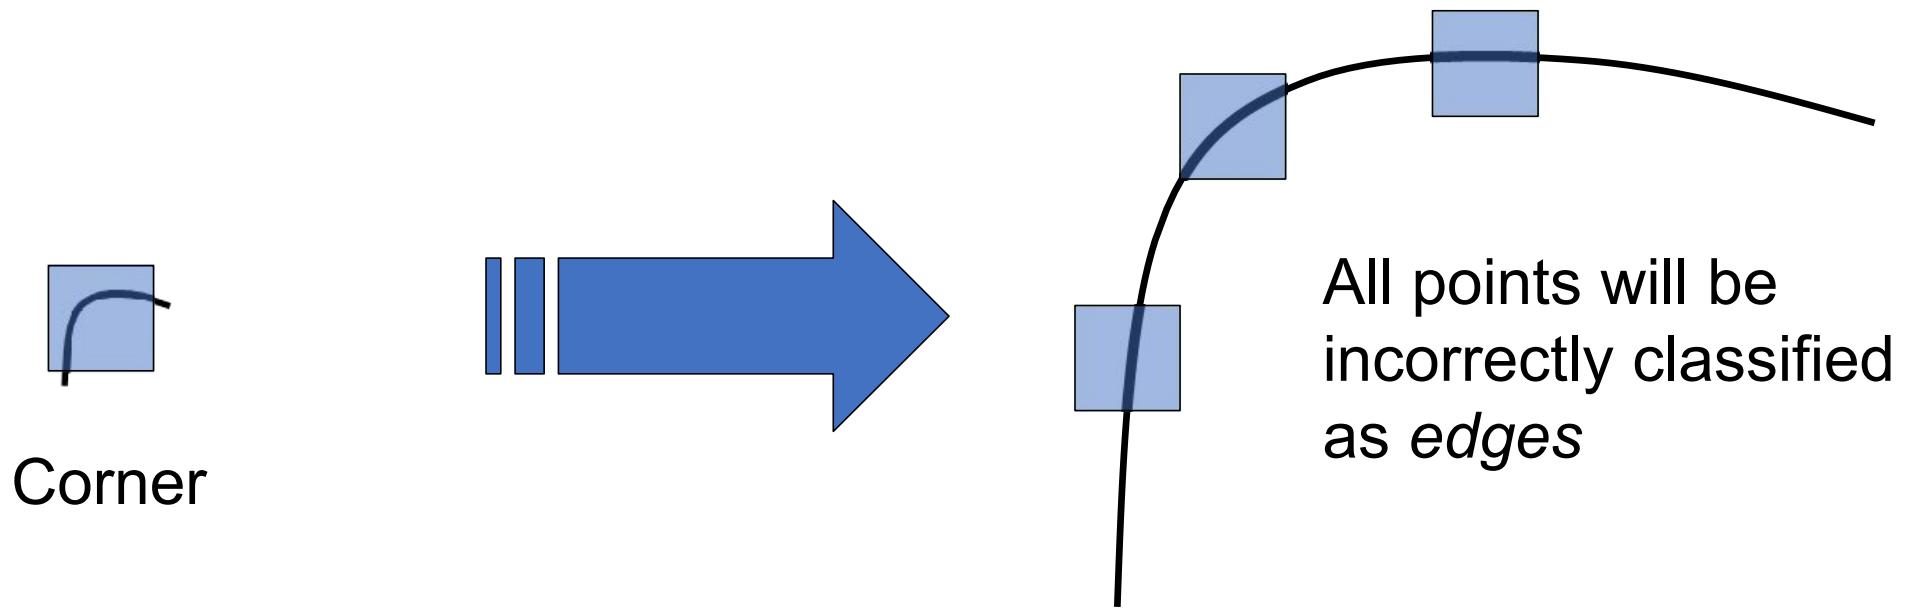
# Image Rotation



- Second moment ellipse rotates but its shape (i.e., eigenvalues) remains the same

**Corner location is covariant w.r.t. rotation**

# Scaling

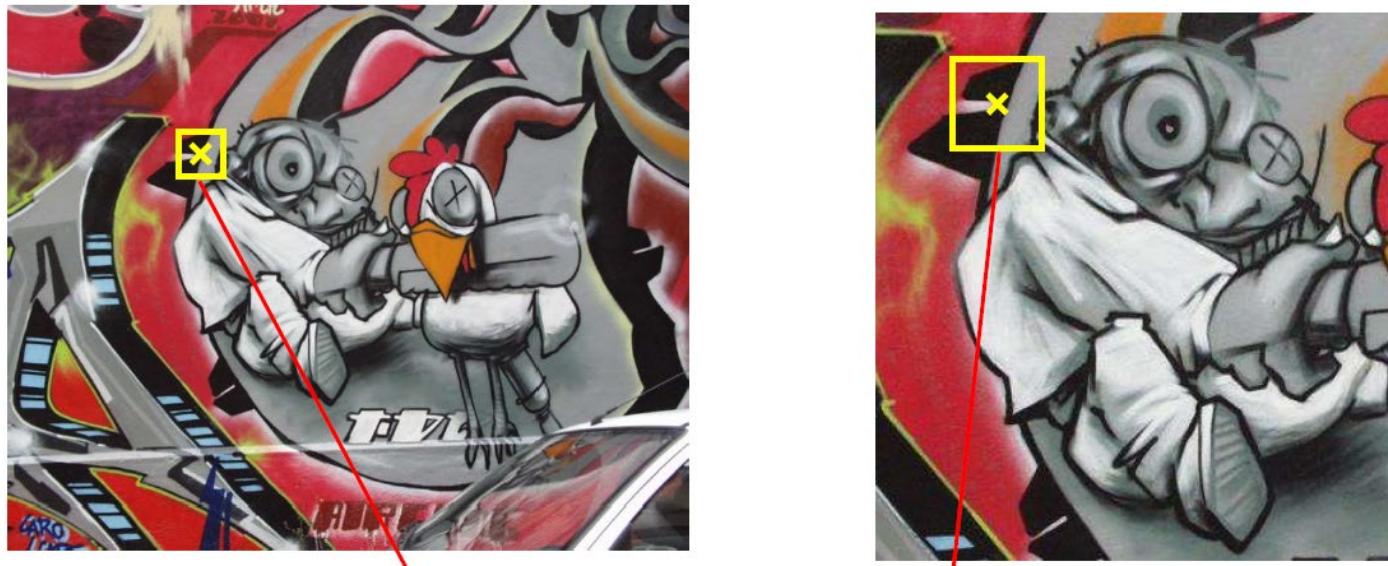


**Corner location is not covariant to scaling!**

# Multi-Scale Detection

# Automatic Scale Selection

- How can we make a feature detector scale-invariant?
- How can we automatically select the scale?



$$f(I_{i_1 \dots i_m}(x, \sigma)) = f(I_{i_1 \dots i_m}(x', \sigma'))$$

- How to find patch sizes at which  $f$  response is equal?
- What is a good  $f$ ?

# Intuition

Find local maxima in both position and scale

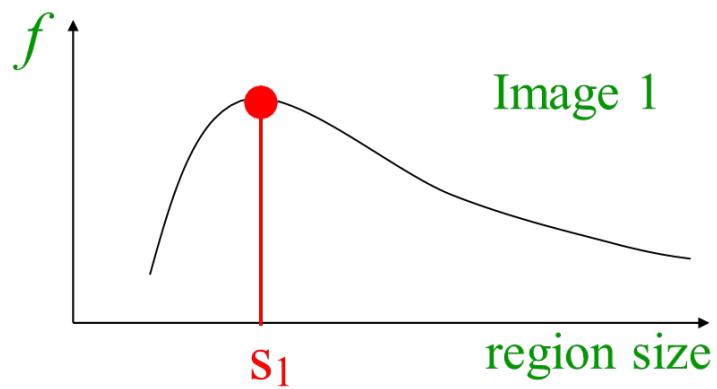
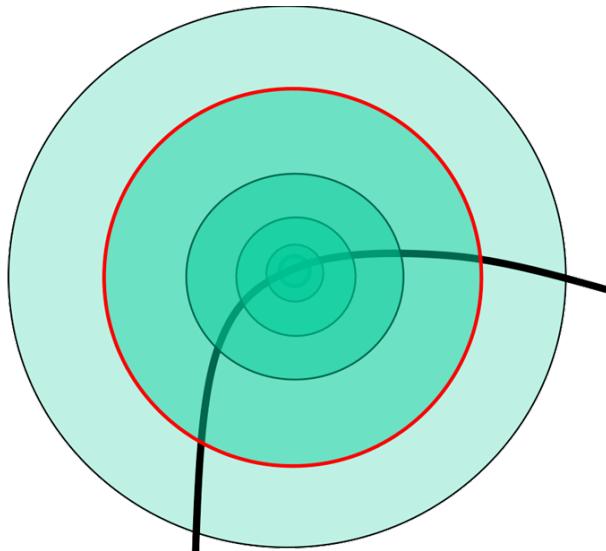


Image 1

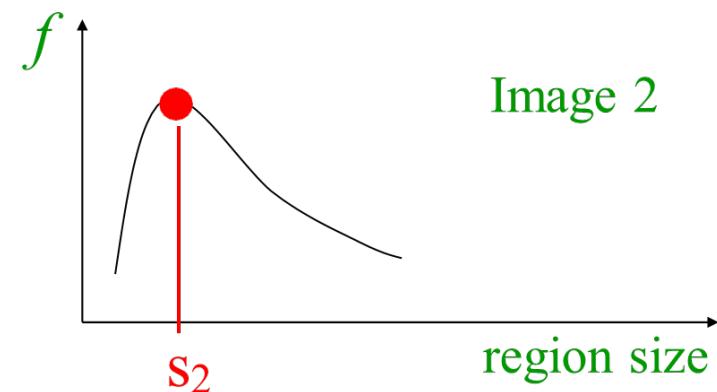
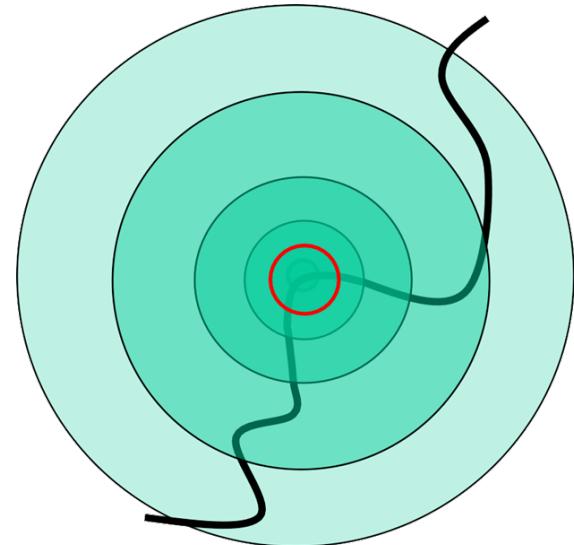
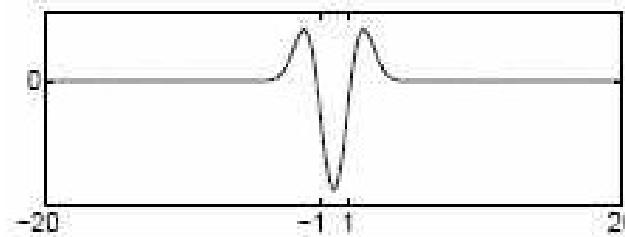


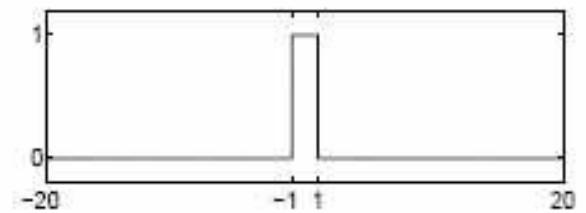
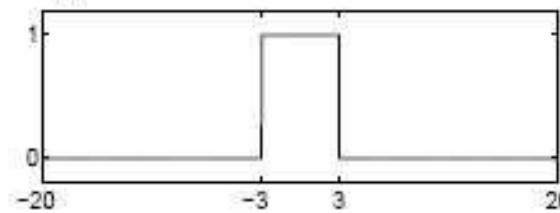
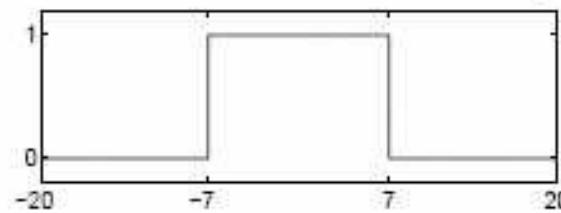
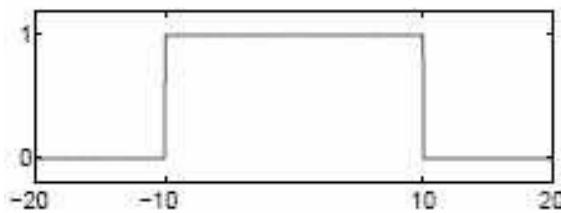
Image 2

# Intuition

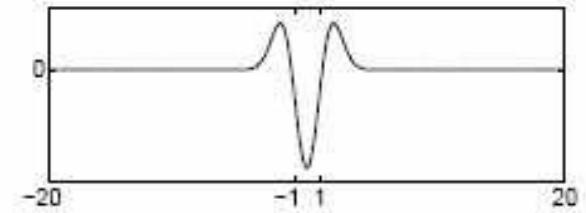
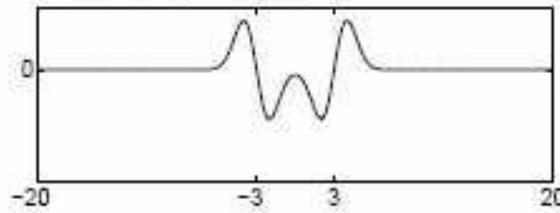
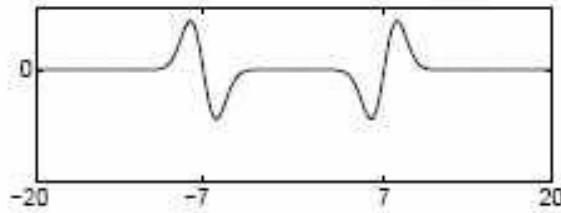
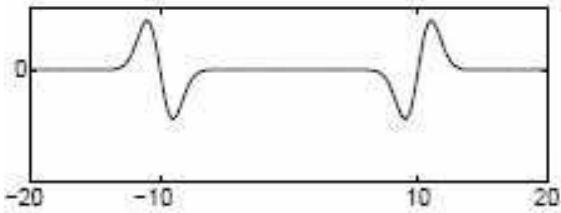


Laplacian filter

Original signal



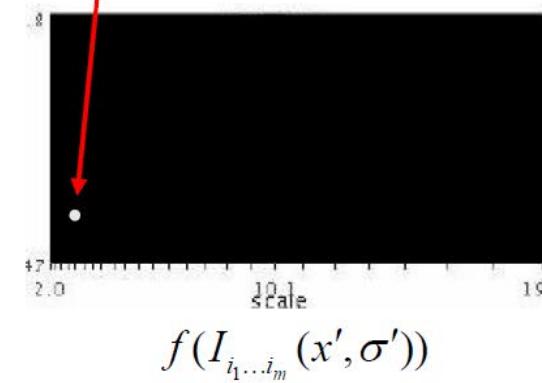
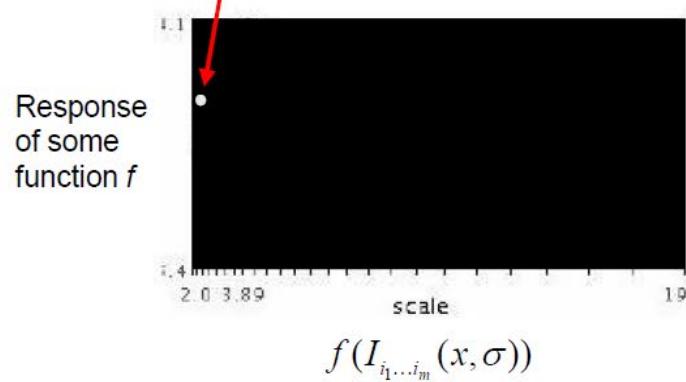
Convolved with Laplacian ( $\sigma = 1$ )



Highest response when the signal has the same characteristic scale as the filter

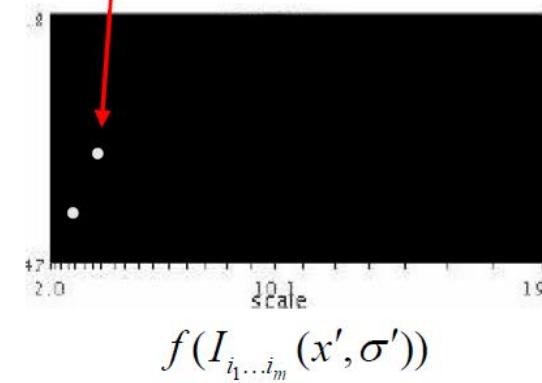
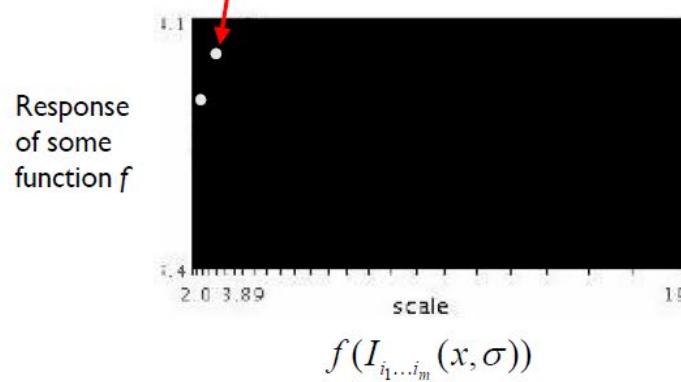
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



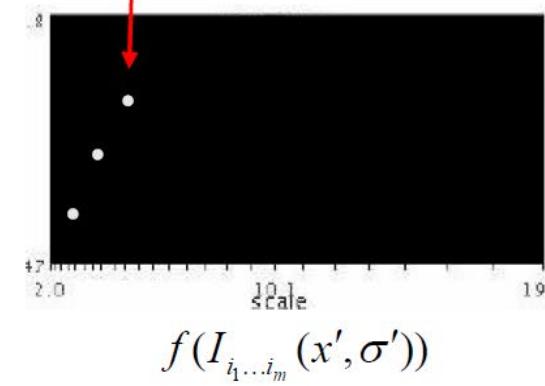
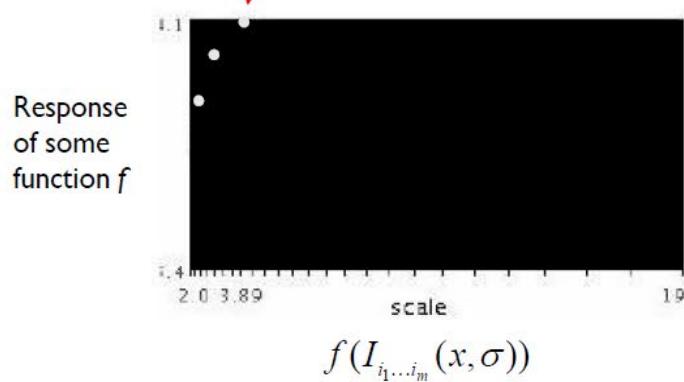
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



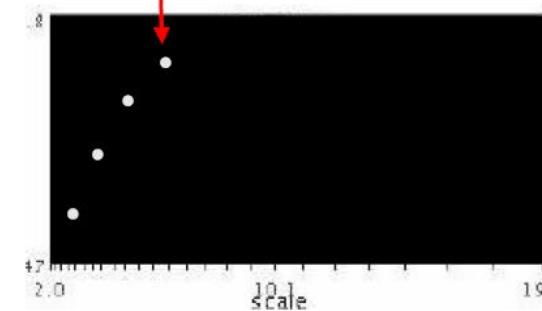
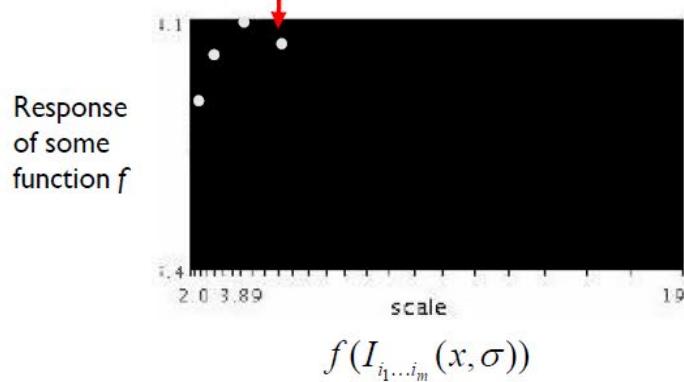
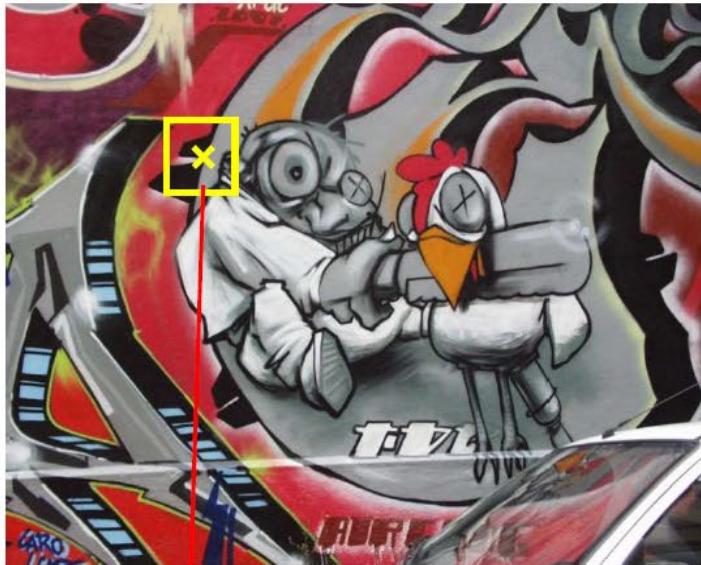
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



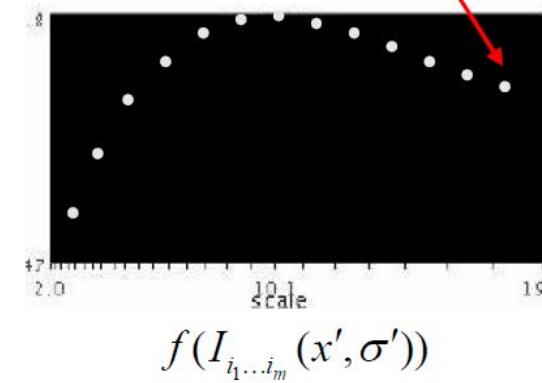
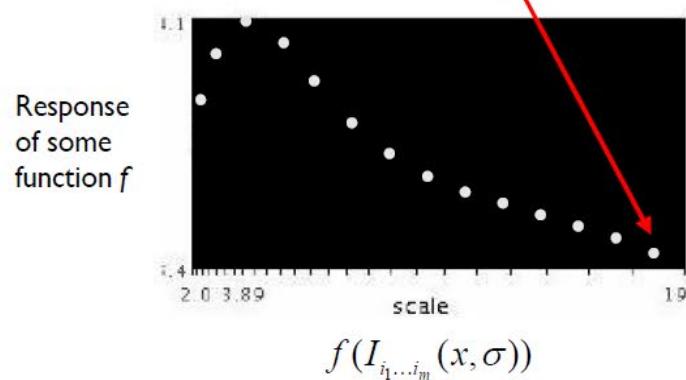
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



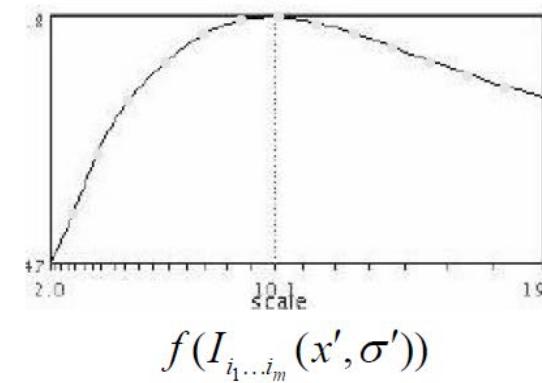
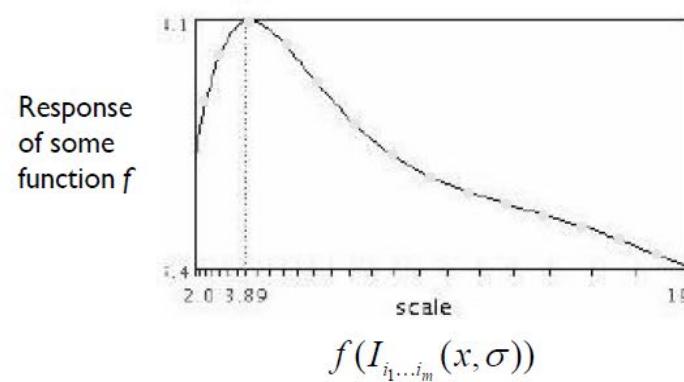
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



# Automatic Scale Selection

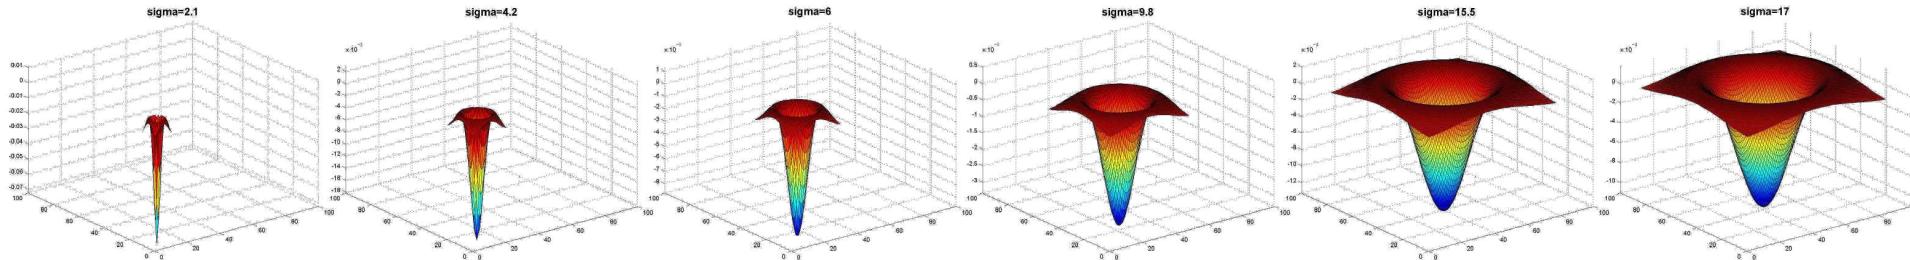
- Function responses for increasing scale (scale signature)



# Multi-scale blob detection



# What happens if you apply different Laplacian filters?



Full size

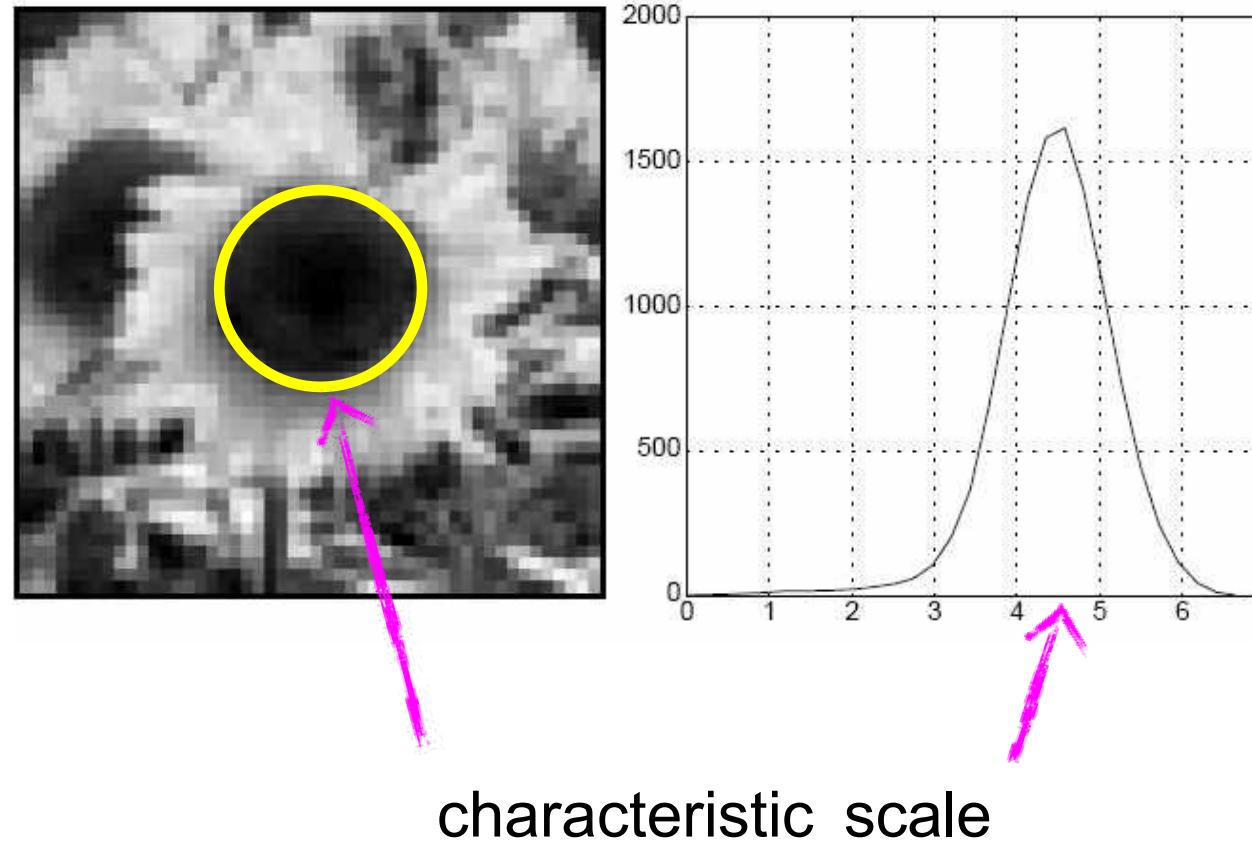


3/4 size

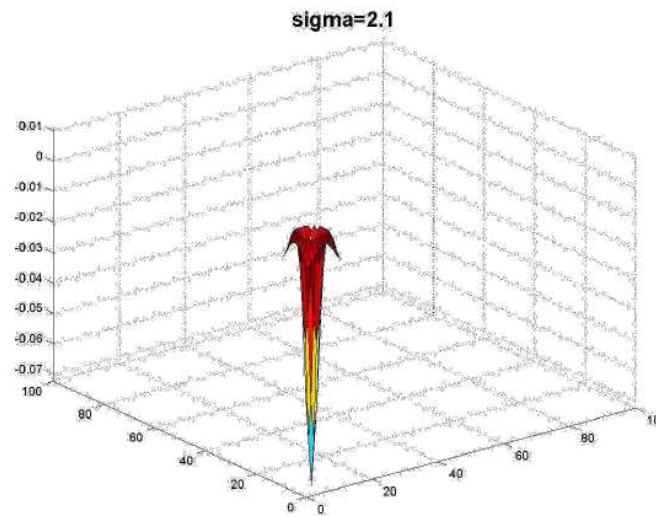


# Characteristic Scale

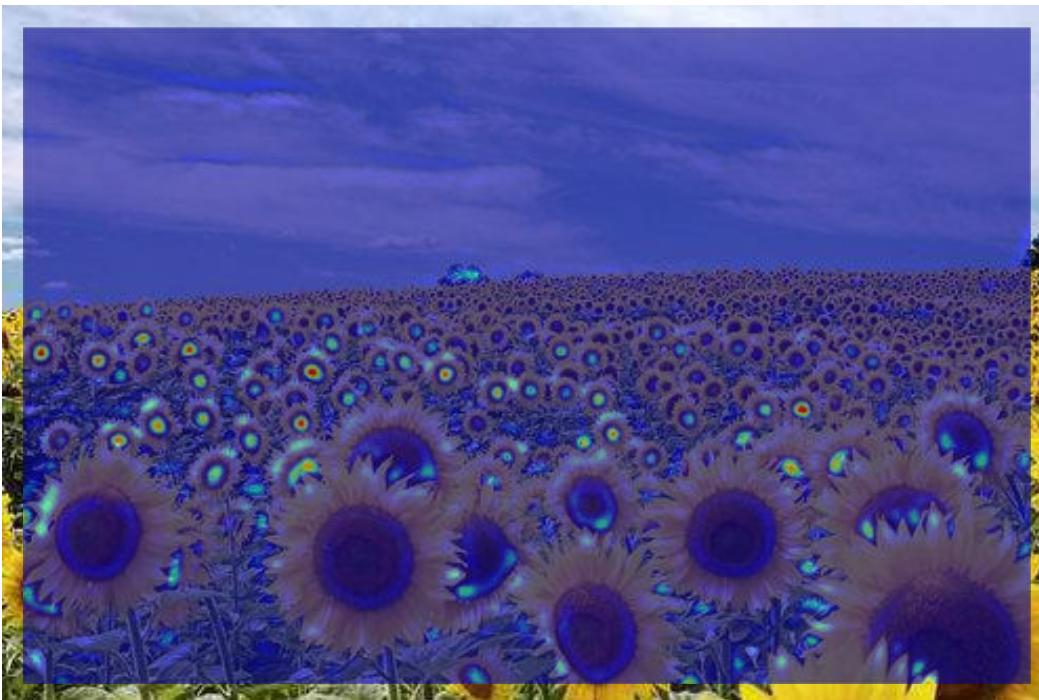
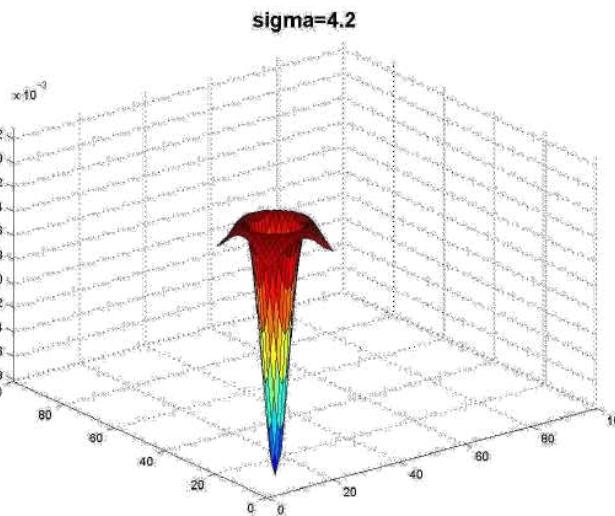
**Characteristic scale:** the scale that produces peak filter response

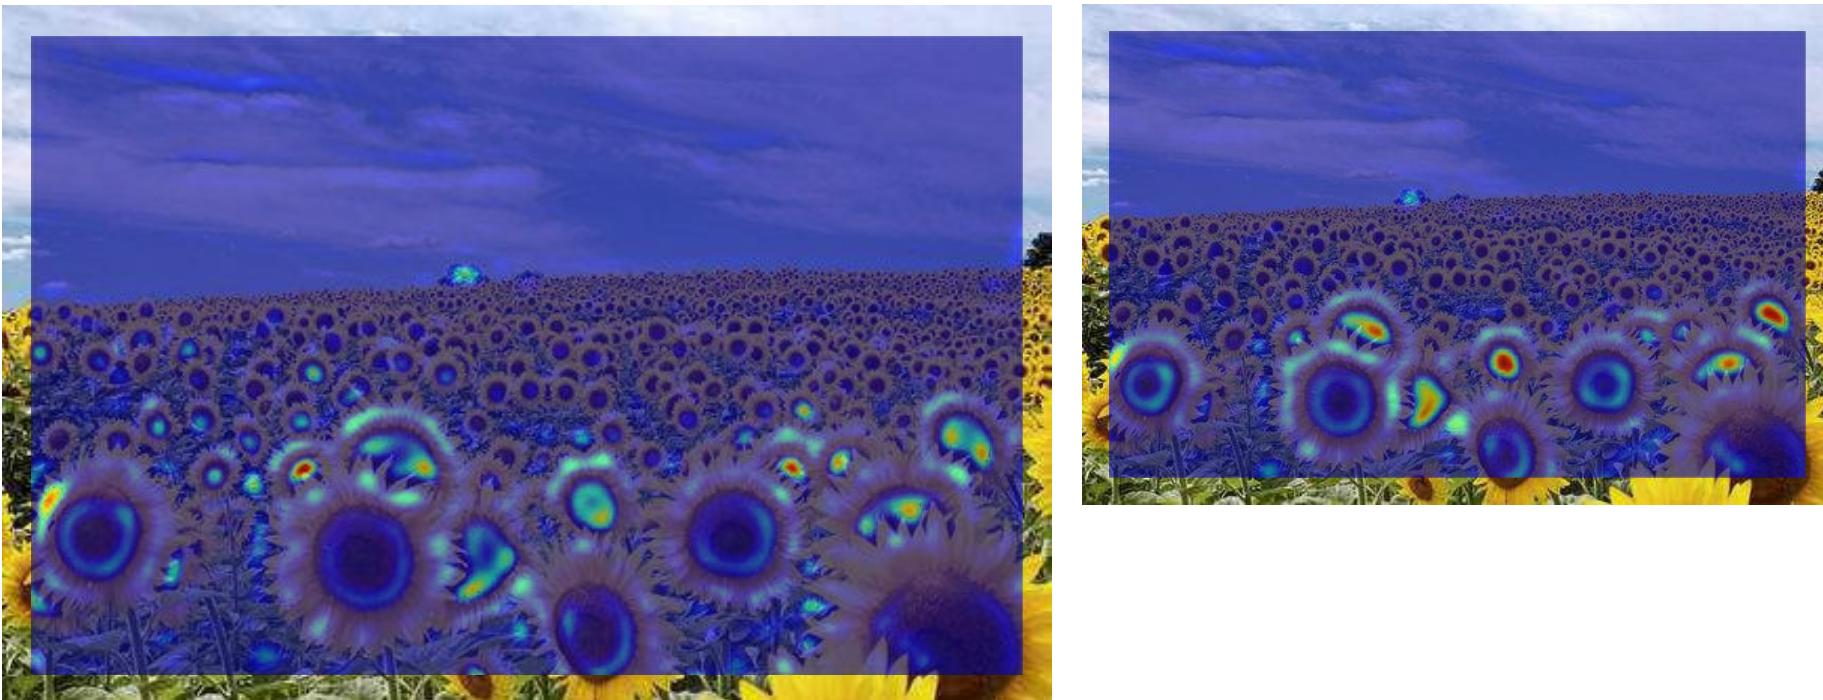
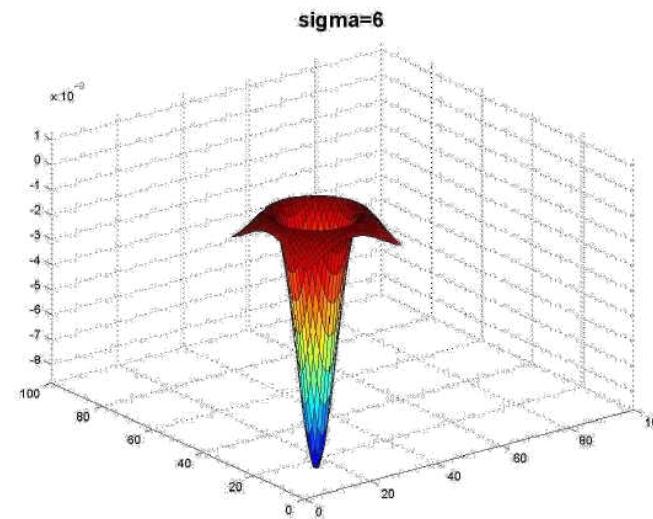


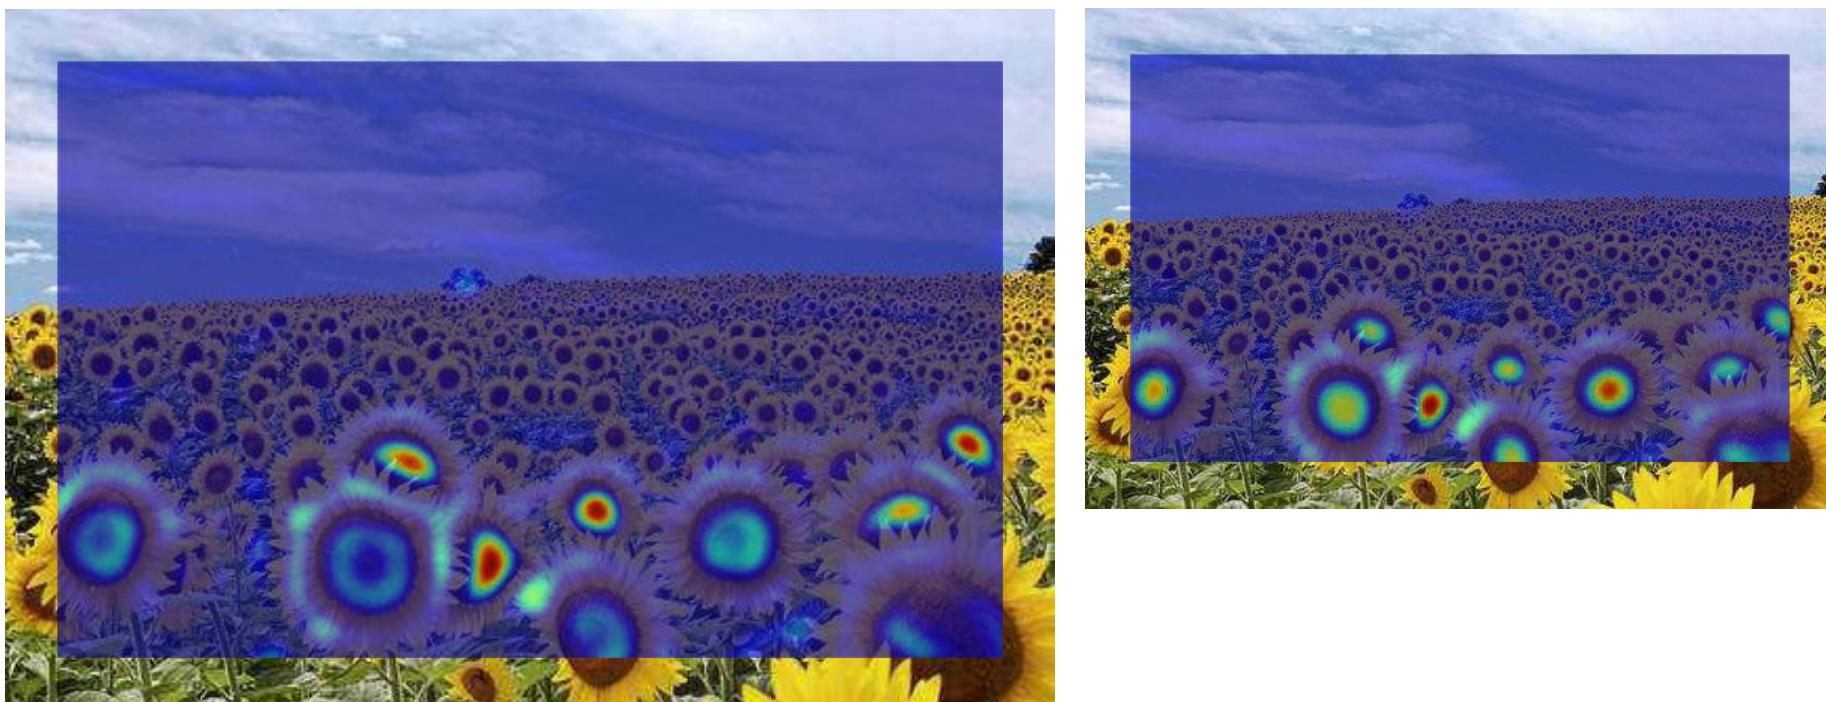
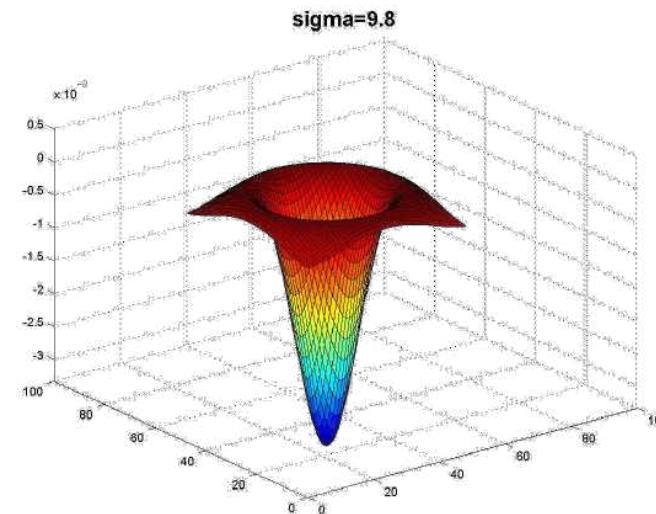
**we need to search over characteristic scales**

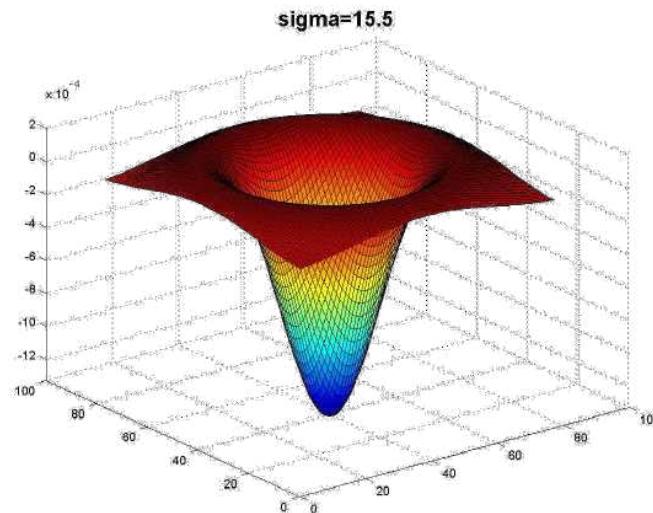


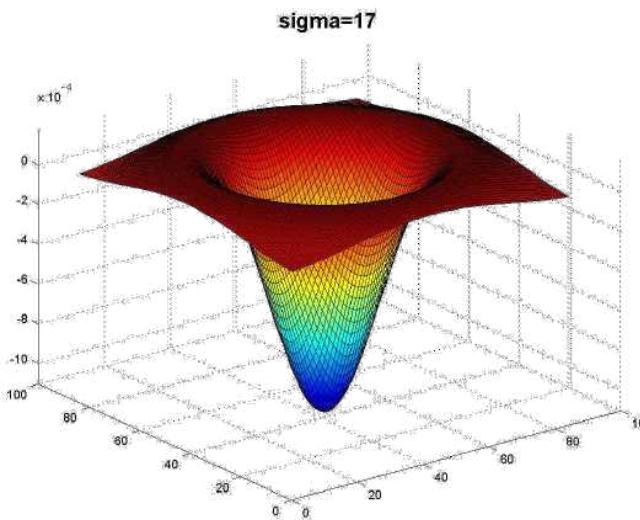
**jet** color scale  
blue: low, red: high





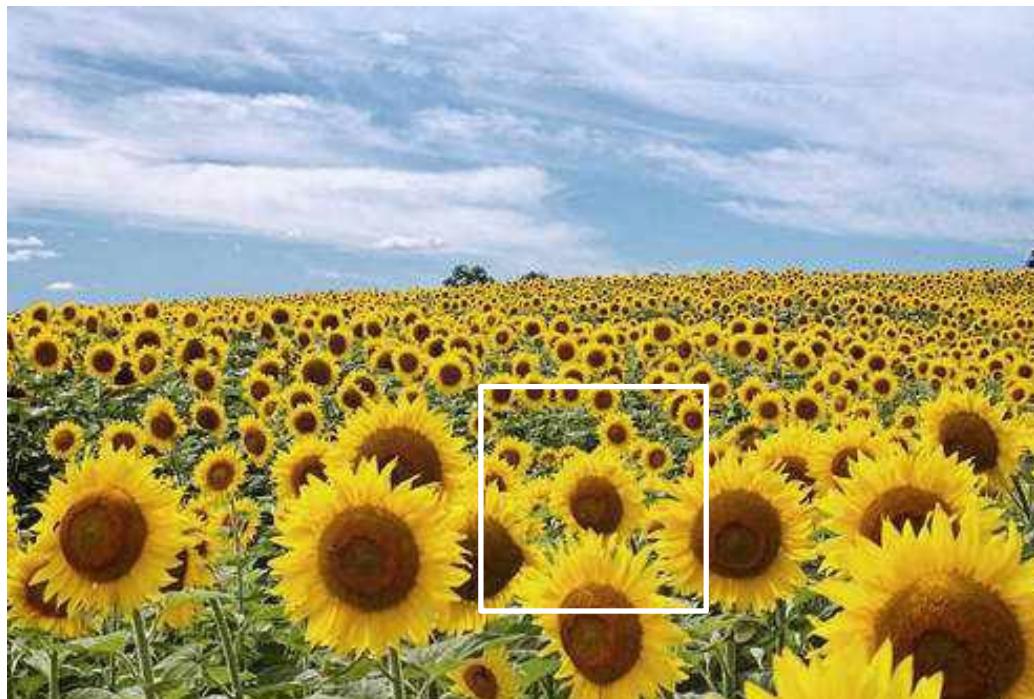






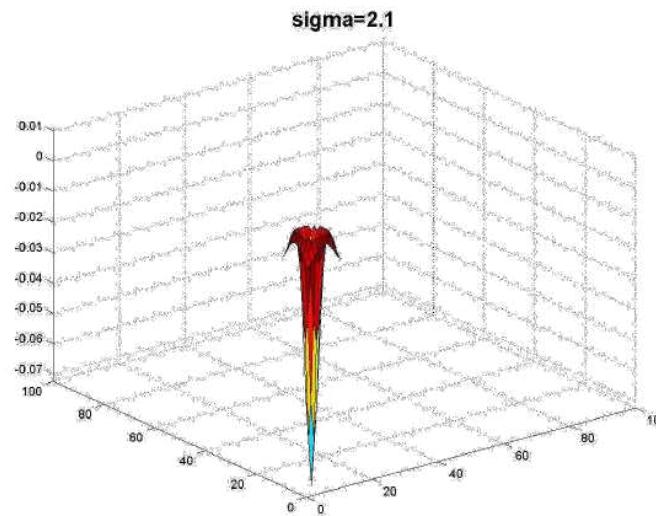
# What happened when you applied different Laplacian filters?

Full size

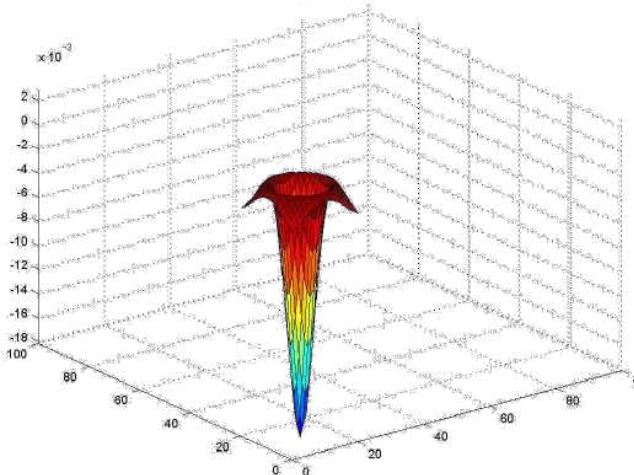


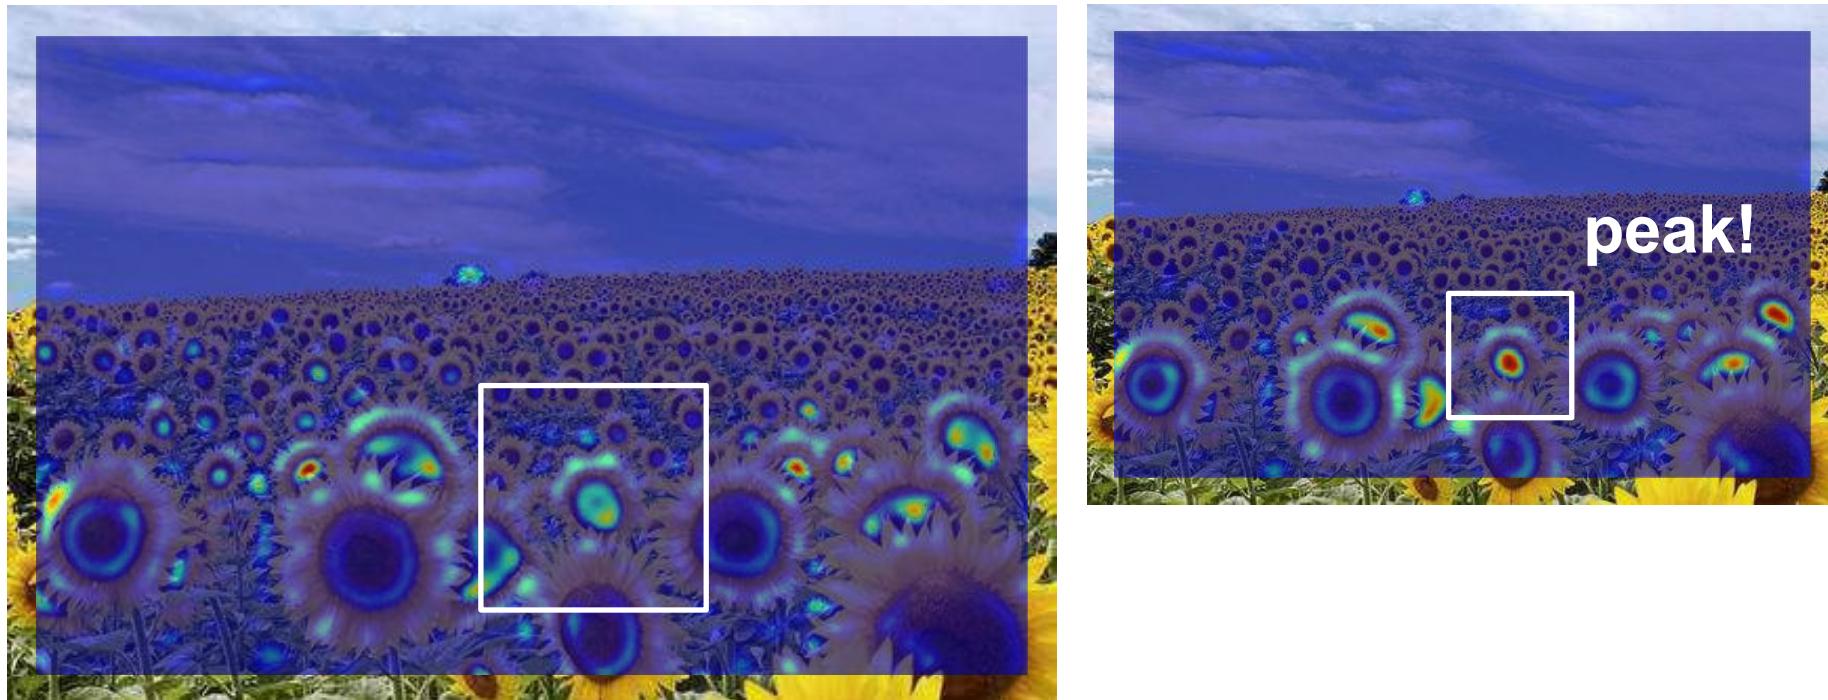
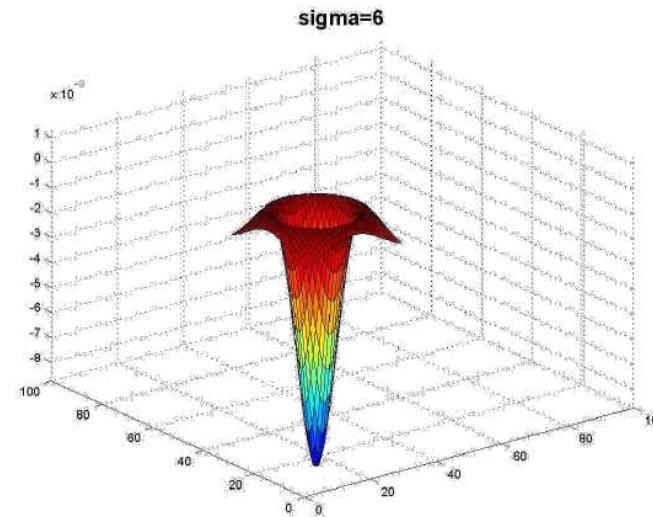
3/4 size

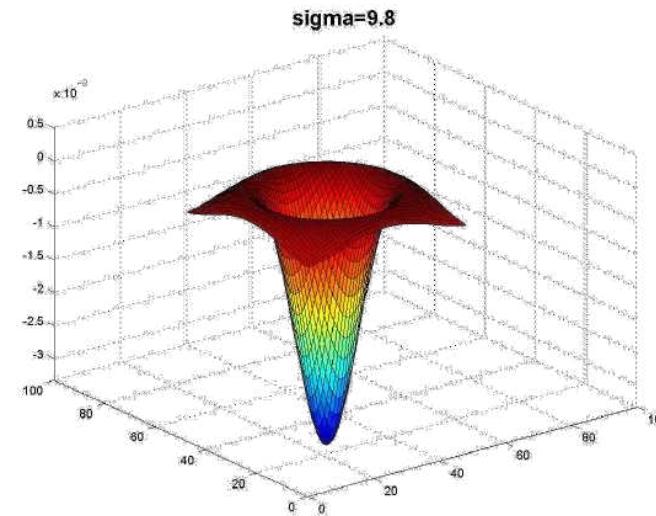


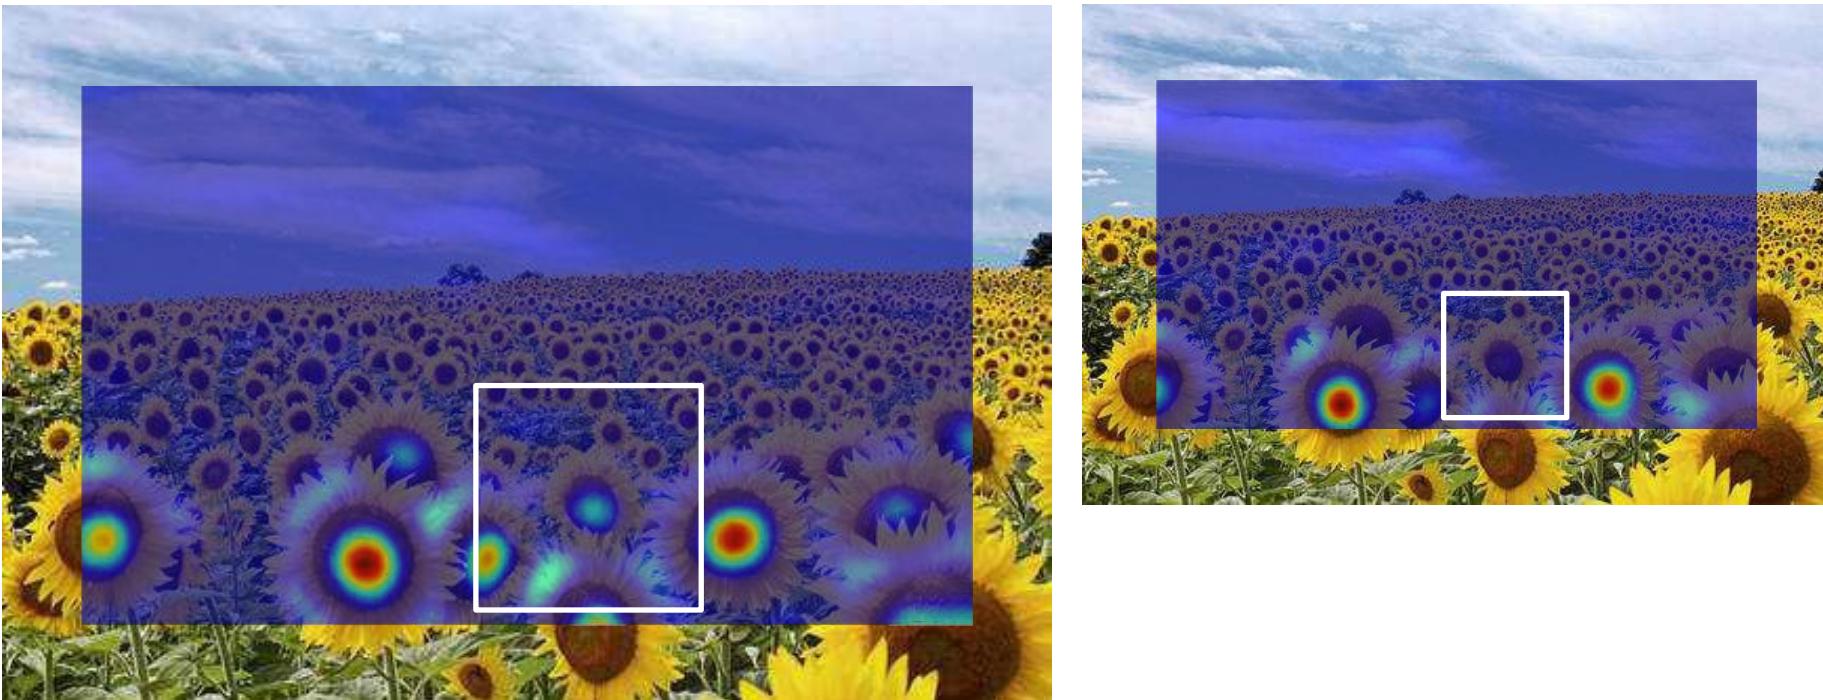
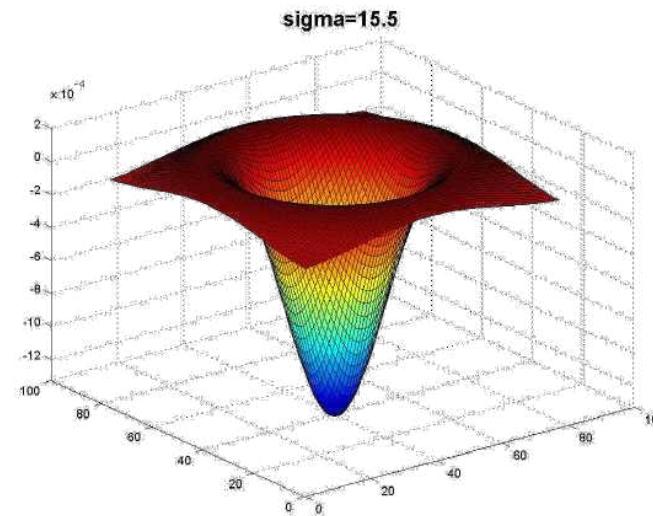


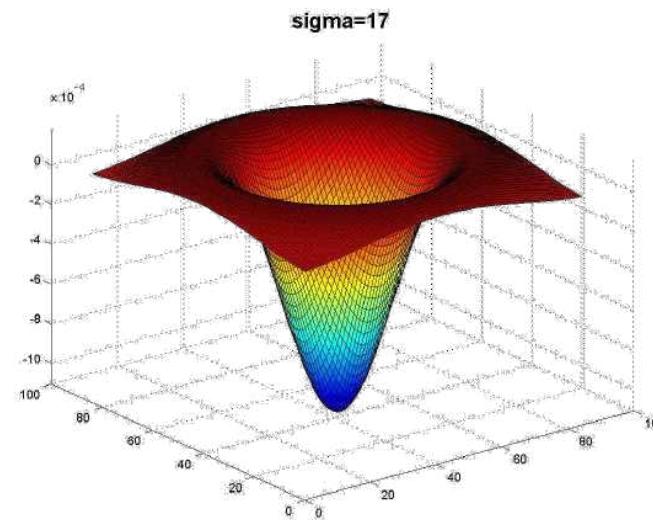
**sigma=4.2**











# What happened when you applied different Laplacian filters?

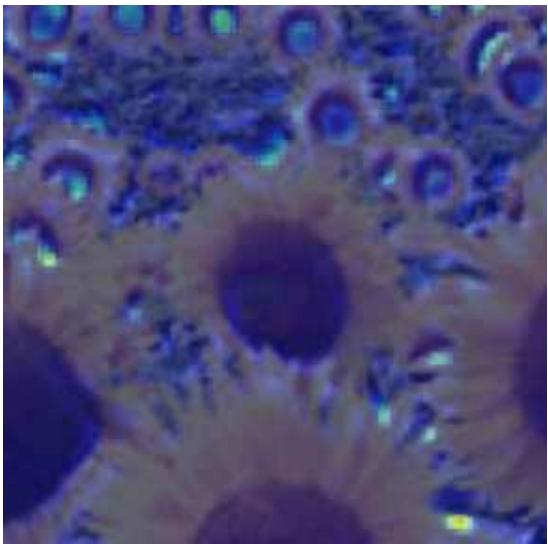
Full size



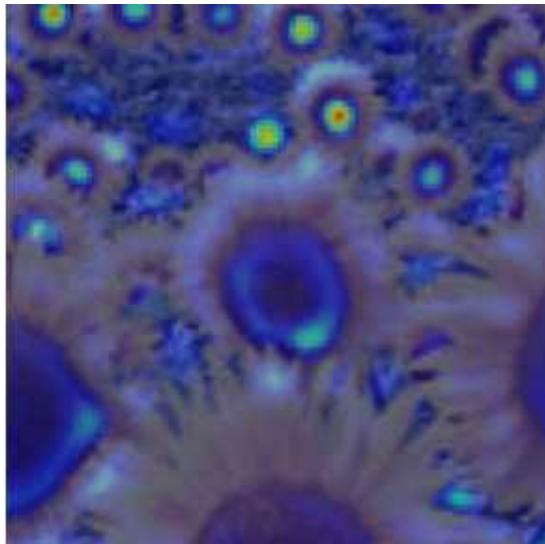
3/4 size



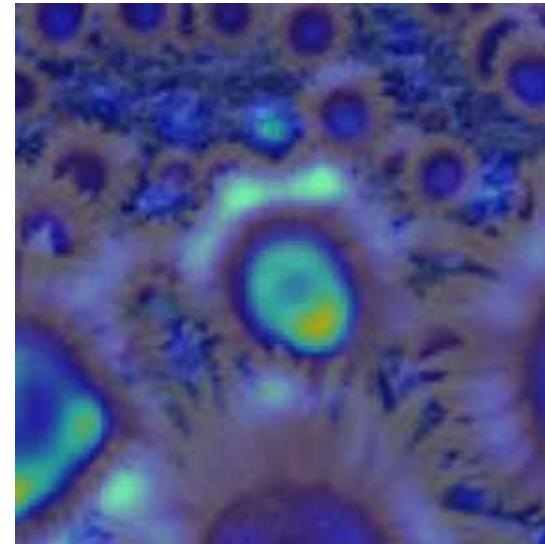
2.1



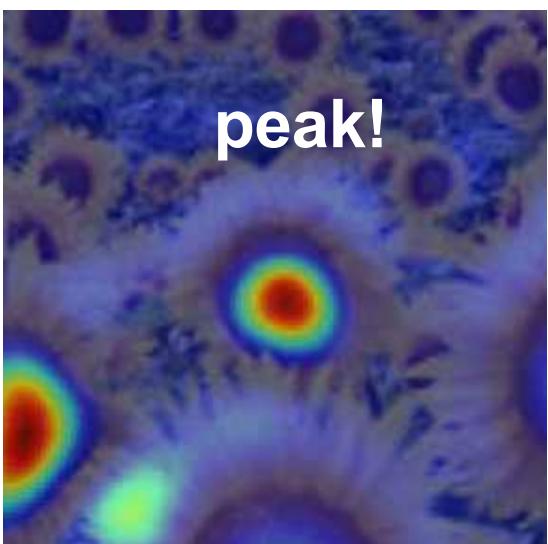
4.2



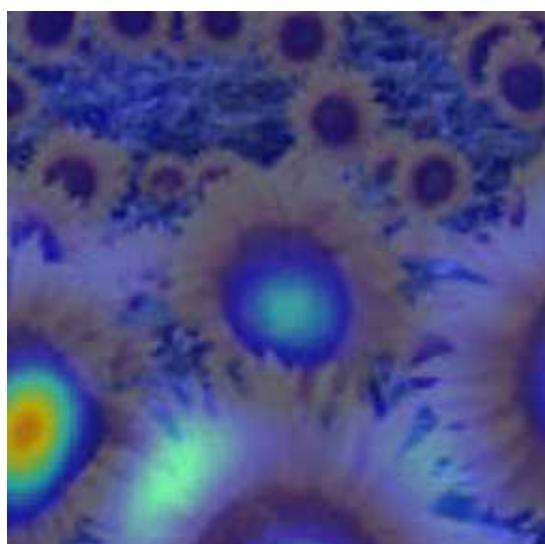
6.0



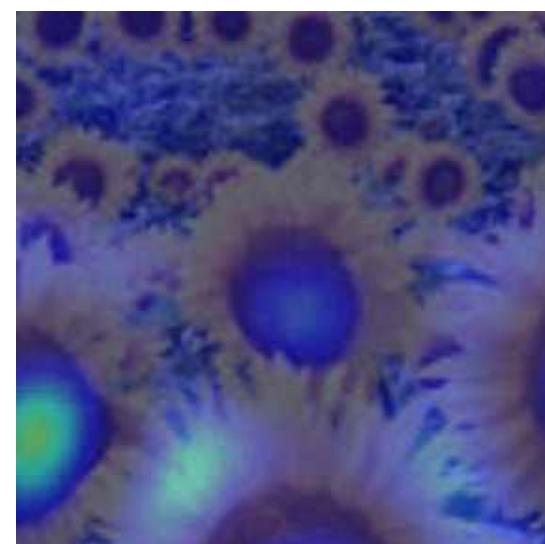
9.8



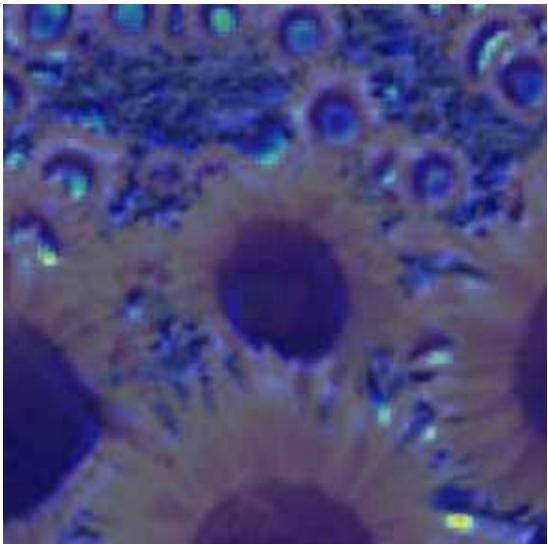
15.5



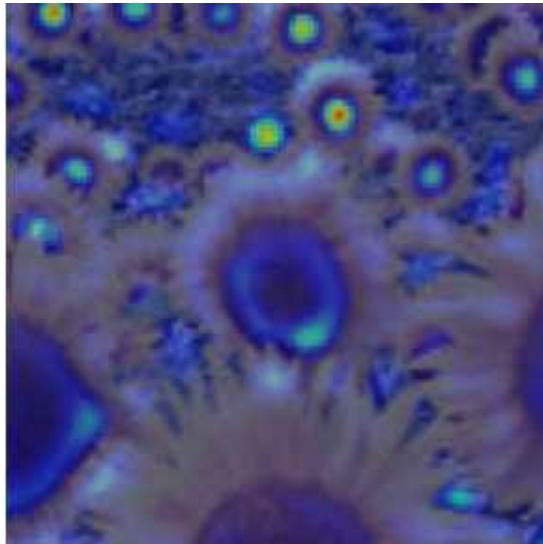
17.0



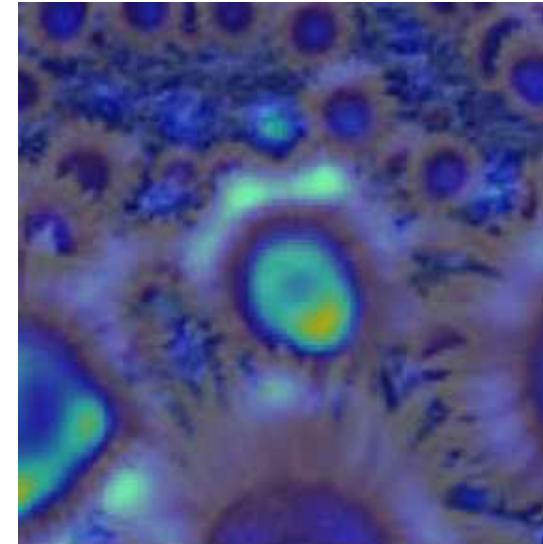
2.1



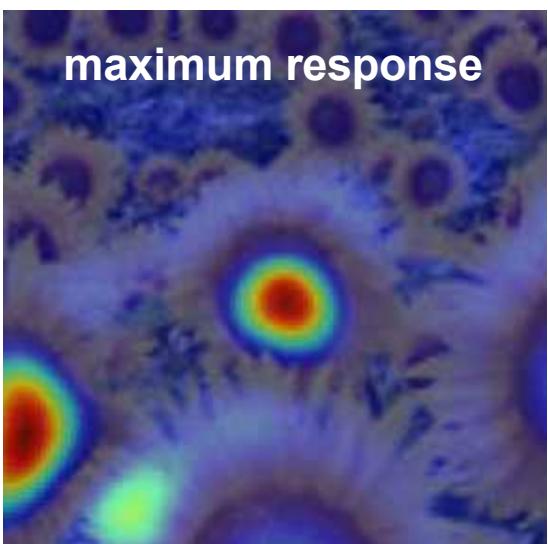
4.2



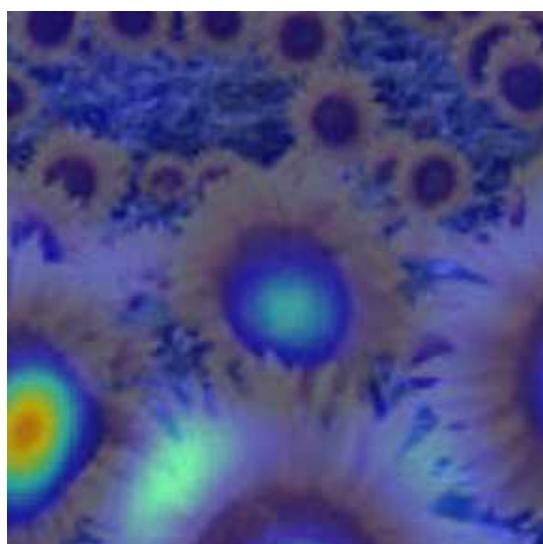
6.0



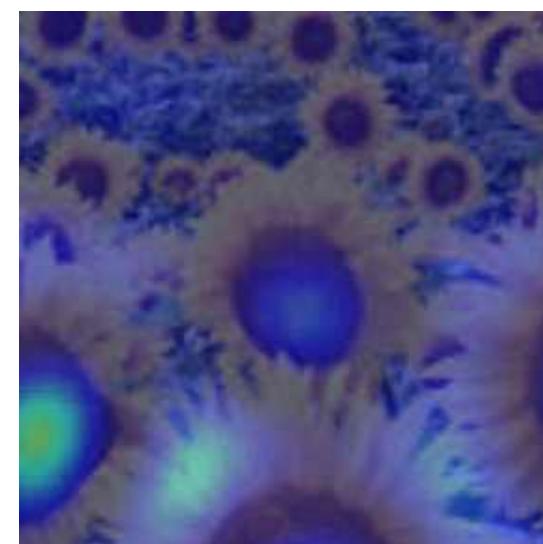
9.8



15.5

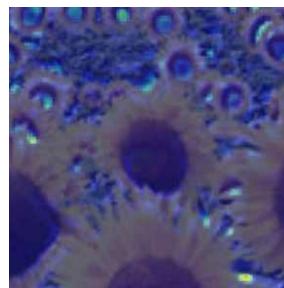


17.0

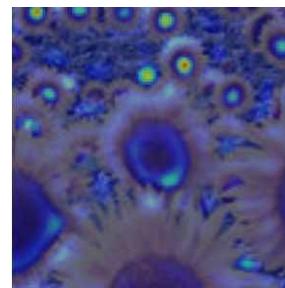


# Optimal Scale

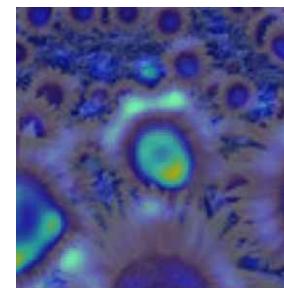
2.1



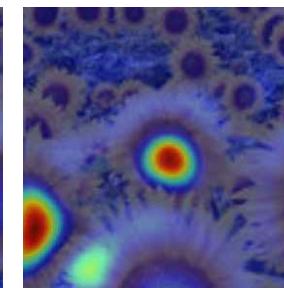
4.2



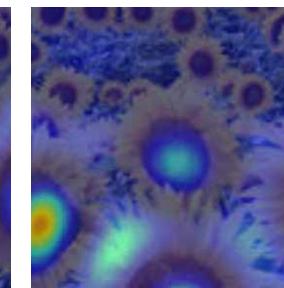
6.0



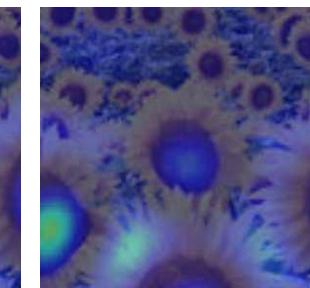
9.8



15.5

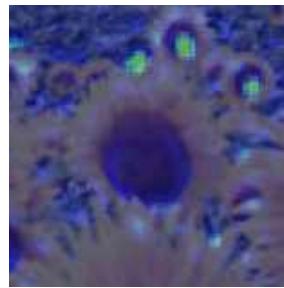


17.0

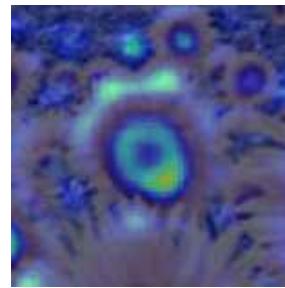


Full size image

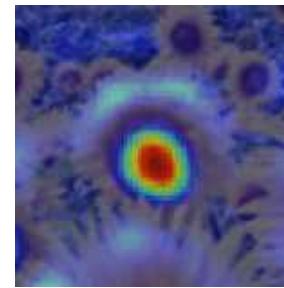
2.1



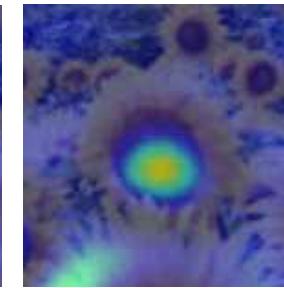
4.2



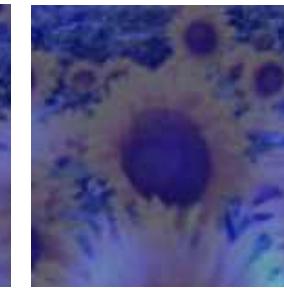
6.0



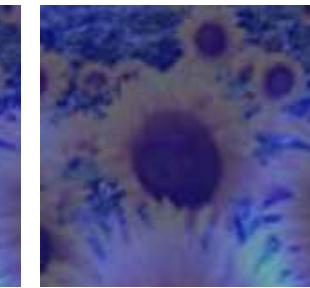
9.8



15.5



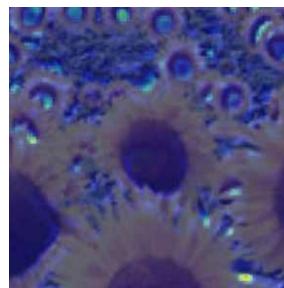
17.0



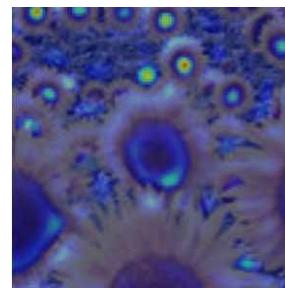
3/4 size image

# Optimal Scale

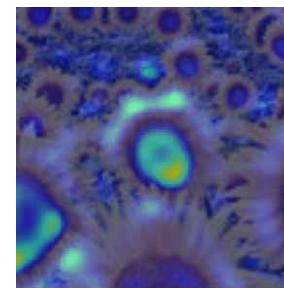
2.1



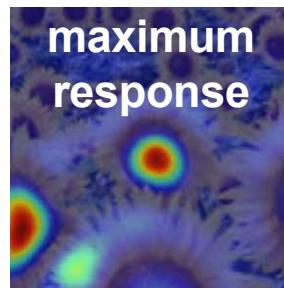
4.2



6.0

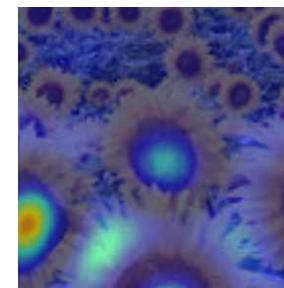


9.8

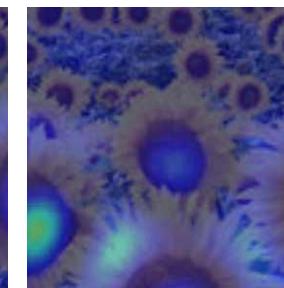


maximum  
response

15.5

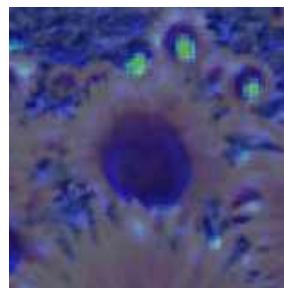


17.0

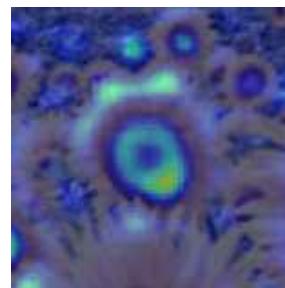


Full size image

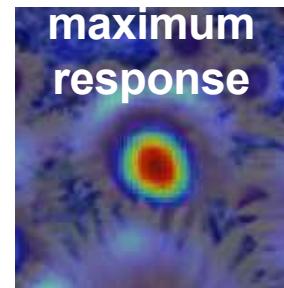
2.1



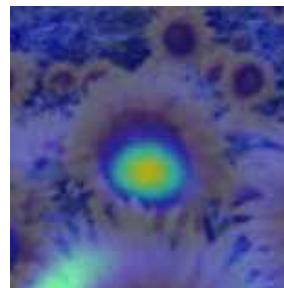
4.2



6.0



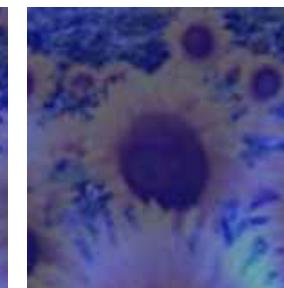
9.8



15.5

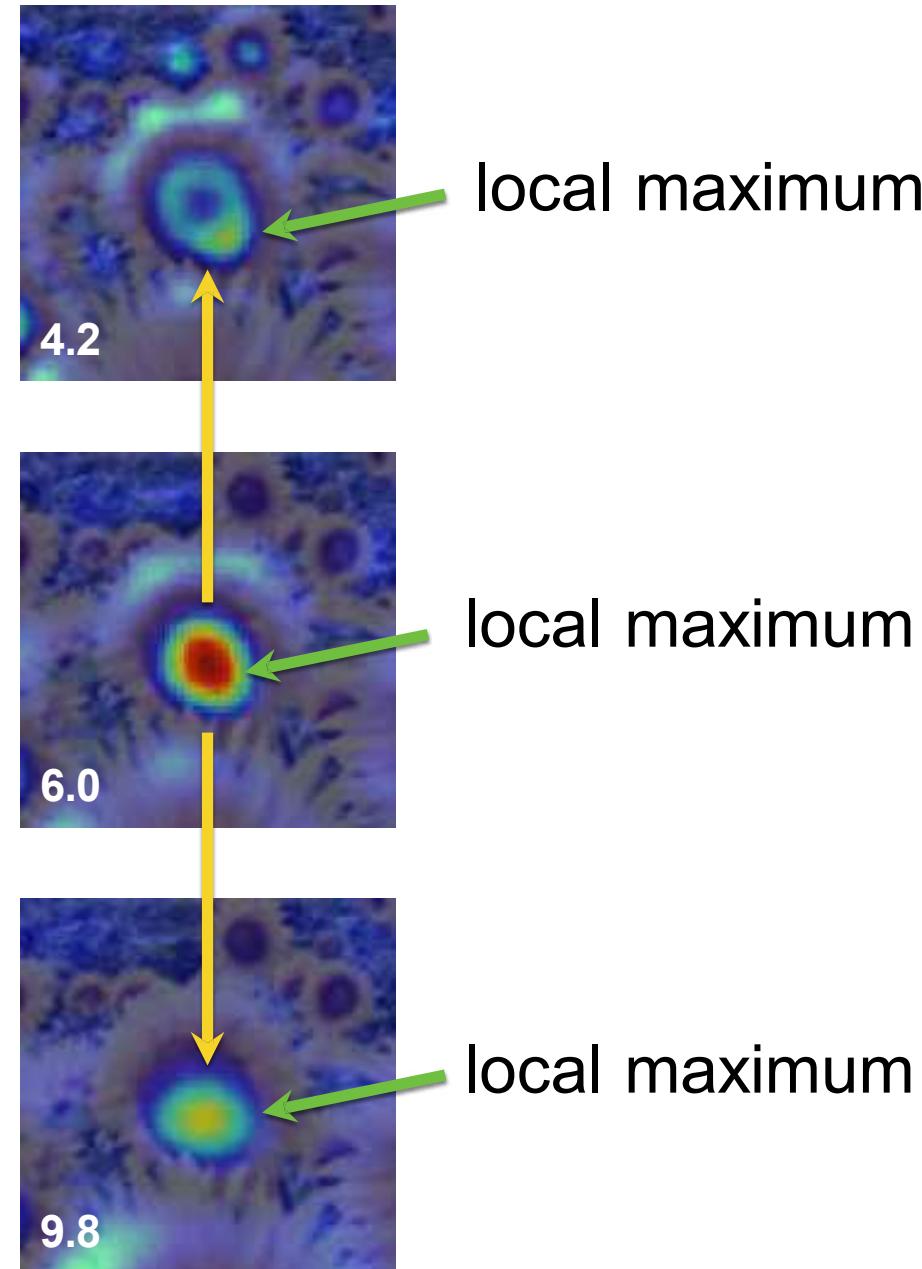


17.0



3/4 size image

cross-scale maximum

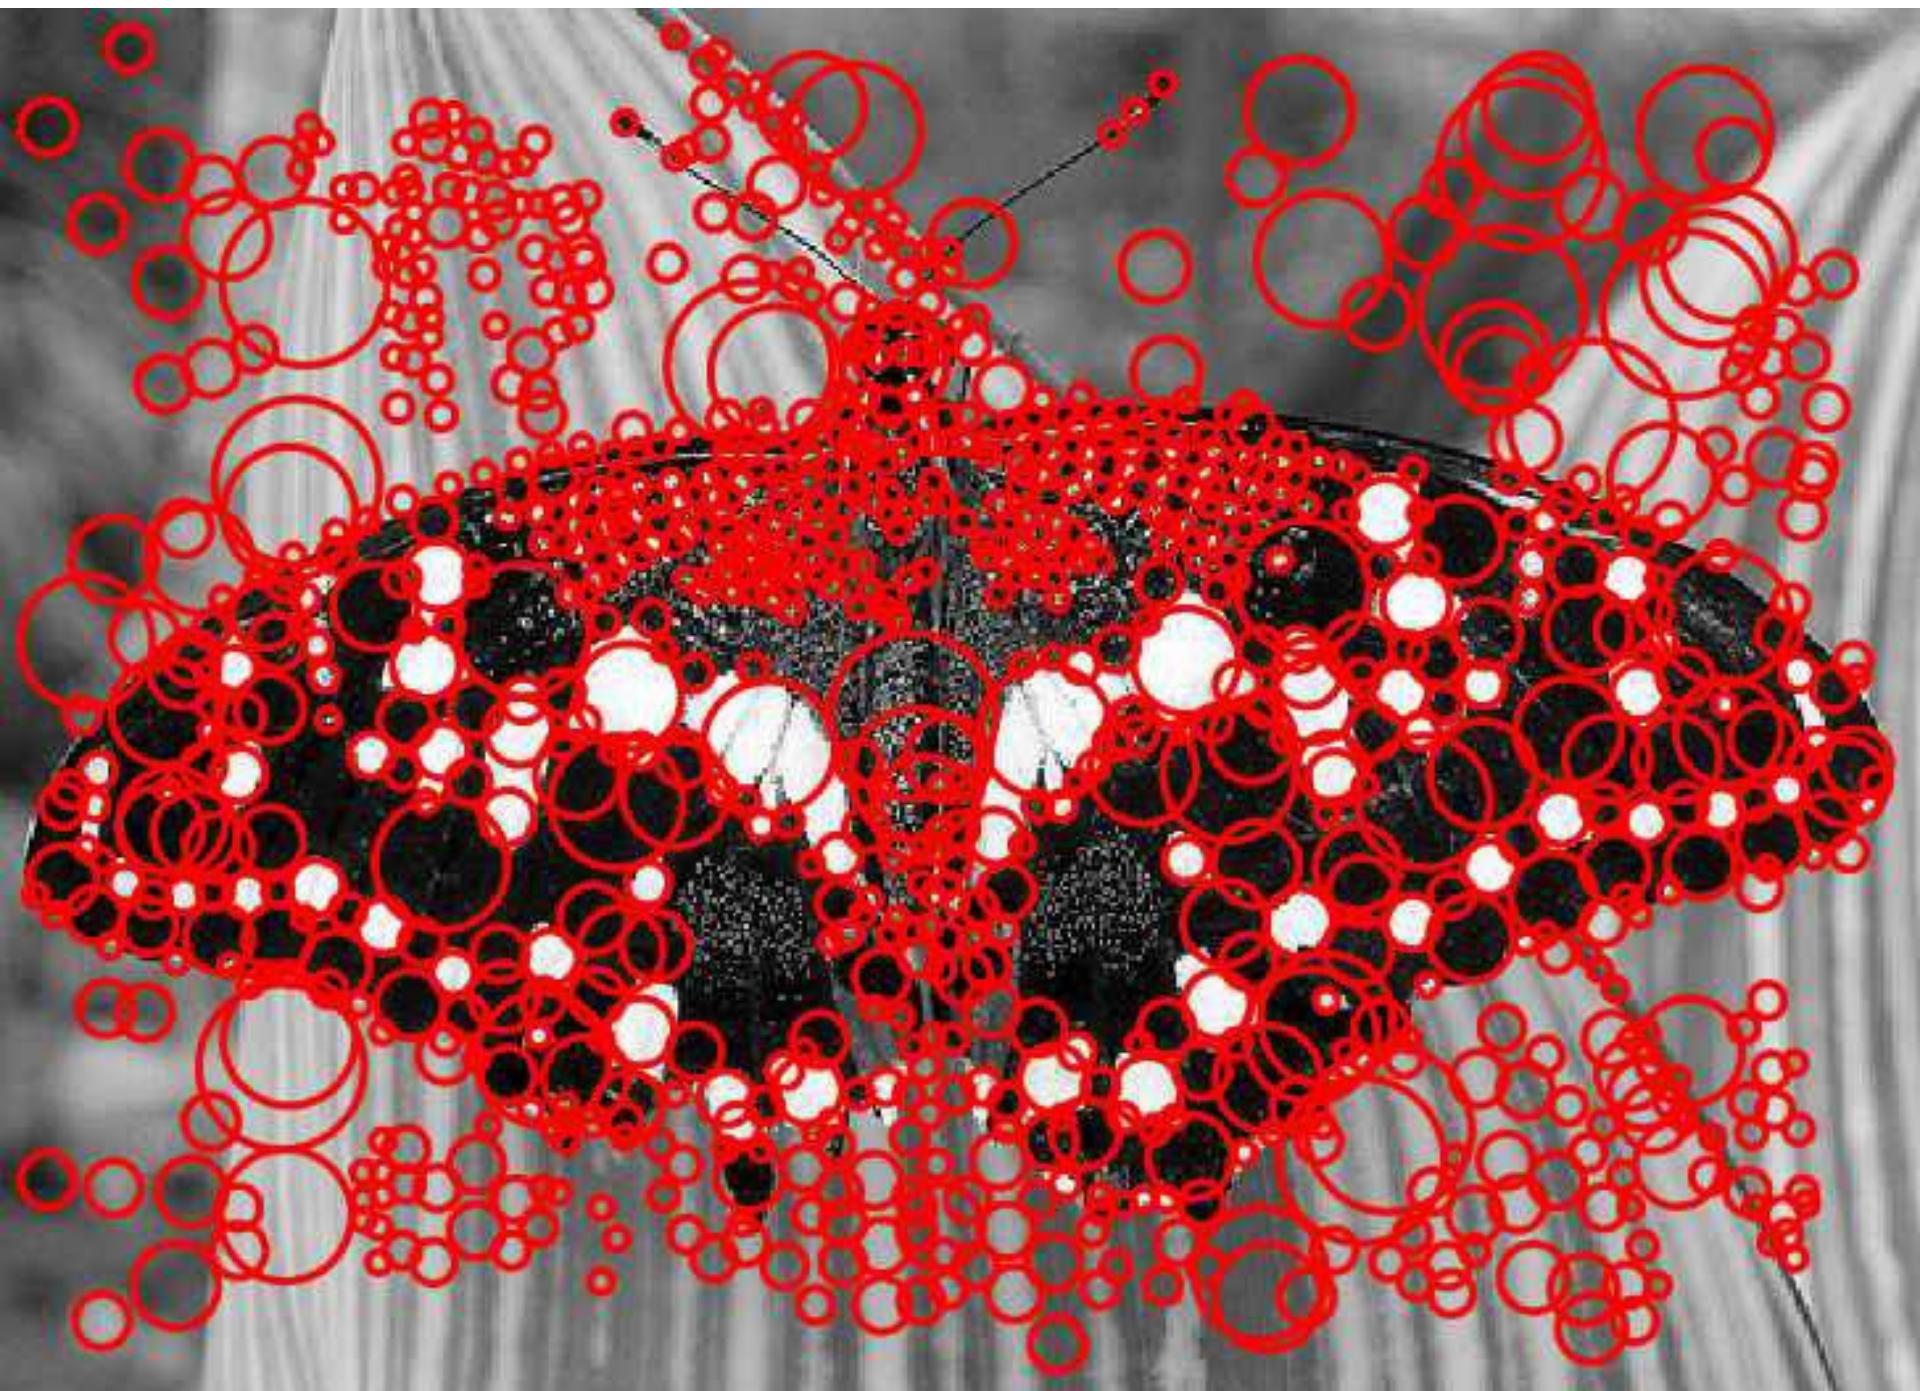


# How would you implement scale selection?

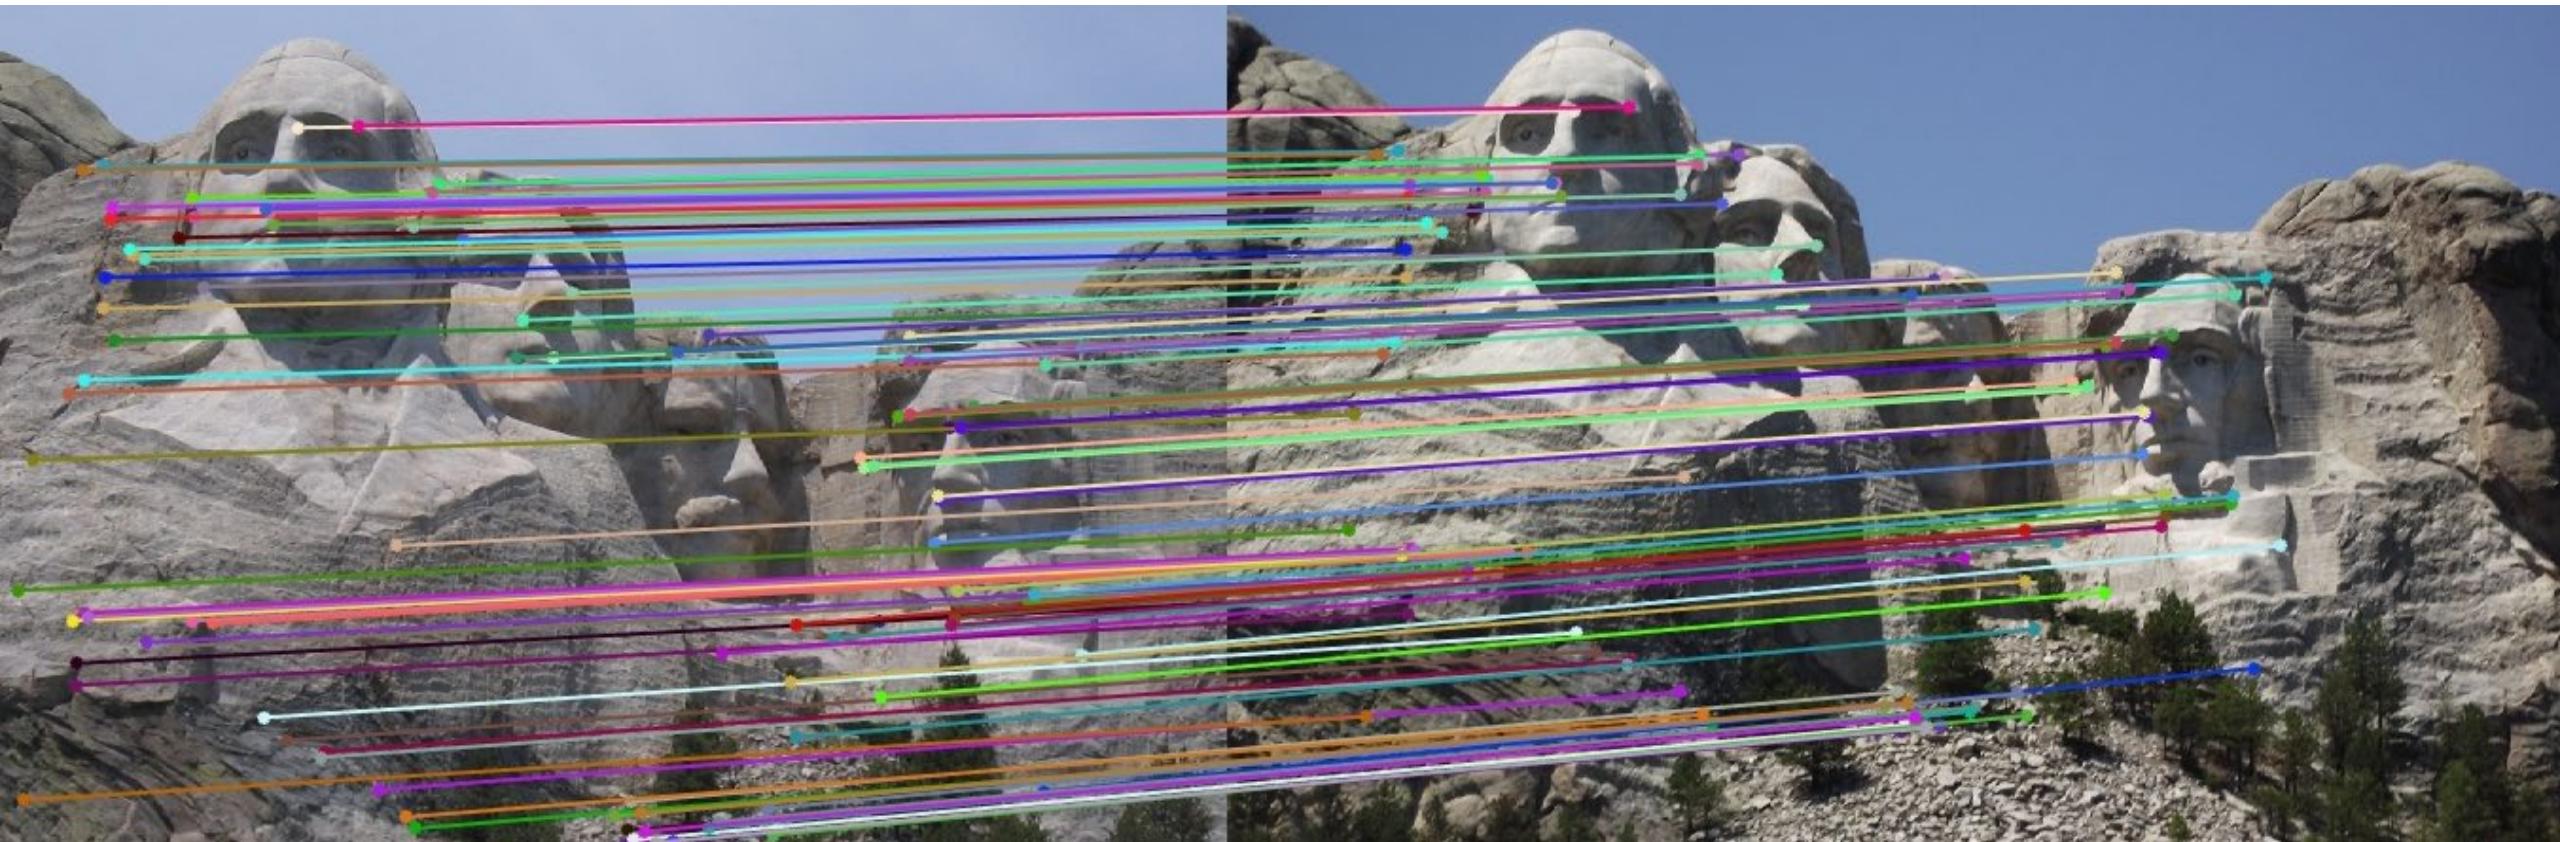
## Algorithm

- For each level of the Gaussian pyramid
  - compute feature response (e.g. Harris, Laplacian)
- For each level of the Gaussian pyramid
  - if local maximum and cross-scale
  - save scale and location of feature  $(x, y, s)$





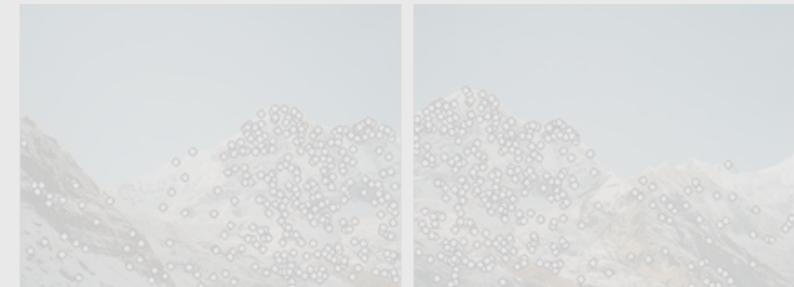
# Feature Matching



# Local Feature Matching: Main Components

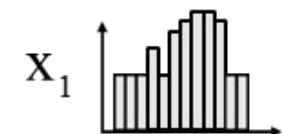
## 1) Detection:

Find a set of distinctive features (e.g., key points)

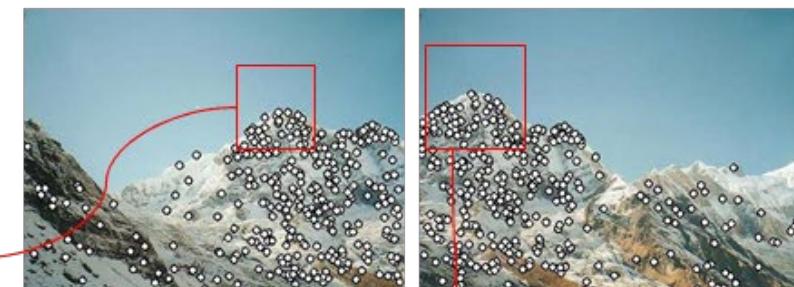


## 2) Description:

Extract feature descriptor around each interest point as vector.



$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$



$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$

## 3) Matching:

Compute distance between feature vectors to find correspondence.

$$d(x_1, x_2) < T$$



K. Grauman, B. Leibe

# Possible Feature Descriptors

- **Templates**
  - Intensity
  - Gradients
  - etc.



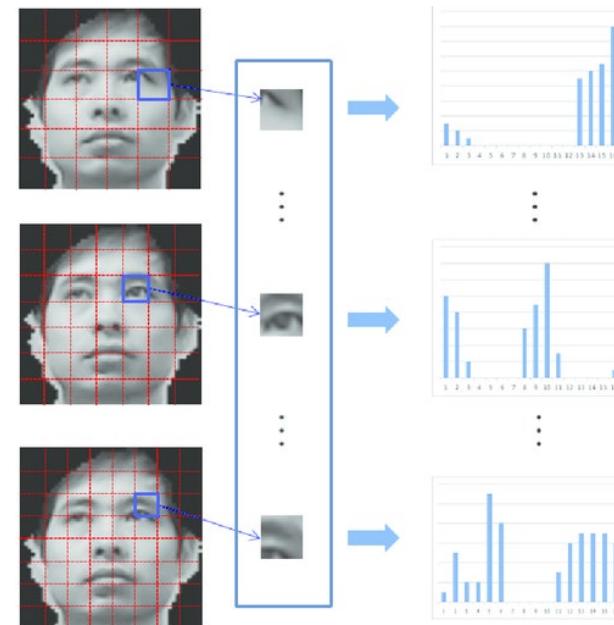
Matching Result



Detected Point

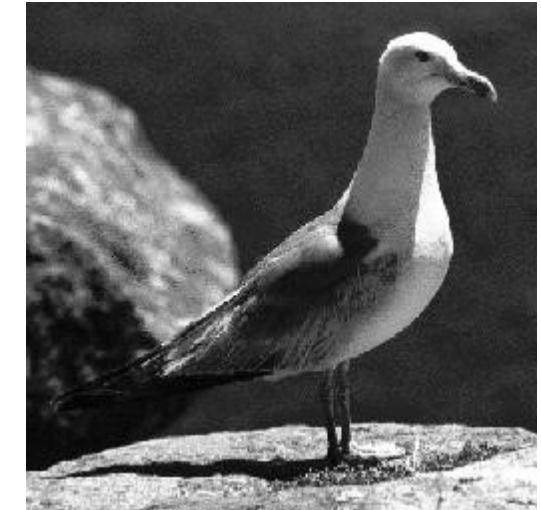
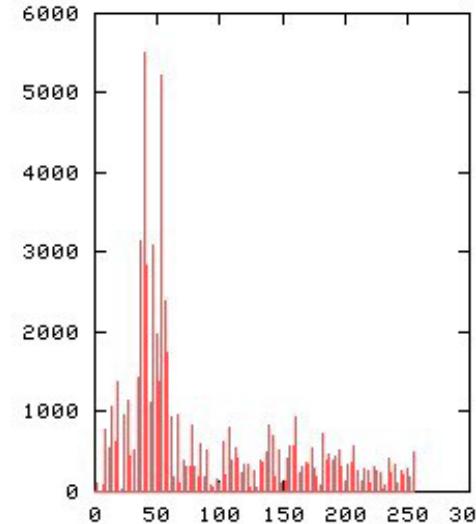


- **Histograms**
  - Color
  - Texture
  - SIFT
  - etc.



# Feature Descriptor: Histograms

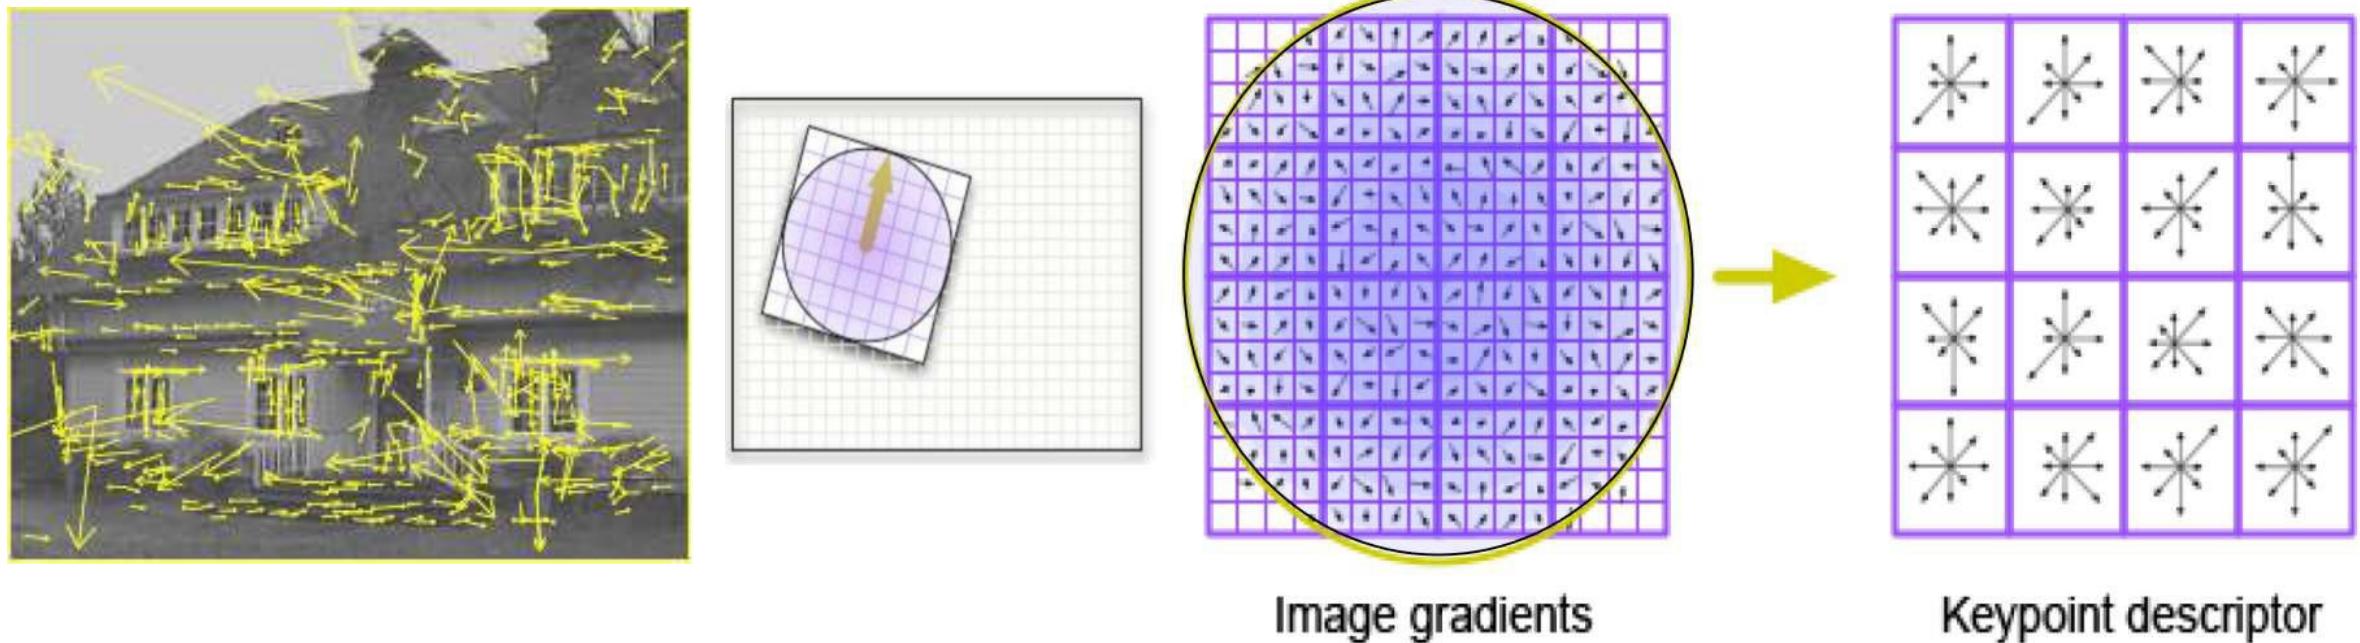
- **Global** histogram to represent distribution of features
  - E.g., how ‘well exposed’ a photo is
- What about a **local** histogram per detected point?



# Feature Descriptor: SIFT

**SIFT:** Scale Invariant Feature Transform

- Paper: “Distinctive image features from scale-invariant keypoints,” DG Lowe
- Extremely popular (76k citations in 2024)

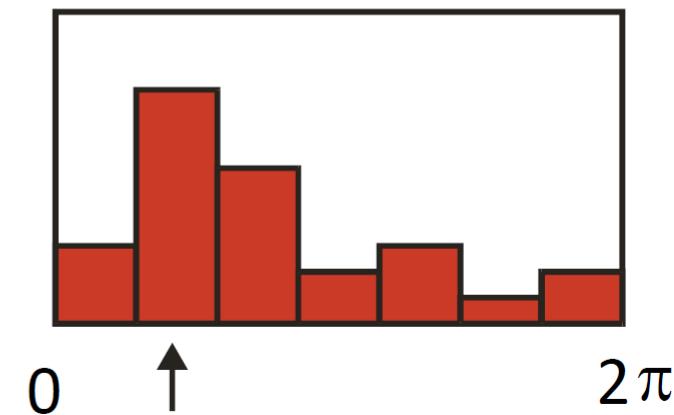
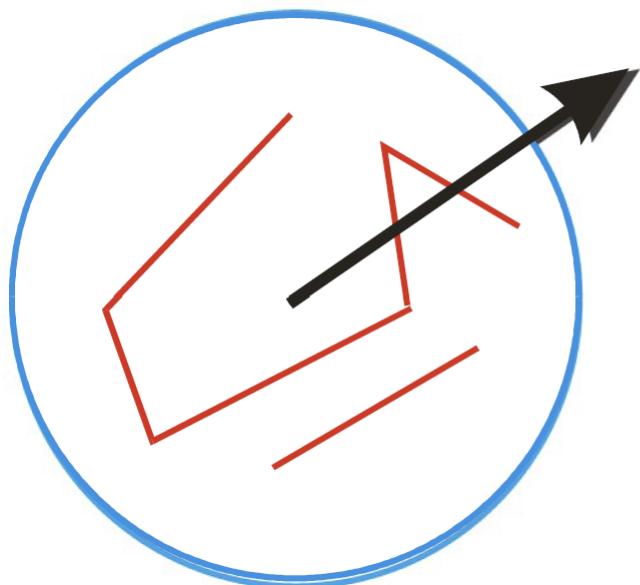


## Algorithm

1. Find Harris corners scale-space extrema as feature point locations
2. Corner post-processing
  - Subpixel position interpolation
  - Discard low-contrast points
  - Eliminate points along edges
3. Orientation estimation per feature point

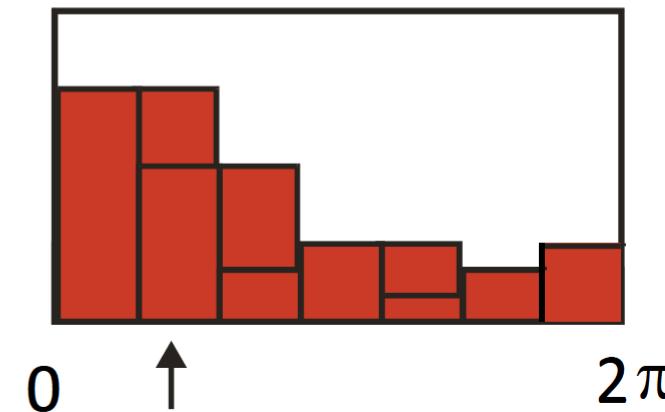
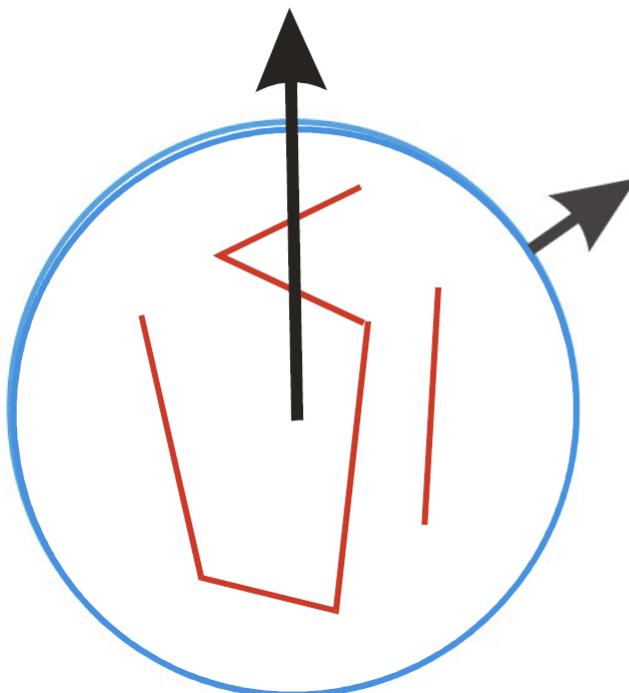
# Orientation Estimation

- Compute gradient orientation histogram
- Select dominant orientation  $\theta$



# Orientation Normalization

- Compute gradient orientation histogram
- Select dominant orientation  $\theta$
- Normalize: rotate to fixed orientation

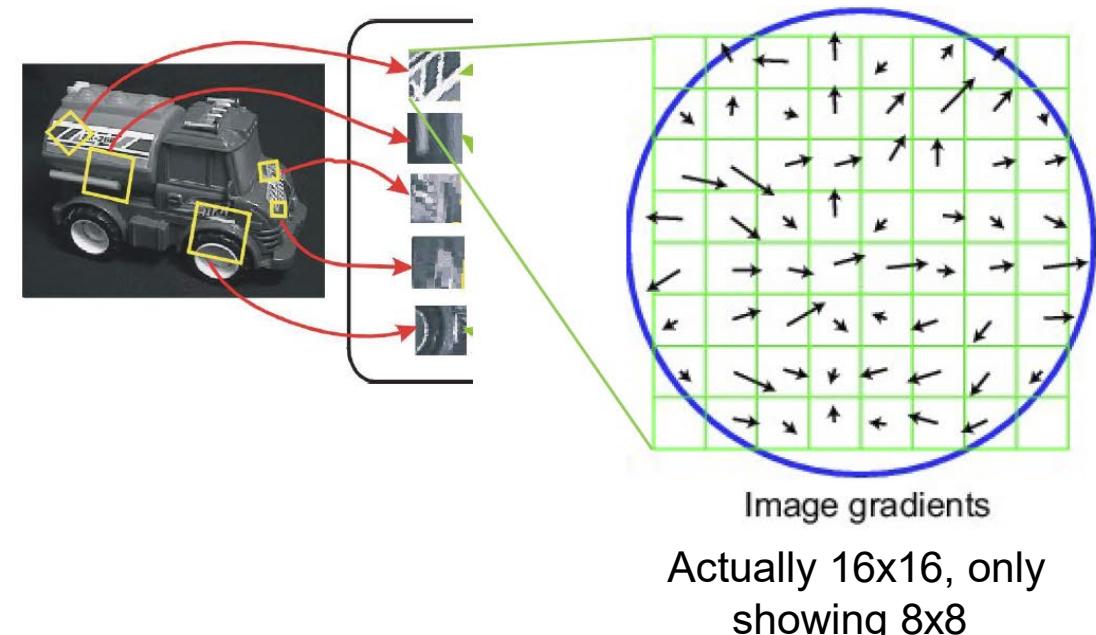
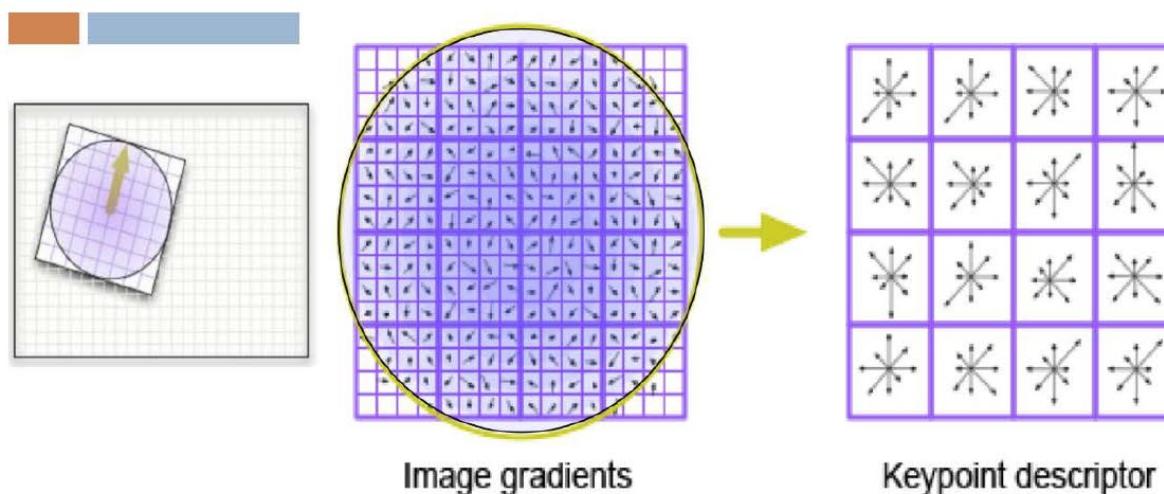


## Algorithm

1. Find Harris corners scale-space extrema as feature point locations
2. Corner post-processing
  - Subpixel position interpolation
  - Discard low-contrast points
  - Eliminate points along edges
3. Orientation estimation per feature point
4. Descriptor extraction
  - Motivation: We want some sensitivity to spatial layout, but not too much, so blocks of histograms give us that

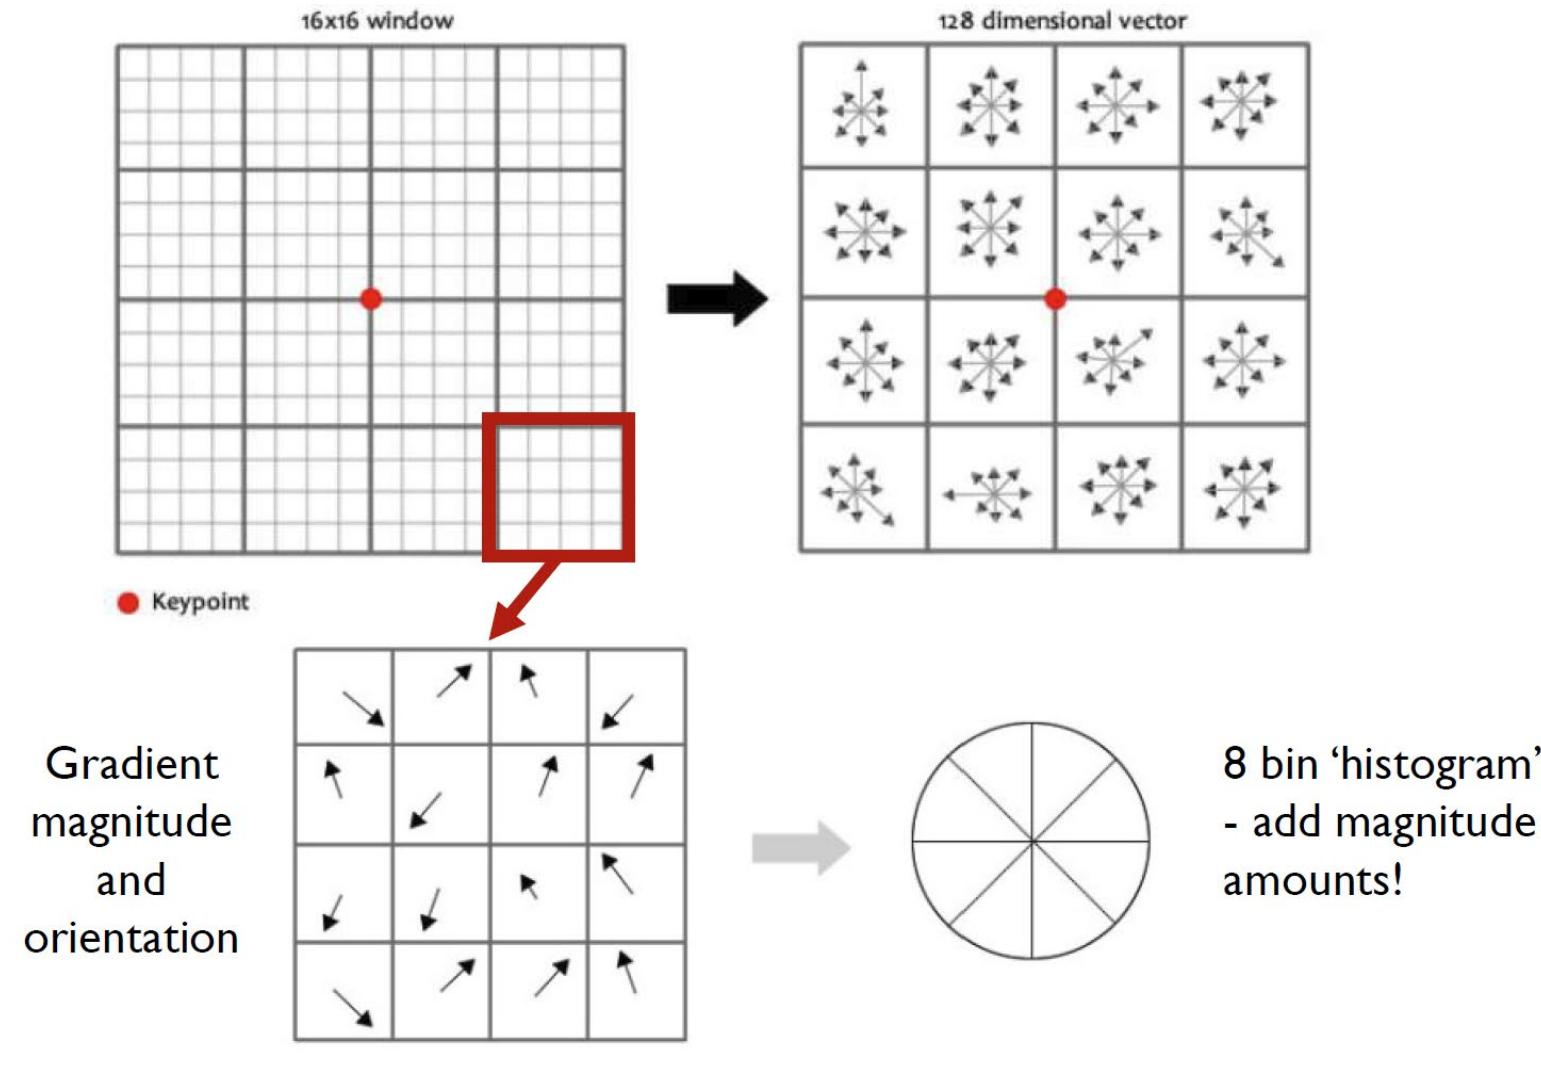
# SIFT Descriptor

- Compute on local  $16 \times 16$  window around detection
- Rotate and scale window according to discovered orientation  $\theta$  (rotation invariance) and scale  $\sigma$  (gain invariance)
  - Resample image to match orientation



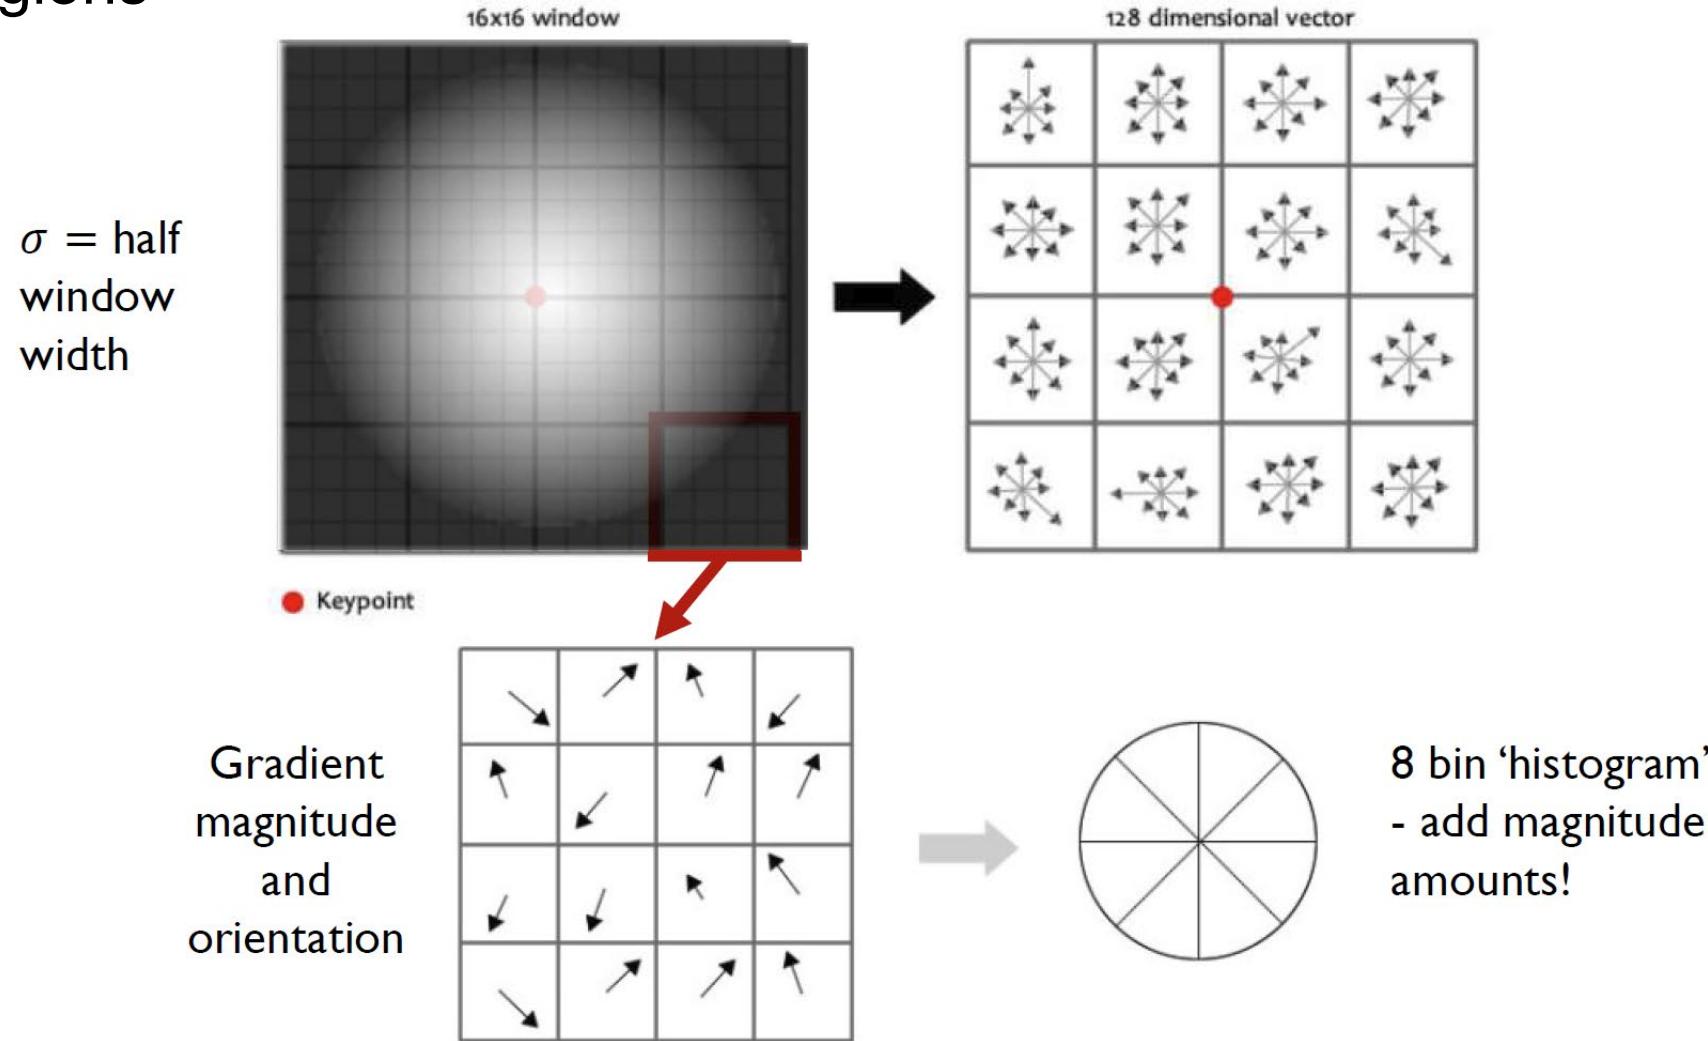
# SIFT Descriptor Extraction

Given a keypoint with scale and orientation

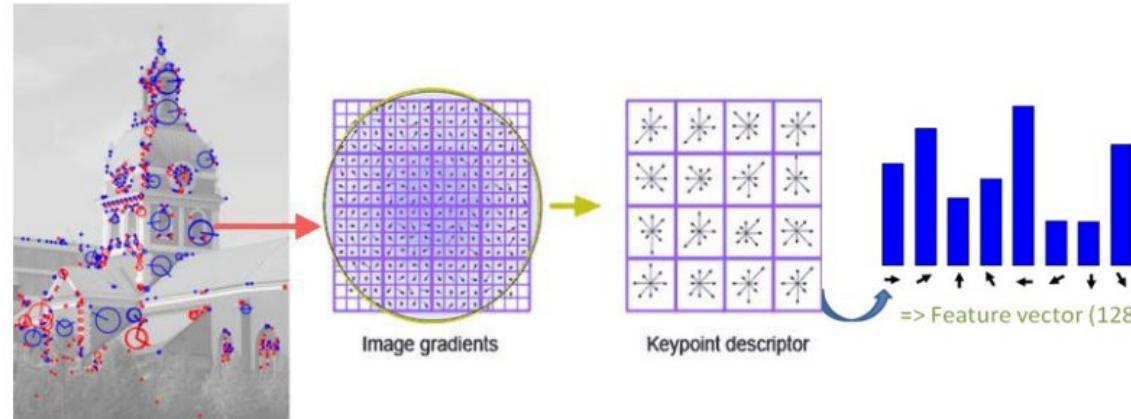


# SIFT Descriptor Extraction

Weight 16x16 grid by Gaussian to add location robustness and reduce effect of outer regions



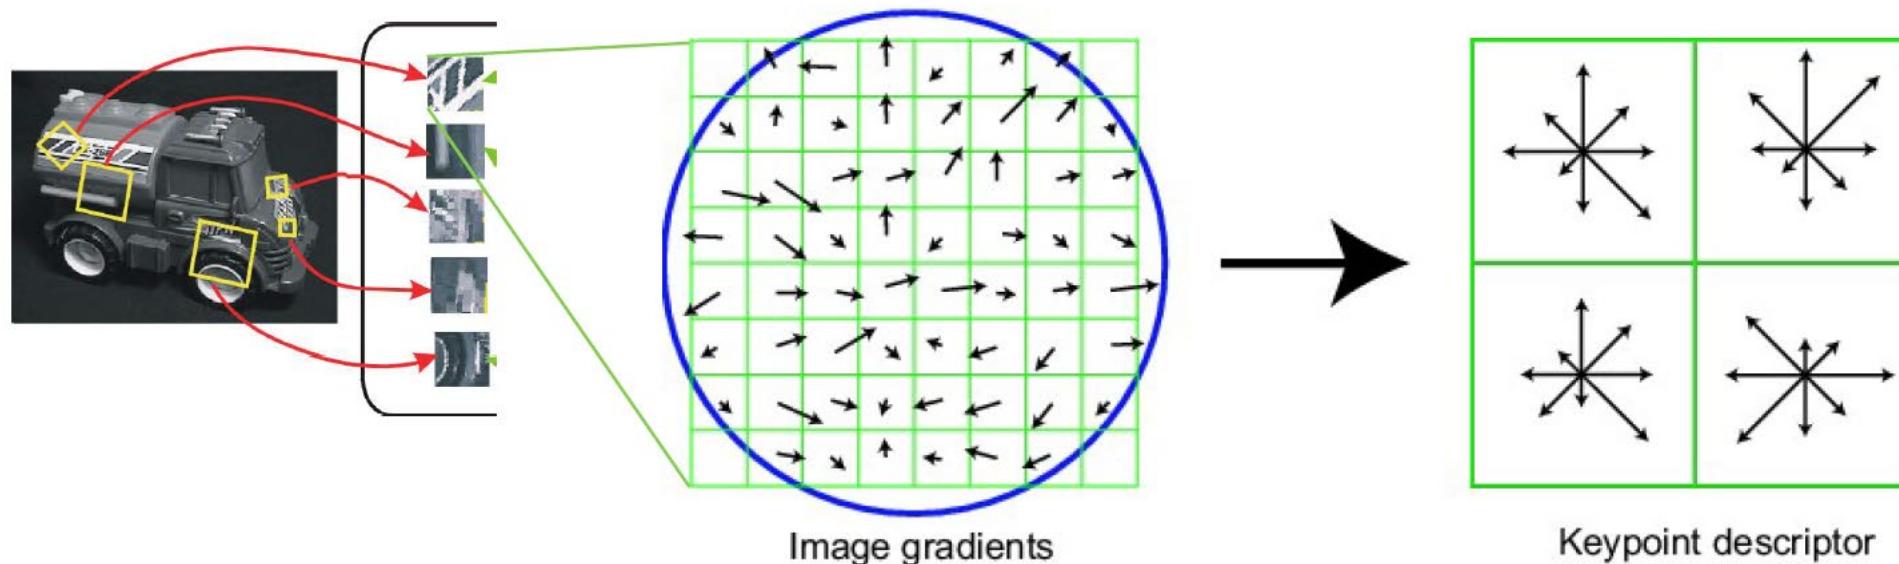
# SIFT Descriptor Extraction



- Extract  $8 \times 16$  values into 128-dim vector
- **Illumination invariance:**
  - Working in gradient space, so robust to  $I = I + b$
  - Normalize vector to  $[0 \dots 1]$ 
    - Robust to  $I = \alpha I$  brightness changes
  - Clamp all vector values  $> 0.2$  to 0.2.
    - Robust to “non-linear illumination effects”
    - Image value saturation / specular highlights
  - Renormalize

# Additional Slide: Reduce effect of illumination

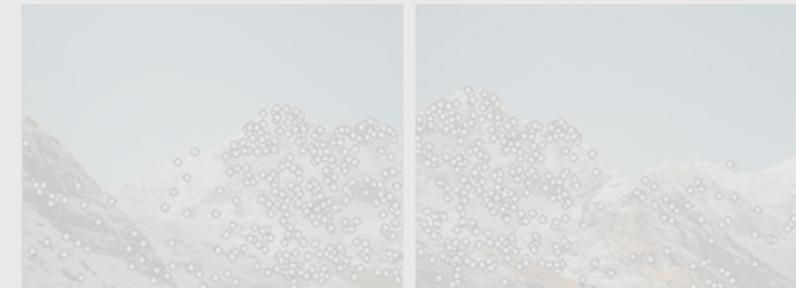
- 128-dim vector normalized to 1
- Threshold gradient magnitudes to avoid excessive influence of high gradients
- After normalization, clamp gradients  $> 0.2$
- Renormalize



# Local Feature Matching: Main Components

## 1) Detection:

Find a set of distinctive features (e.g., key points)

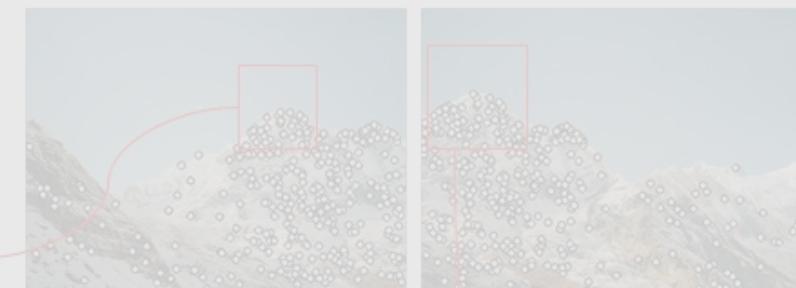


## 2) Description:

Extract feature descriptor around each interest point as vector.



$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$



$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$

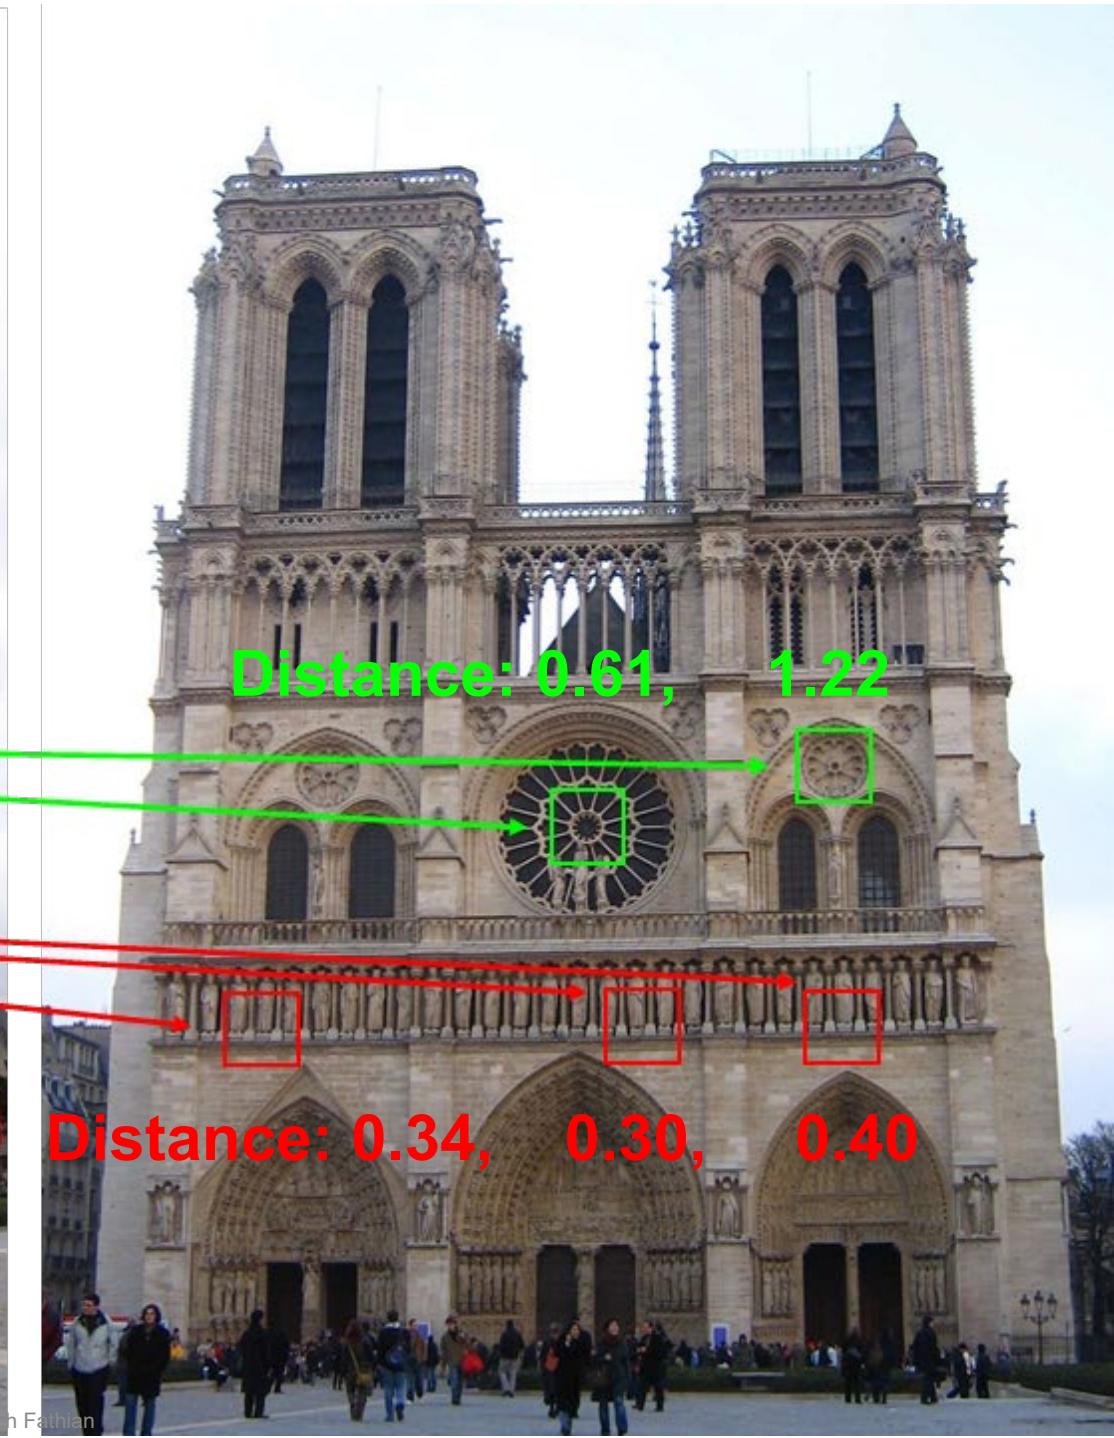
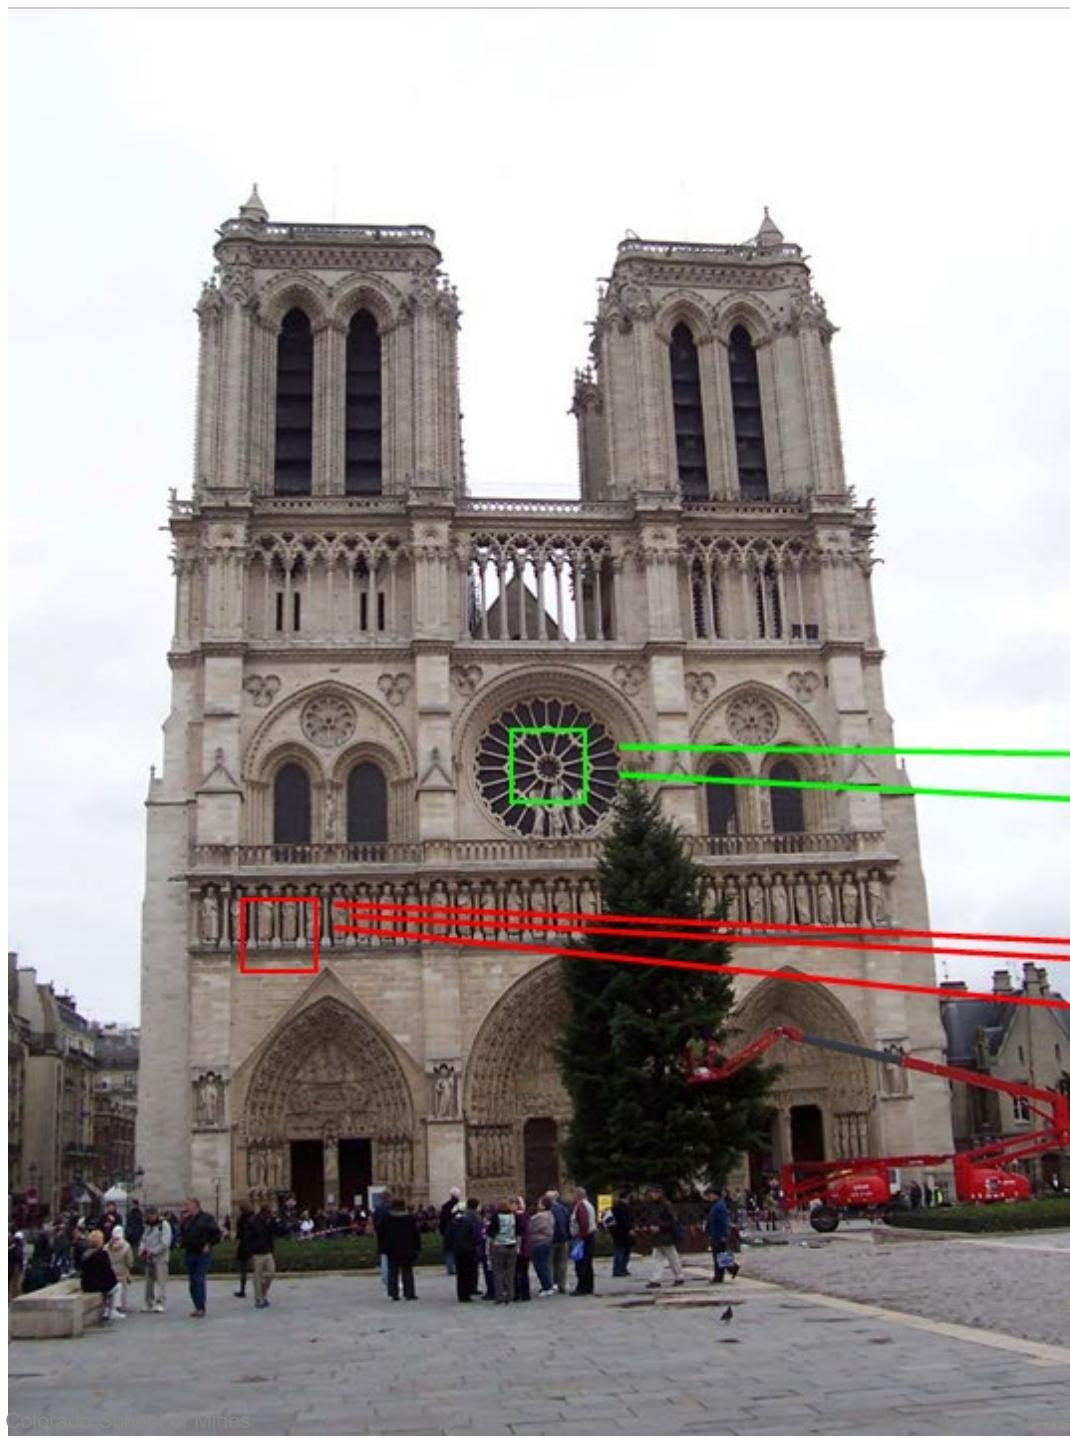
## 3) Matching:

Compute distance between feature vectors to find correspondence.

$$d(x_1, x_2) < T$$



K. Grauman, B. Leibe



# Euclidean Distance vs. Cosine Similarity

- Euclidean distance:

$$d(\mathbf{p}, \mathbf{q}) = d(\mathbf{q}, \mathbf{p}) = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \cdots + (q_n - p_n)^2}$$

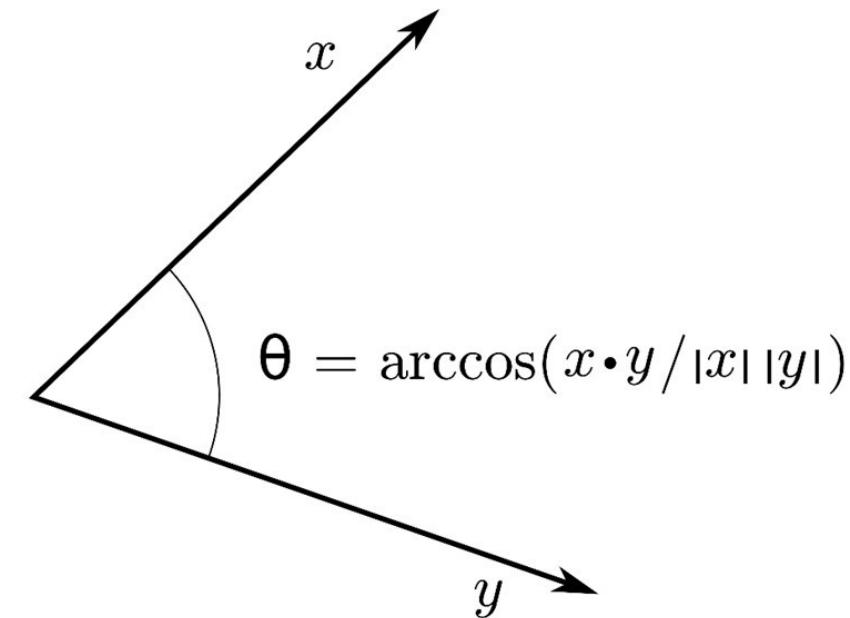
$$= \sqrt{\sum_{i=1}^n (q_i - p_i)^2}.$$

$$\|\mathbf{q} - \mathbf{p}\| = \sqrt{(\mathbf{q} - \mathbf{p}) \cdot (\mathbf{q} - \mathbf{p})}.$$

- Cosine similarity:

$$\mathbf{a} \cdot \mathbf{b} = \|\mathbf{a}\|_2 \|\mathbf{b}\|_2 \cos \theta$$

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\|_2 \|\mathbf{B}\|_2}$$



$$\theta = \arccos(x \cdot y / |x| |y|)$$

# Efficiency Consideration

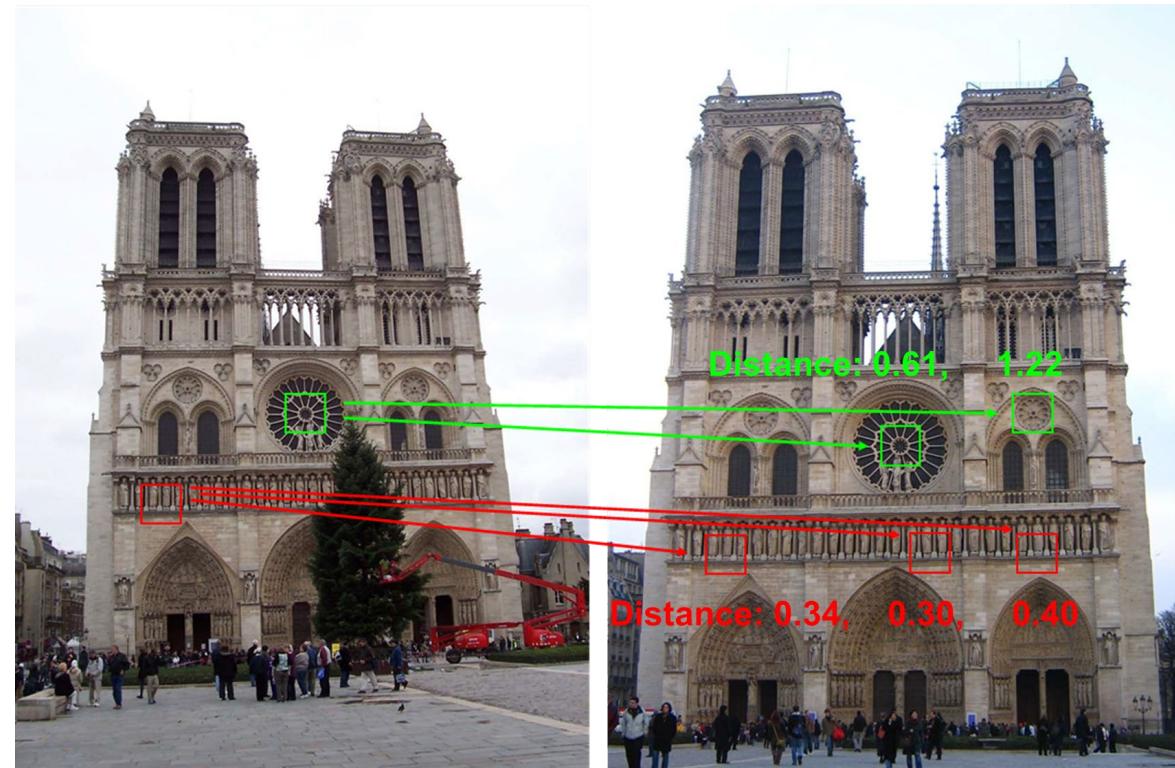
- Naïve looping: Expensive
- Operate on matrices of descriptors
  - E.g., for row vectors,

`features_image1 * features_image2T`

produces matrix of dot product results for all pairs of features

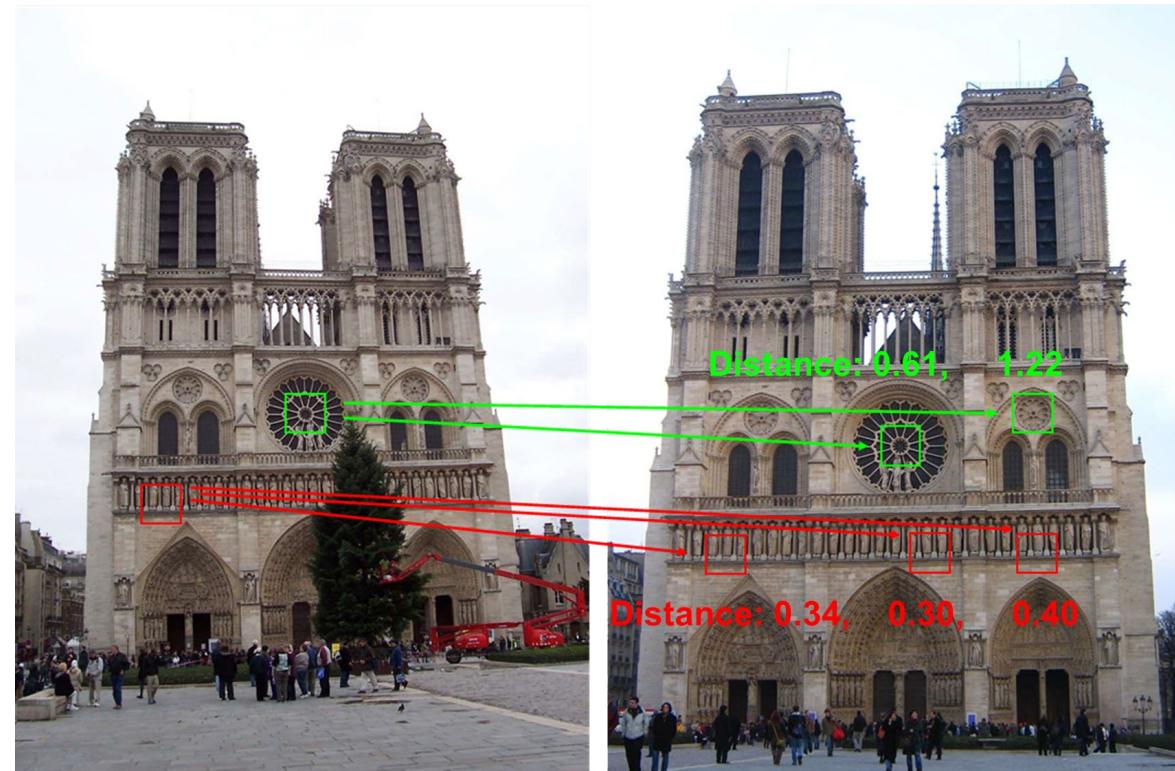
# Feature Matching

- **Criteria 1:**
  - Compute distance in feature space, e.g., Euclidean distance between 128-dim SIFT descriptors
  - Match point to lowest distance (nearest neighbor)
- **Problems:**
  - Does everything have a match?



# Feature Matching

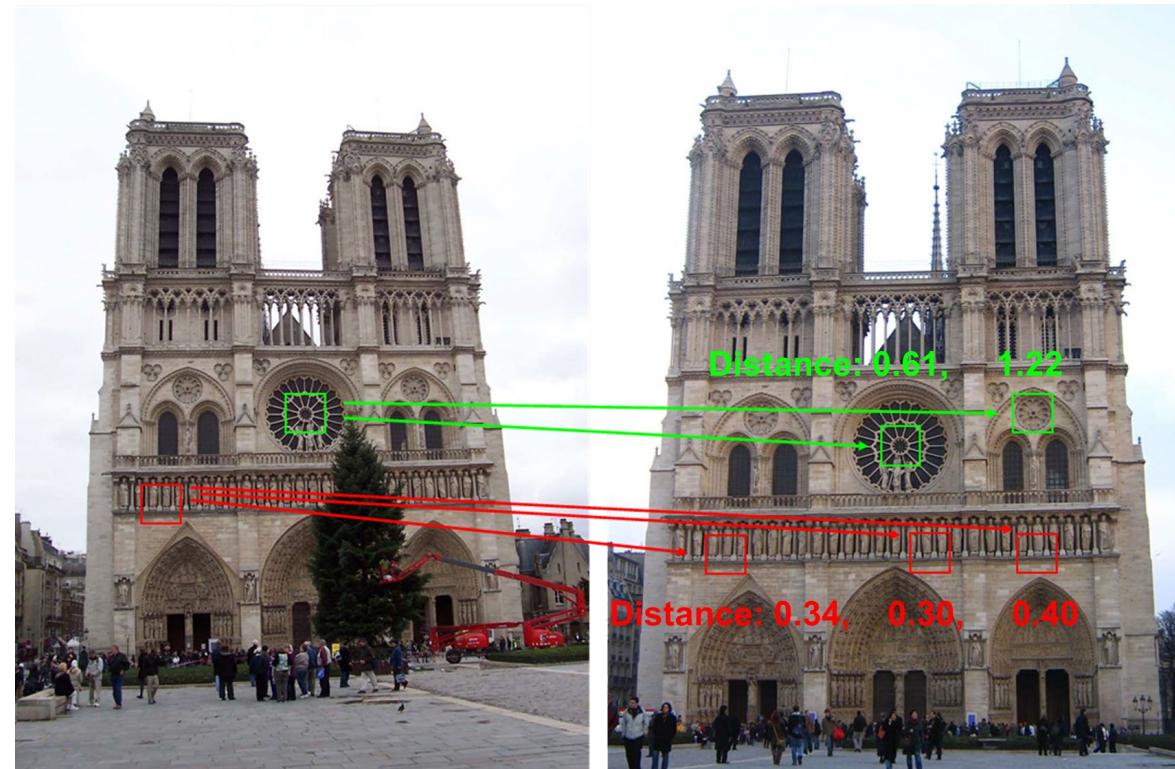
- **Criteria 2:**
  - Compute distance in feature space, e.g., Euclidean distance between 128-dim SIFT descriptors
  - Match point to lowest distance (nearest neighbor)
  - Ignore anything higher than threshold (no match!)



- **Problems:**
  - Threshold is hard to pick
  - Non-distinctive features could have lots of close matches, only one of which is correct

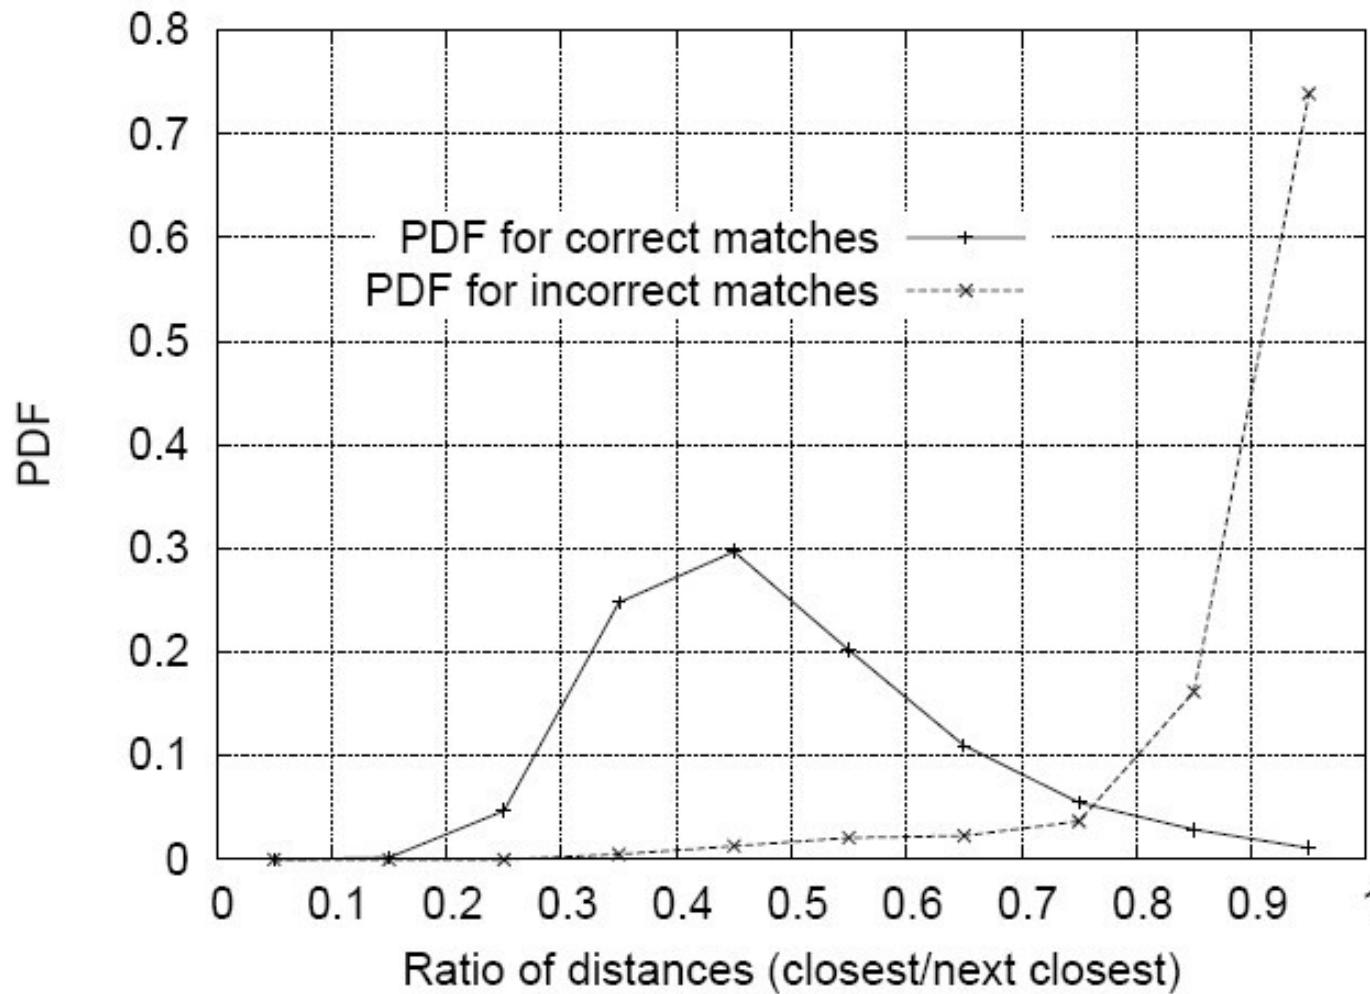
# Nearest Neighbor Distance Ratio

- Compare distance of closest (NN1) and second-closest (NN2) feature vector neighbor:
  - If  $NN1 \approx NN2$ , ratio  $NN1/NN2$  will be  $\approx 1 \Rightarrow$  matches too close
  - As  $NN1 \ll NN2$ , ratio  $NN1/NN2$  tends to 0
- Sorting by this ratio puts matches in order of confidence
- Threshold ratio – but how to choose?



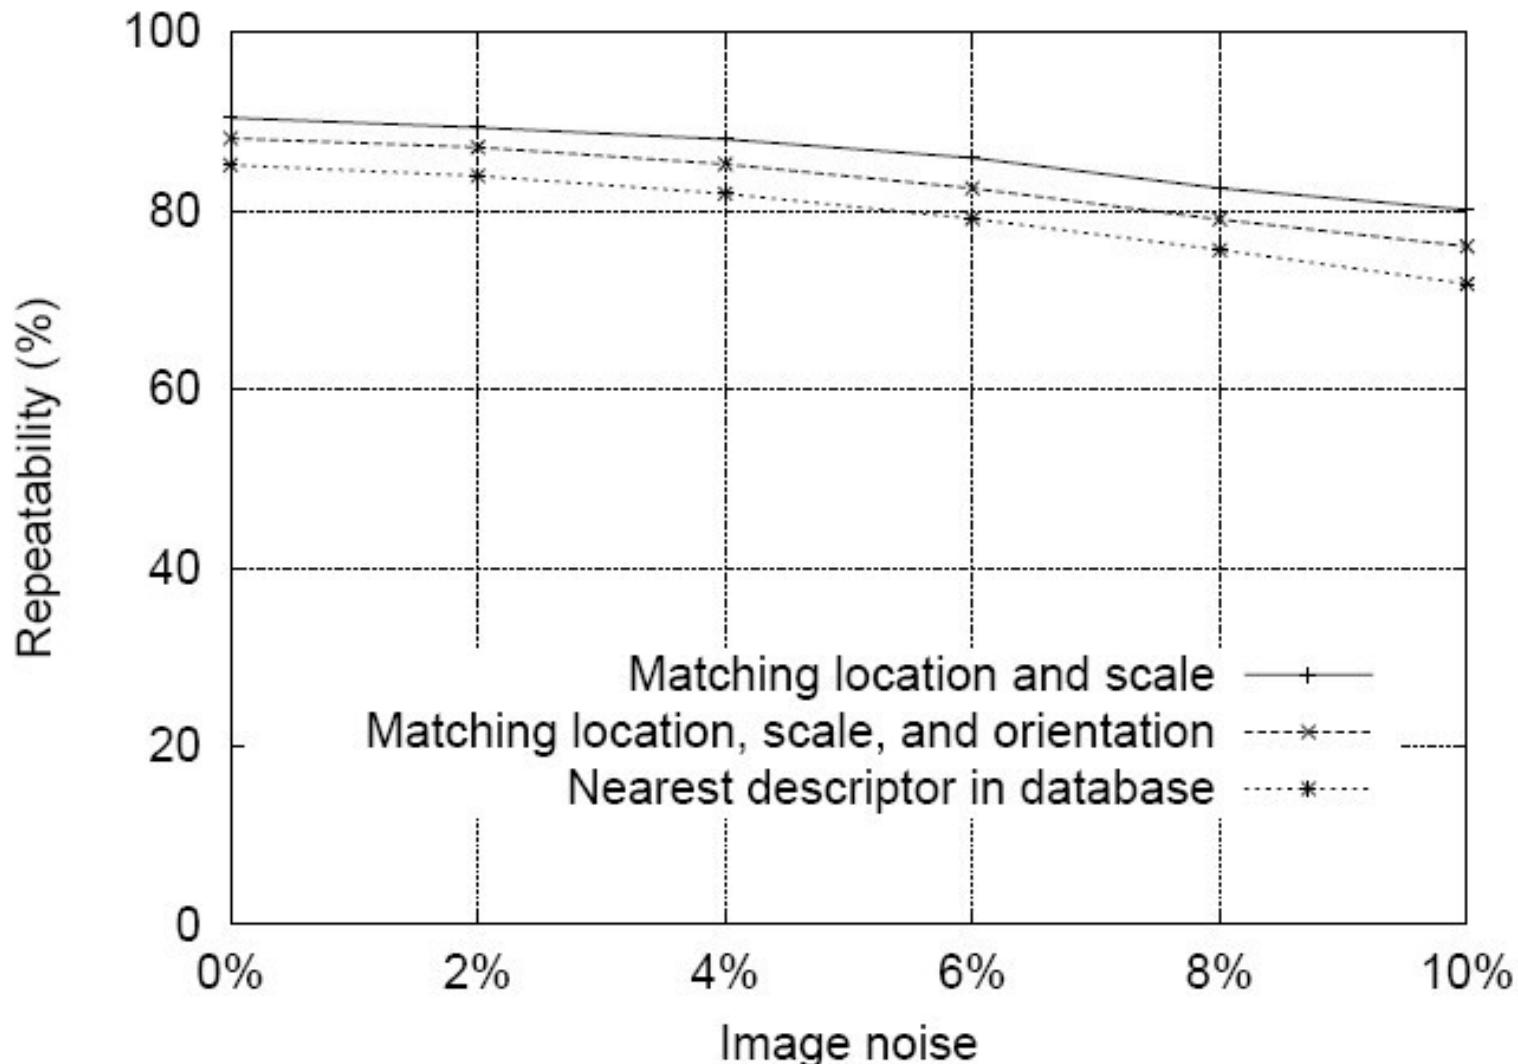
# Nearest Neighbor Distance Ratio

- Lowe computed a probability distribution functions of ratios
- 40,000 keypoints with hand-labeled ground truth

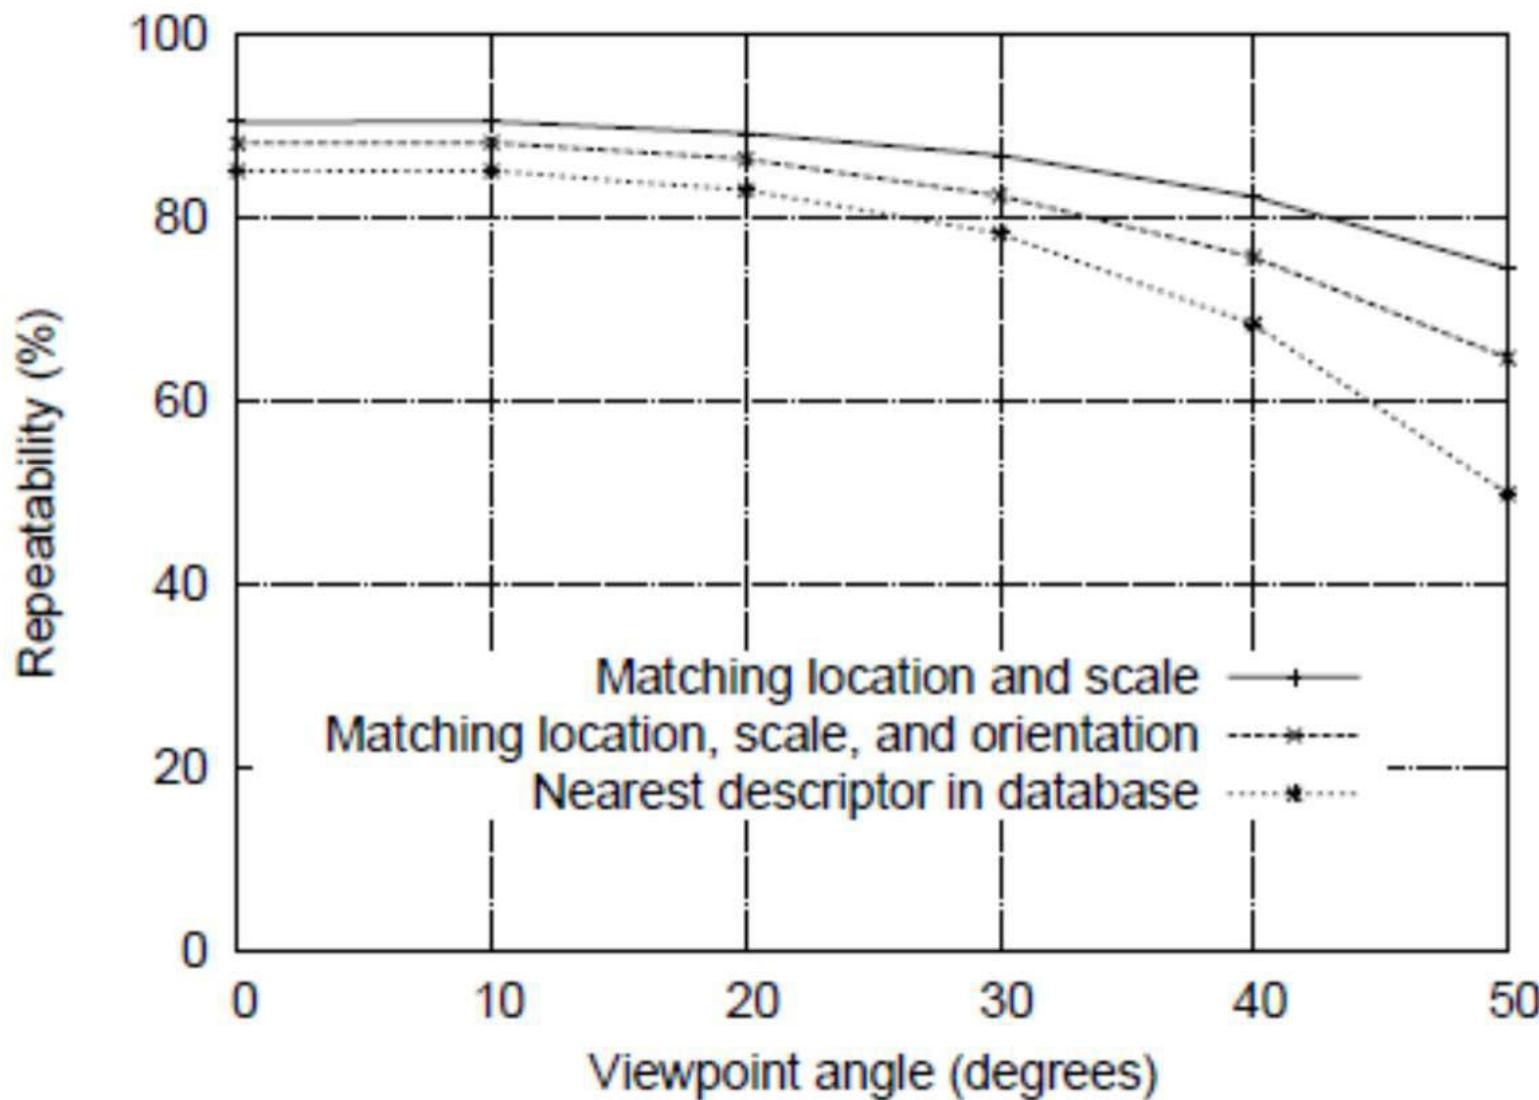


Ratio threshold depends on your application's view on the trade-off between the number of false positives and true positives!

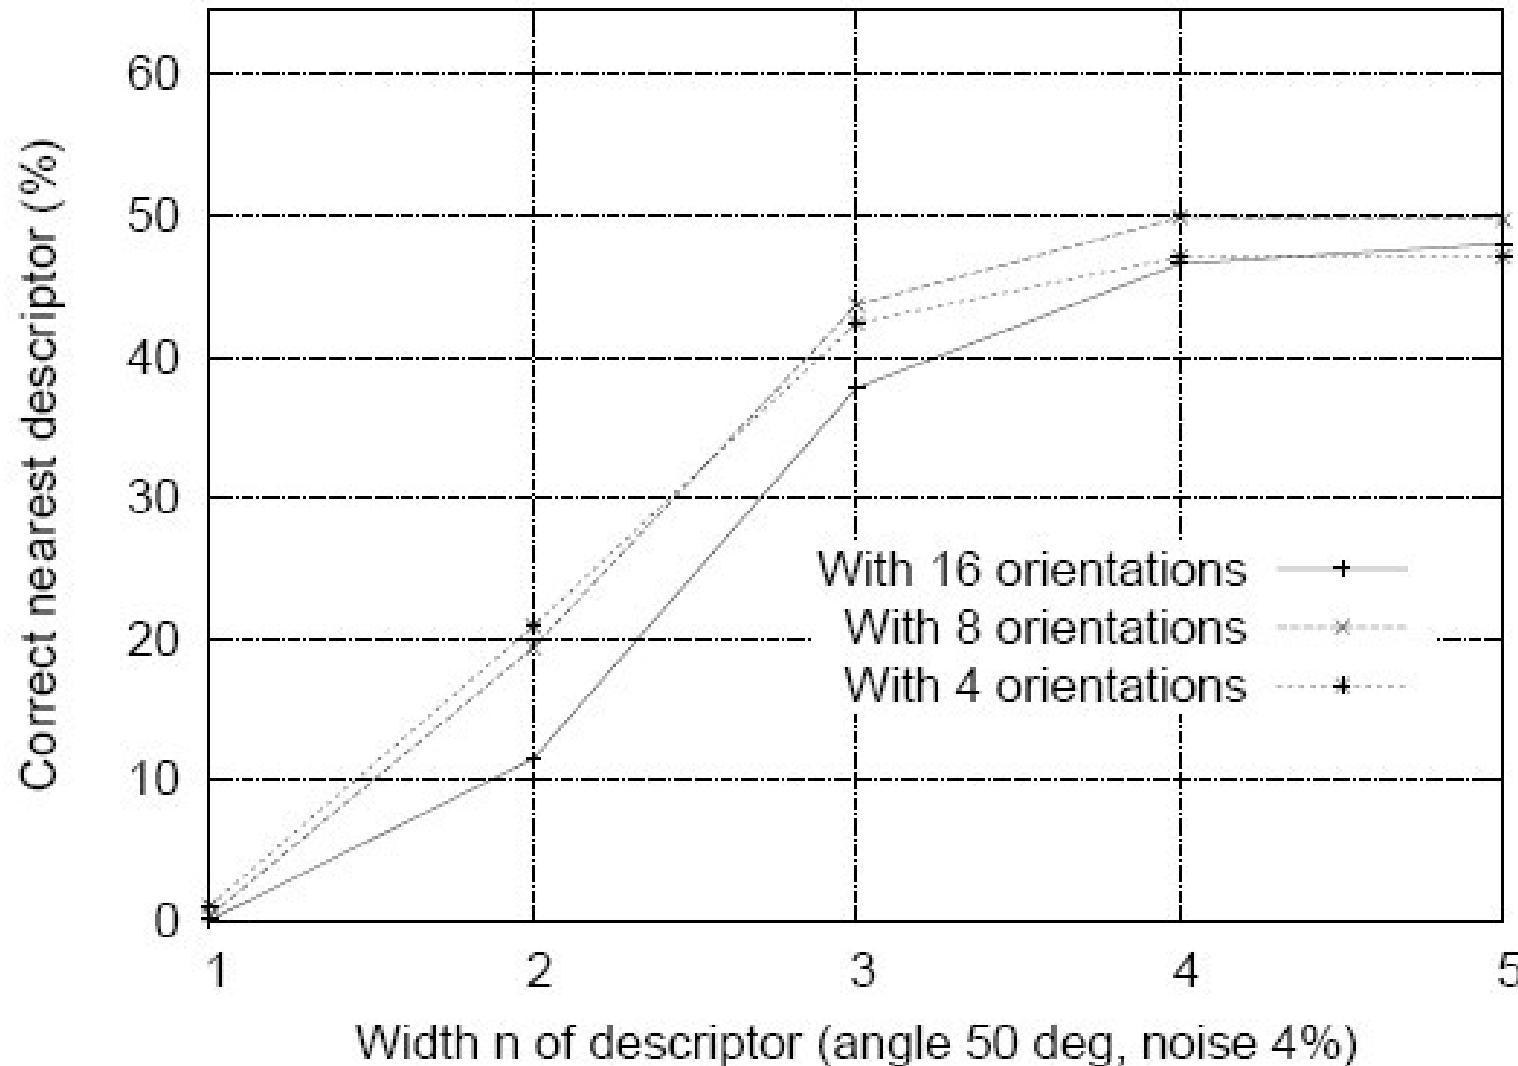
# Additional Slide: SIFT Repeatability



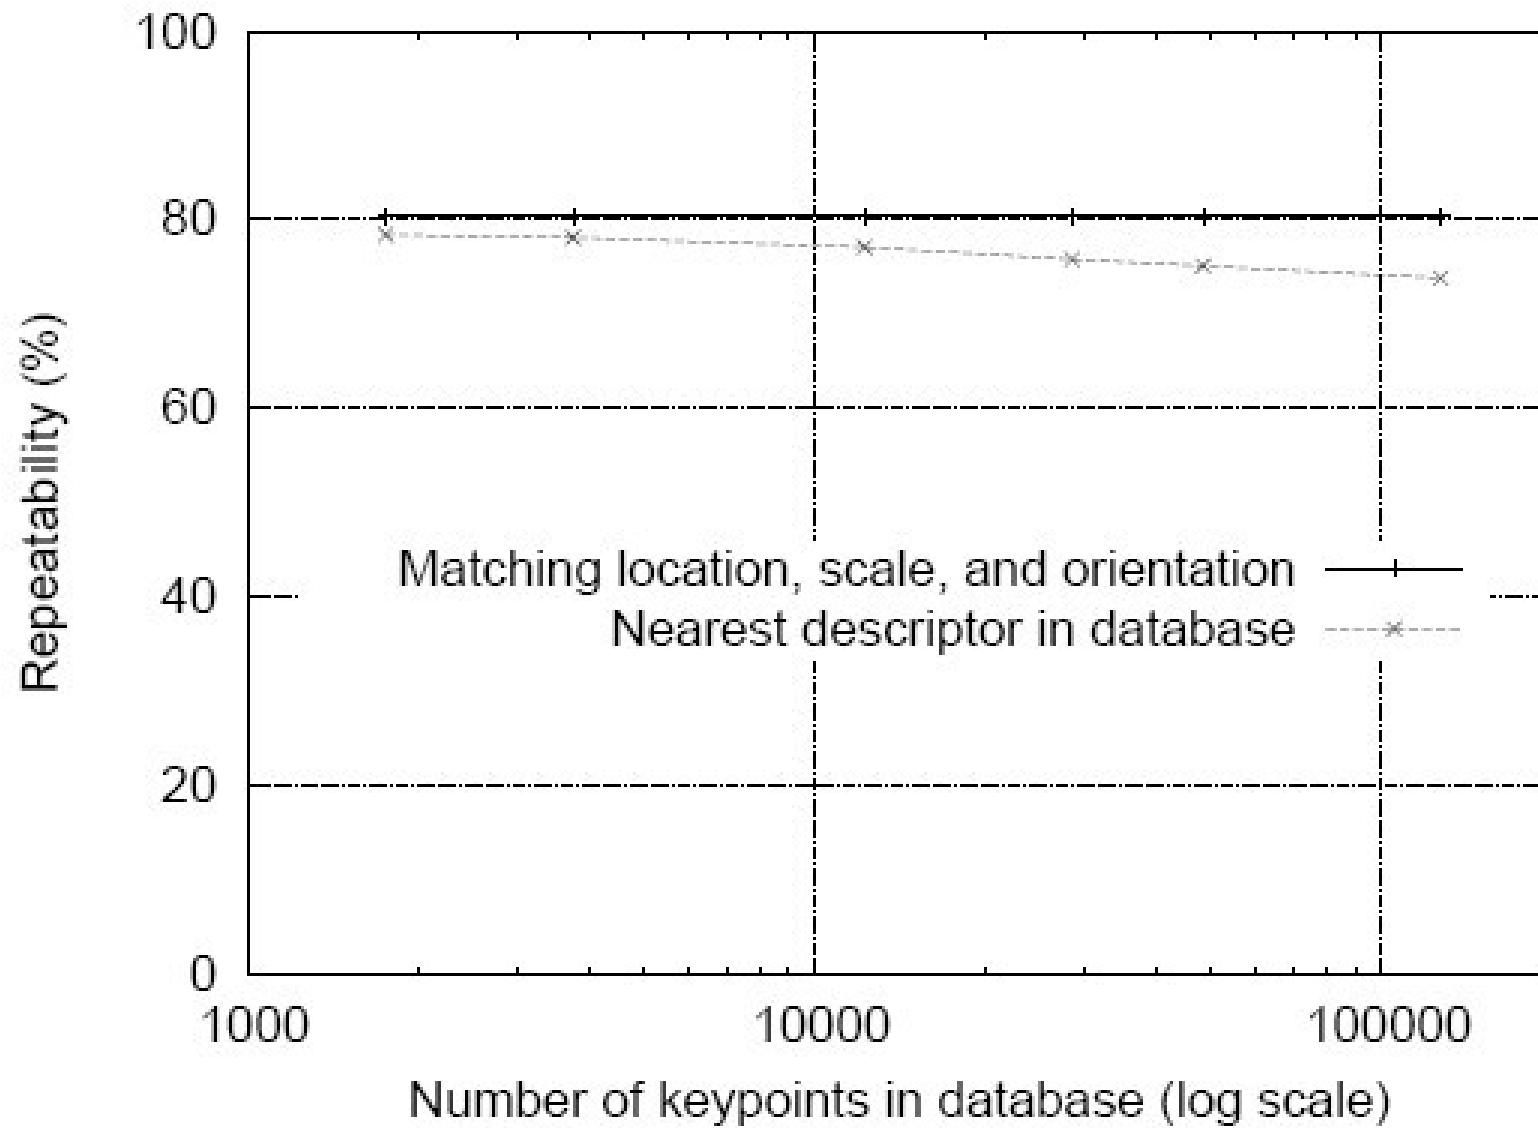
# Additional Slide: SIFT Repeatability



# Additional Slide: SIFT Repeatability

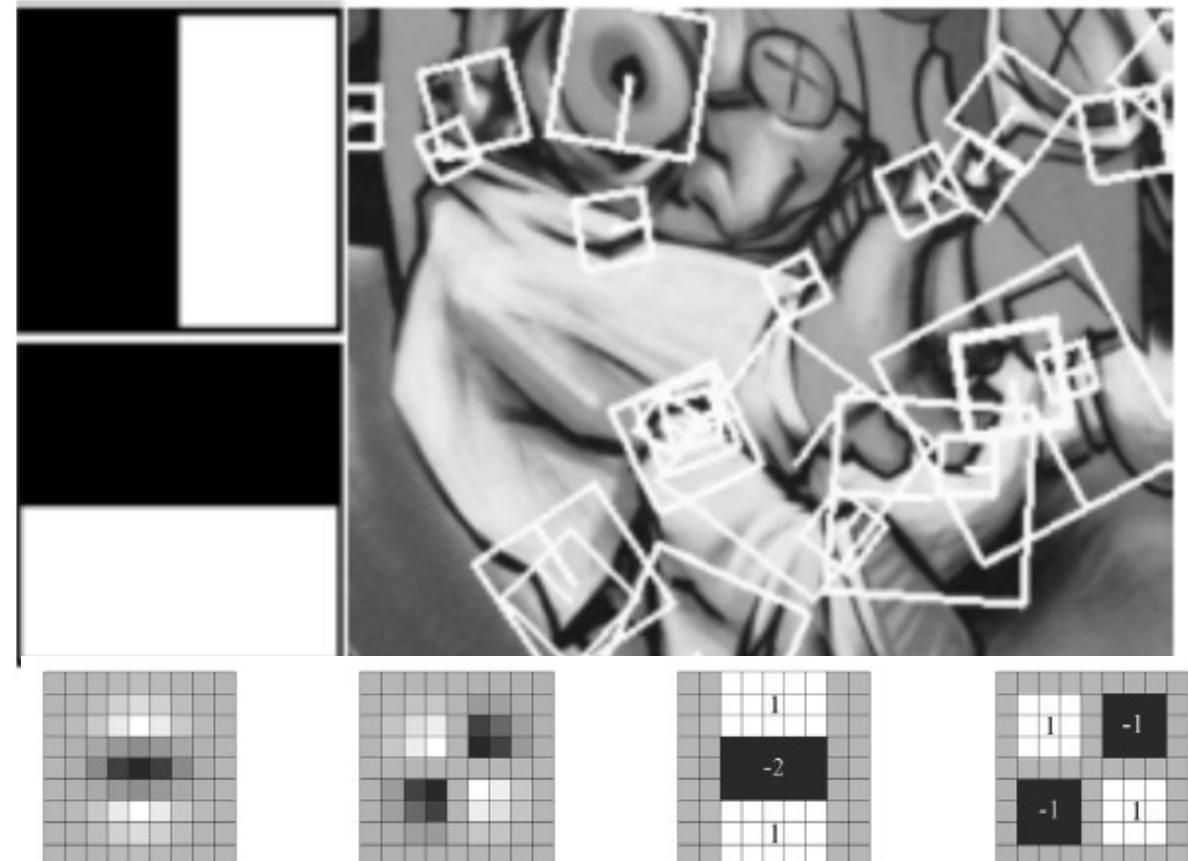


# Additional Slide: SIFT Repeatability

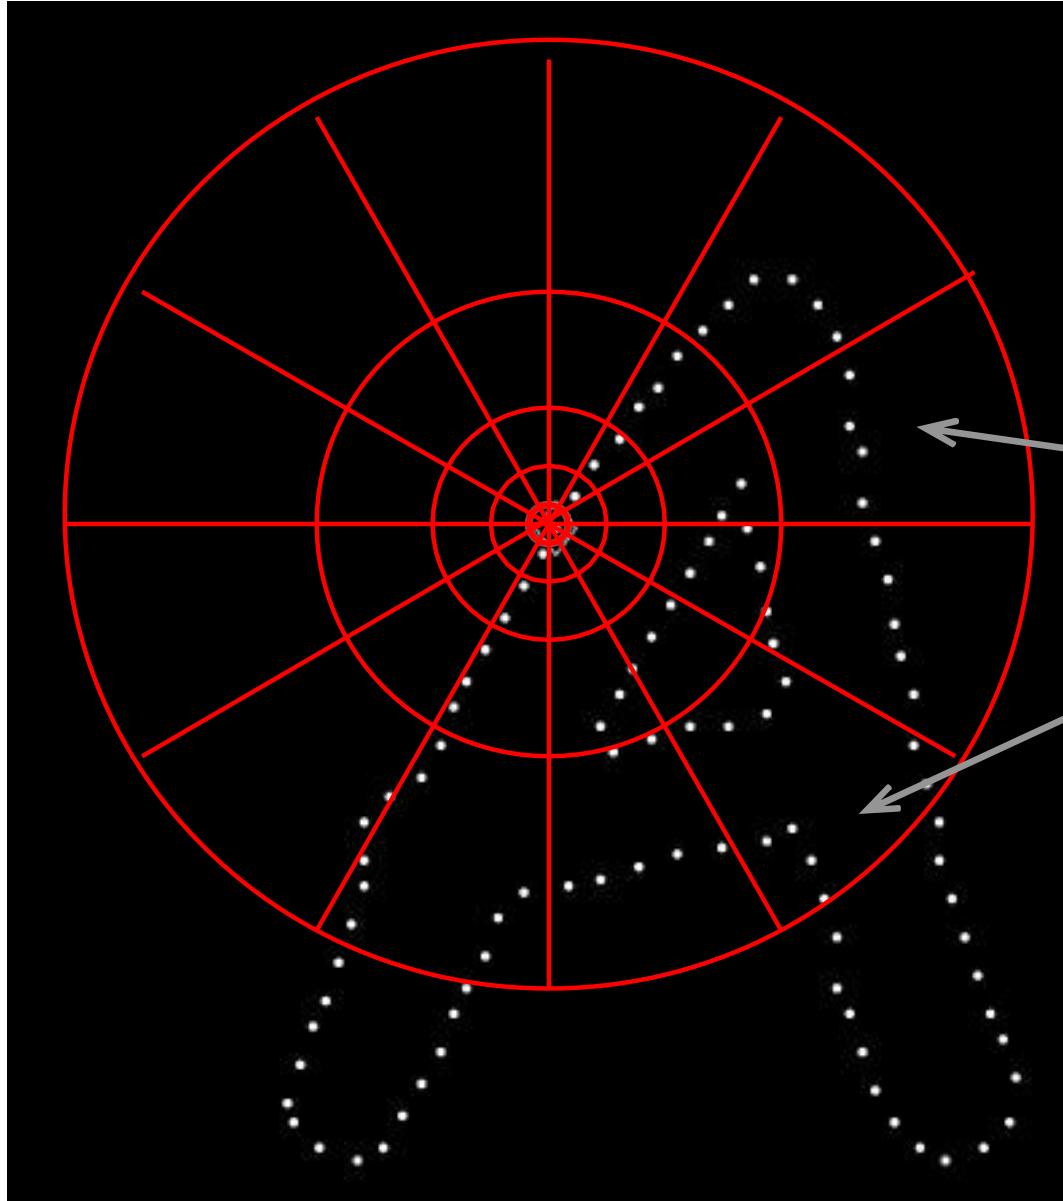


# Local Descriptors: SURF

- SURF is a fast approximation of SIFT
  - Efficient computation by 2D box filters & integral images  
⇒ 6 times faster than SIFT
  - Equivalent quality for object identification
- GPU implementation available
  - Feature extraction @ 200Hz  
(detector + descriptor, 640×480 img)  
<http://www.vision.ee.ethz.ch/~surf>



# Local Descriptors: Shape Context



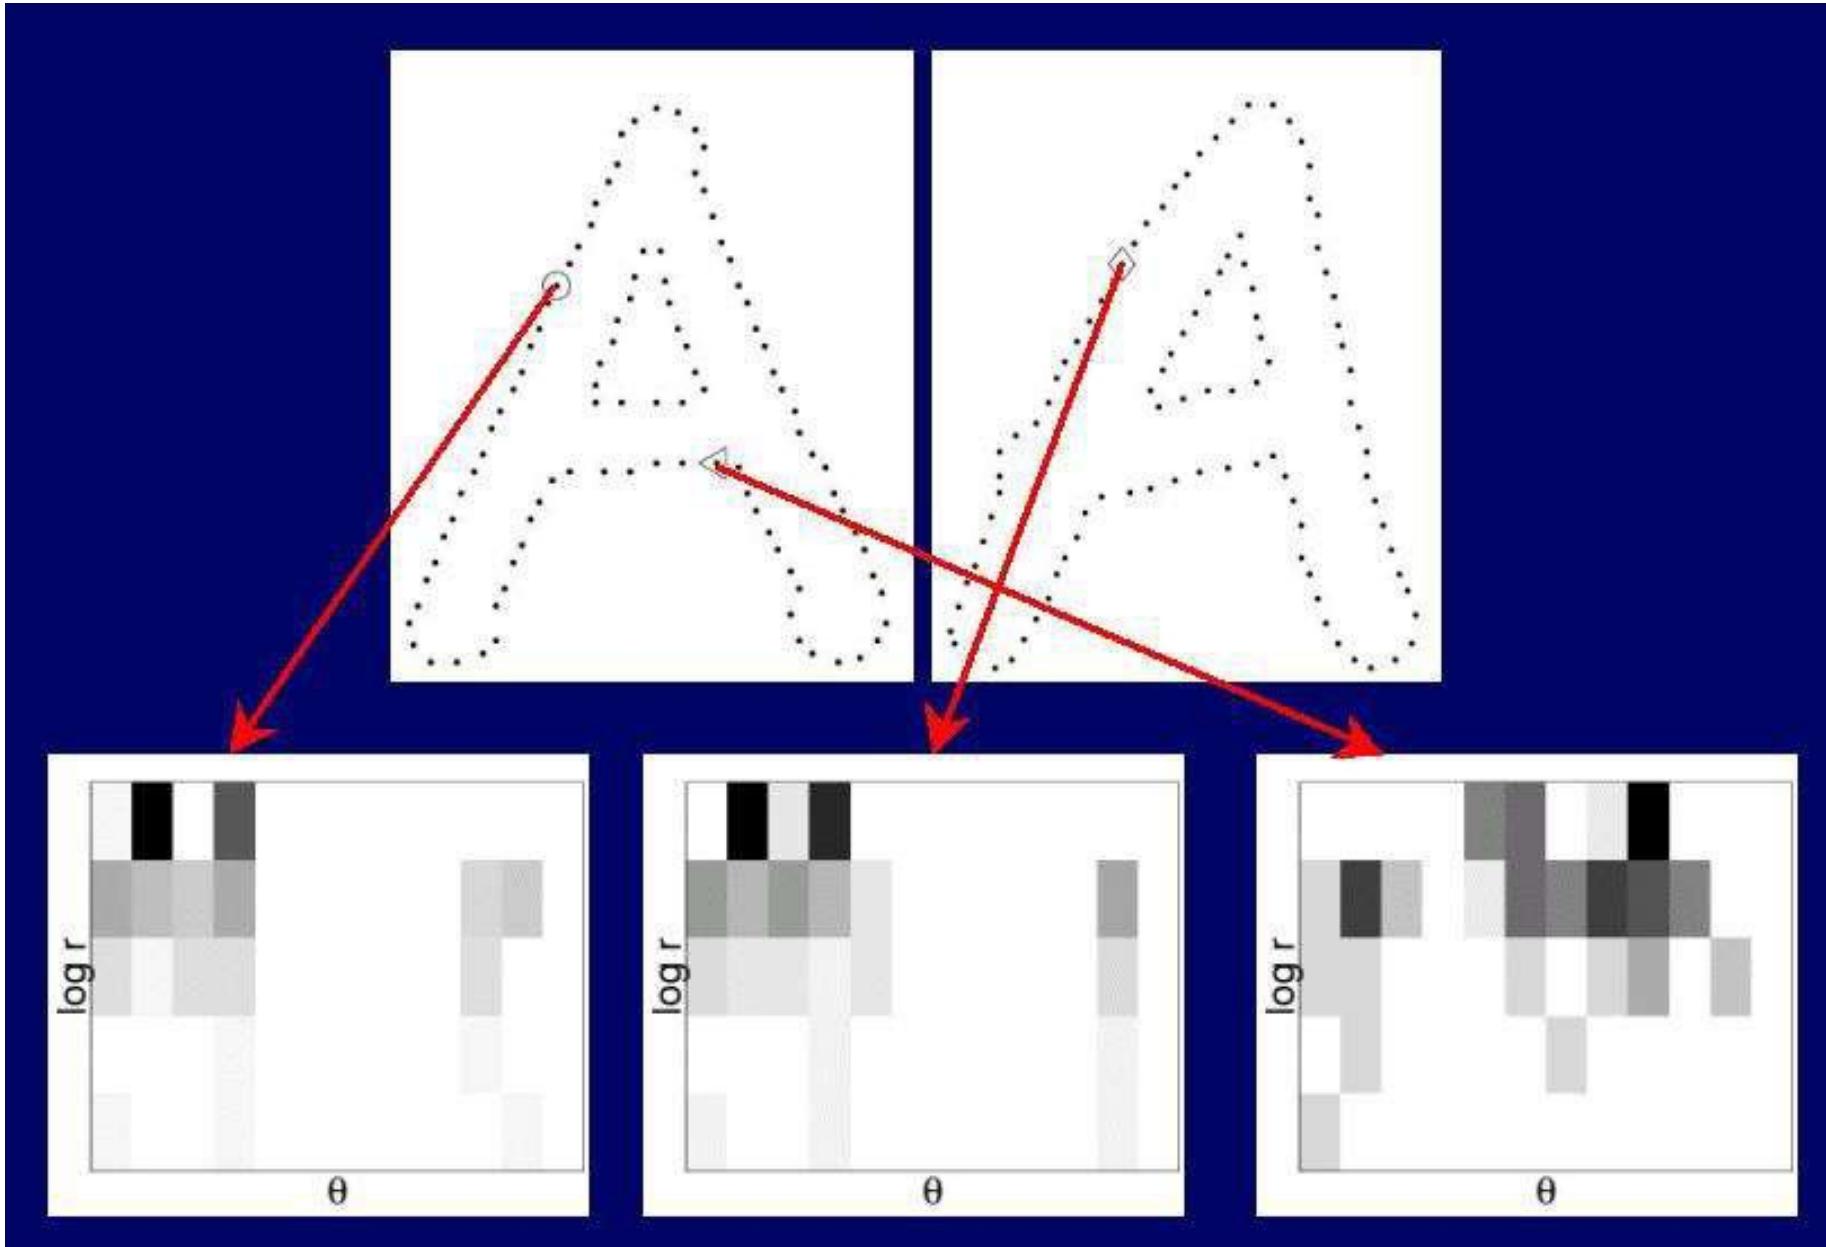
Count the number of points inside each bin:

Count = 4

Count = 10

**Log-polar** binning:  
More precision for nearby points,  
more flexibility for farther points.

# Shape Context Descriptor



# Self-similarity Descriptor

## Application:

- Paper: “Matching Local Self-Similarities across Images and Videos,” Shechtman & Irani, 2007
- Correlate image patch centered at  $q$  with a larger surrounding image region (e.g., of radius 40 pixels), resulting in a local internal “correlation surface”
- Transform correlation surface into a binned log-polar representation
- This results in a compact descriptor  $d_q$  for every pixel  $q$

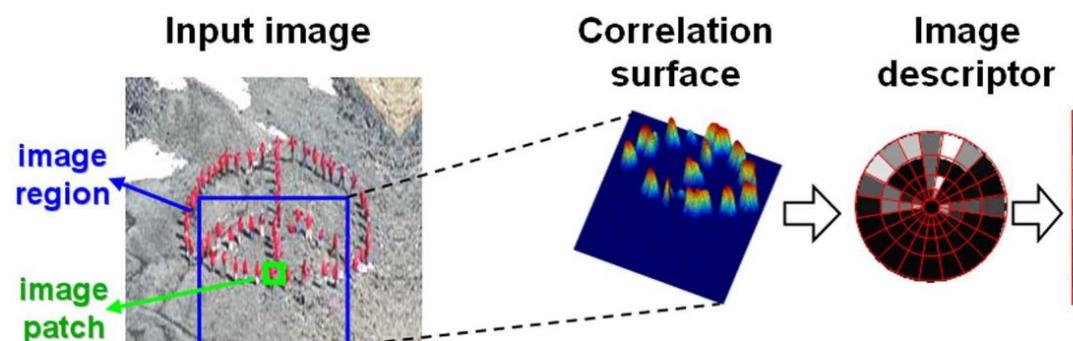


Figure 1. These images of the same object (a heart) do *NOT* share common image properties (colors, textures, edges), but *DO* share a similar geometric layout of local internal self-similarities.

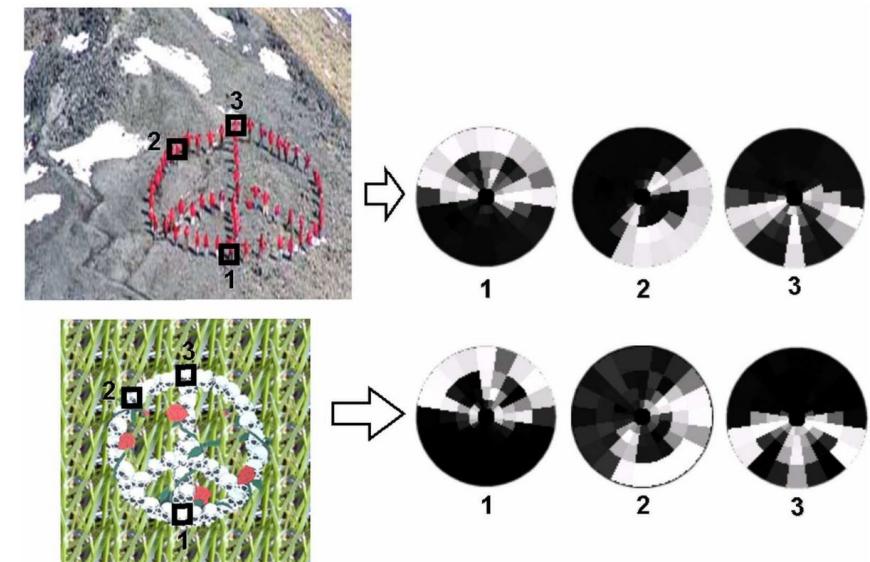


Figure 3. Corresponding “Self-similarity descriptors”. We show a few corresponding points (1,2,3) across two images of the same object, with their “self-similarity” descriptors. Despite the large difference in photometric properties between the two images, their corresponding “self-similarity” descriptors are quite similar.