

# **Introduction to Computer Vision**

---

**Kaveh Fathian**

Assistant Professor

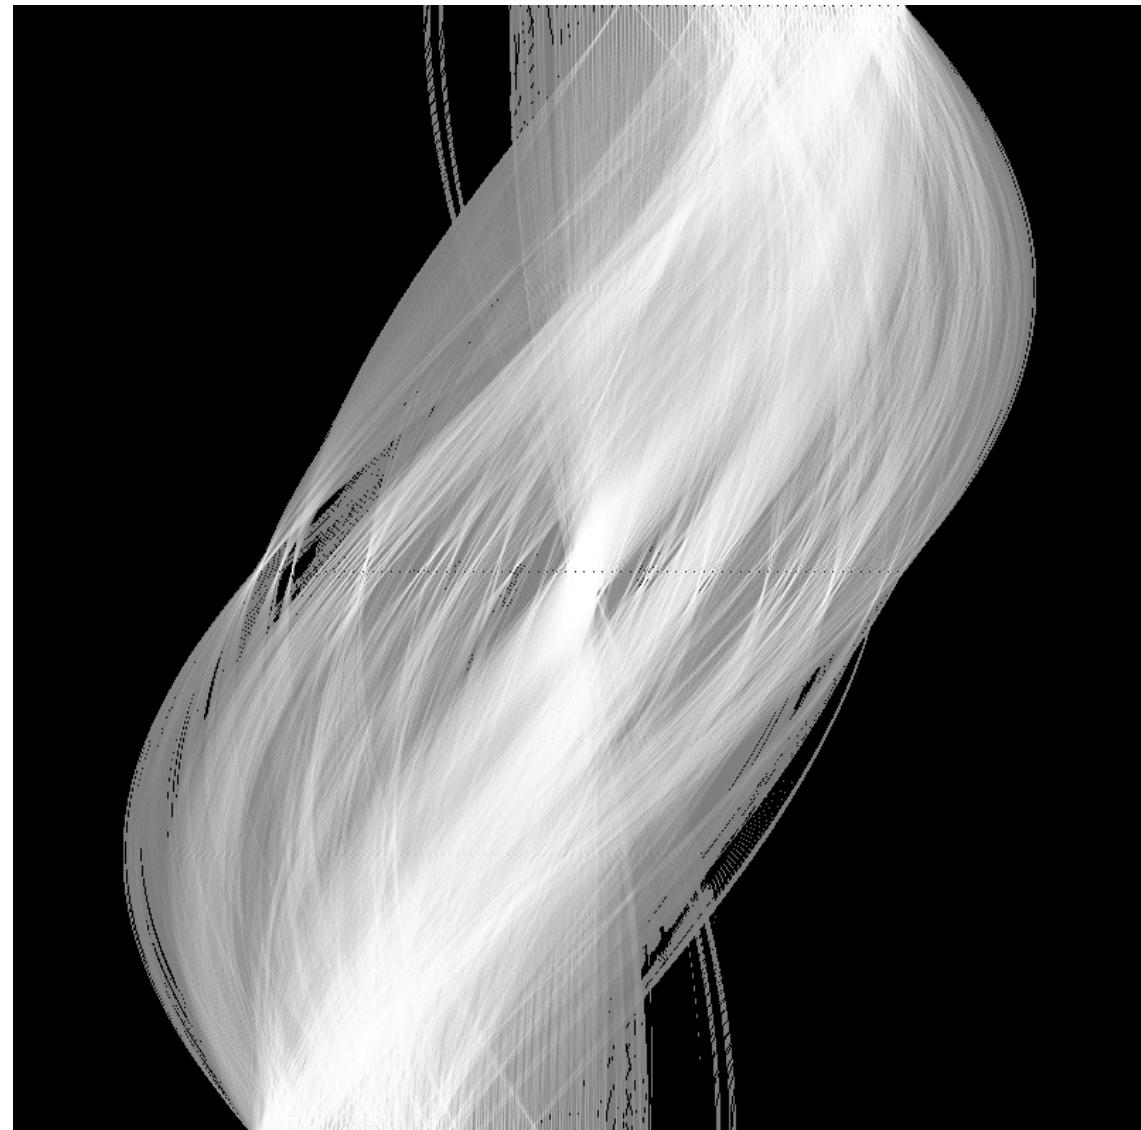
Computer Science Department

Colorado School of Mines

**Lecture 7**

# Learning Outcomes

- Hough Transform
- Line Detection
- Optical Flow



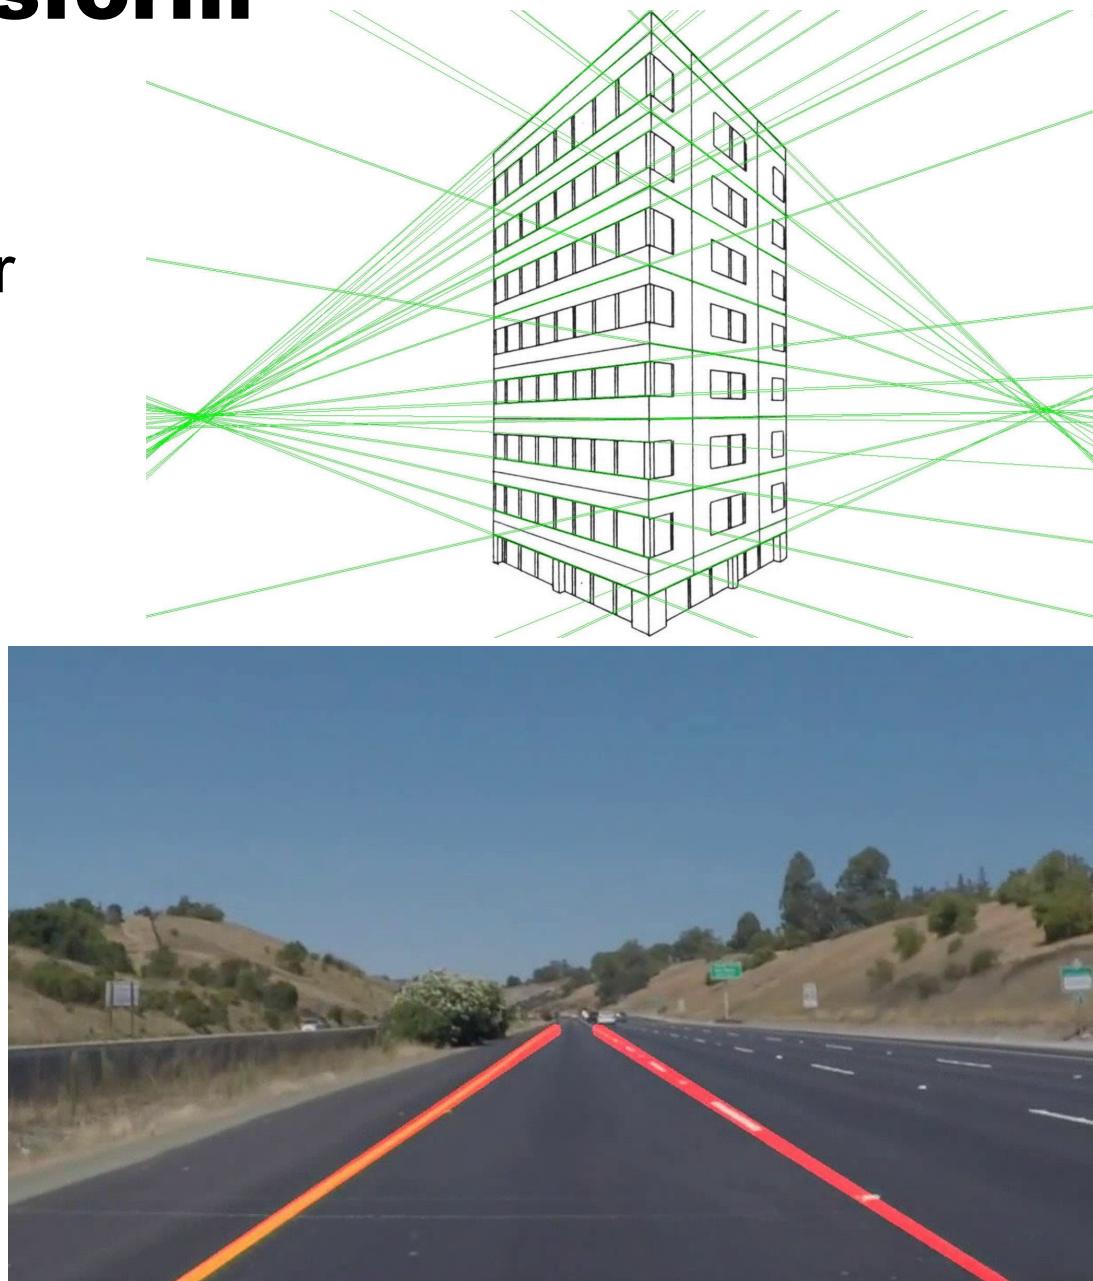
# Hough transform

Hough transform is a **generic framework** for detecting a parametric model

Example: Hough transform can detect lines

Why Hough transform? E.g., vs. Canny edge detection?

- Edges don't have to be connected
- Lines can be occluded
- Key idea: edges **vote** for the possible models

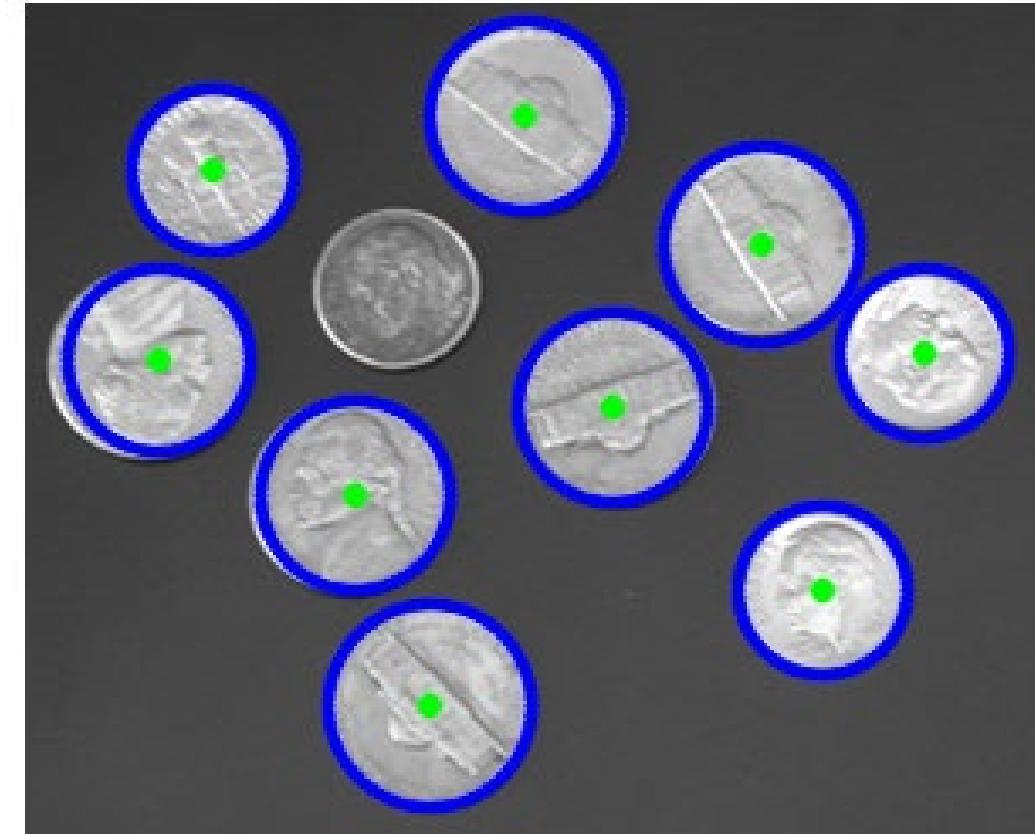


# Circles

source



detected circles



Example: Hough transform for circle detection  
What is its application?

# Image and parameter space

variables  
 $y = mx + b$   
parameters

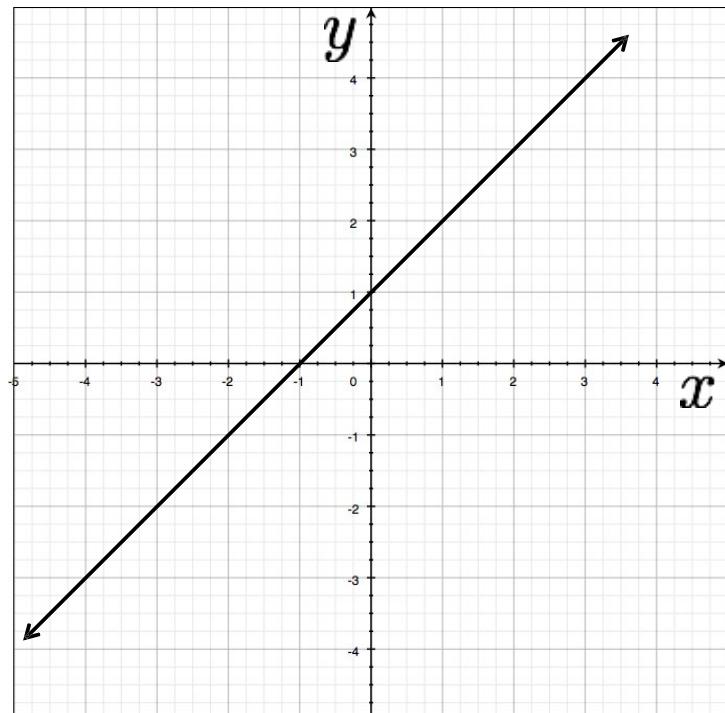
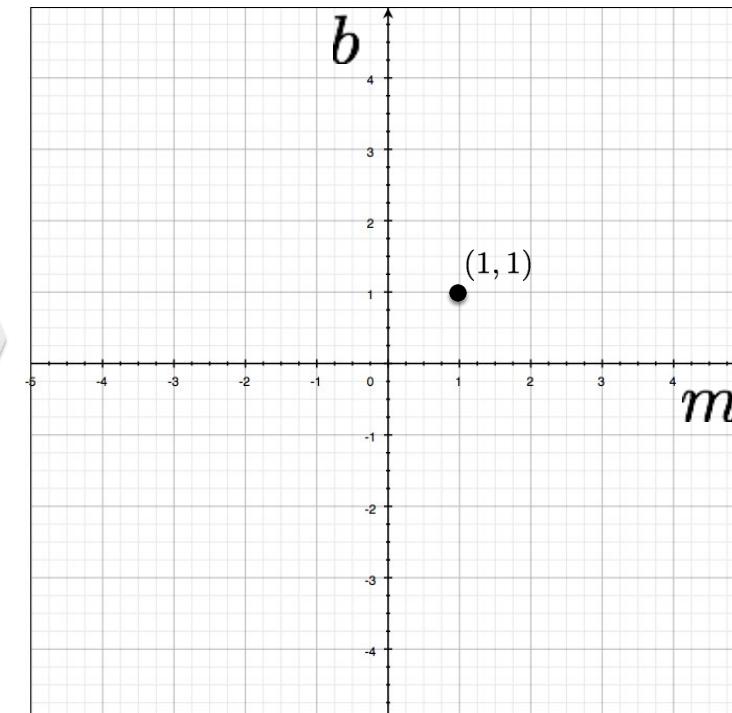


Image space

variables  
 $y - mx = b$   
parameters

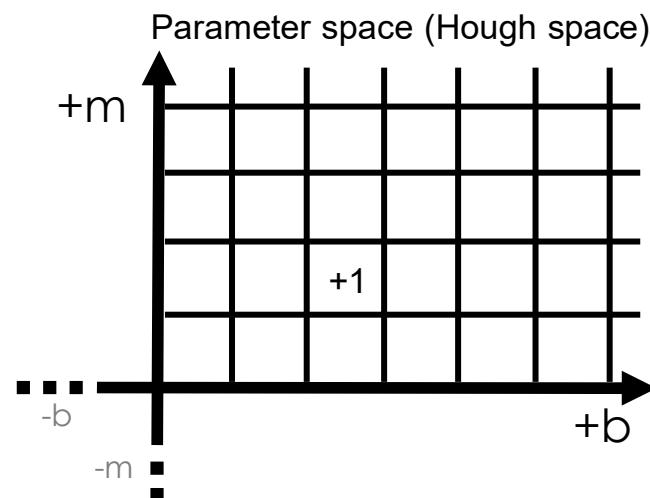
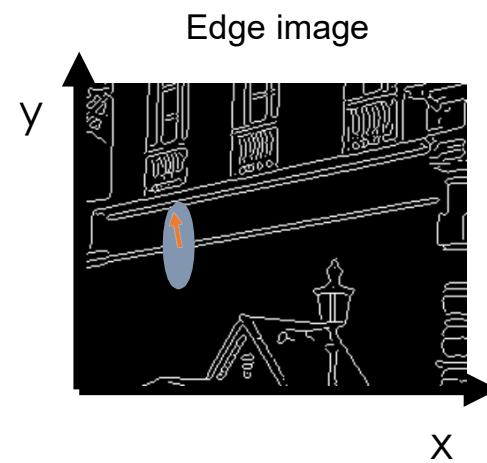


a line  
becomes a  
point

Parameter space

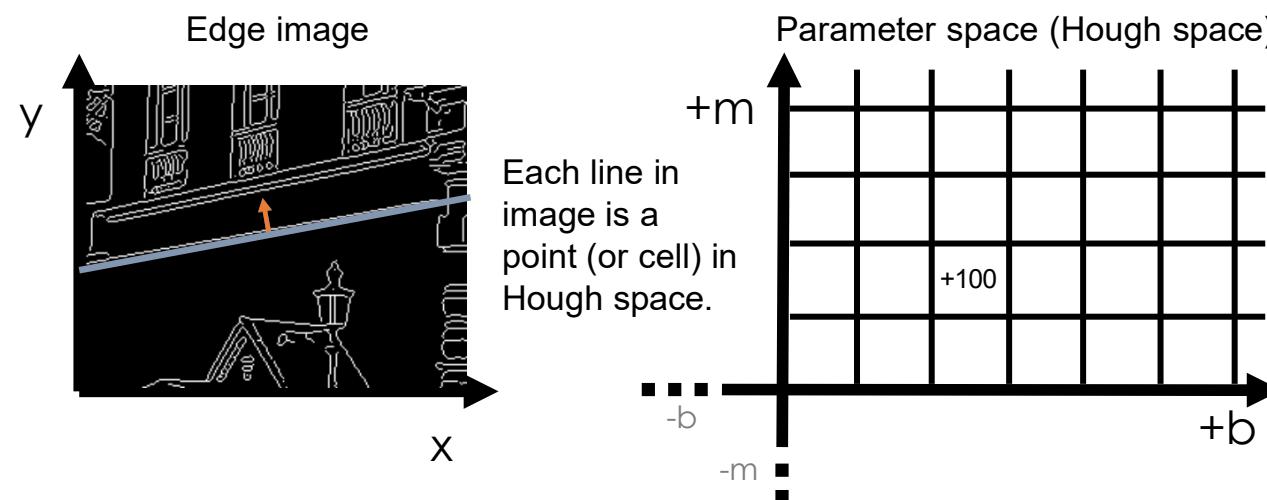
# Hough Transform: Outline

- Create a *grid* of candidate  $m, b$  parameter values: **accumulator**
  - Why a grid?
  - $m, b$  are continuous; grid discretizes into hypotheses.
- Each edge pixel votes for a set of parameters, which increments those values in grid



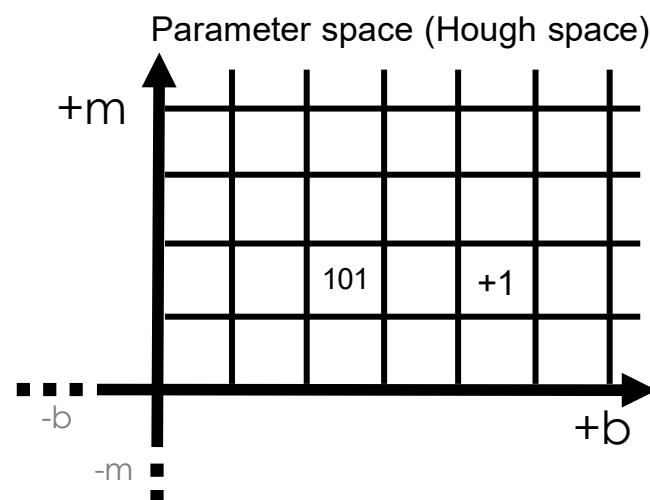
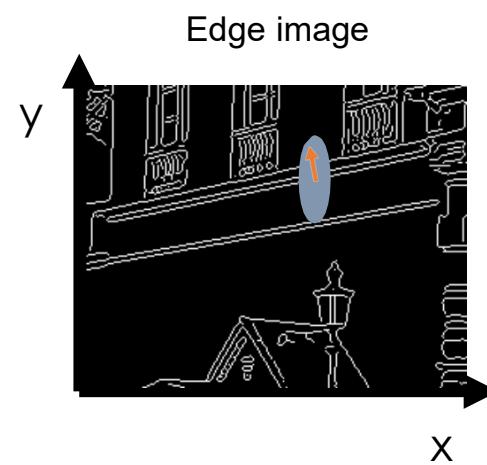
# Hough Transform: Outline

- Create a *grid* of candidate  $m, b$  parameter values: **accumulator**
  - Why a grid?
  - $m, b$  are continuous; grid discretizes into hypotheses.
- Each edge pixel votes for a set of parameters, which increments those values in grid



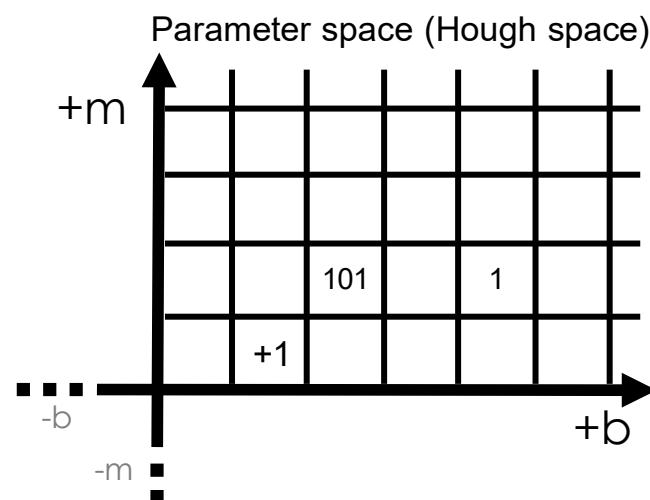
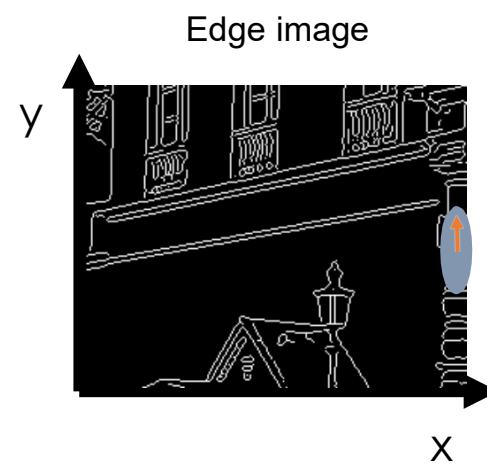
# Hough Transform: Outline

- Create a *grid* of candidate  $m, b$  parameter values: **accumulator**
  - Why a grid?
  - $m, b$  are continuous; grid discretizes into hypotheses.
- Each edge pixel votes for a set of parameters, which increments those values in grid



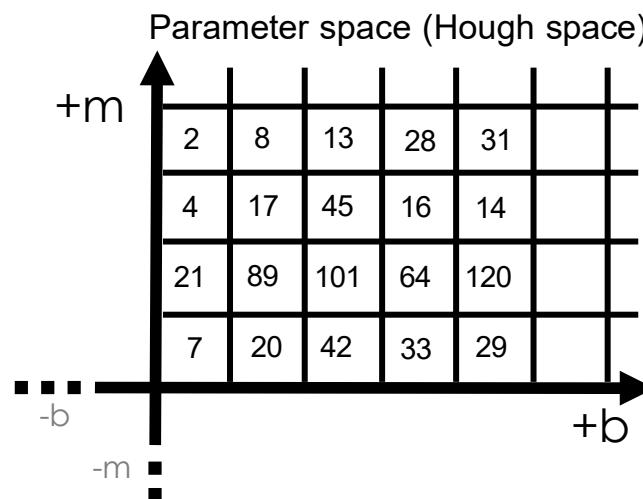
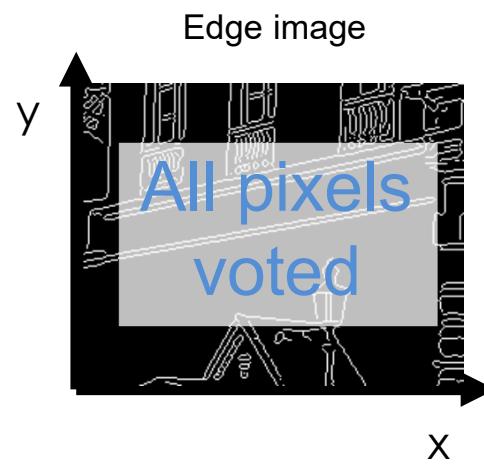
# Hough Transform: Outline

- Create a *grid* of candidate  $m, b$  parameter values: **accumulator**
  - Why a grid?
  - $m, b$  are continuous; grid discretizes into hypotheses.
- Each edge pixel votes for a set of parameters, which increments those values in grid



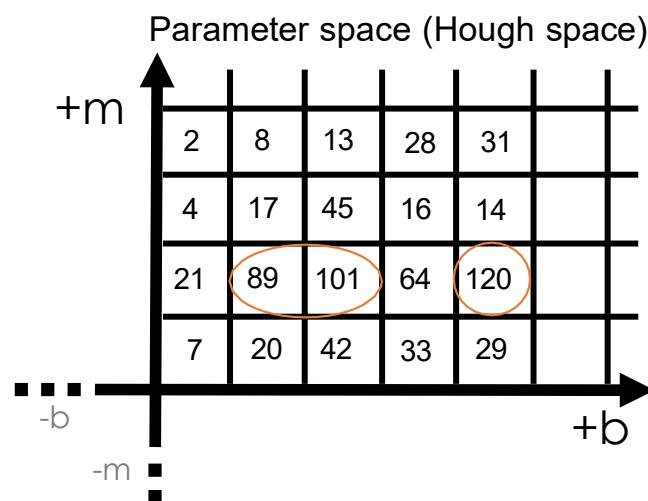
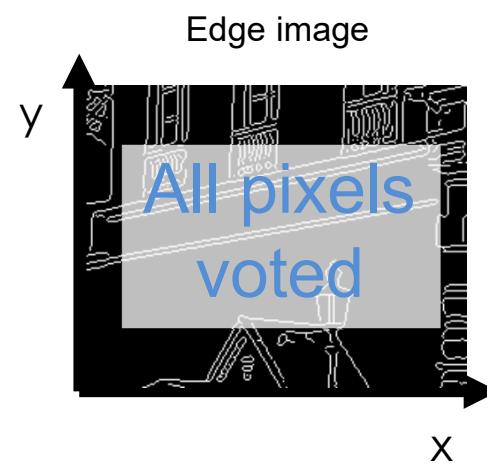
# Hough Transform: Outline

- Create a *grid* of candidate  $m, b$  parameter values: **accumulator**
  - Why a grid?
  - $m, b$  are continuous; grid discretizes into hypotheses.
- Each edge pixel votes for a set of parameters, which increments those values in grid



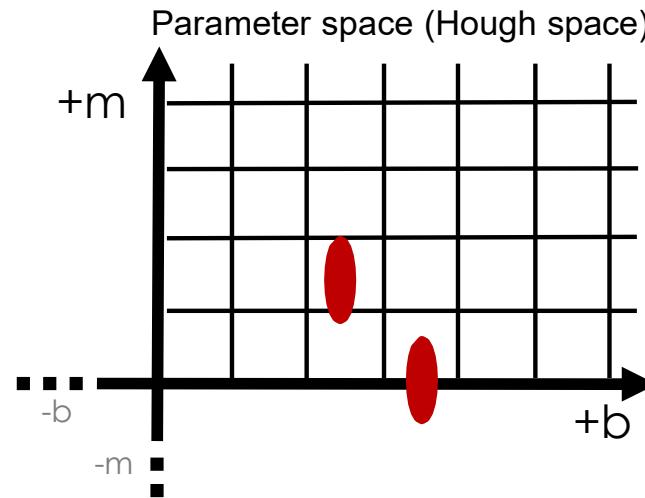
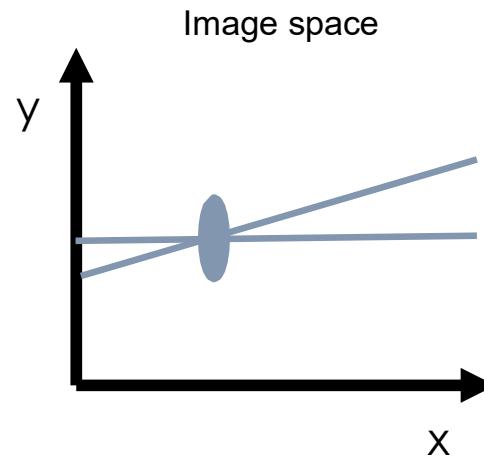
# Hough Transform: Outline

- Create a *grid* of candidate  $m, b$  parameter values: **accumulator**
  - Why a grid?
  - $m, b$  are continuous; grid discretizes into hypotheses.
- Each edge pixel votes for a set of parameters, which increments those values in grid
- Find maxima – our line candidates.



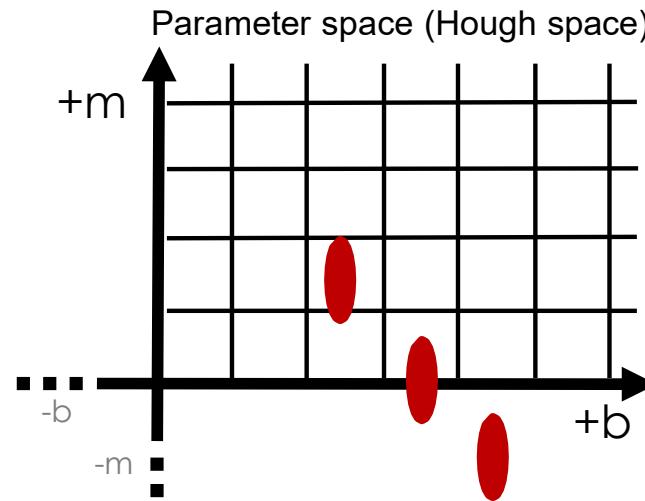
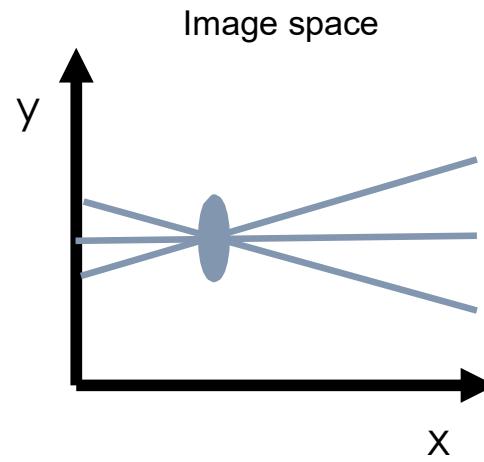
# Hough Transform: Step back

- Hough space represents all possible lines.
- With gradient information constraint:
  - Edge is single point in Hough space.
- Without gradient orientation information?



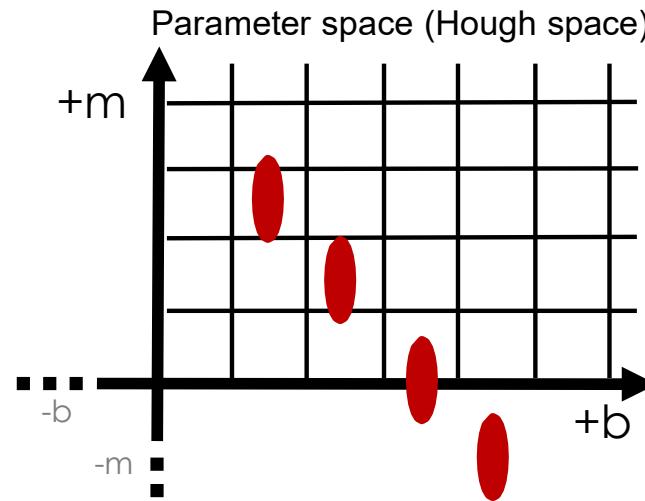
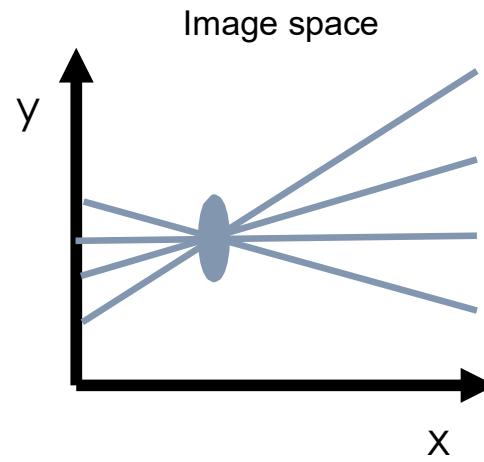
# Hough Transform: Step back

- Hough space represents all possible lines.
- With gradient information constraint:
  - Edge is single point in Hough space.
- Without gradient orientation information?



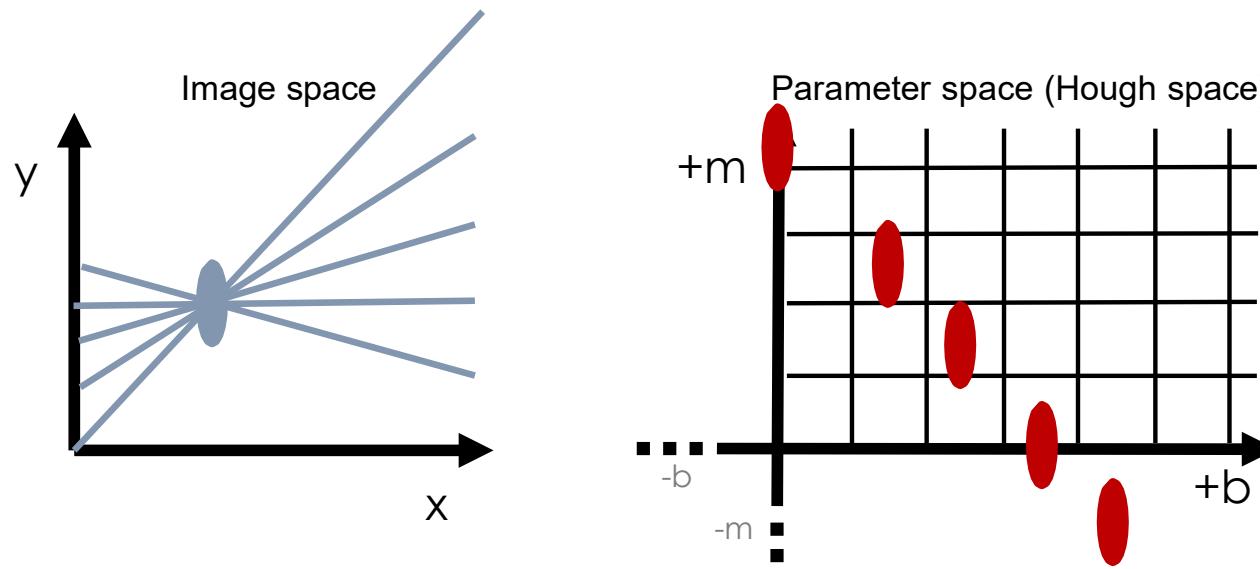
# Hough Transform: Step back

- Hough space represents all possible lines.
- With gradient information constraint:
  - Edge is single point in Hough space.
- Without gradient orientation information?



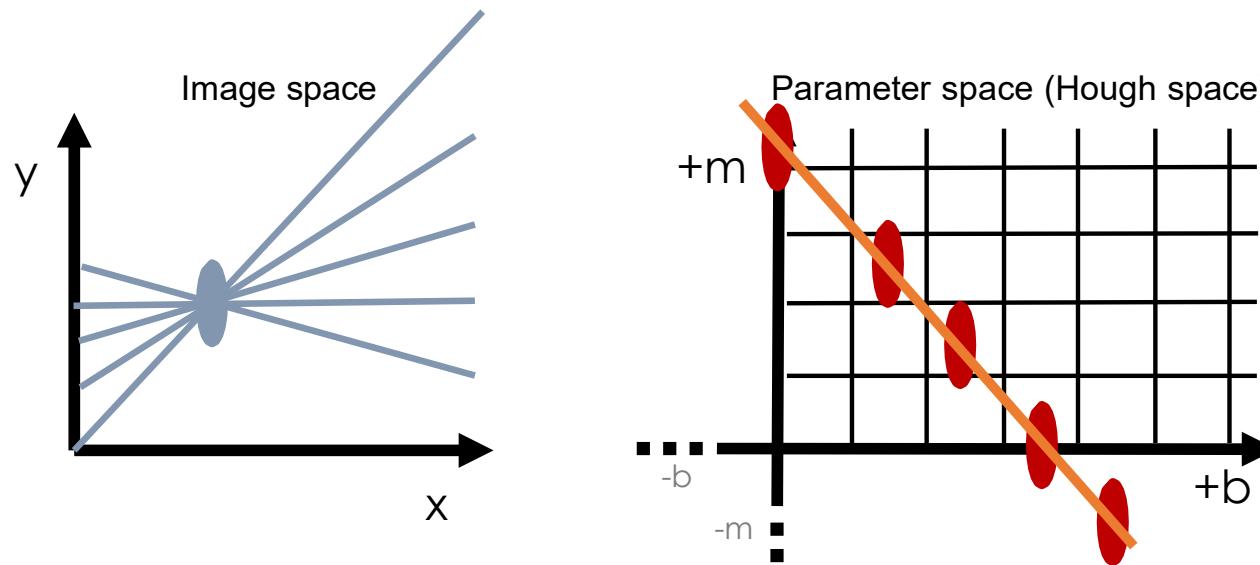
# Hough Transform: Step back

- Hough space represents all possible lines.
- With gradient information constraint:
  - Edge is single point in Hough space.
- Without gradient orientation information?



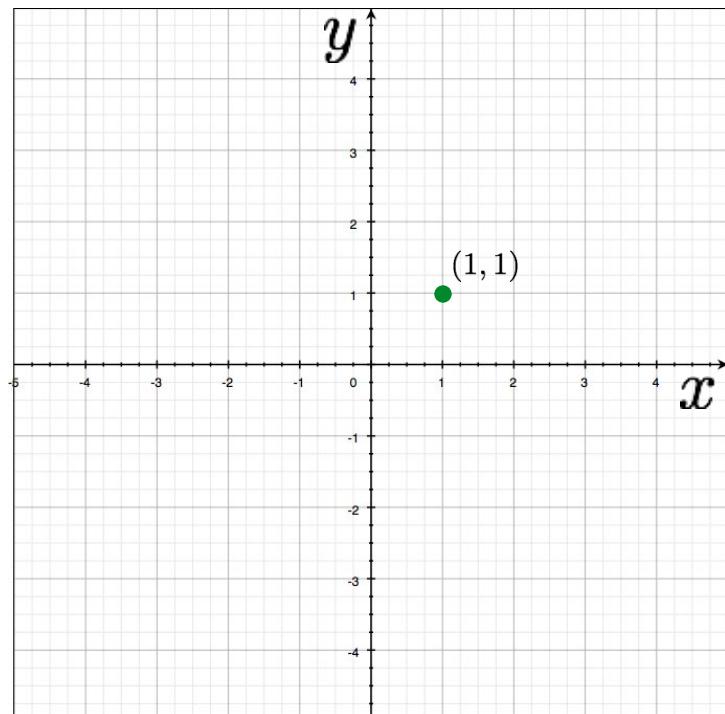
# Hough Transform: Step back

- Hough space represents all possible lines.
- With gradient information constraint:
  - Edge is single point in Hough space.
- Without gradient orientation information?
  - Point is line in Hough space



# Image and parameter space

variables  
 $y = mx + b$   
parameters



a point becomes a line

variables  
 $y - mx = b$   
parameters

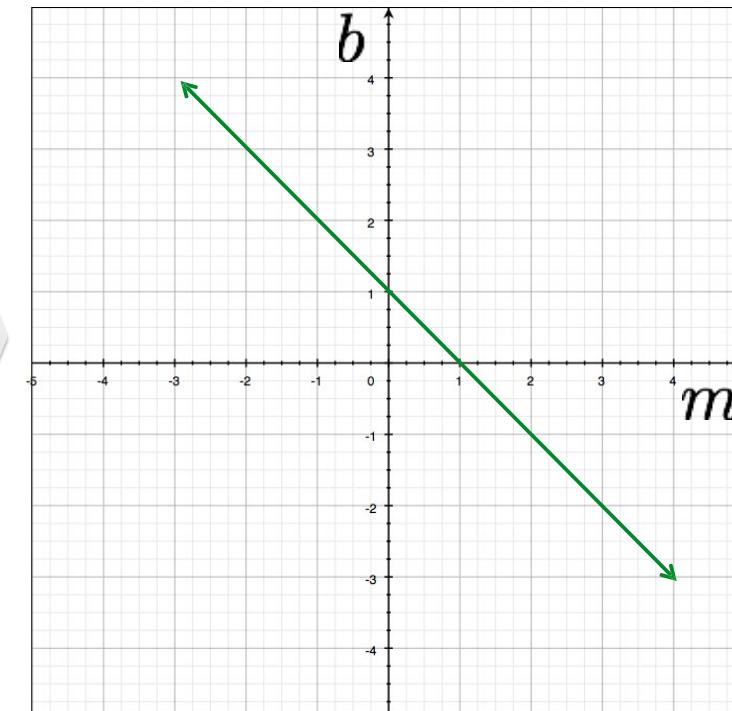
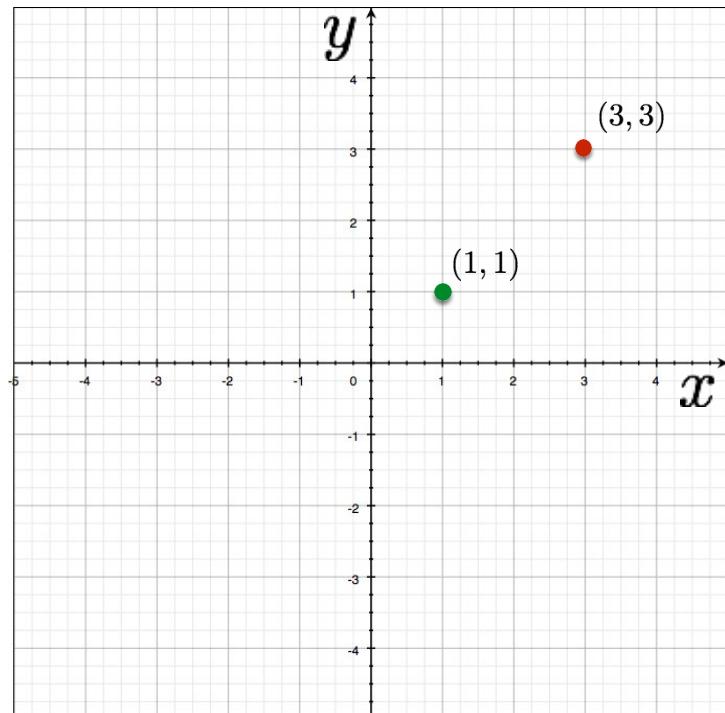


Image space

Parameter space

# Image and parameter space

variables  
 $y = mx + b$   
parameters



two points  
become  
?

variables  
 $y - mx = b$   
parameters

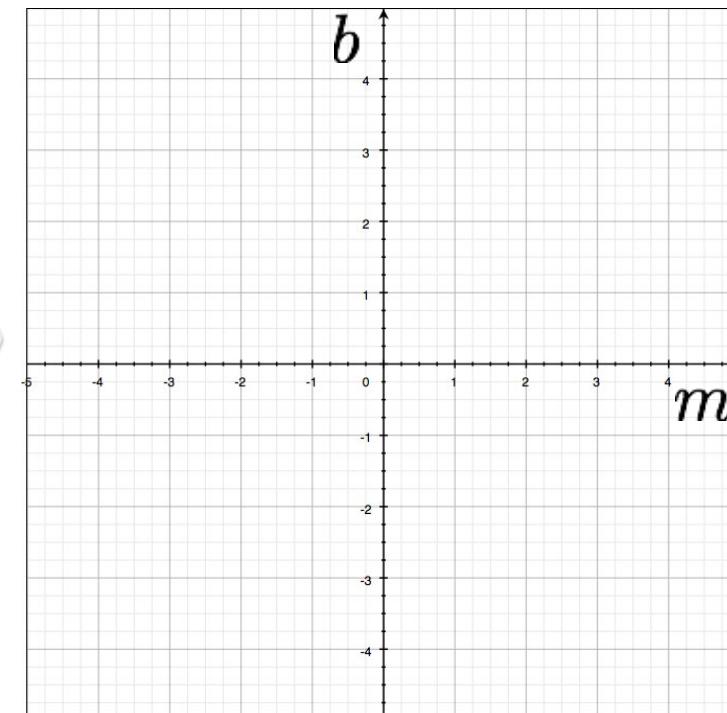
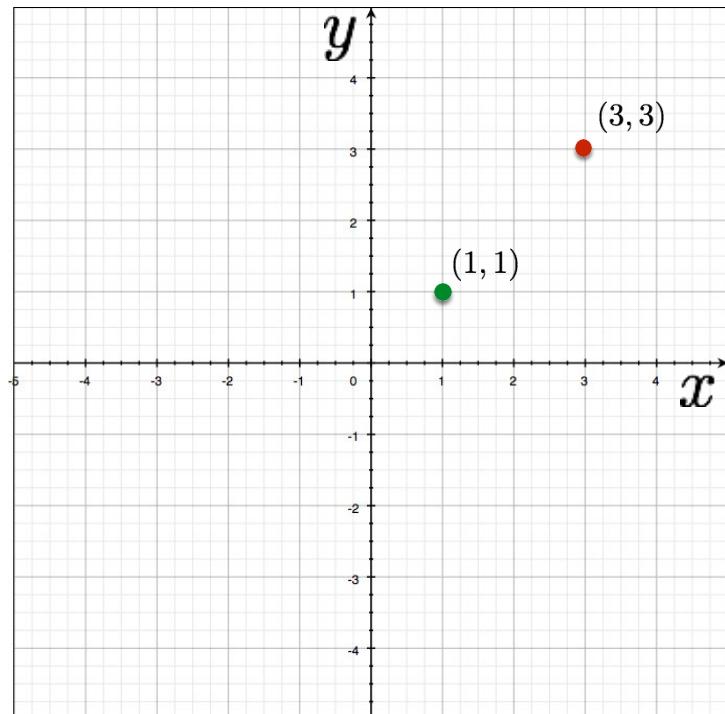


Image space

Parameter space

# Image and parameter space

variables  
 $y = mx + b$   
parameters



two points  
become  
?

variables  
 $y - mx = b$   
parameters

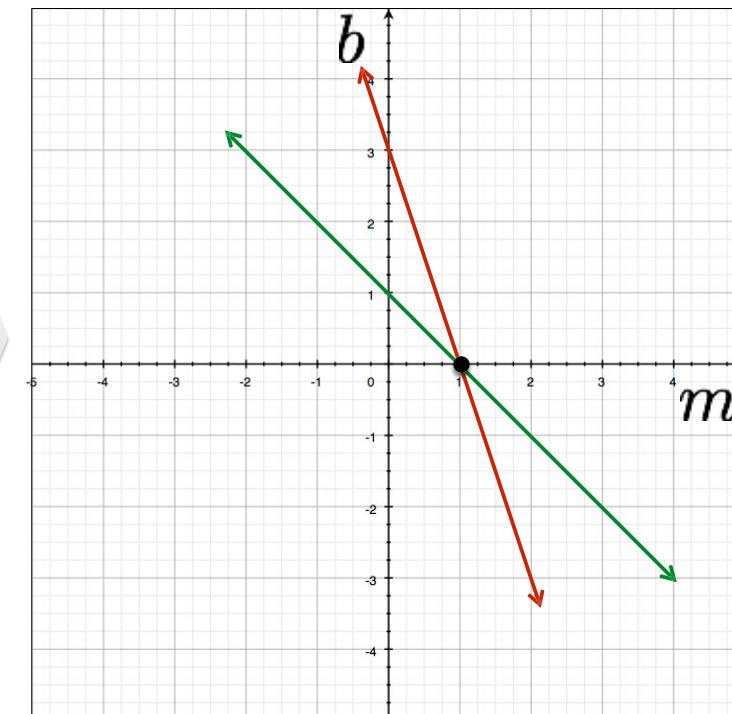


Image space

Parameter space

# Image and parameter space

variables  
 $y = mx + b$   
parameters

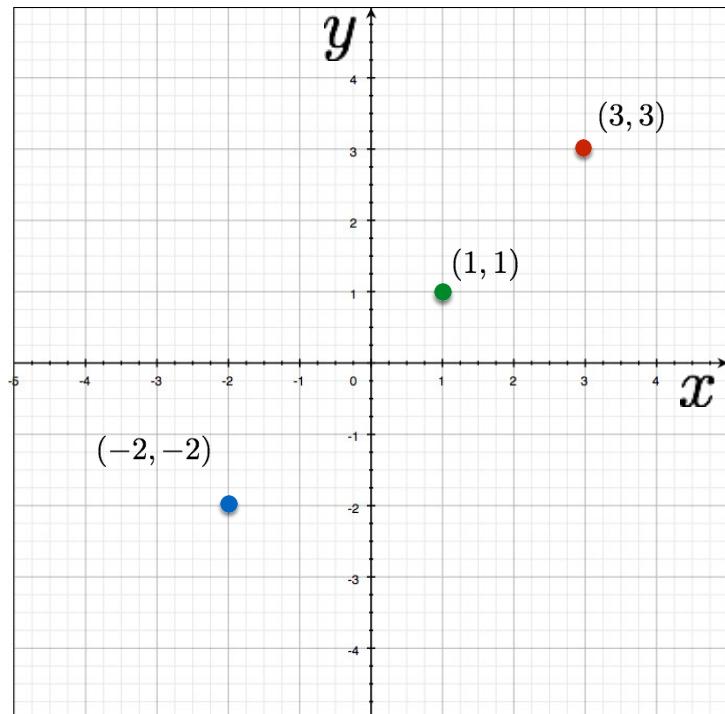
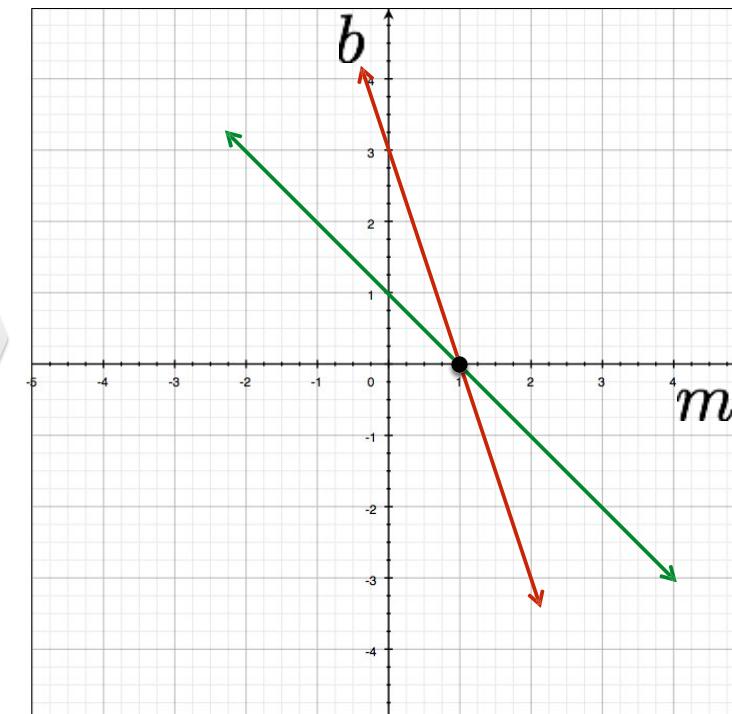


Image space

variables  
 $y - mx = b$   
parameters

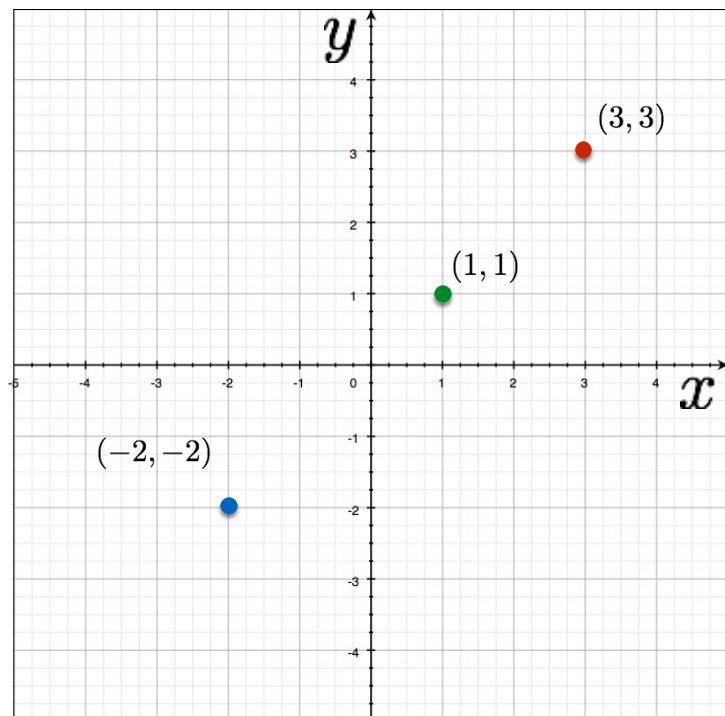


Parameter space

three points  
become  
?

# Image and parameter space

variables  
 $y = mx + b$   
parameters



three points  
become  
?

variables  
 $y - mx = b$   
parameters

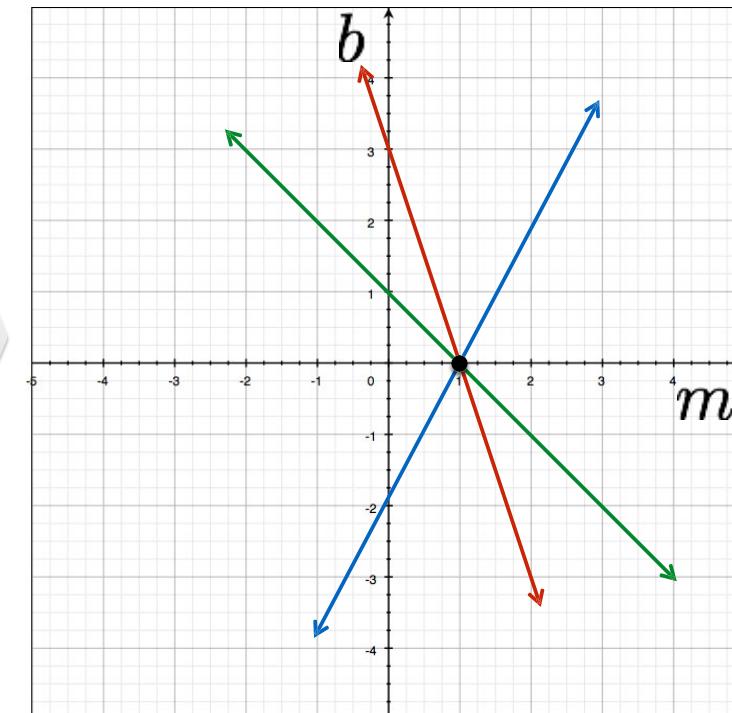


Image space

Parameter space

# Image and parameter space

variables  
 $y = mx + b$   
parameters

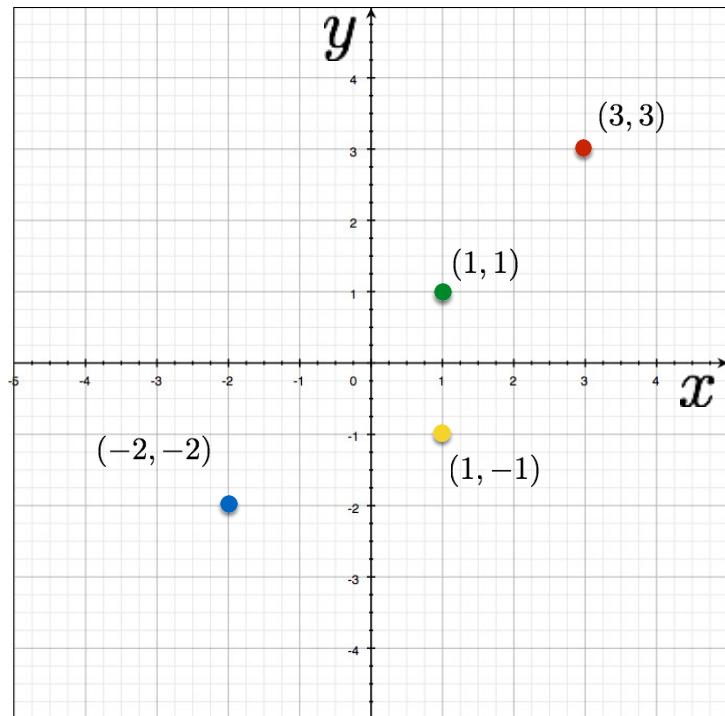
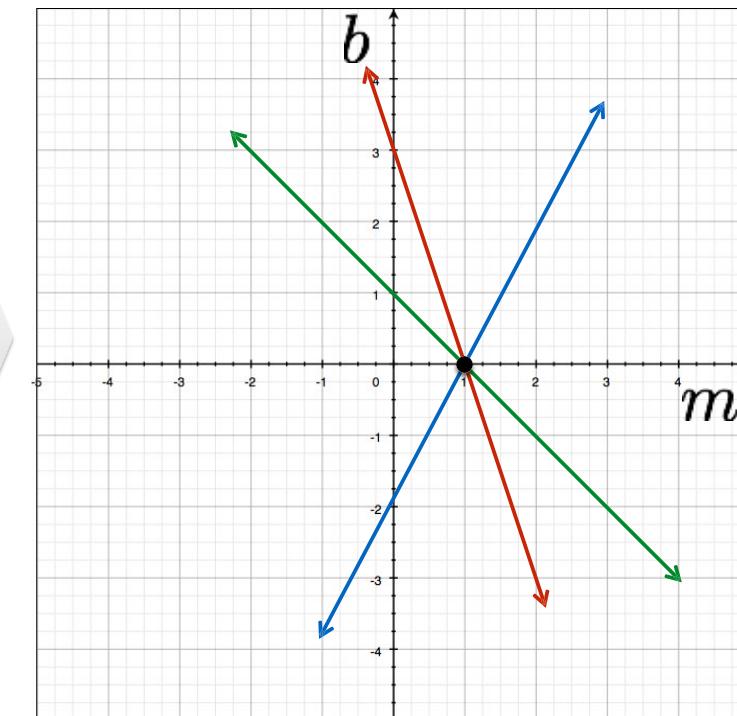


Image space

variables  
 $y - mx = b$   
parameters



Parameter space

# Image and parameter space

variables  
 $y = mx + b$   
parameters

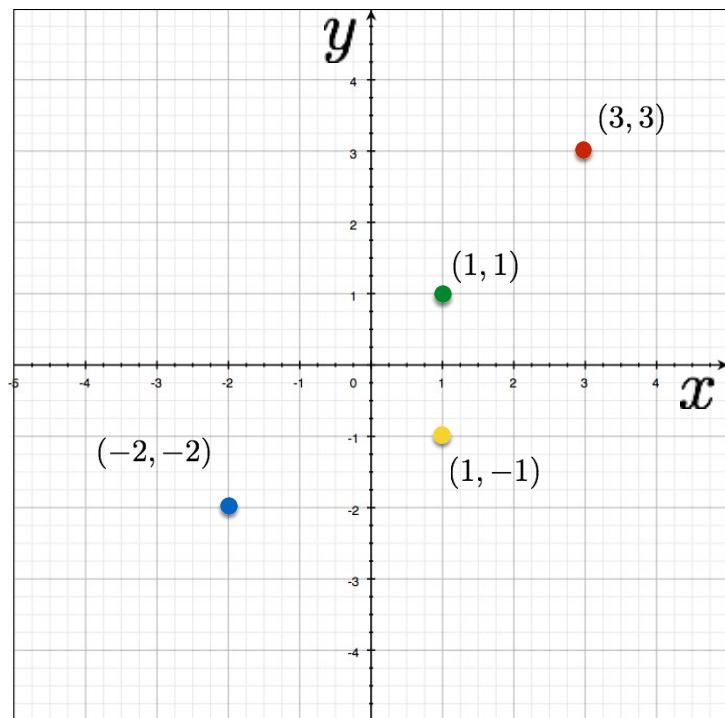
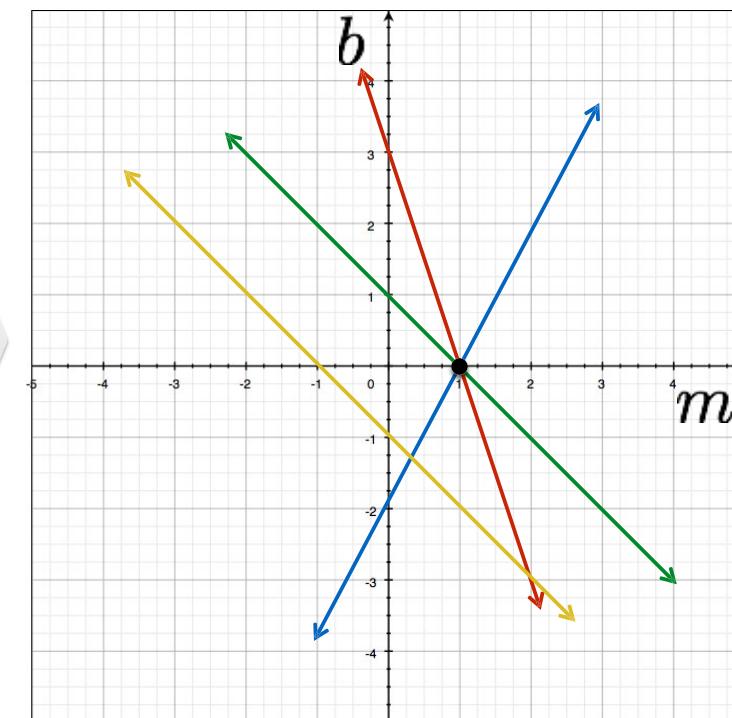


Image space

variables  
 $y - mx = b$   
parameters

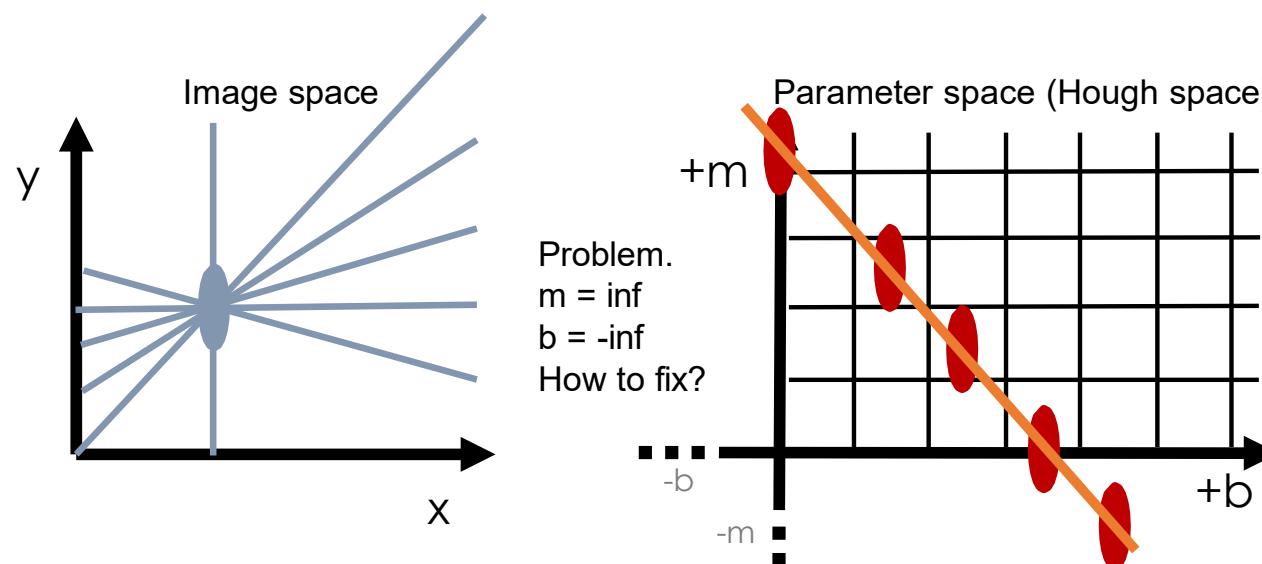


Parameter space

four points  
become  
?

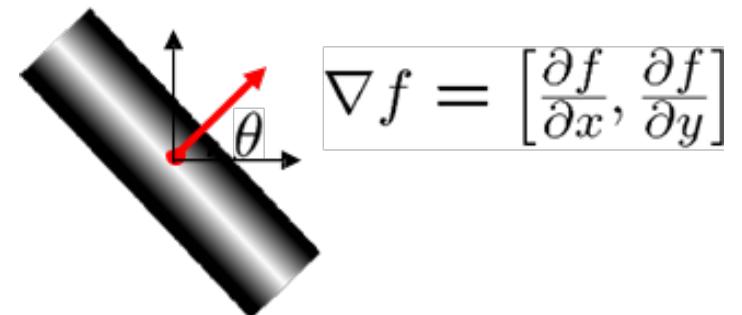
# Hough Transform: Step back

- Hough space represents all possible lines
- With gradient information constraint:
  - Edge is single point in Hough space
- Without gradient orientation information?
  - How big is Hough space?

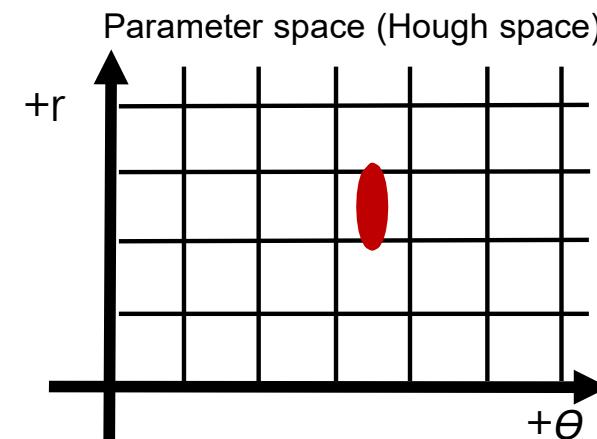
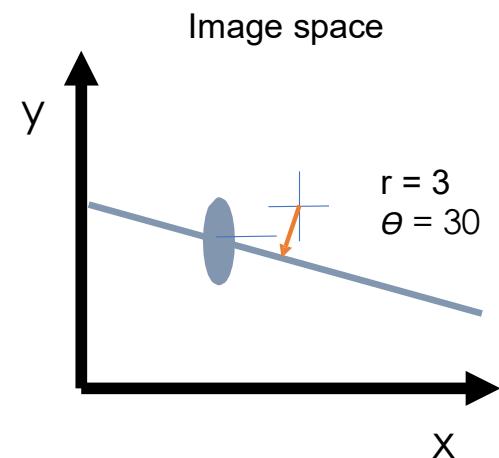


# Hough Transform: Line Normal Form

- Use  $\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$ 
  - Space is 0 to 360
- Use  $r = \text{distance to line from some origin}$ 
  - $r_i = x_i \cos \theta_i + y_i \sin \theta_i$
  - Space is  $\pm \sqrt{\max_x^2 + \max_y^2}$

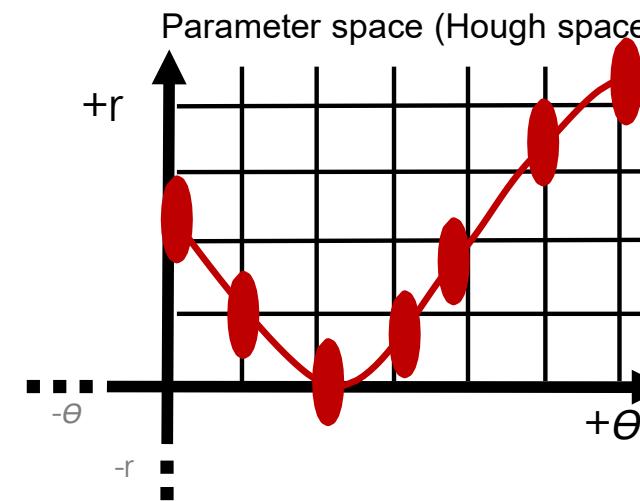
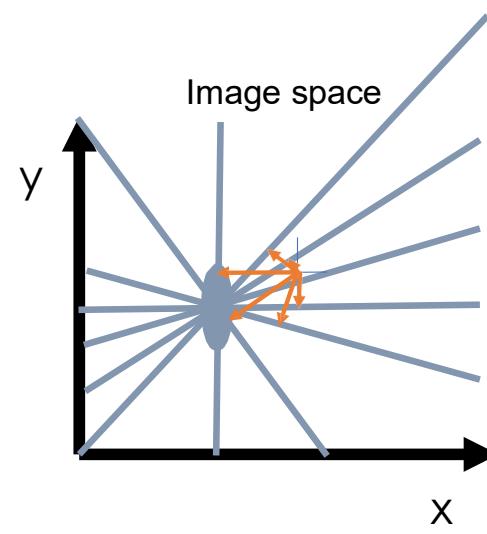


$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$



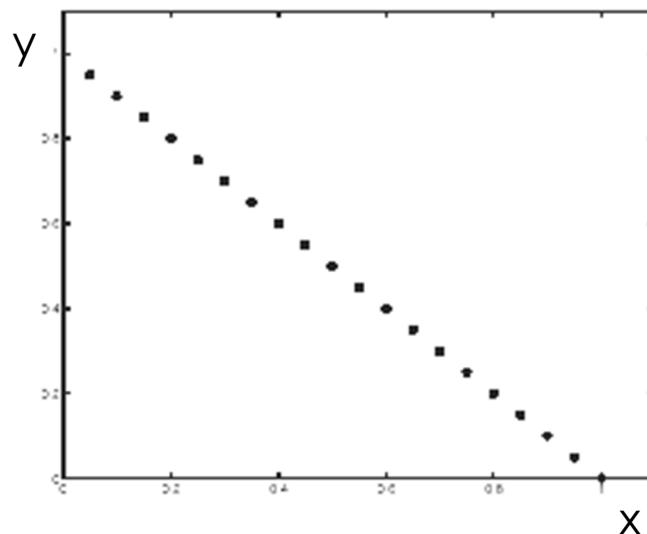
# Hough Transform: Line Normal Form

- In this line form, unoriented edge draws a sinusoid in Hough space:

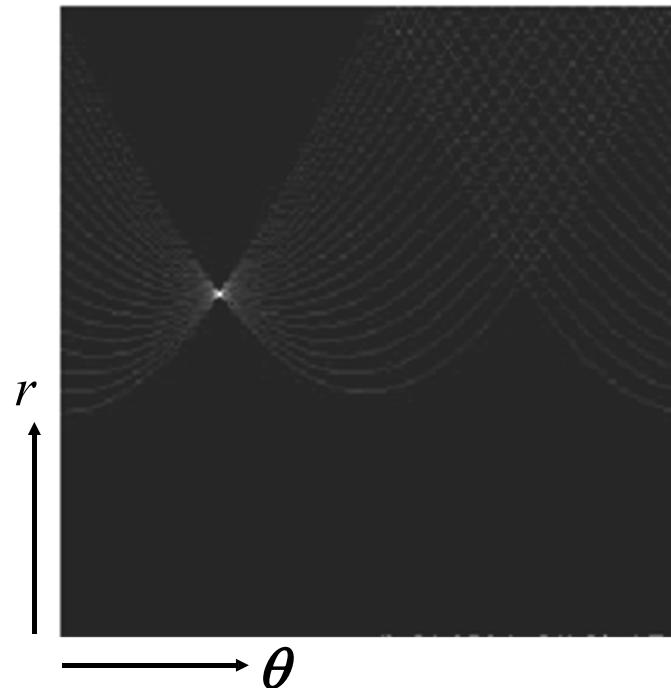


# Hough transform - experiment

Next few images *ignore* edge orientation.  
Each point is one sinusoid.

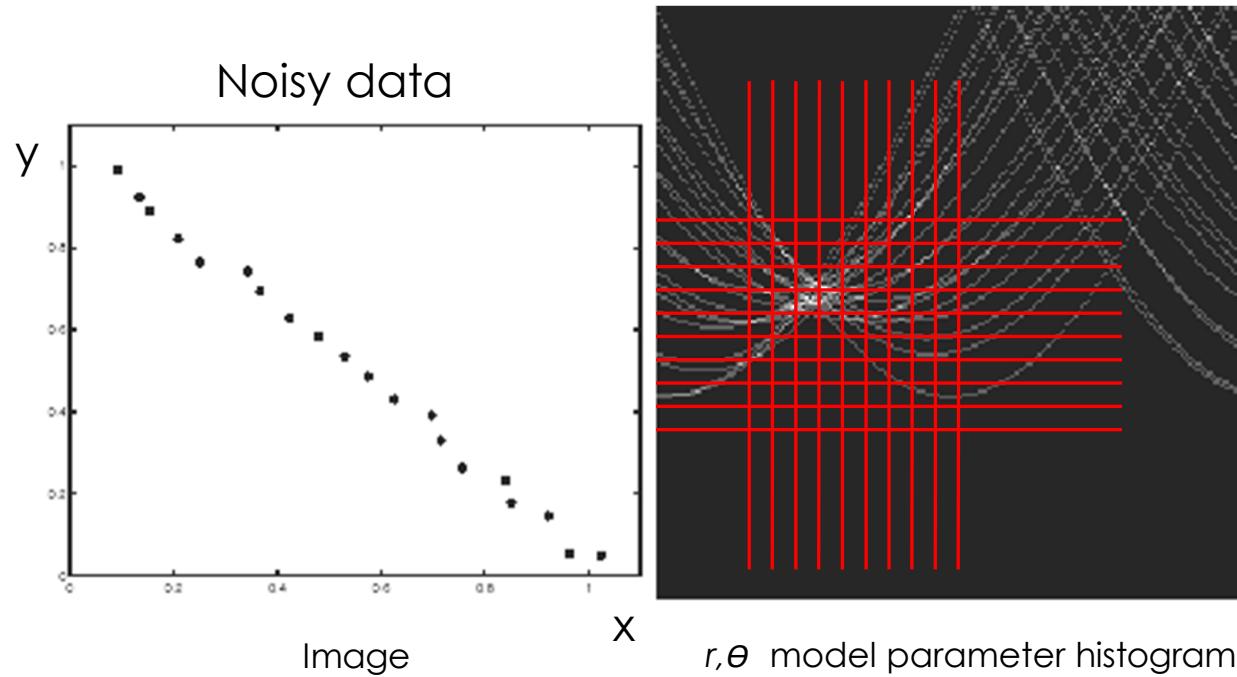


Image



$r, \theta$  model parameter histogram

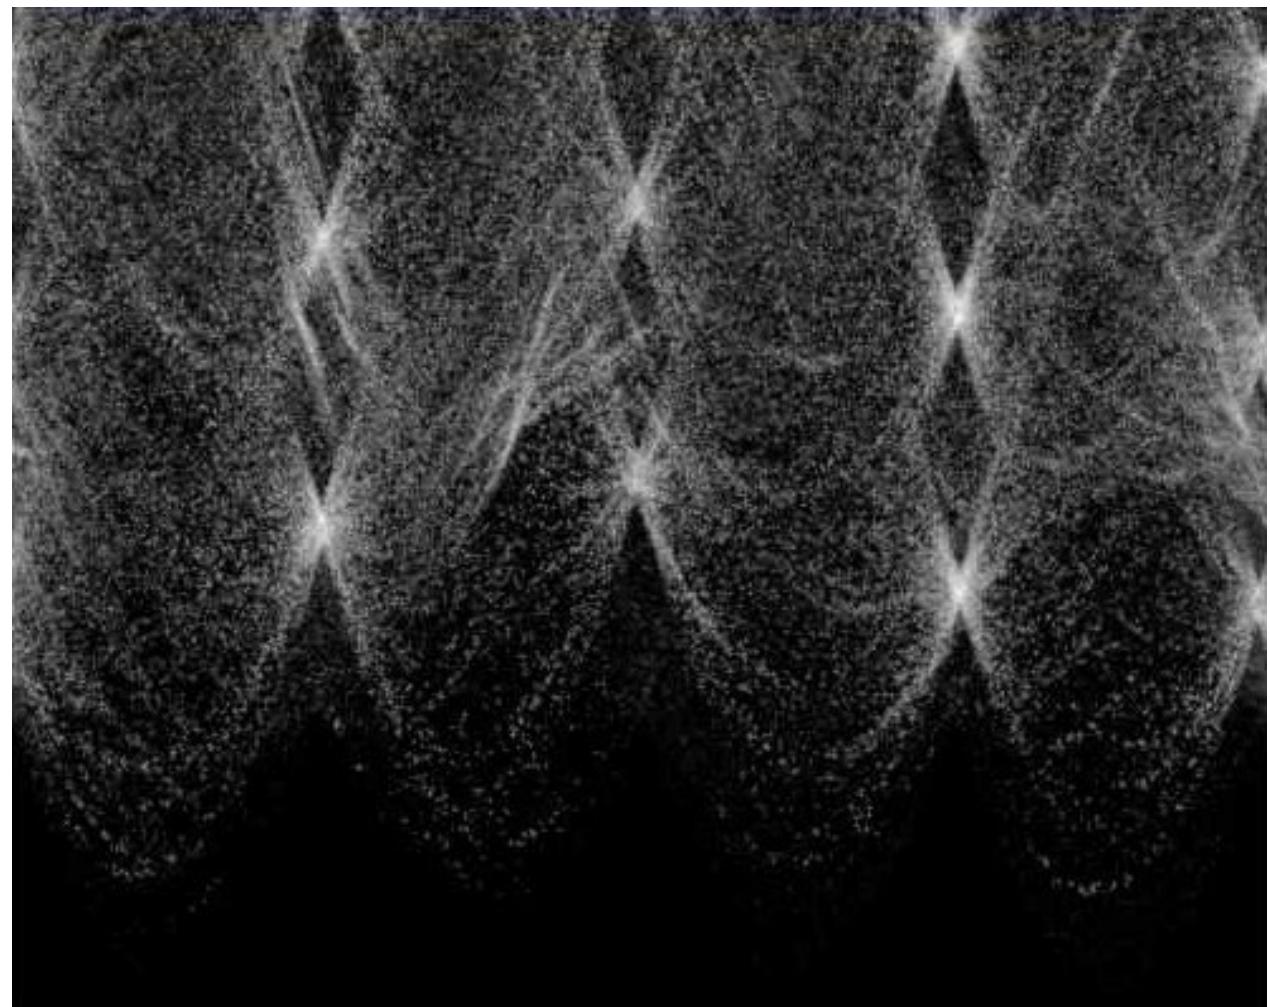
# Hough transform - experiment



Need to adjust grid size or smooth

- Practical considerations
  - Bin size
  - Smoothing
  - Finding multiple lines
  - Finding line segments

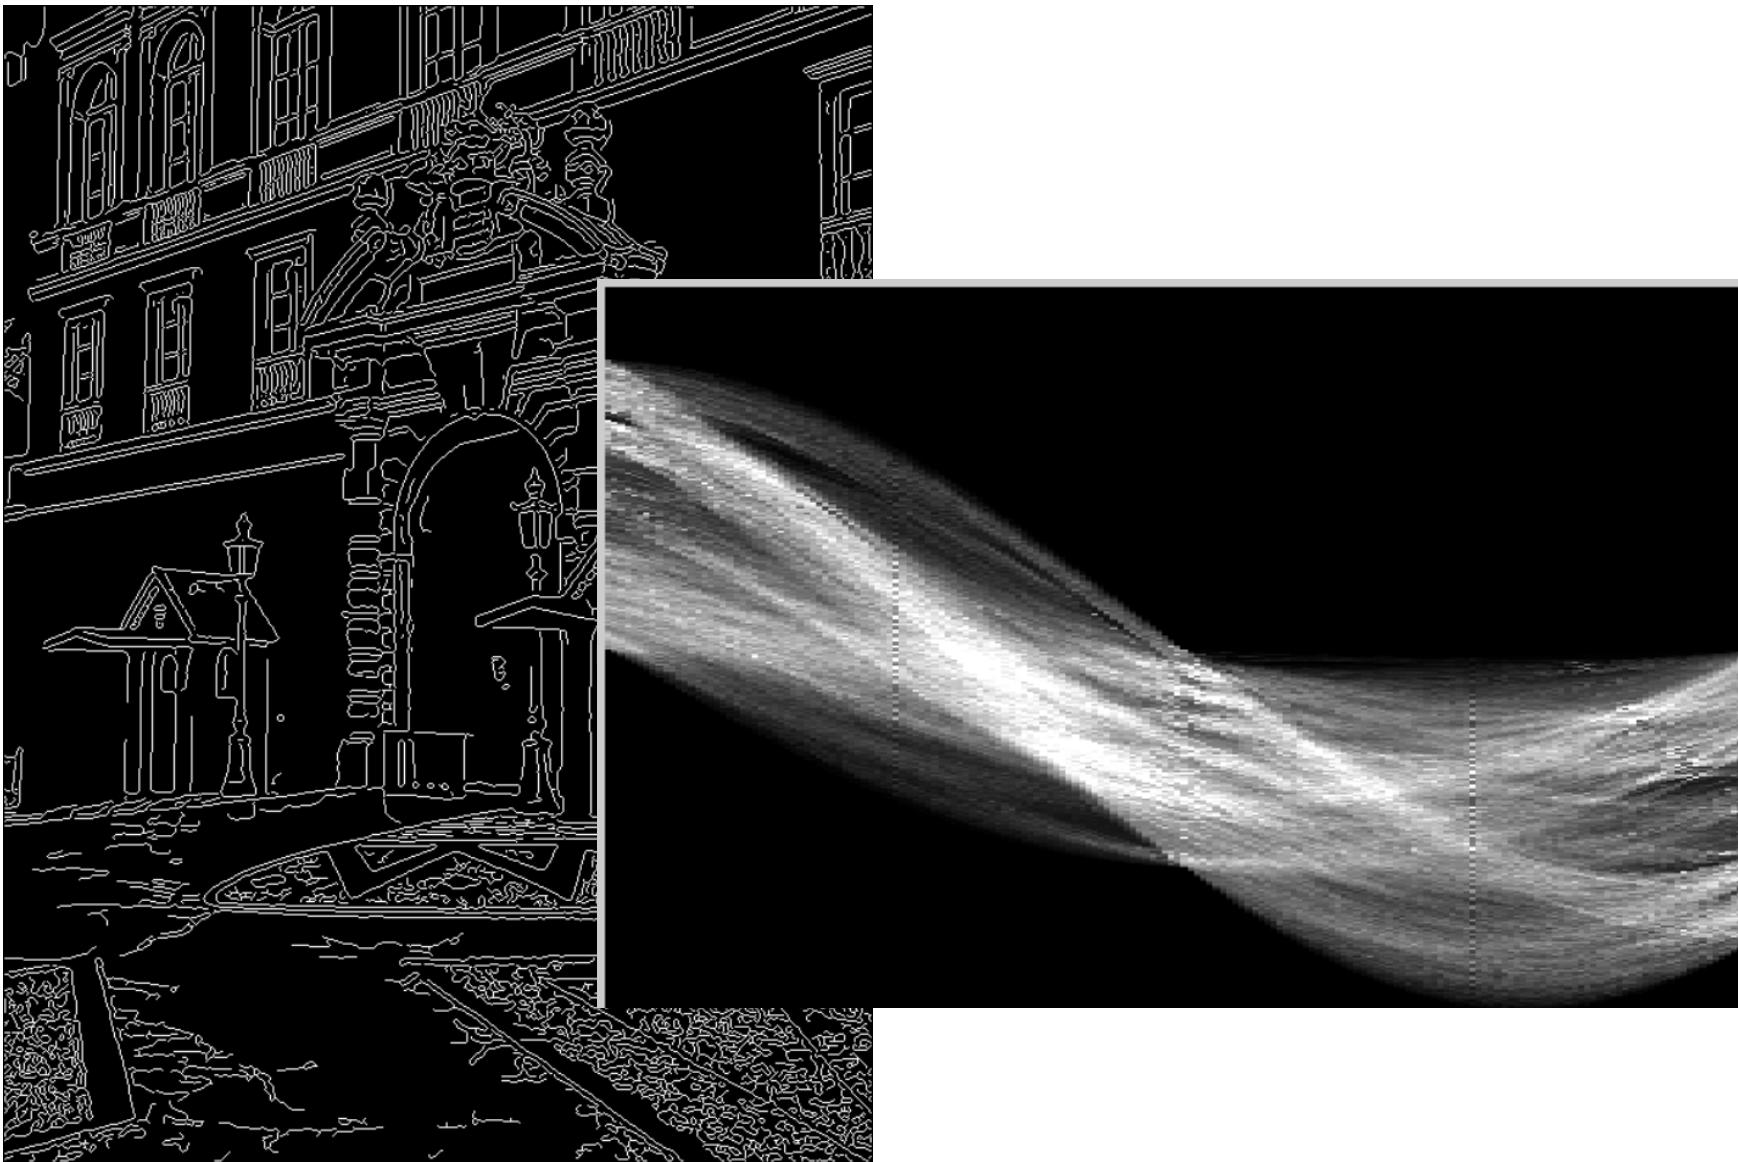
# Hough transform example



# Image → Canny

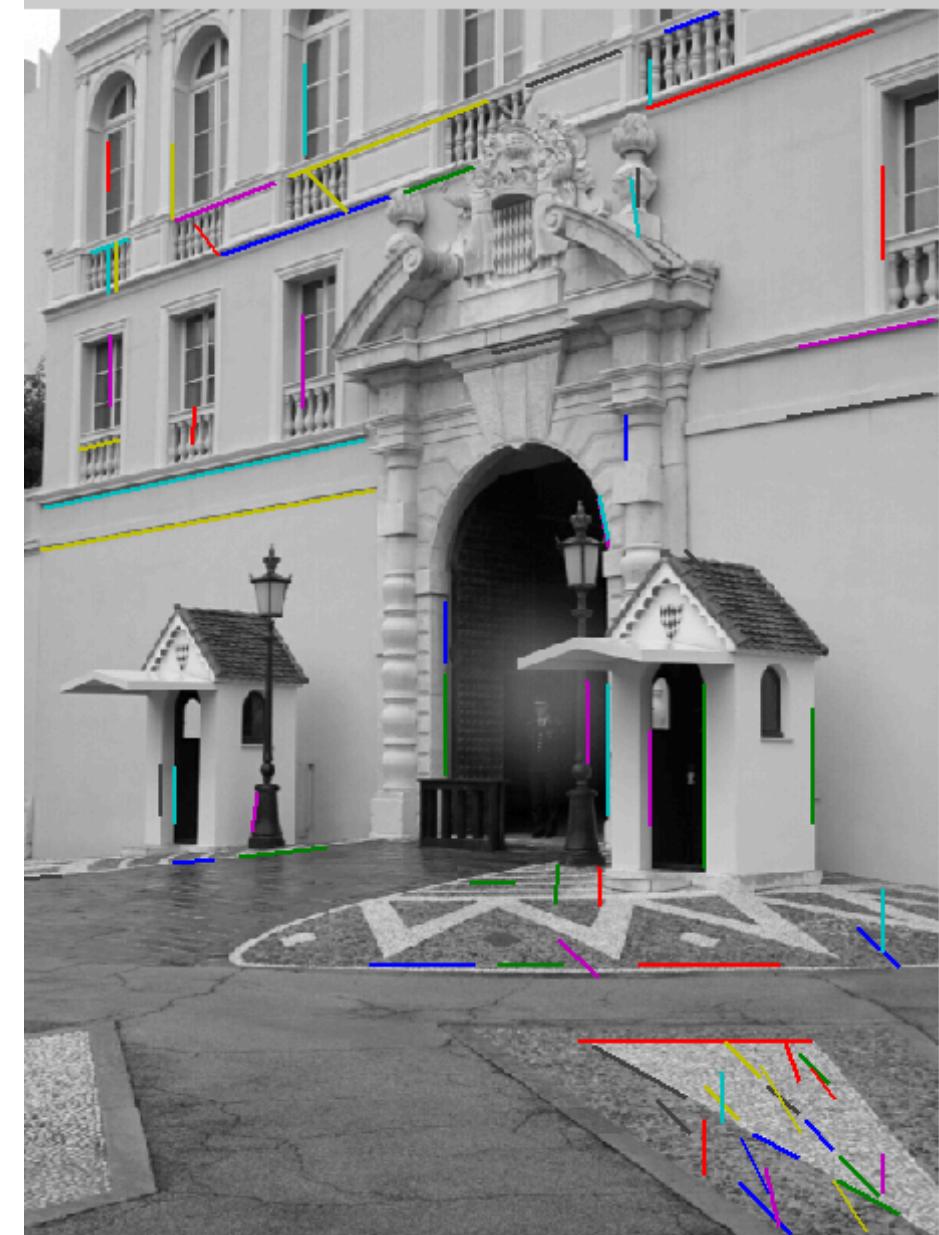
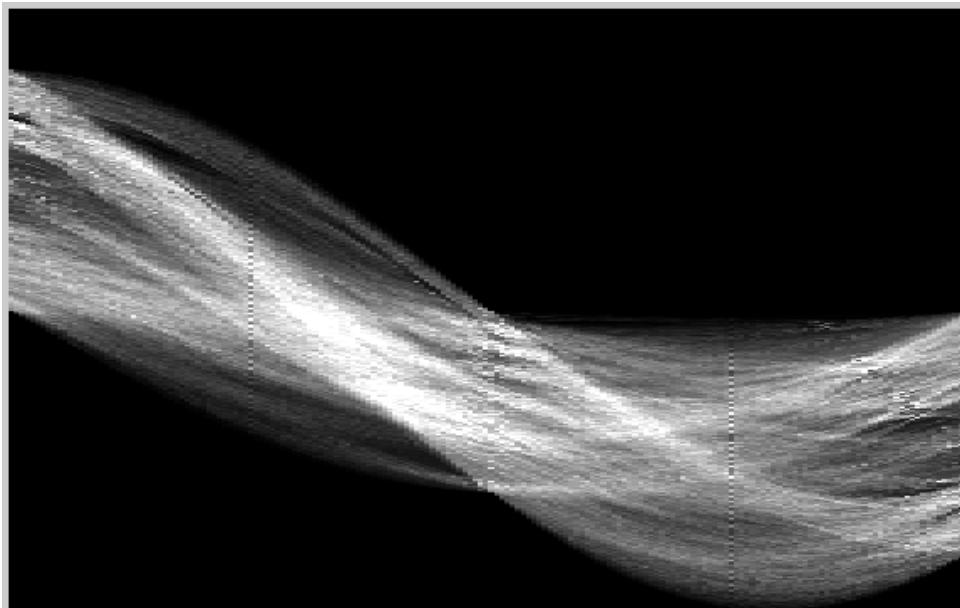


# Canny → Hough votes

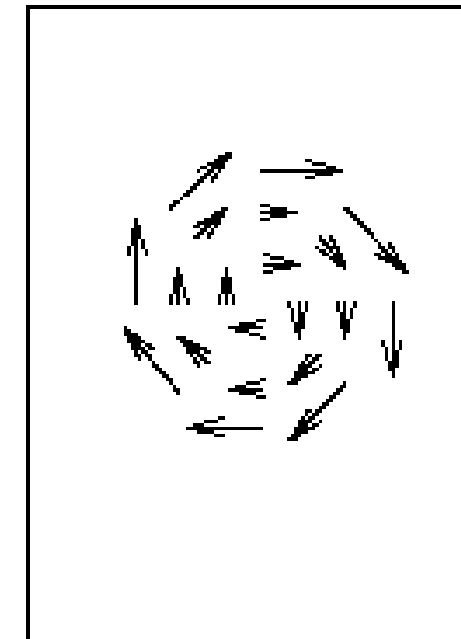
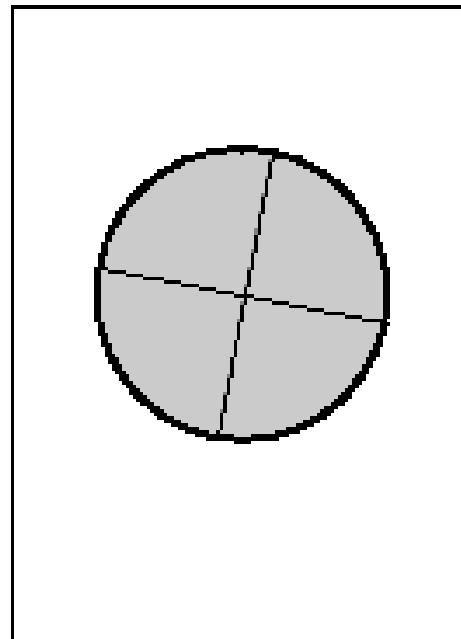
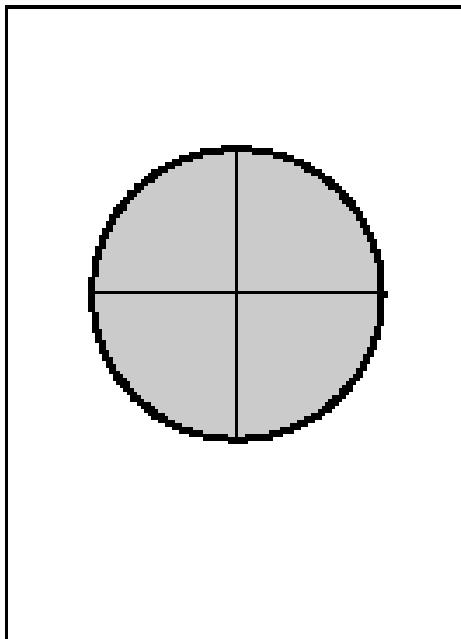


# Hough votes → Edges

Find peaks and post-process.



# Optical Flow



# Fundamental Equations of Computer Vision

I. Image Filtering

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l]$$

2. Optical Flow

$$I(x, y, t) = I(x + \Delta x, y + \Delta y, t + \Delta t)$$

3. Camera Geometry

$$\mathbf{x} = \mathbf{K}[\mathbf{R} \ \mathbf{t}] \mathbf{X} \quad x^T F x' = 0$$

4. Machine Learning

$$\arg \min_{\mathbf{S}} \sum_{i=1}^k \sum_{\mathbf{x} \in S_i} \|\mathbf{x} - \boldsymbol{\mu}_i\|^2. \quad y = \varphi(\sum_{i=1}^n w_i x_i + b) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

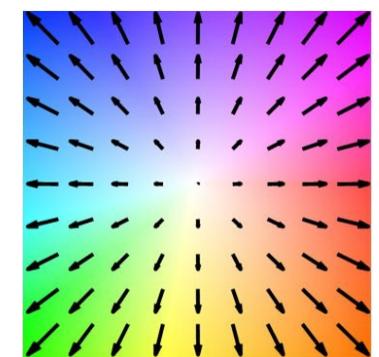
# Optical Flow: Dense Correspondence Over Time



Input [Liu *et al.* CVPR'08]

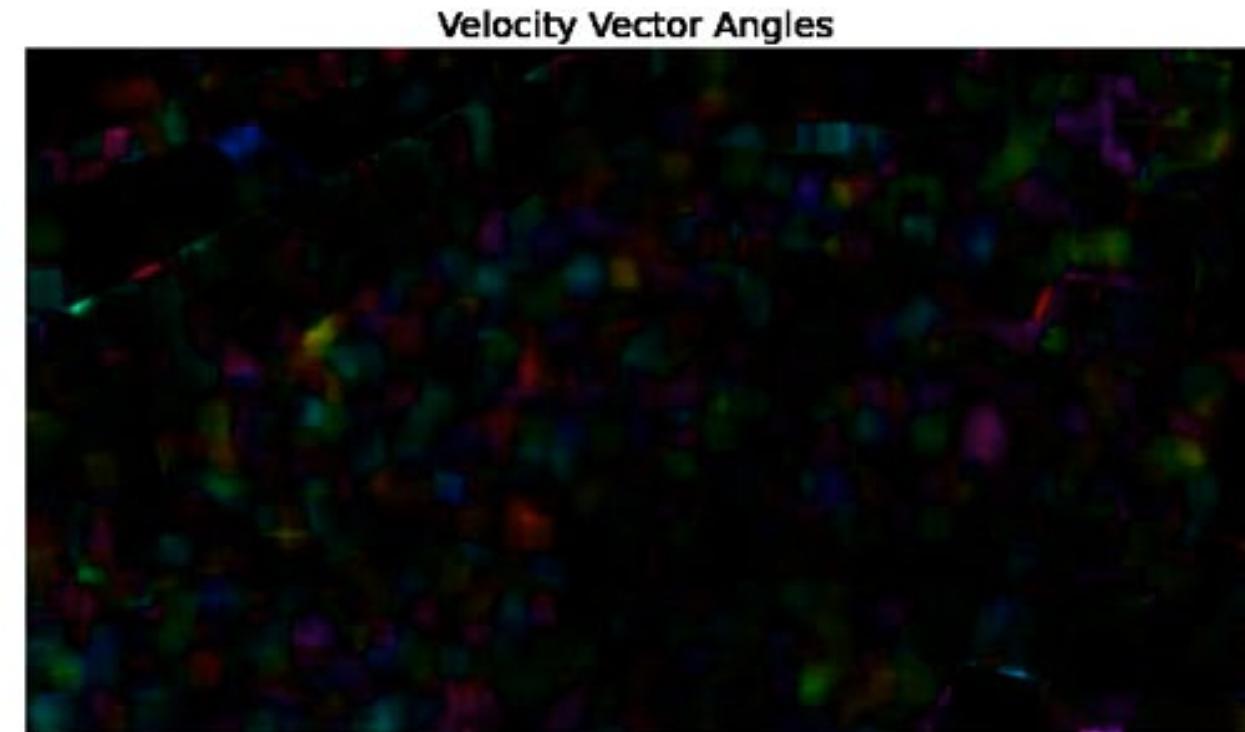


Optical flow (2D motion vector)

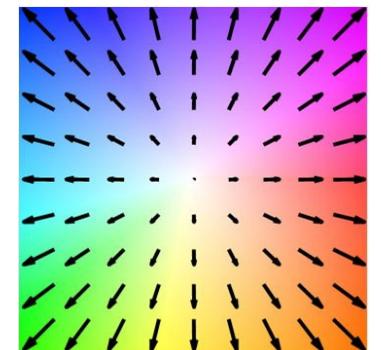


Color key  
[Baker *et al.* IJCV'11]

## Optical Flow to find the Walk Direction of people in a Video with OpenCV-python



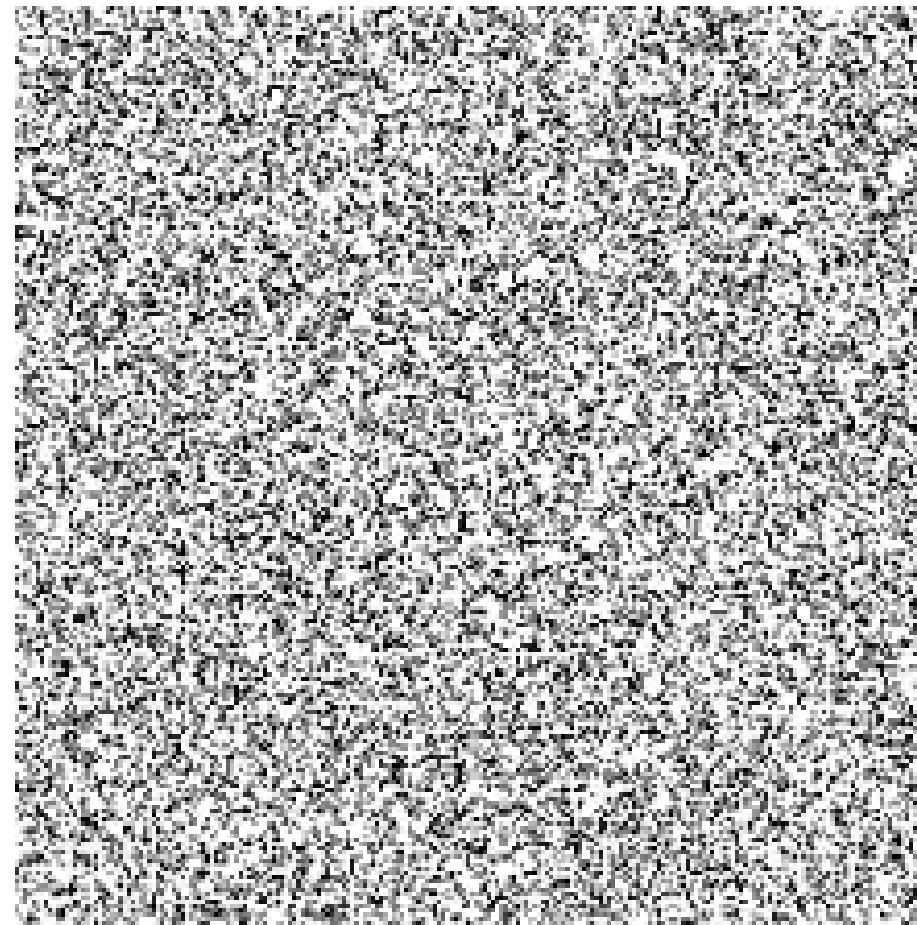
[https://www.youtube.com/watch?v=e\\_TeY6QRp4c](https://www.youtube.com/watch?v=e_TeY6QRp4c)



Color key

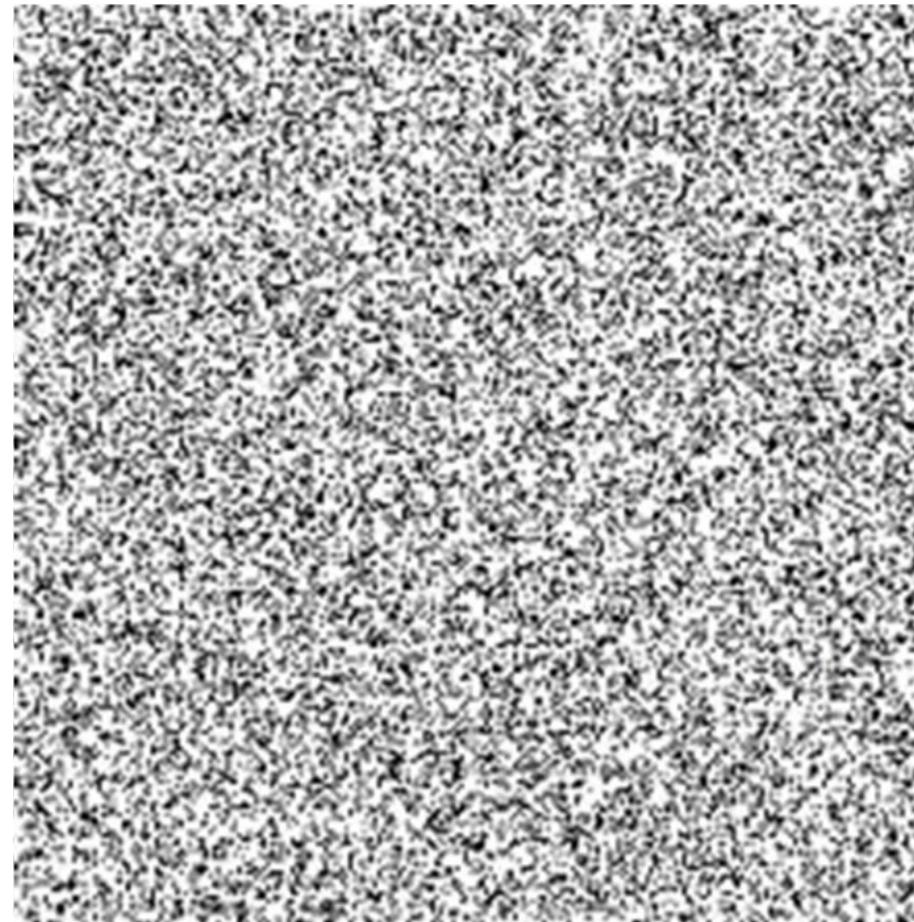
# **Correspondence and perceptual organization**

Sometimes correspondence/motion is the only cue



# **Correspondence and perceptual organization**

Sometimes correspondence/motion is the only cue





<https://www.youtube.com/watch?v=mNrqvcS0oI0>

# Applications of optical flow



# Applications of optical flow



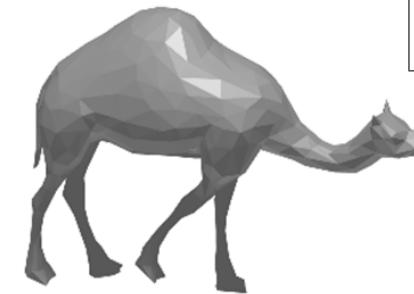
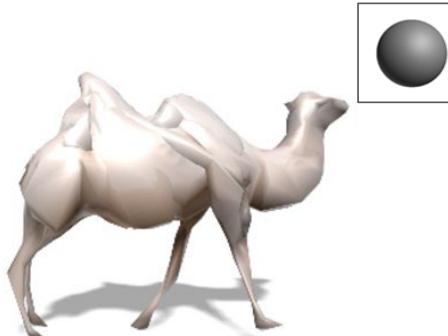
<https://youtu.be/MjViy6kyiqs?si=zXenEX6O6VJkEaQF>

Super SloMo [Jiang, Sun, et al. CVPR 2018 Spotlight] Incorporated into **NVIDIA NGX** SDK for the Turing GPU.

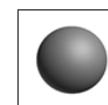
# Face Unblur for Pixel 6



# Articulated shape reconstruction from a monocular video



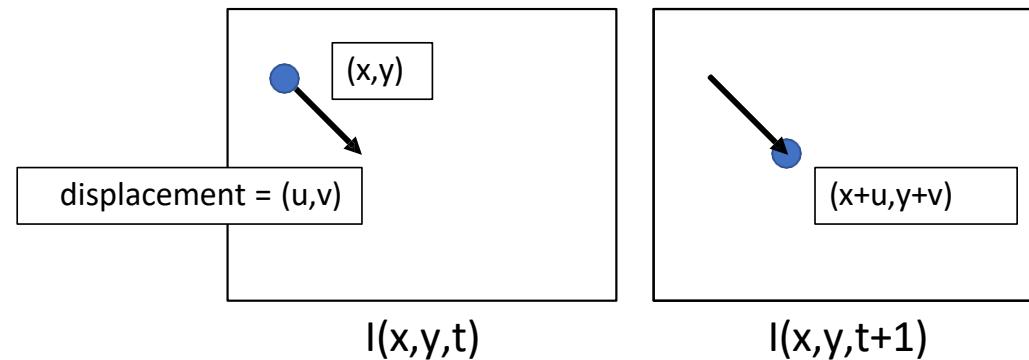
A-CSM: Kulkarni et al. CVPR 2020



VIBE: Kocabas et al. CVPR 2020

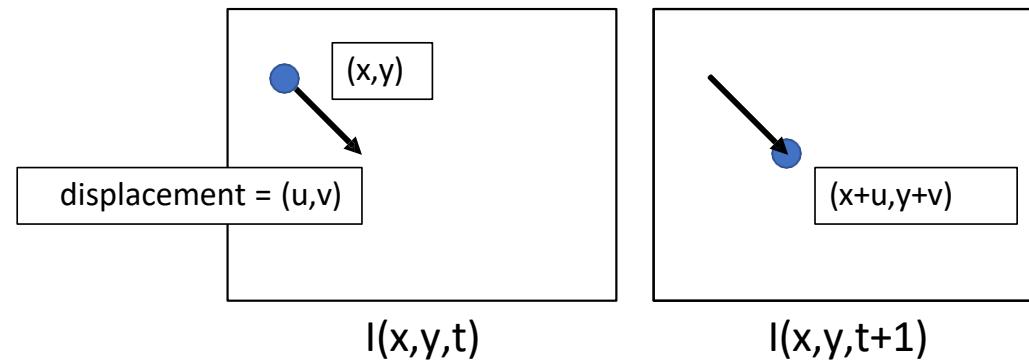
**Challenge:** Solving non-rigid 3D shape from 2D measurements without template or category prior is highly *under-constrained*

# Optical Flow



Brightness constancy:  $I(x, y, t) = I(x + u, y + v, t + 1)$

# Optical Flow



Brightness constancy:  $I(x, y, t) = I(x + u, y + v, t + 1)$

Recall Taylor Expansion:  $I(x + u, y + v, t + 1) = I(x, y, t) + I_x u + I_y v + I_t \dots$

# Optical Flow Equation

$$\begin{aligned} I(x + u, y + v, t + 1) &= I(x, y, t) \\ 0 &= I(x + u, y + v, t + 1) - I(x, y, t) \quad \text{Taylor} \\ &\approx I(x, y, t) + I_t + I_x u + I_y v - I(x, y, t) \quad \text{Expansion} \\ &= I_t + I_x u + I_y v \\ &= I_t + \nabla I \cdot [u, v] \end{aligned}$$

When is this approximation **bad?**  
**u or v big**

# Optical Flow Equation

Brightness constancy equation

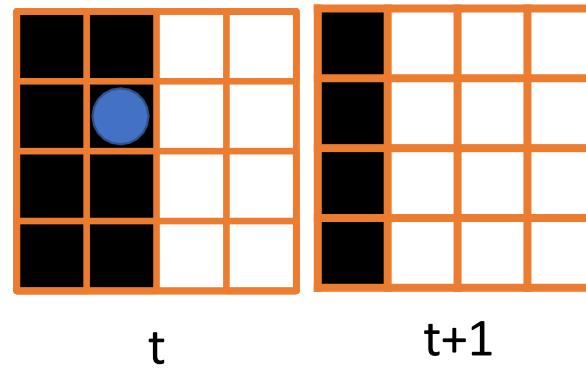
$$I_x u + I_y v + I_t = 0$$

What do static image gradients have  
to do with motion estimation?



# Brightness Constancy Example

$$I_x u + I_y v + I_t = 0$$



t

t+1

$$I_t = 1 - 0 = 1$$

@  $I_y = 0$

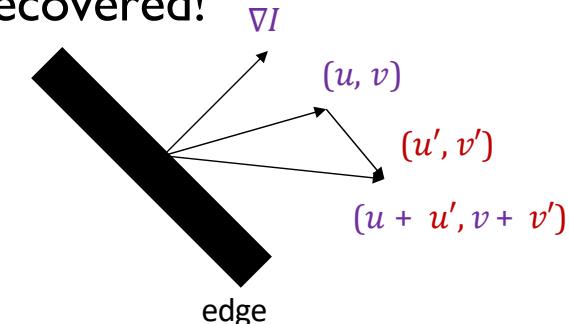
$$I_x = 1 - 0 = 1$$

**What's u?   What's v?**

# The Brightness Constancy Constraint

$$I_x u + I_y v + I_t = 0$$

- Given the gradients  $I_x$ ,  $I_y$  and  $I_t$ , can we uniquely recover the motion  $(u, v)$ ?
  - One equation, two unknowns
  - Suppose  $(u, v)$  satisfies the constraint:  $\nabla I \cdot (u, v) + I_t = 0$
  - Then  $\nabla I \cdot (u + u', v + v') + I_t = 0$  for any  $(u', v')$  s.t.  $\nabla I \cdot (u', v') = 0$
  - Interpretation: the component of the flow perpendicular to the gradient (i.e., parallel to the edge) cannot be recovered!



# Solving Ambiguity – Lucas Kanade

- 2 unknowns  $[u, v]$ , 1 eqn per pixel How do we get more equations?
- Assume **spatial coherence**: pixel's neighbors have move together / have same  $[u, v]$
- 5x5 window gives 25 new equations

$$I_t + I_x u + I_y v = 0$$
$$\begin{bmatrix} I_x(p_1) & I_y(p_1) \\ \vdots & \vdots \\ I_x(p_{25}) & I_y(p_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(p_1) \\ \vdots \\ I_t(p_{25}) \end{bmatrix}$$

B. Lucas and T. Kanade. An iterative image registration technique with an application to stereo vision. In *Proceedings of the International Joint Conference on Artificial Intelligence*, pp. 674–679, 1981.

# Solving for $\mathbf{u}, \mathbf{v}$

$$\begin{bmatrix} I_x(p_1) & I_y(p_1) \\ \vdots & \vdots \\ I_x(p_{25}) & I_y(p_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(p_1) \\ \vdots \\ I_t(p_{25}) \end{bmatrix}$$
$$\underset{25 \times 2}{A} \underset{2 \times 1}{d} = \underset{25 \times 1}{b}$$

What's the solution?

$$(A^T A) \mathbf{d} = A^T \mathbf{b} \rightarrow \mathbf{d} = (A^T A)^{-1} A^T \mathbf{b}$$

Intuitively, need to solve (sum over pixels in window)

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$
$$A^T A \quad A^T b$$

# Challenges for traditional methods



Large motion, motion blur



Textureless regions

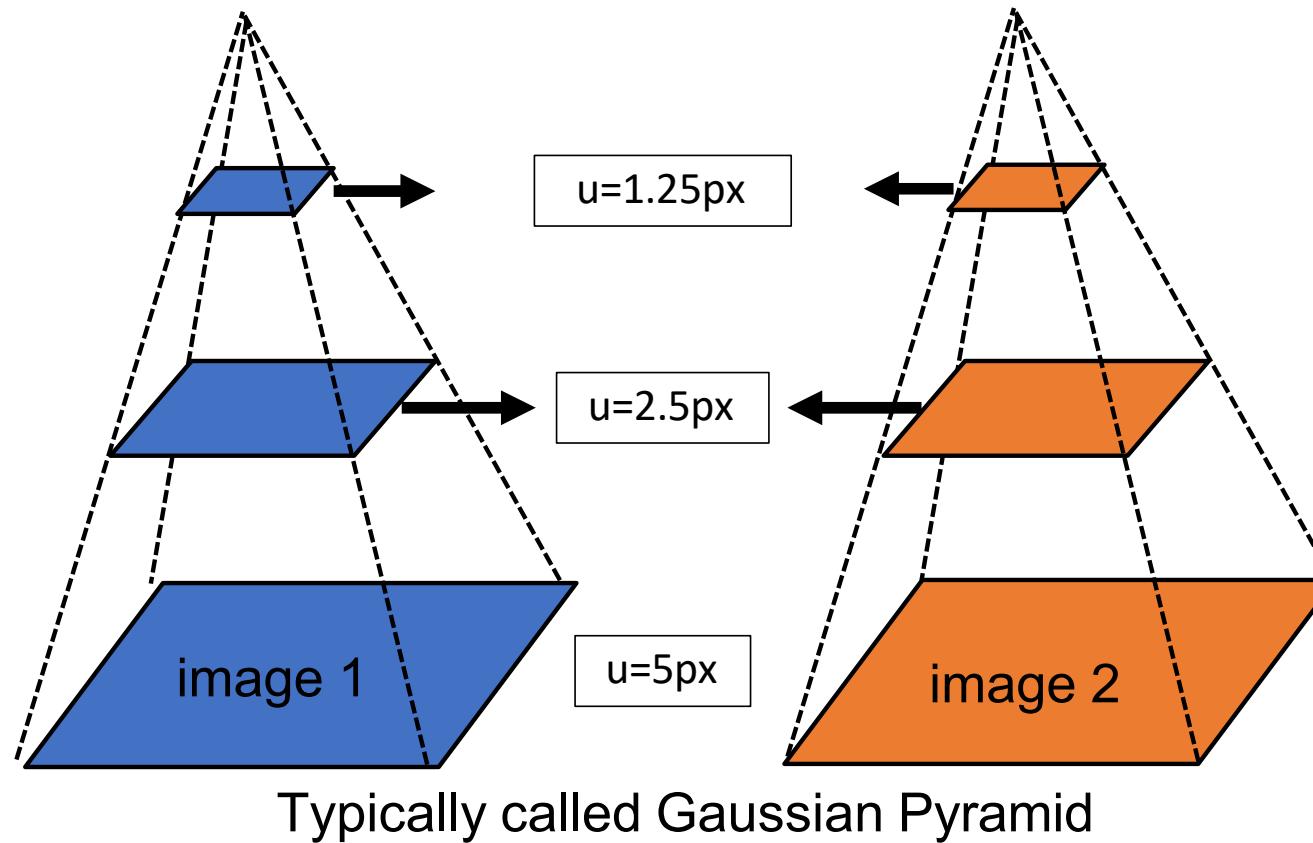


Occlusions

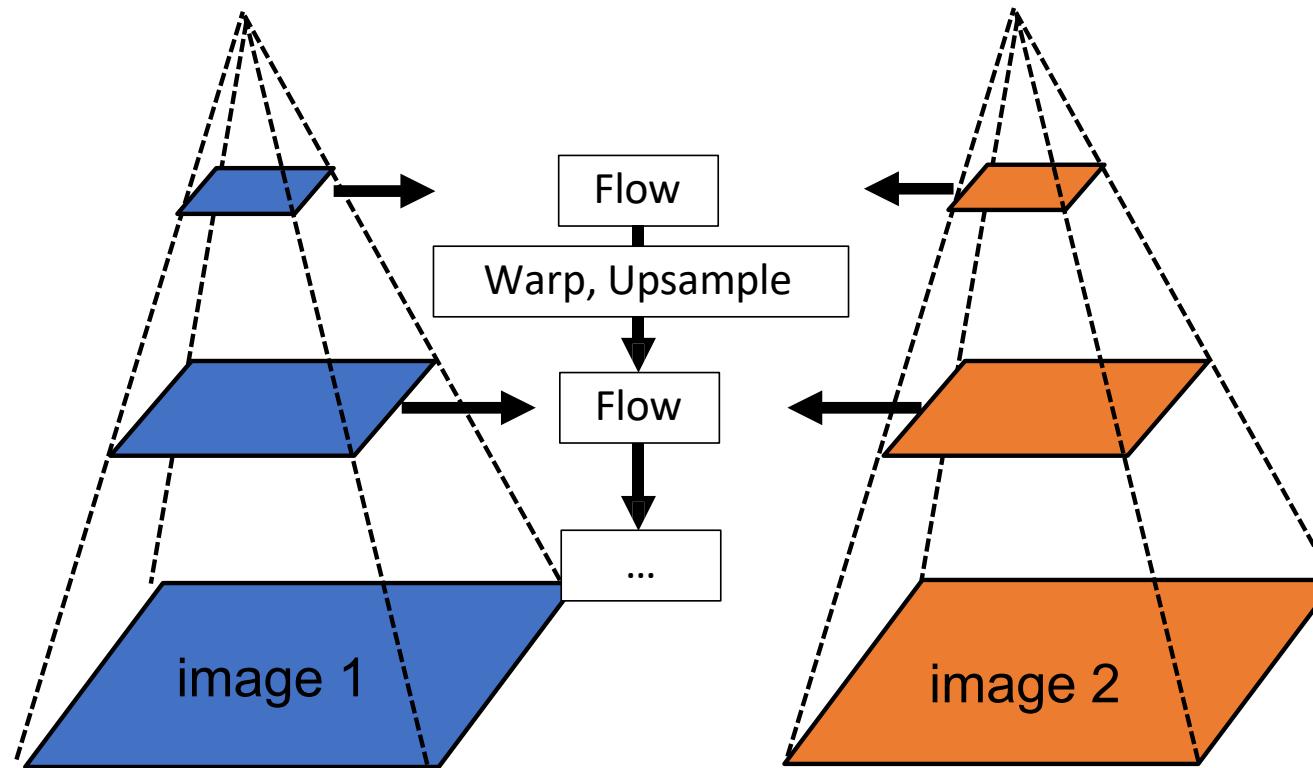


Lighting changes, noise ...

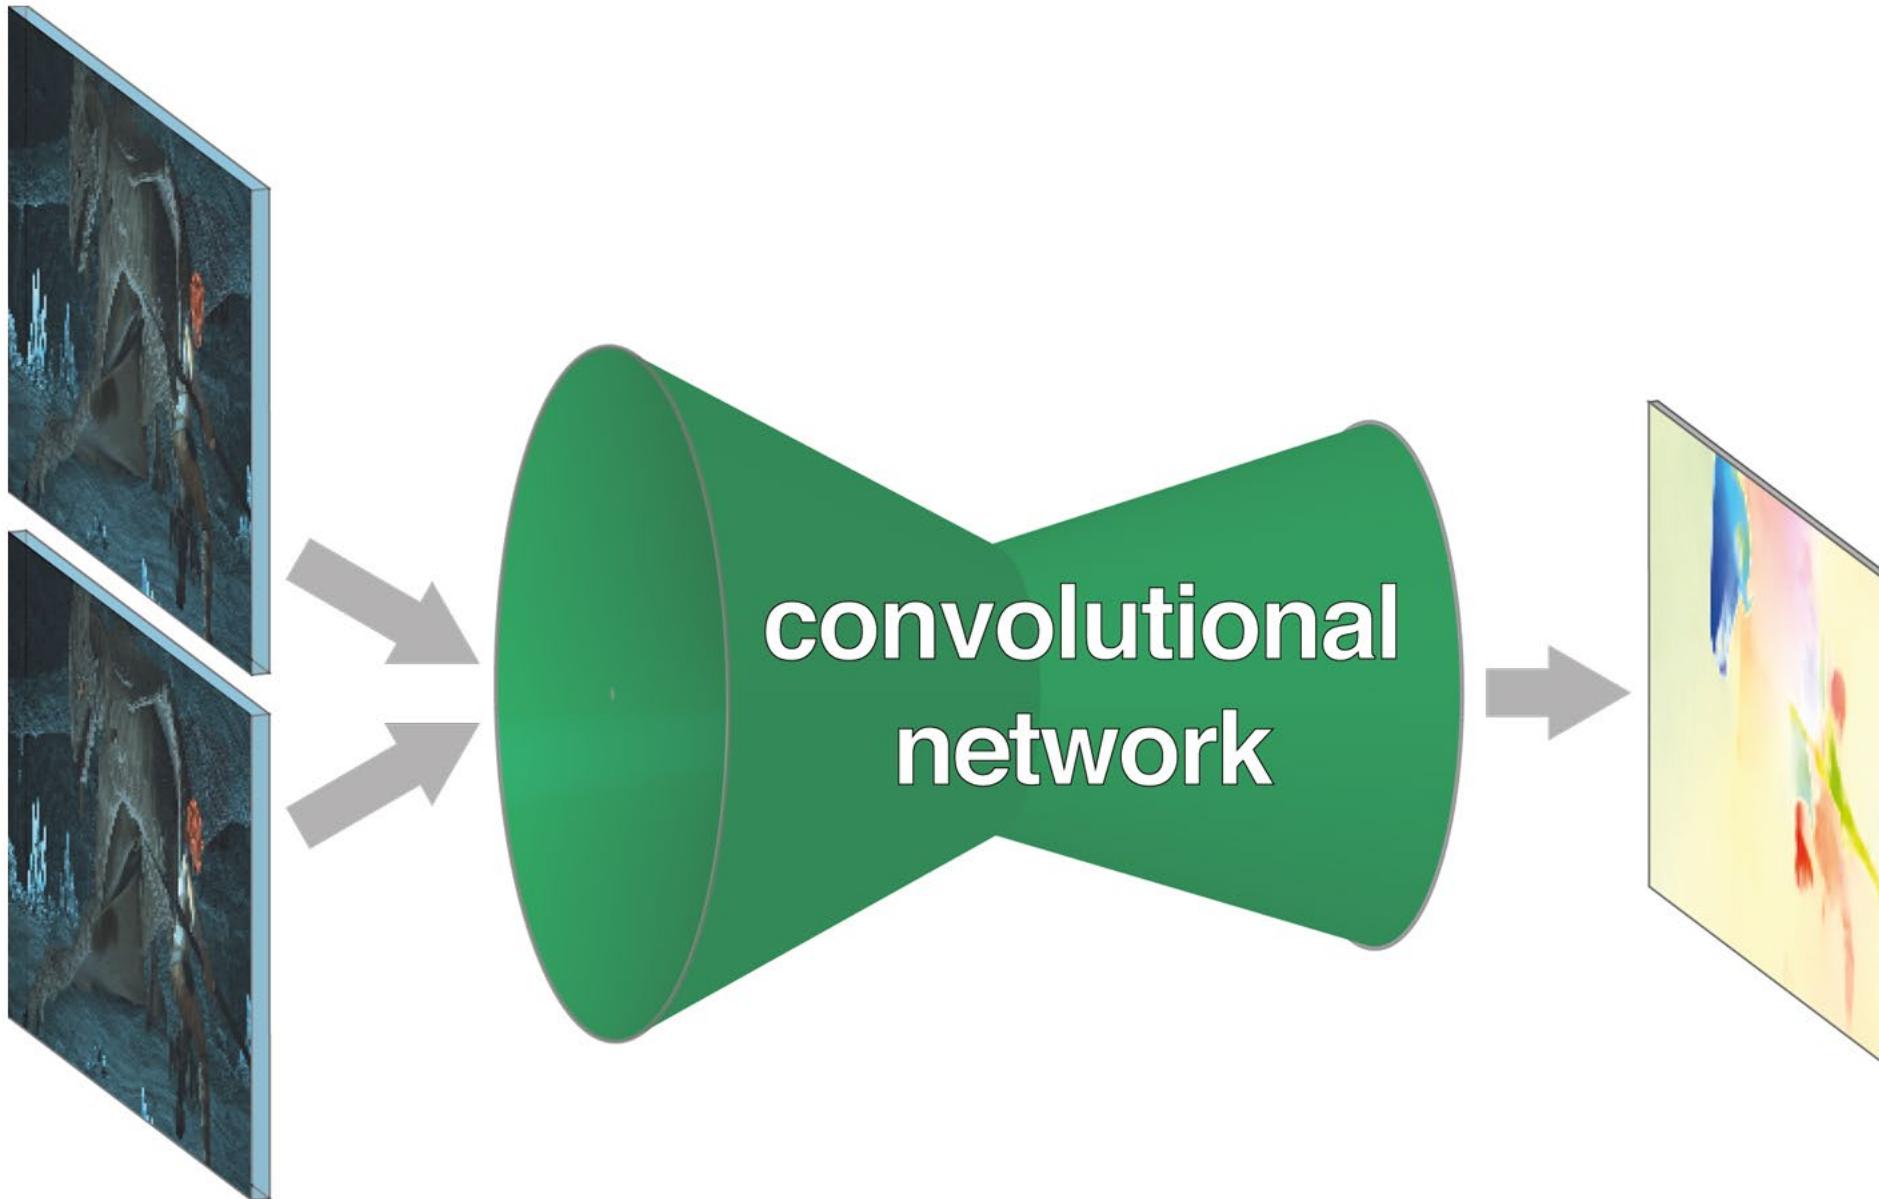
# Coarse-to-fine optical flow estimation



# Coarse-to-fine optical flow estimation



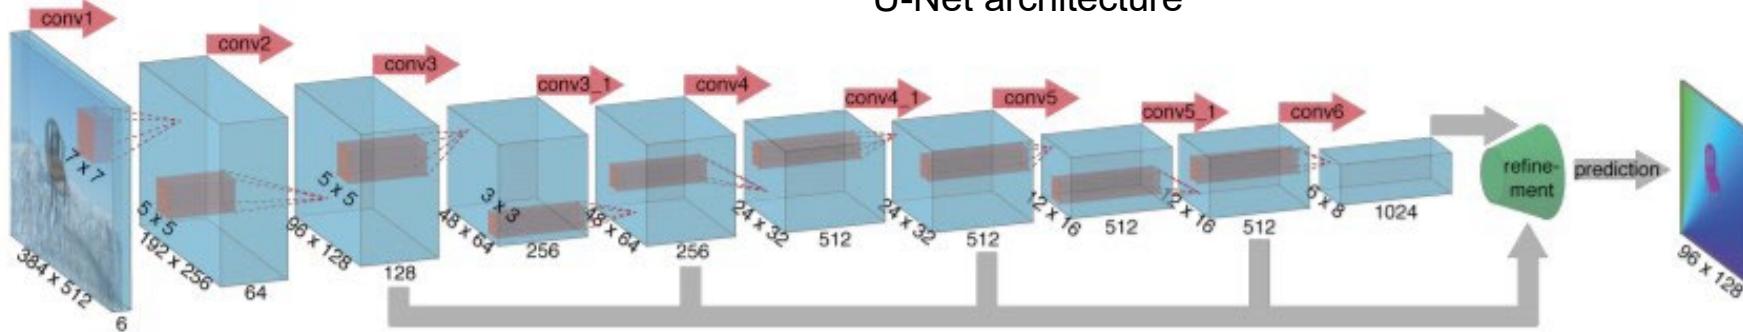
# Typical DL Pipeline for Optical Flow



# FlowNetS: mapping from images to flow

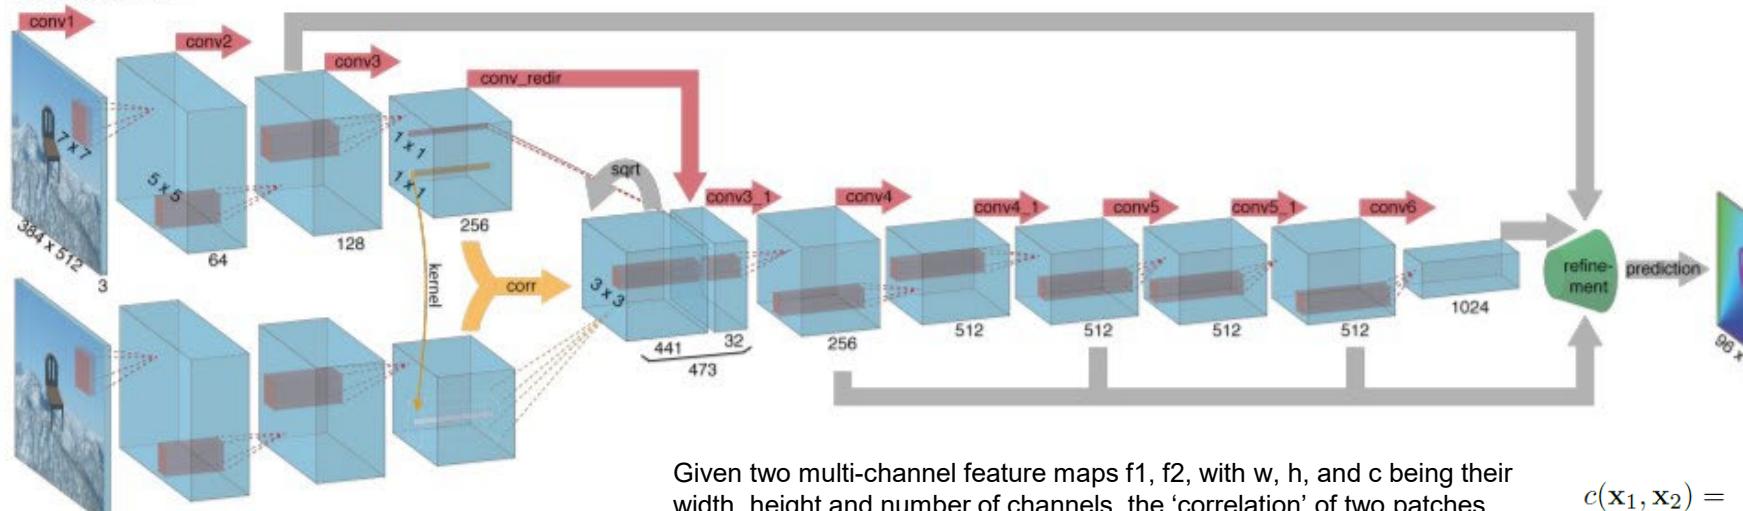
[Dosovitskiy et al. ICCV'15]

FlowNetSimple



U-Net architecture

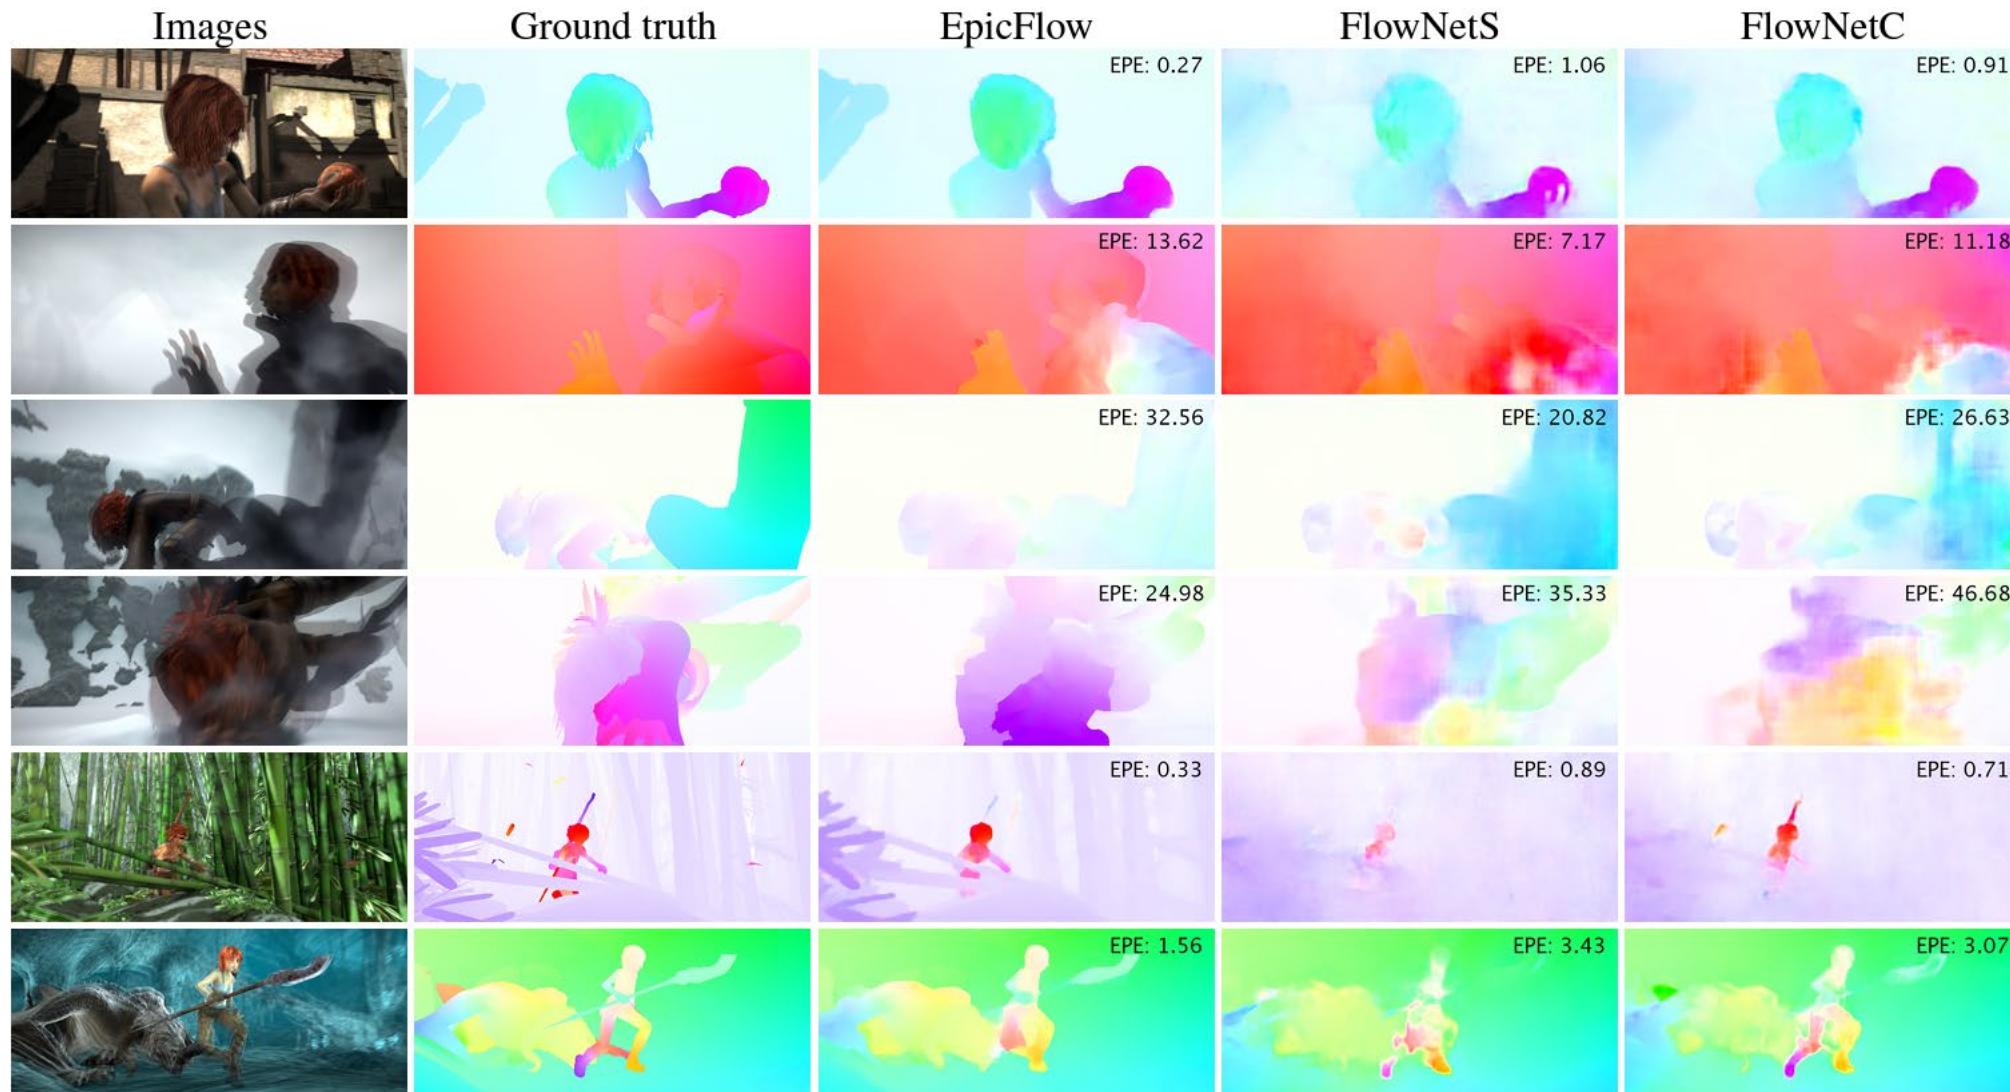
FlowNetCorr



Given two multi-channel feature maps  $f_1, f_2$ , with  $w, h$ , and  $c$  being their width, height and number of channels, the 'correlation' of two patches centered at  $x_1$  in the first map and  $x_2$  in the second map is then defined as:

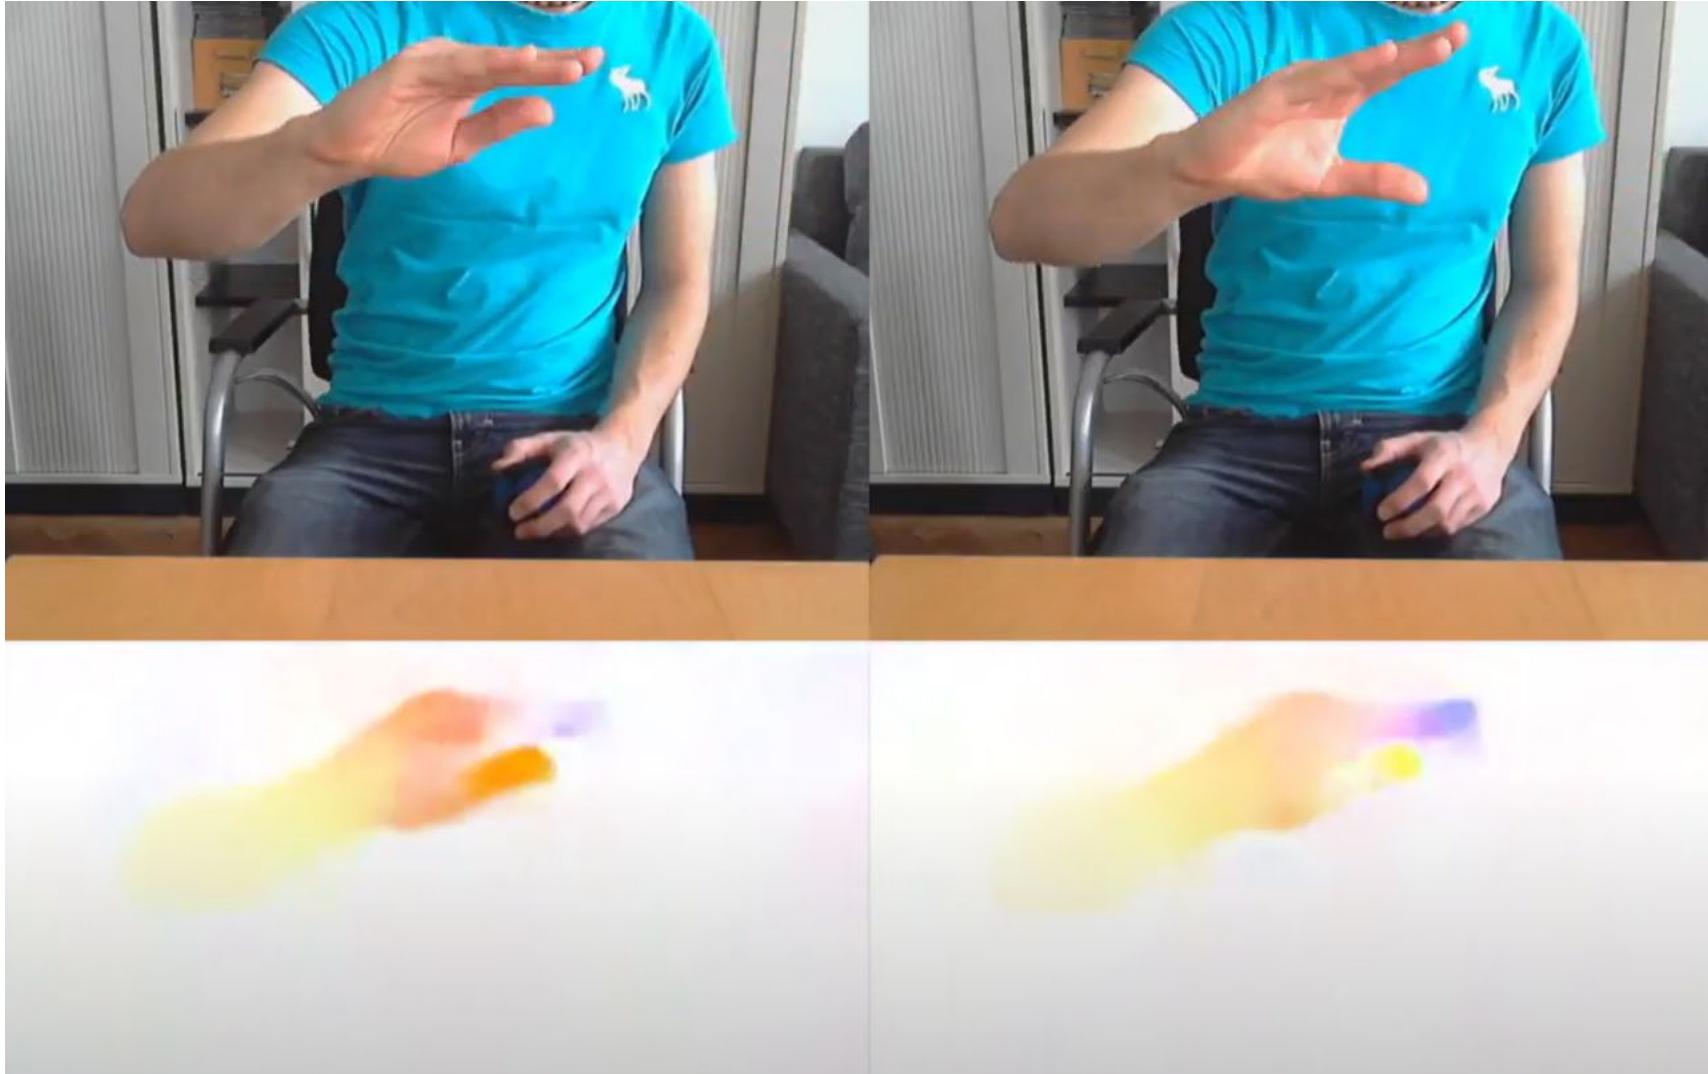
$$c(x_1, x_2) = \sum_{o \in [-k, k] \times [-k, k]} \langle f_1(x_1 + o), f_2(x_2 + o) \rangle$$

# FlowNetS: mapping from images to flow



**EPE:** endpoint error

# FlowNetS: mapping from images to flow



[https://youtu.be/k\\_wkDLJ8lJE?si=wJj1RePbw-NA75lU](https://youtu.be/k_wkDLJ8lJE?si=wJj1RePbw-NA75lU)